

Towards Ultra-Reliable CPS: Reliability Analysis of Distributed Real-Time Systems

Arpan Gujarati



MAX PLANCK INSTITUTE
FOR SOFTWARE SYSTEMS

Cyber-Physical Systems (CPS)

*Integration of **computation**, **networking**,
and **physical processes****

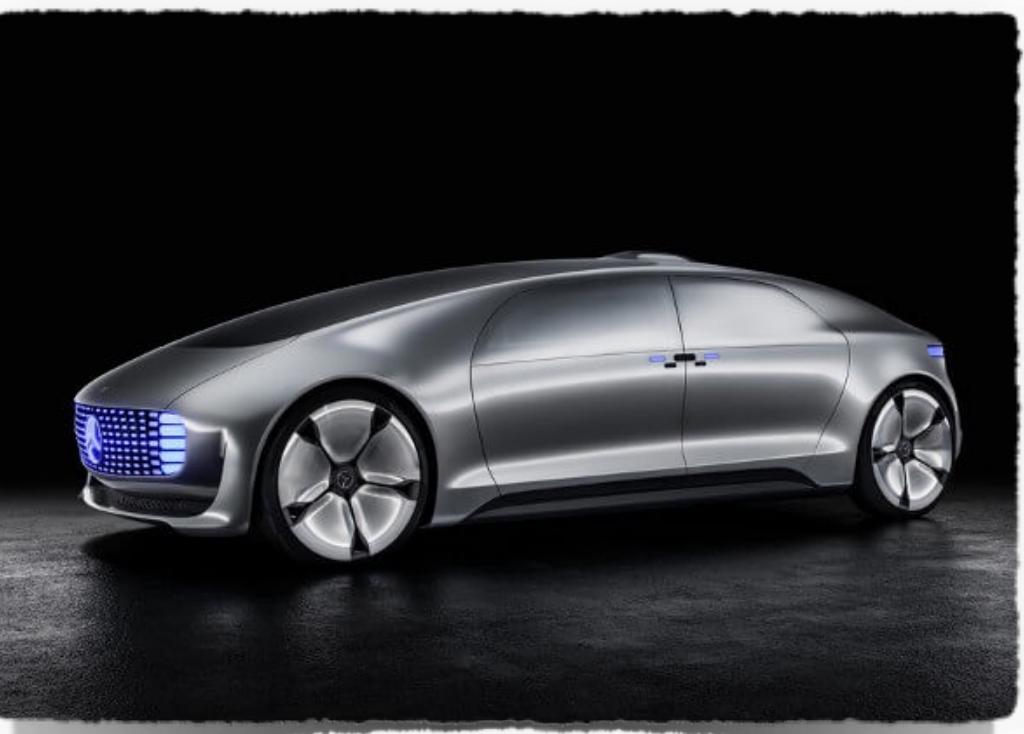
- ▶ Computers control physical processes with **feedback loops**
- ▶ Physical processes affect computations, and vice versa

* Lee. "The past, present, and future of cyber-physical systems: A focus on models." Sensors 15.3 (2015)

Cyber-Physical Systems (CPS)

*Integration of **computation**, **networking**,
and **physical processes****

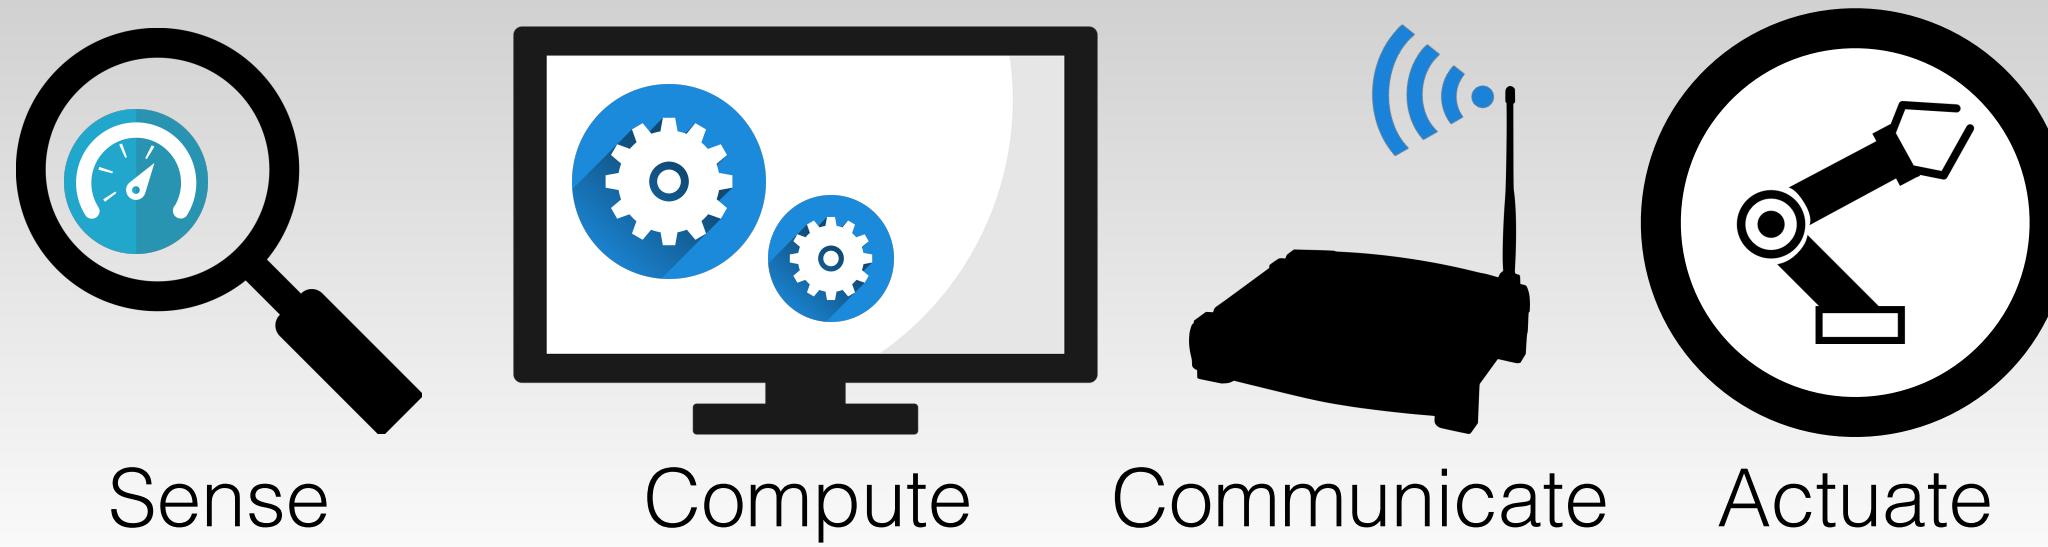
- ▶ Computers control physical processes with **feedback loops**
- ▶ Physical processes affect computations, and vice versa



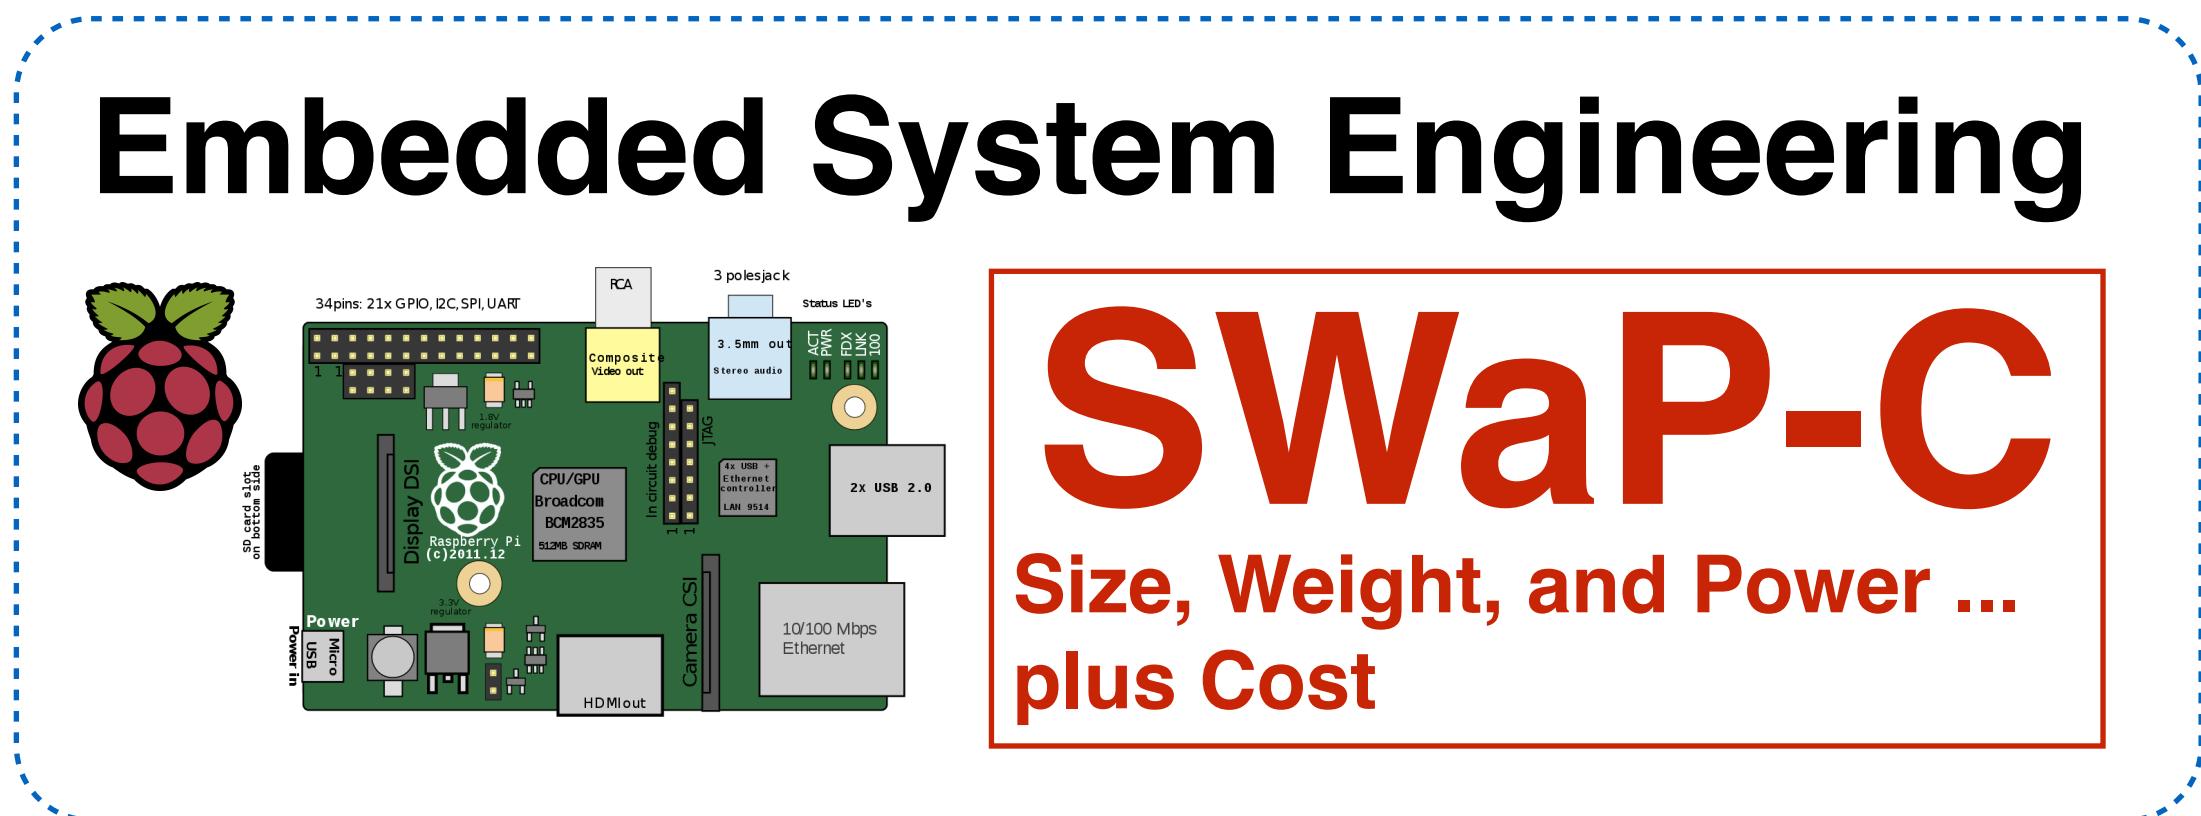
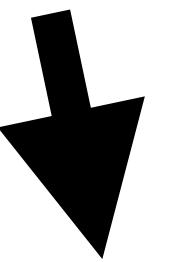
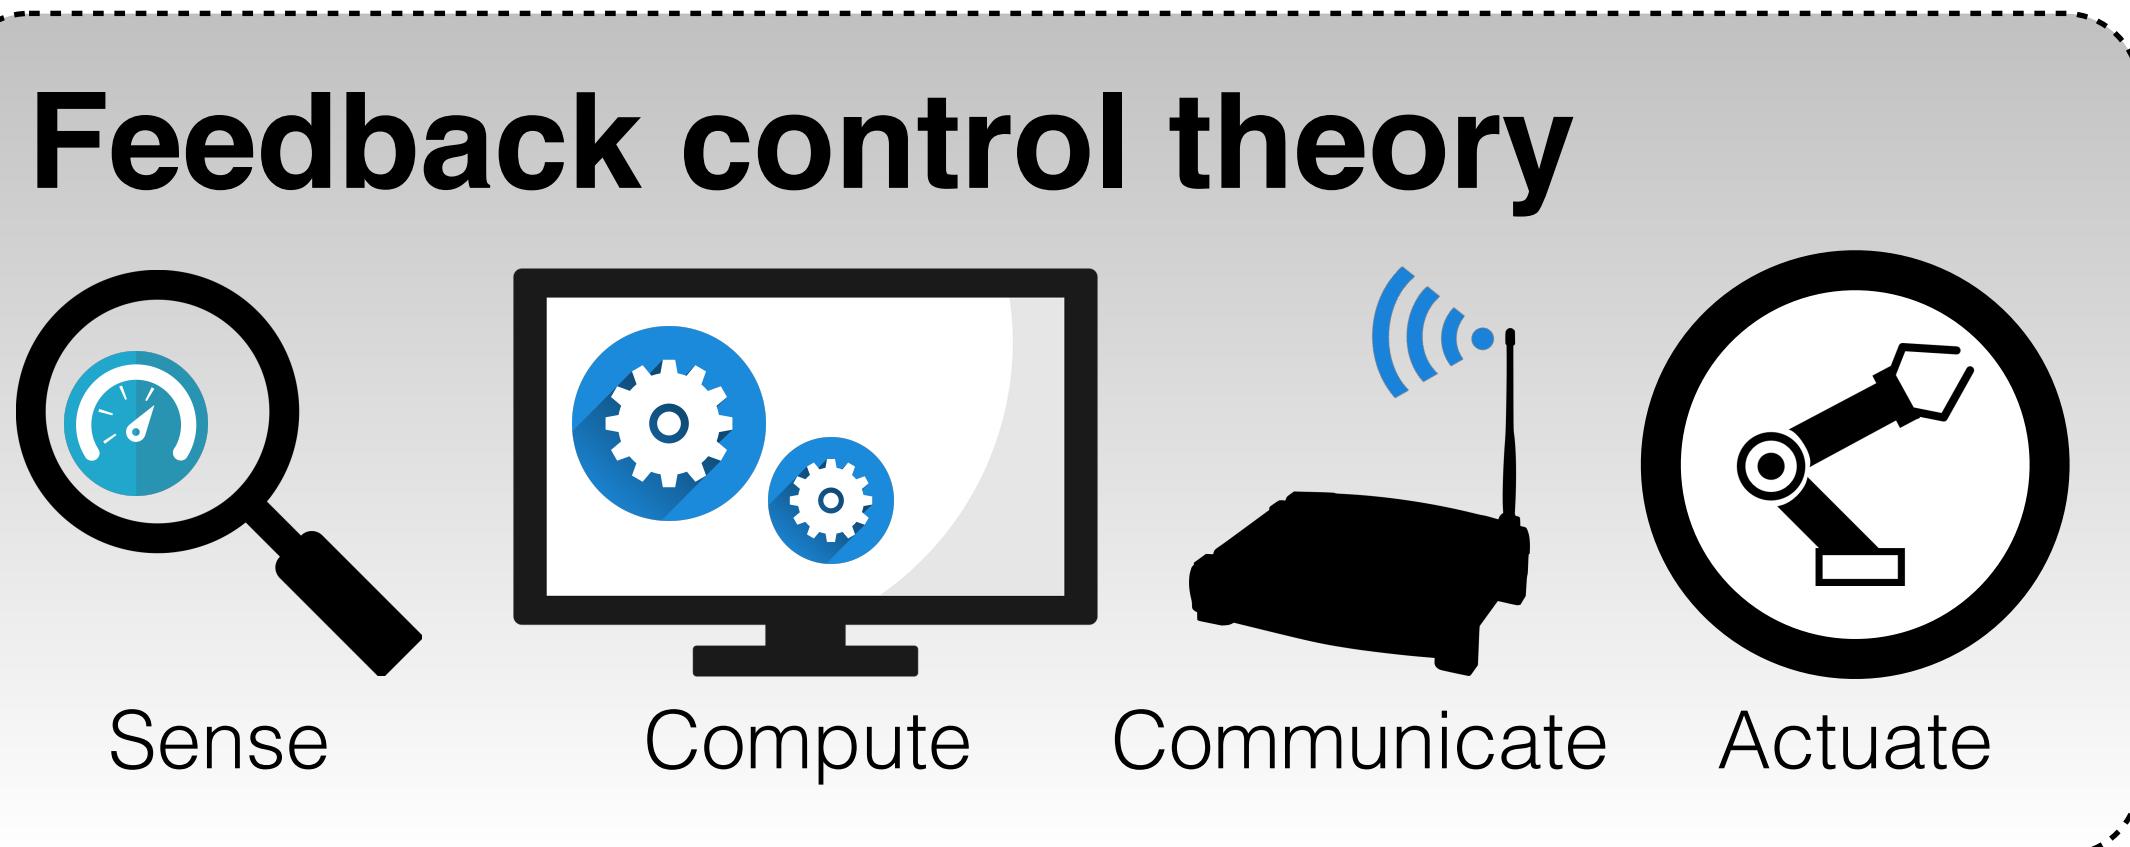
* Lee. "The past, present, and future of cyber-physical systems: A focus on models." Sensors 15.3 (2015)

Developing safety-critical CPS

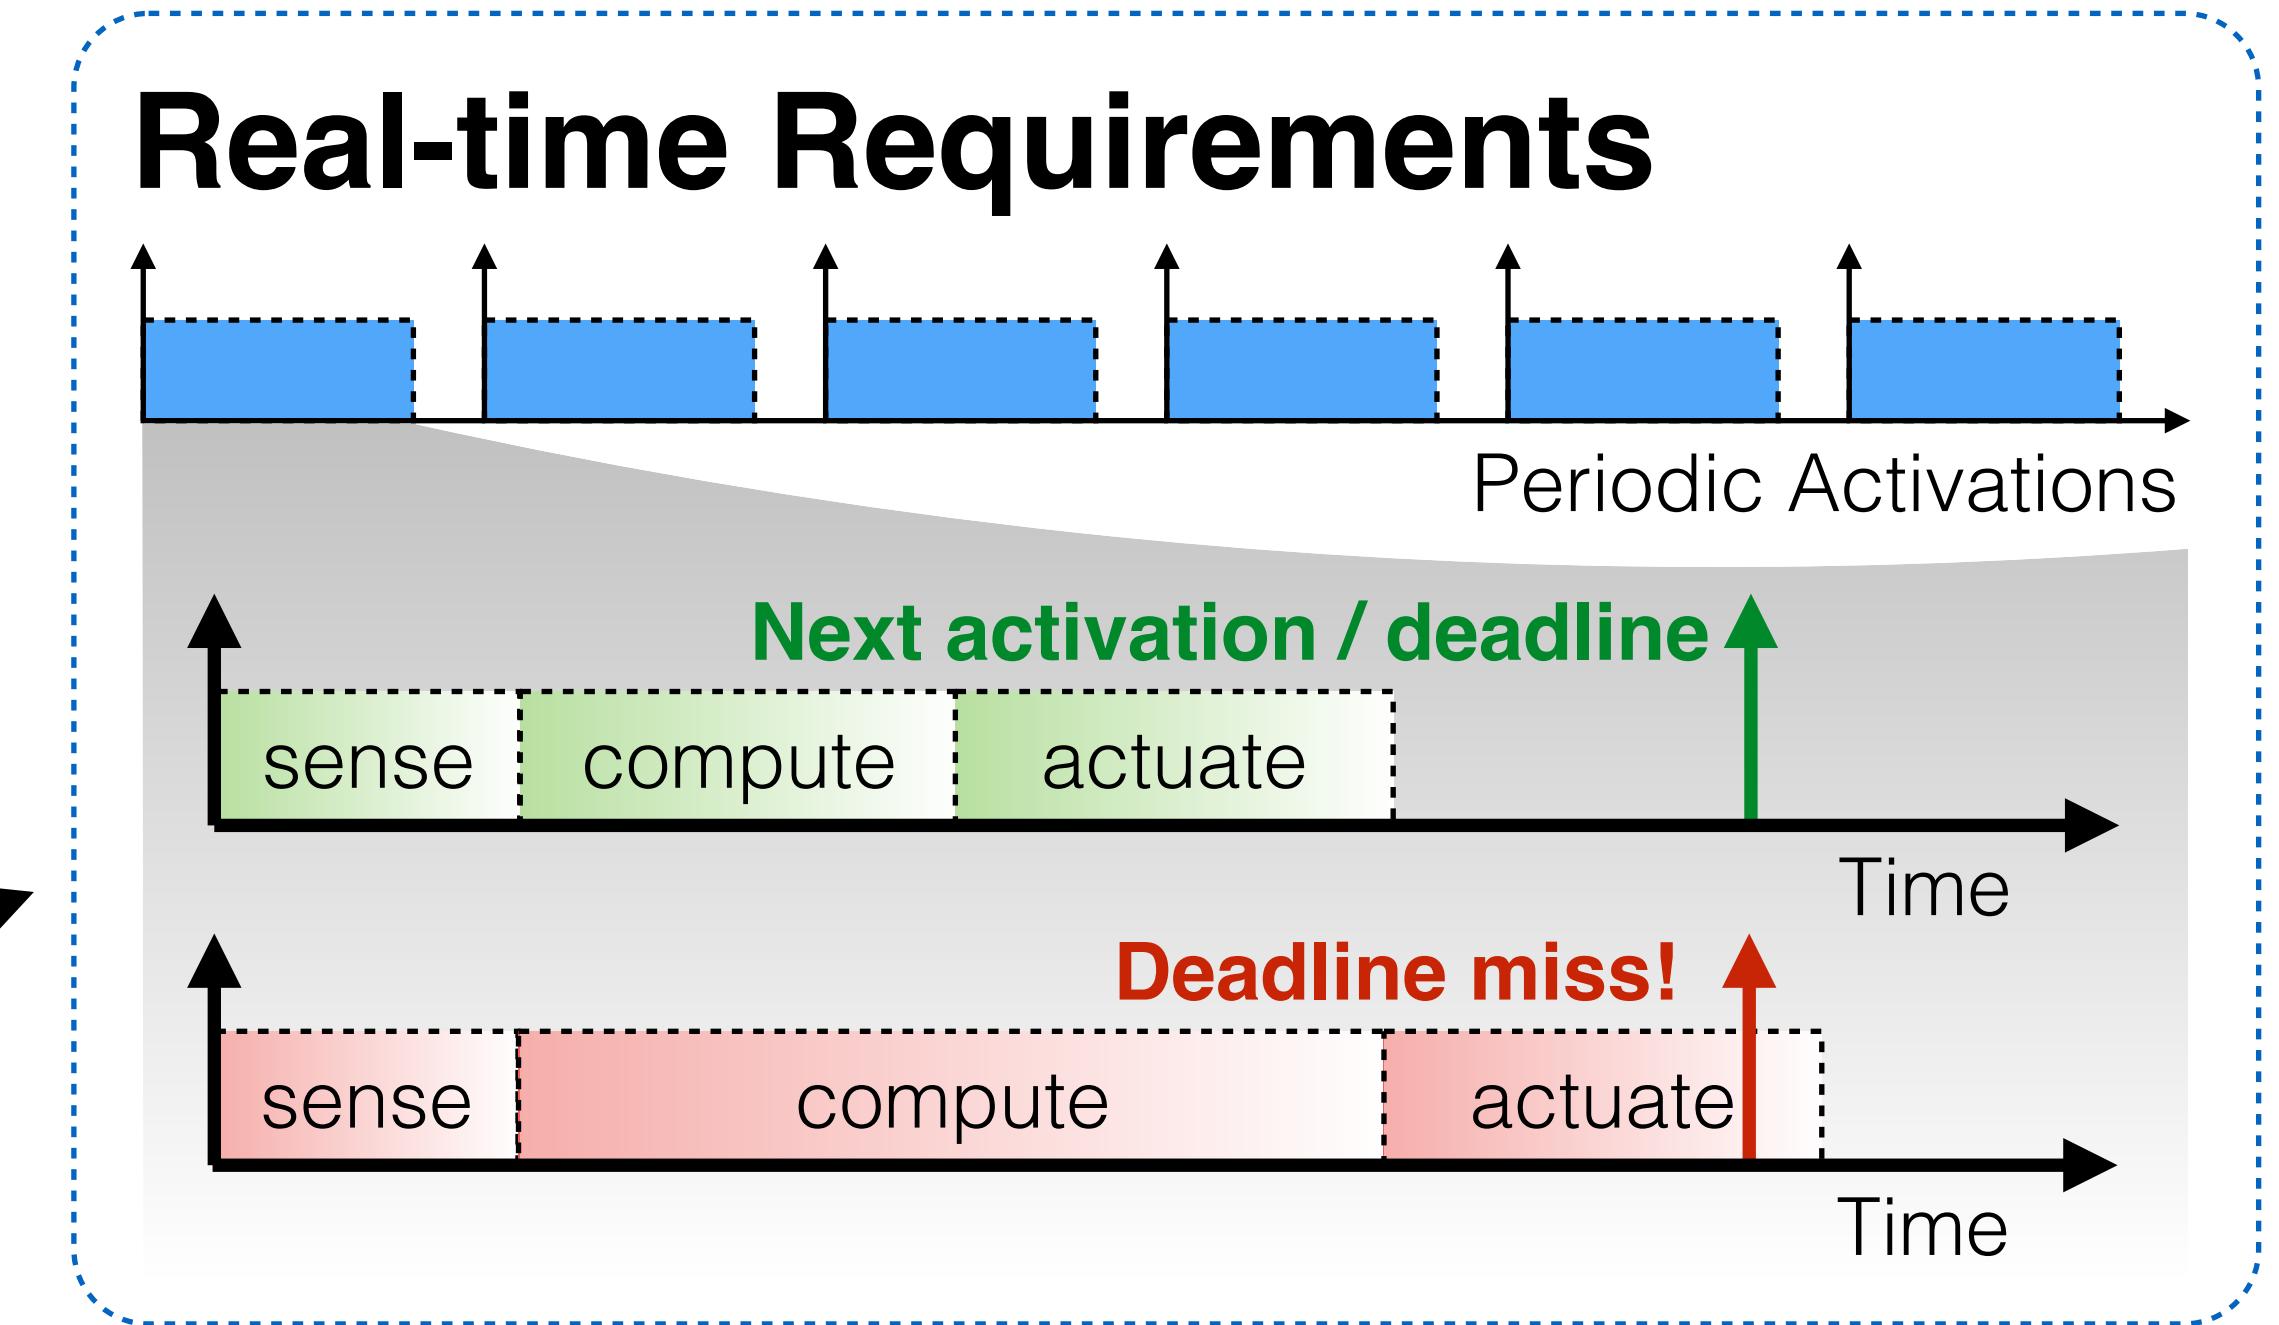
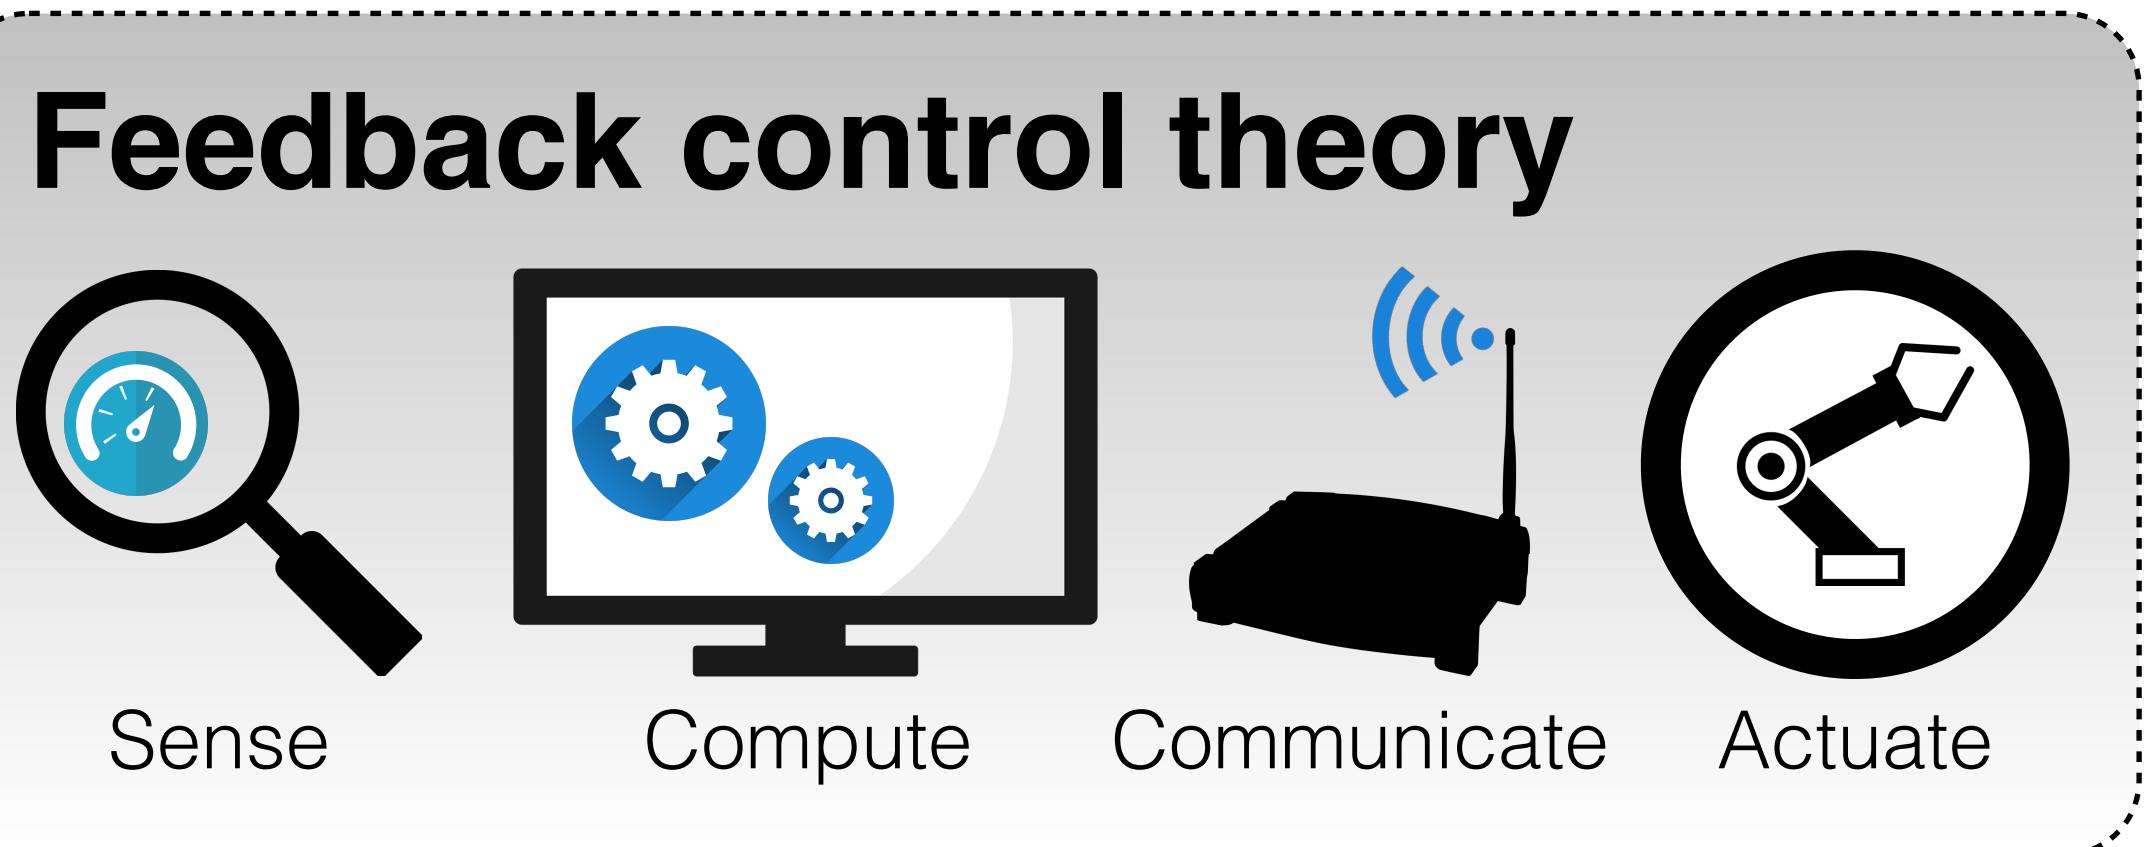
Feedback control theory



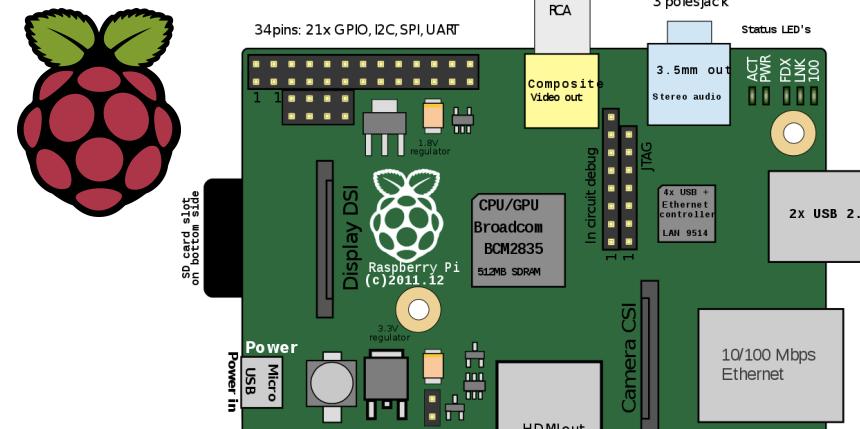
Developing safety-critical CPS



Developing safety-critical CPS

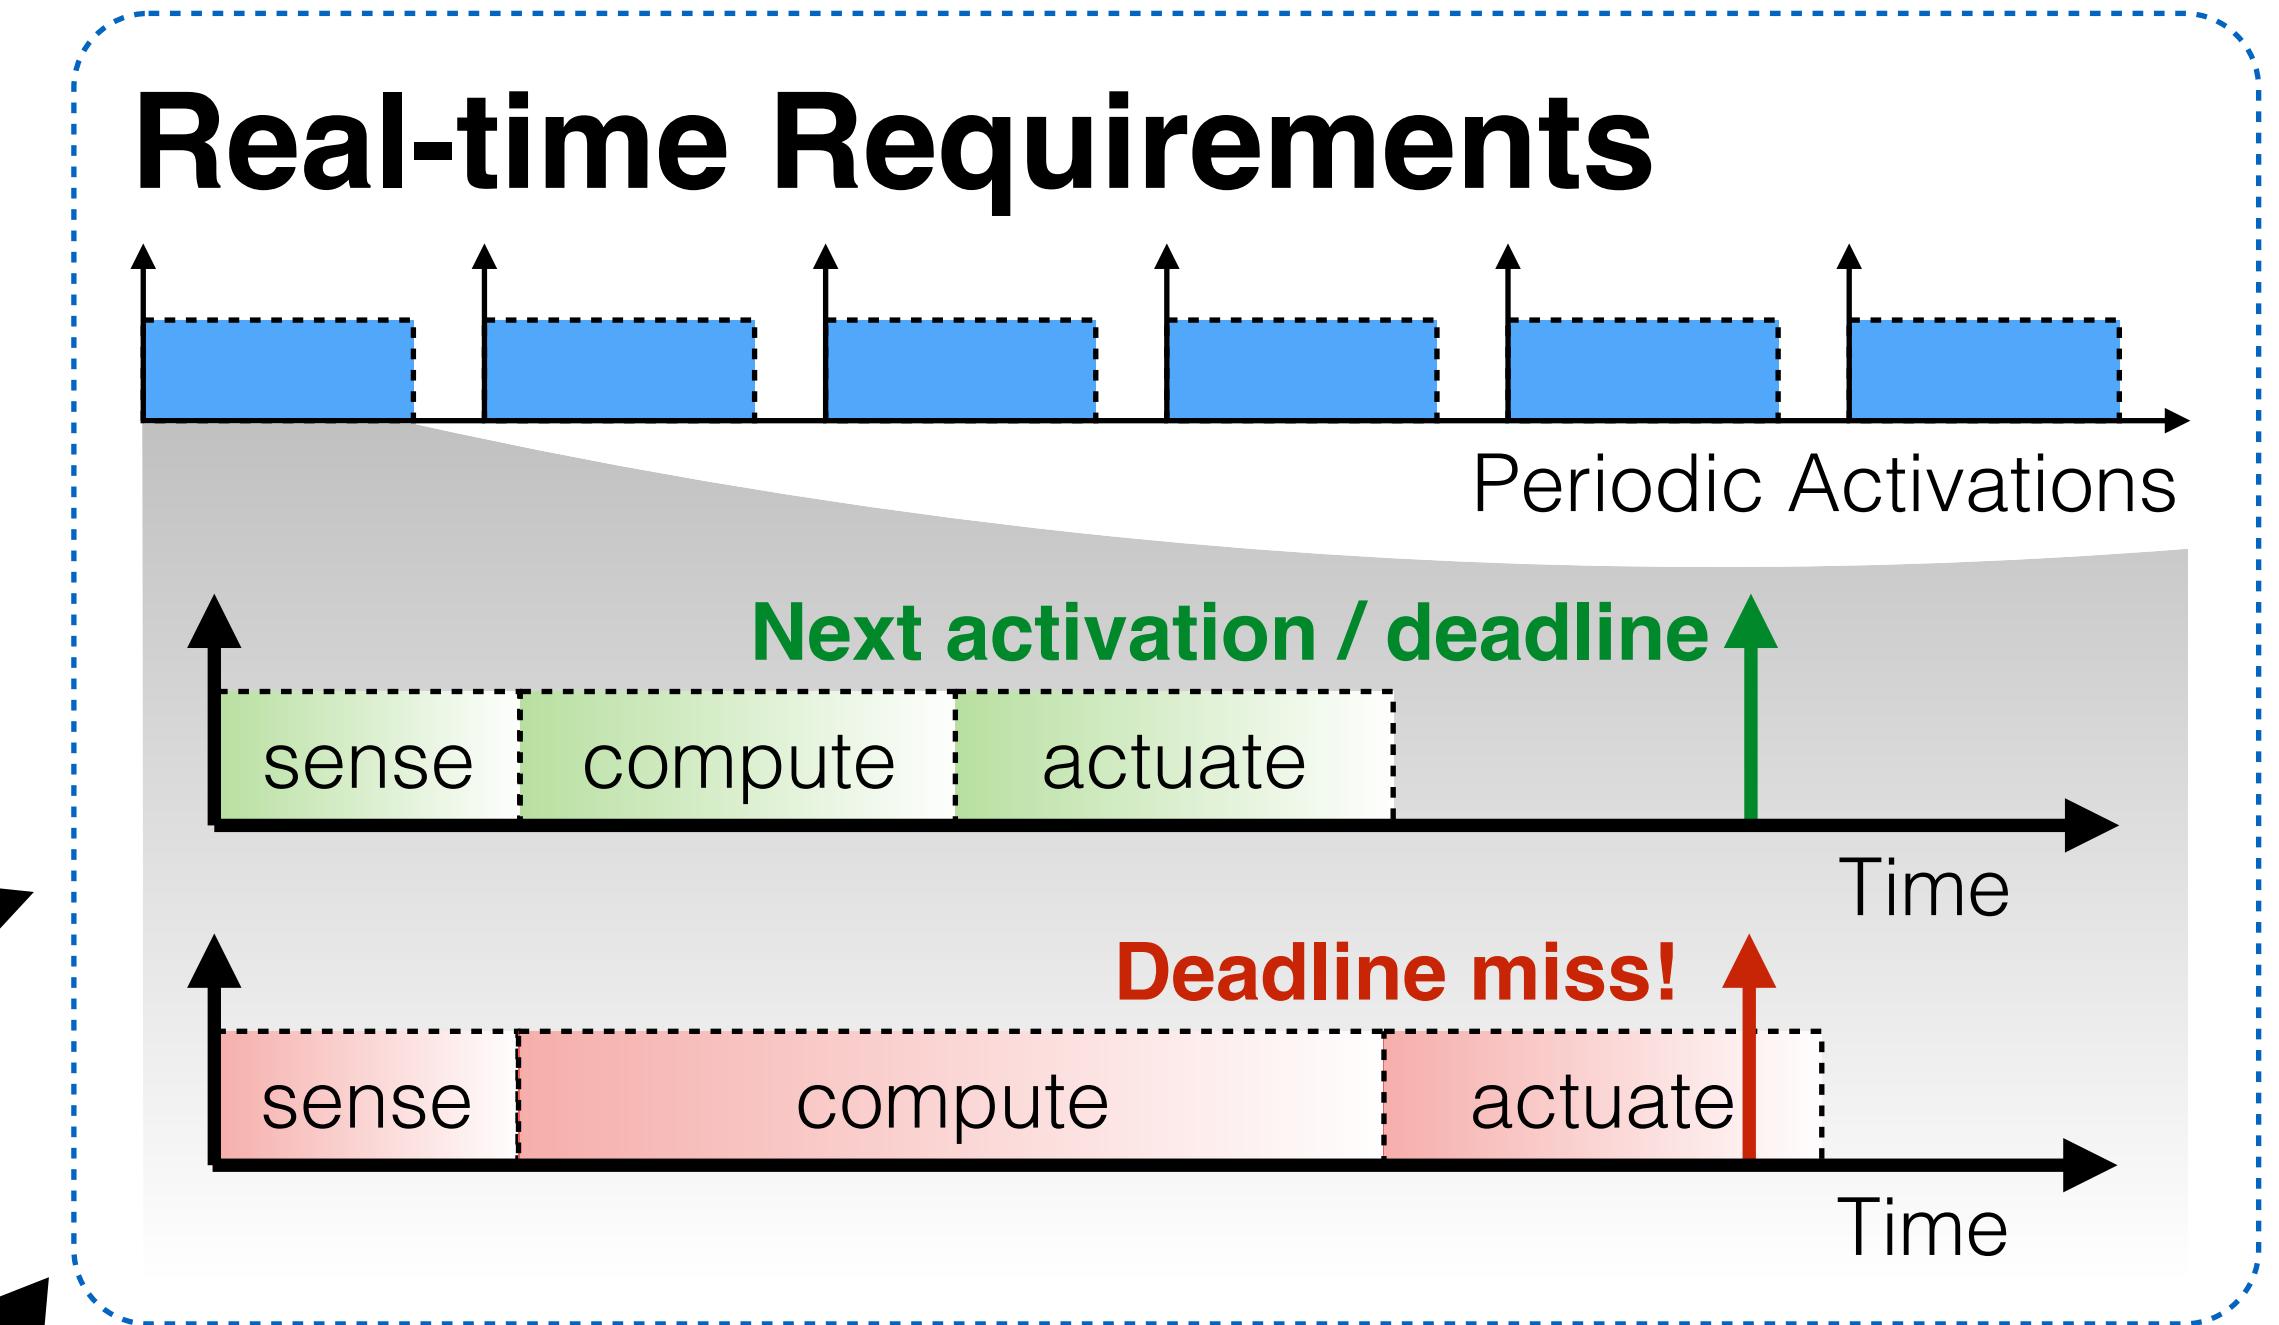
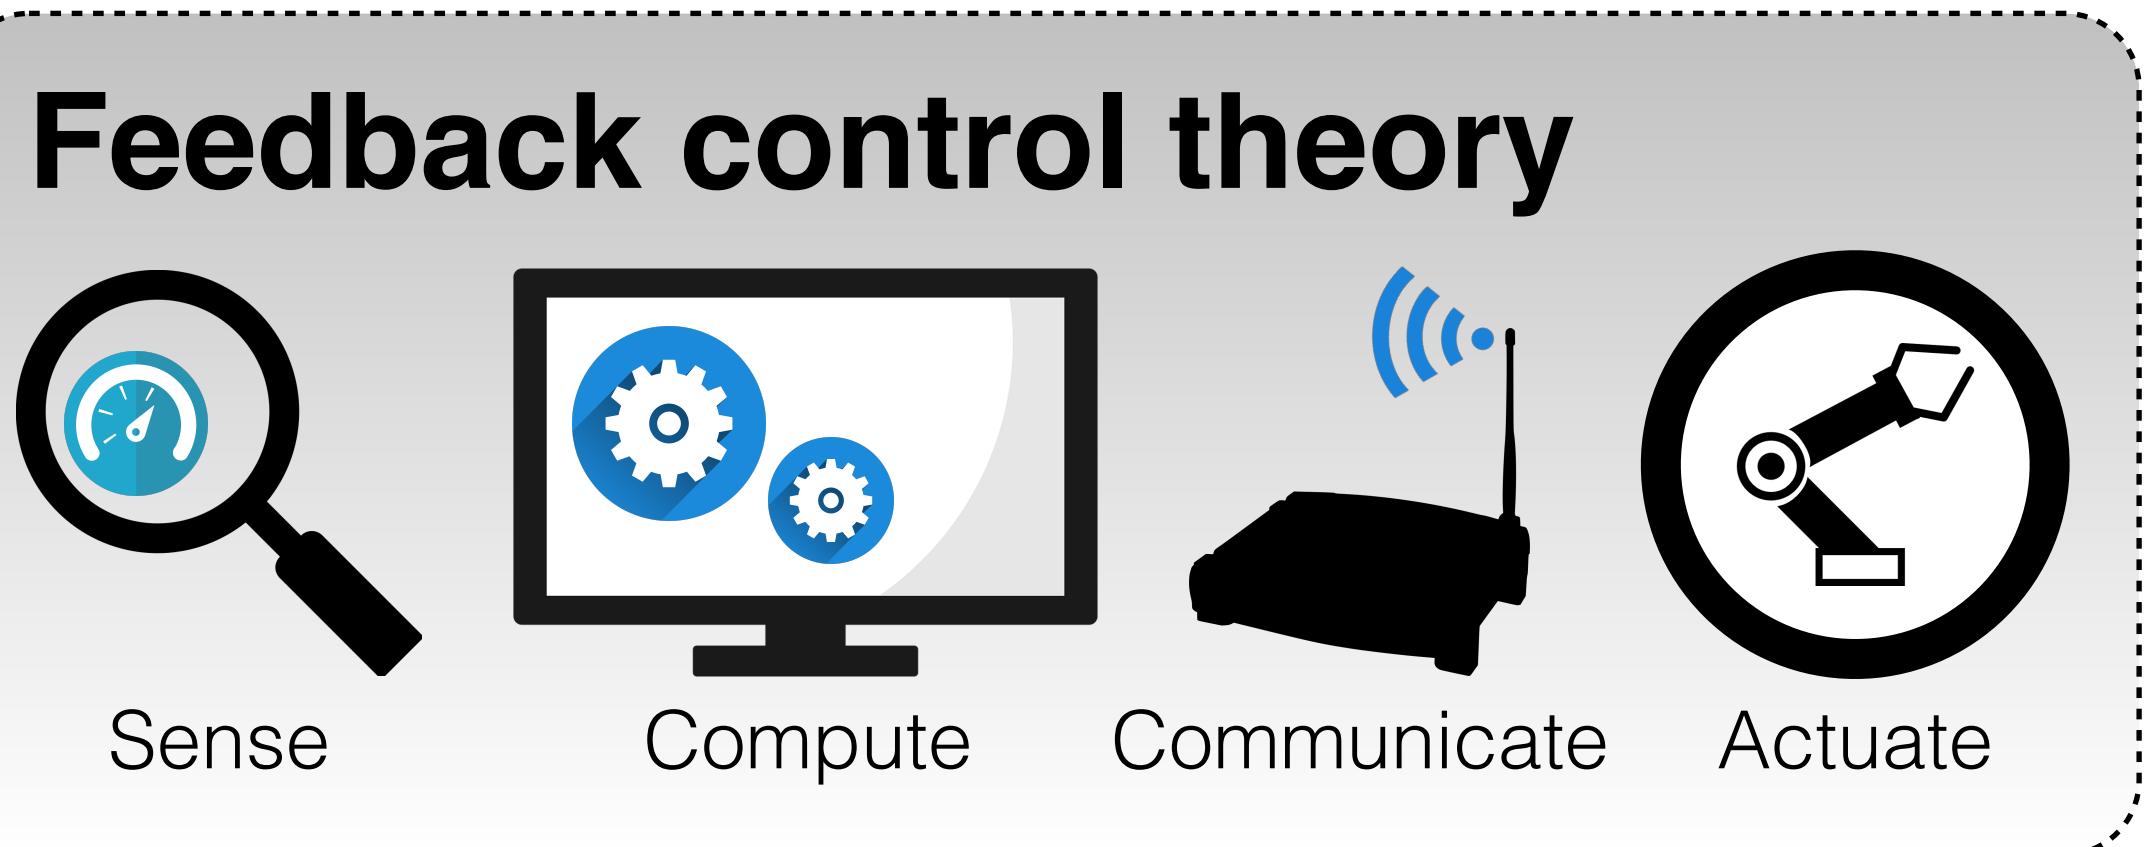


Embedded System Engineering

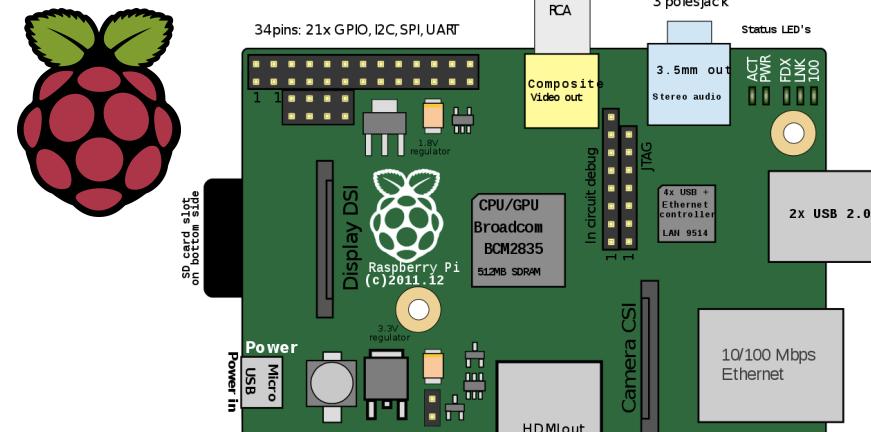


SWaP-C
Size, Weight, and Power ...
plus Cost

Developing safety-critical CPS



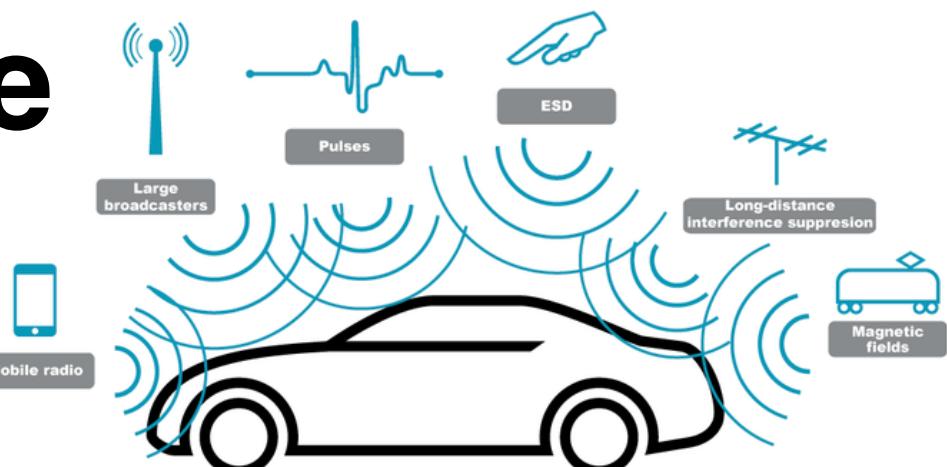
Embedded System Engineering



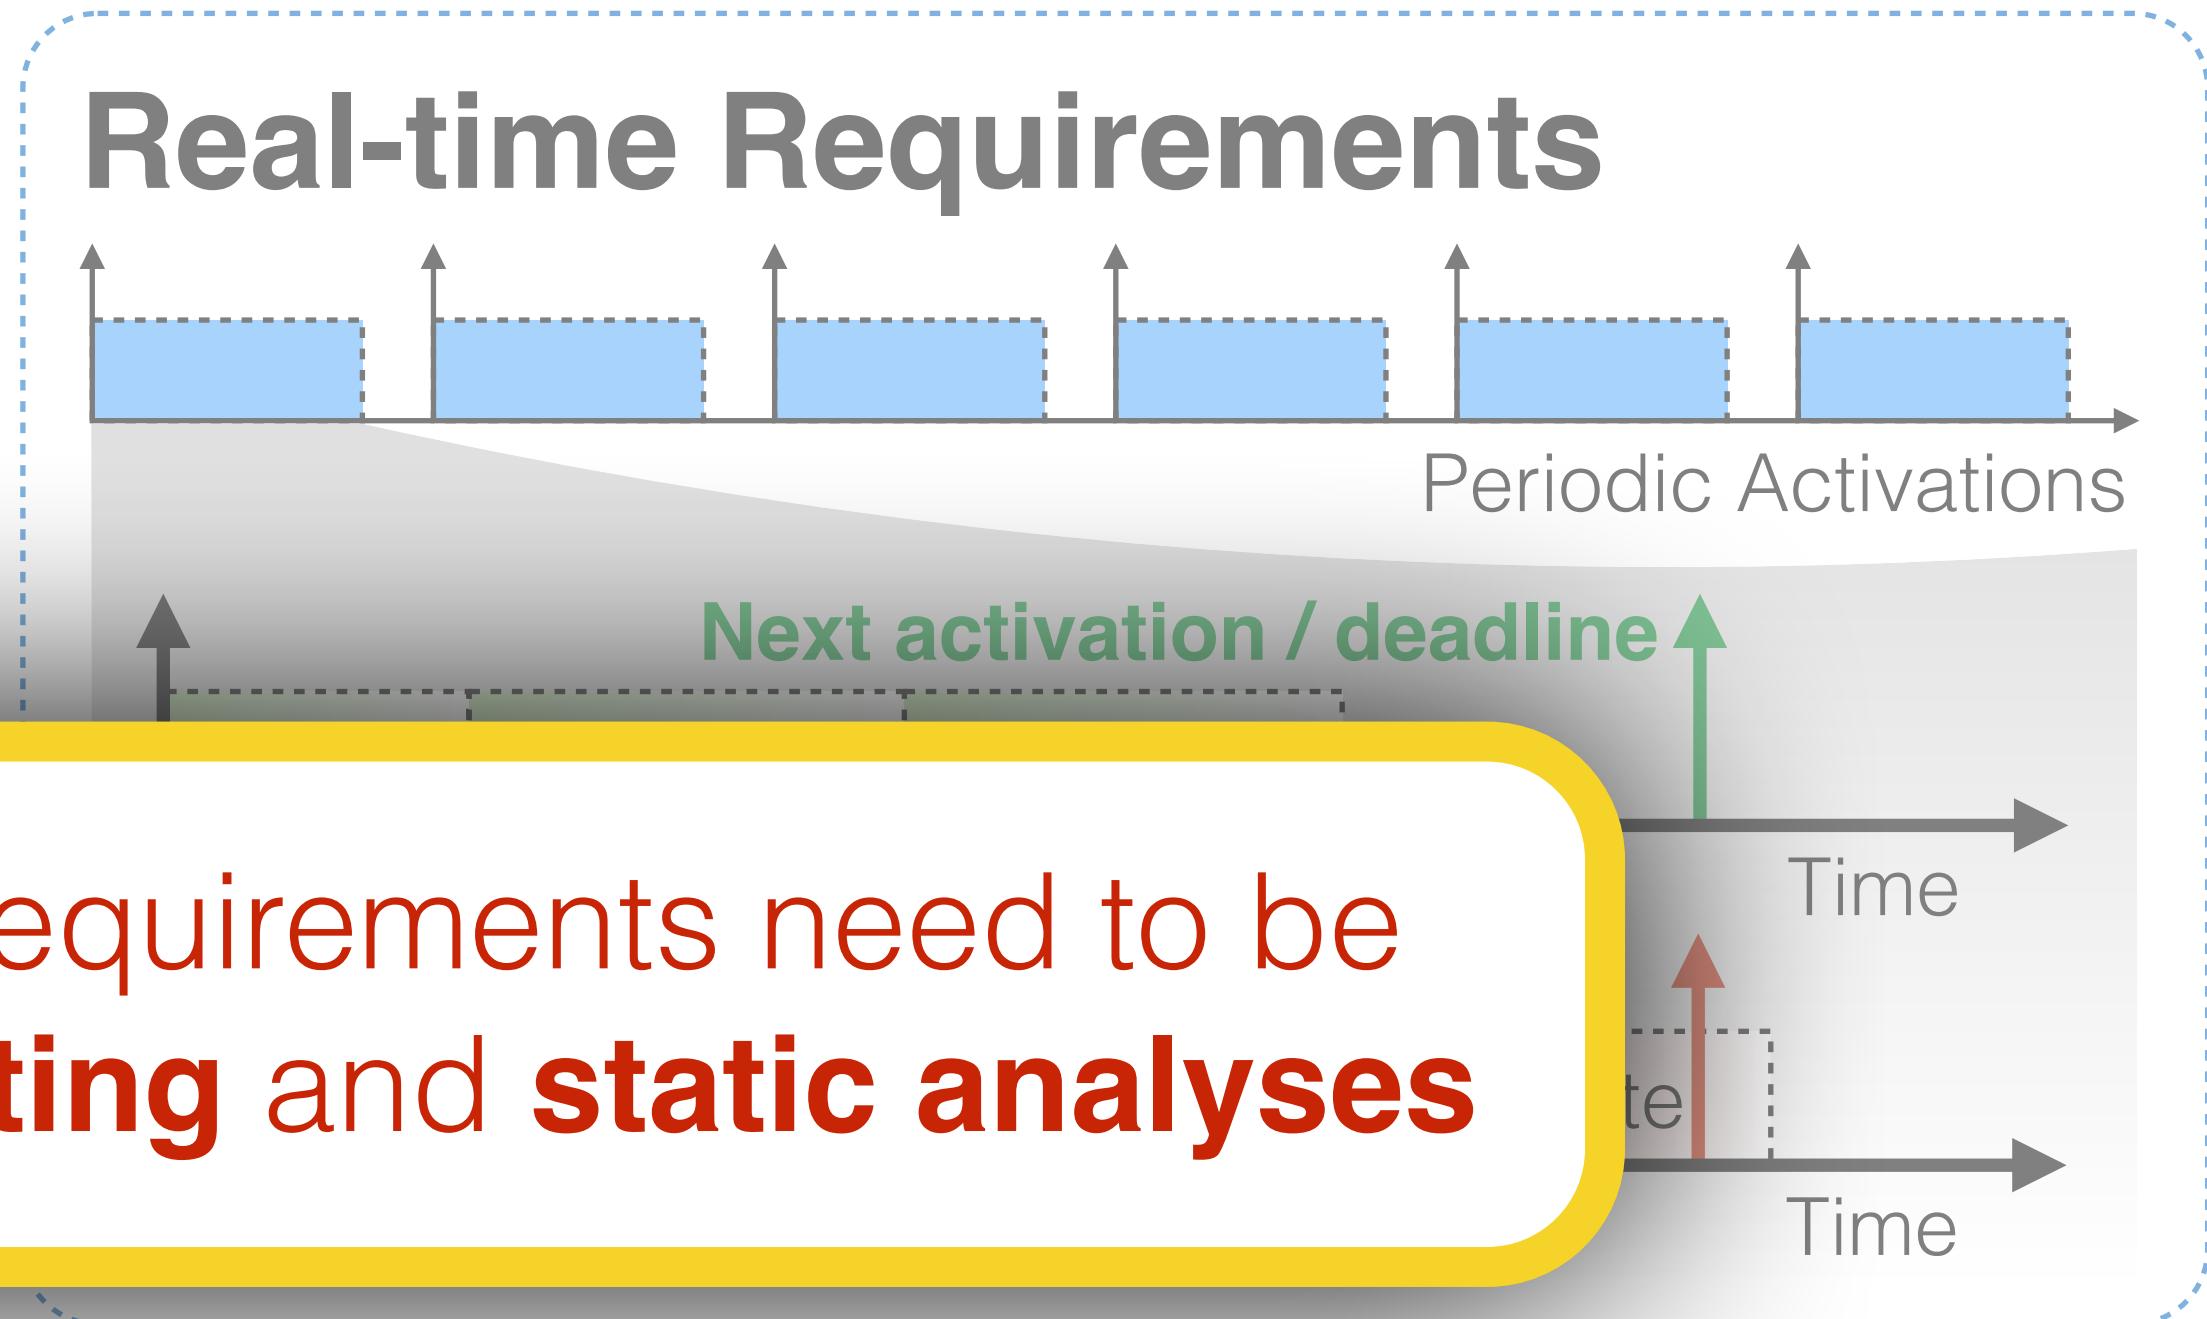
SWaP-C
Size, Weight, and Power ...
plus Cost

Fault Tolerance

Harsh environment
can result in
hardware faults

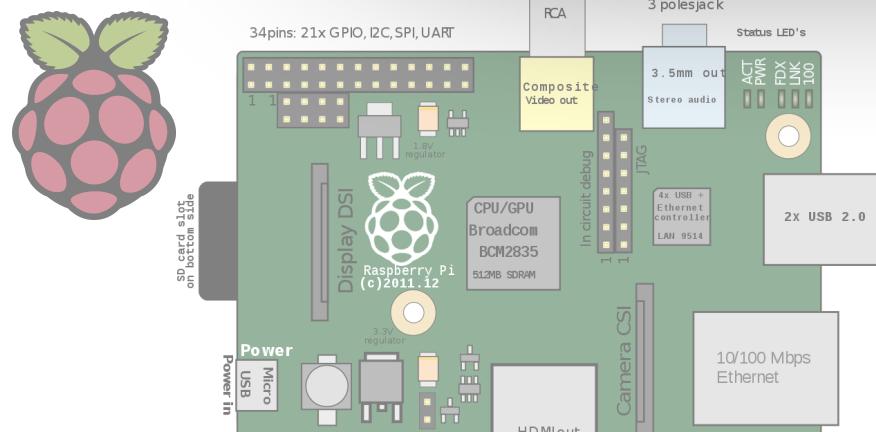


Developing safety-critical CPS



For safety certification, these requirements need to be validated **in advance** using **testing** and **static analyses**

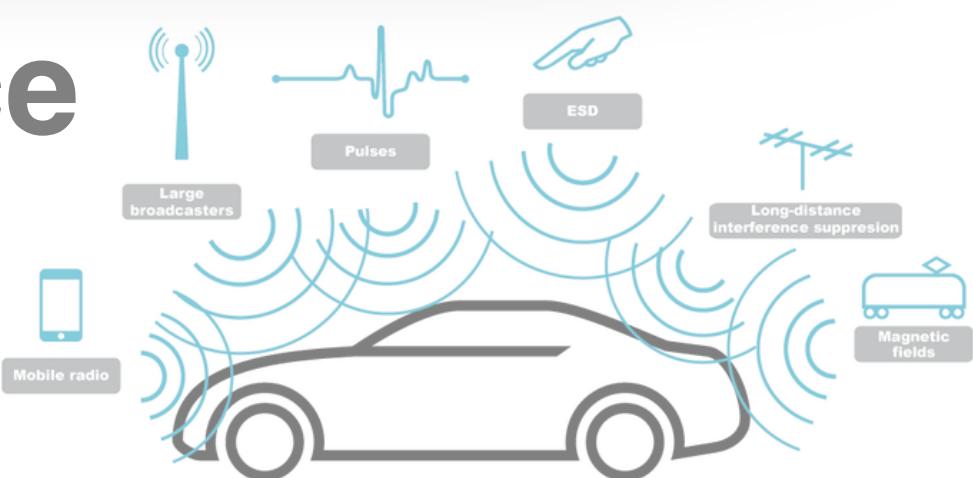
Embedded System Engineering



SWaP-C
Size, Weight, and Power ...
plus Cost

Fault Tolerance

Harsh environment
can result in
hardware faults



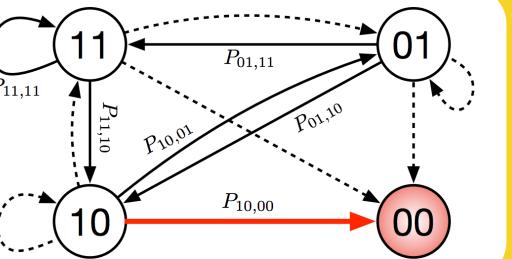
Reliability Analysis of Distributed Real-Time Systems

Best Presentation Award

Controller Area
Network [ECRTS '18]

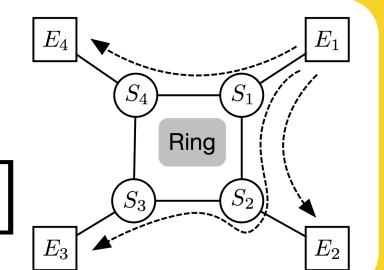


Periodic Weakly-Hard
Systems [ECRTS '19]



Distinguished Paper Award

Replica Consistency
over Ethernet [RTAS '20]



Reliability Analysis of Distributed Real-Time Systems

Best Presentation Award

Controller Area
Network [ECRTS '18]

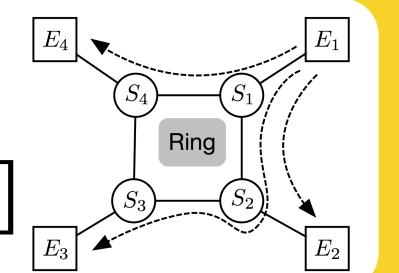


Periodic Weakly-Hard
Systems [ECRTS '19]



Distinguished Paper Award

Replica Consistency
over Ethernet [RTAS '20]



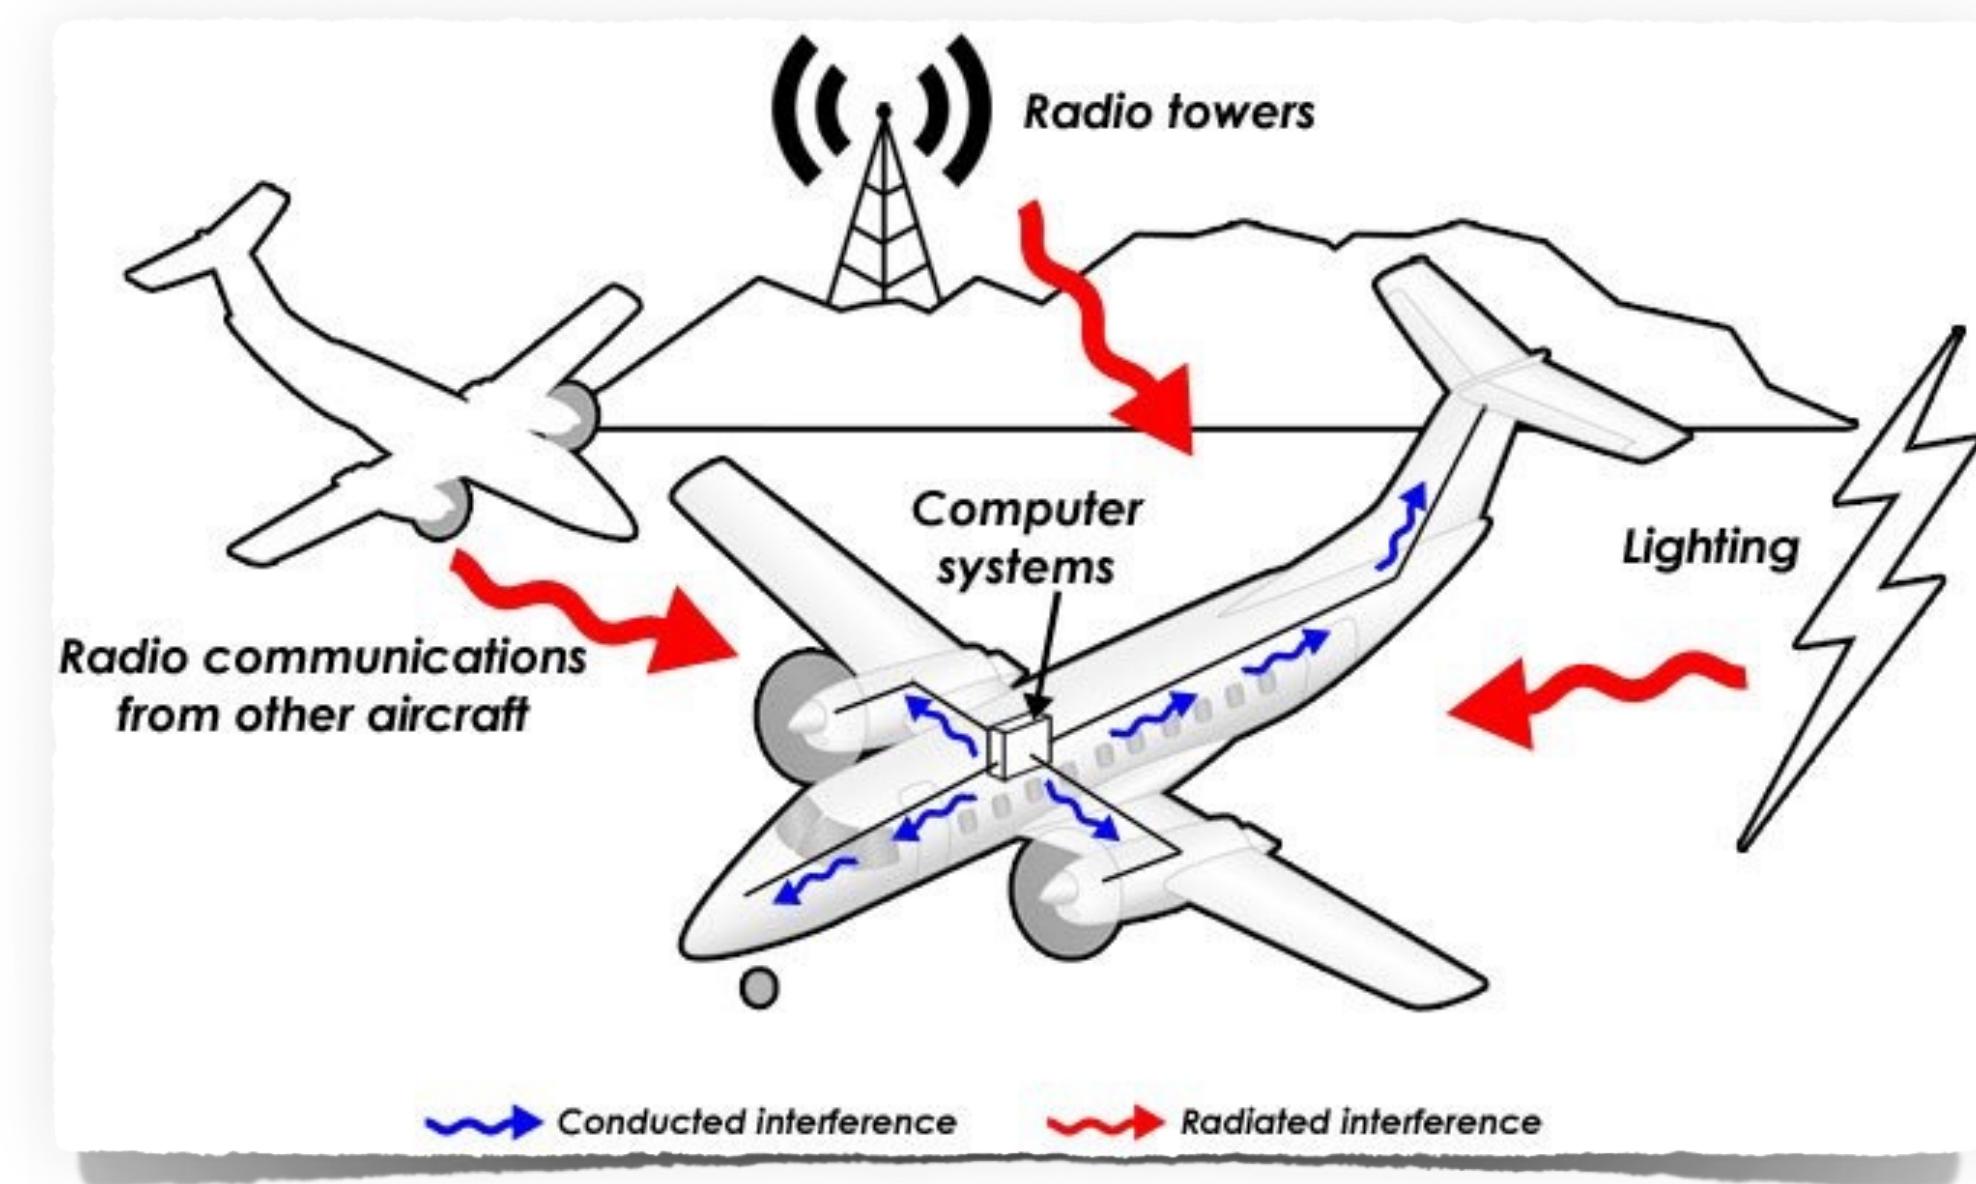
↑
This Talk

↑
This Talk

Why do we need **reliability analyses**?

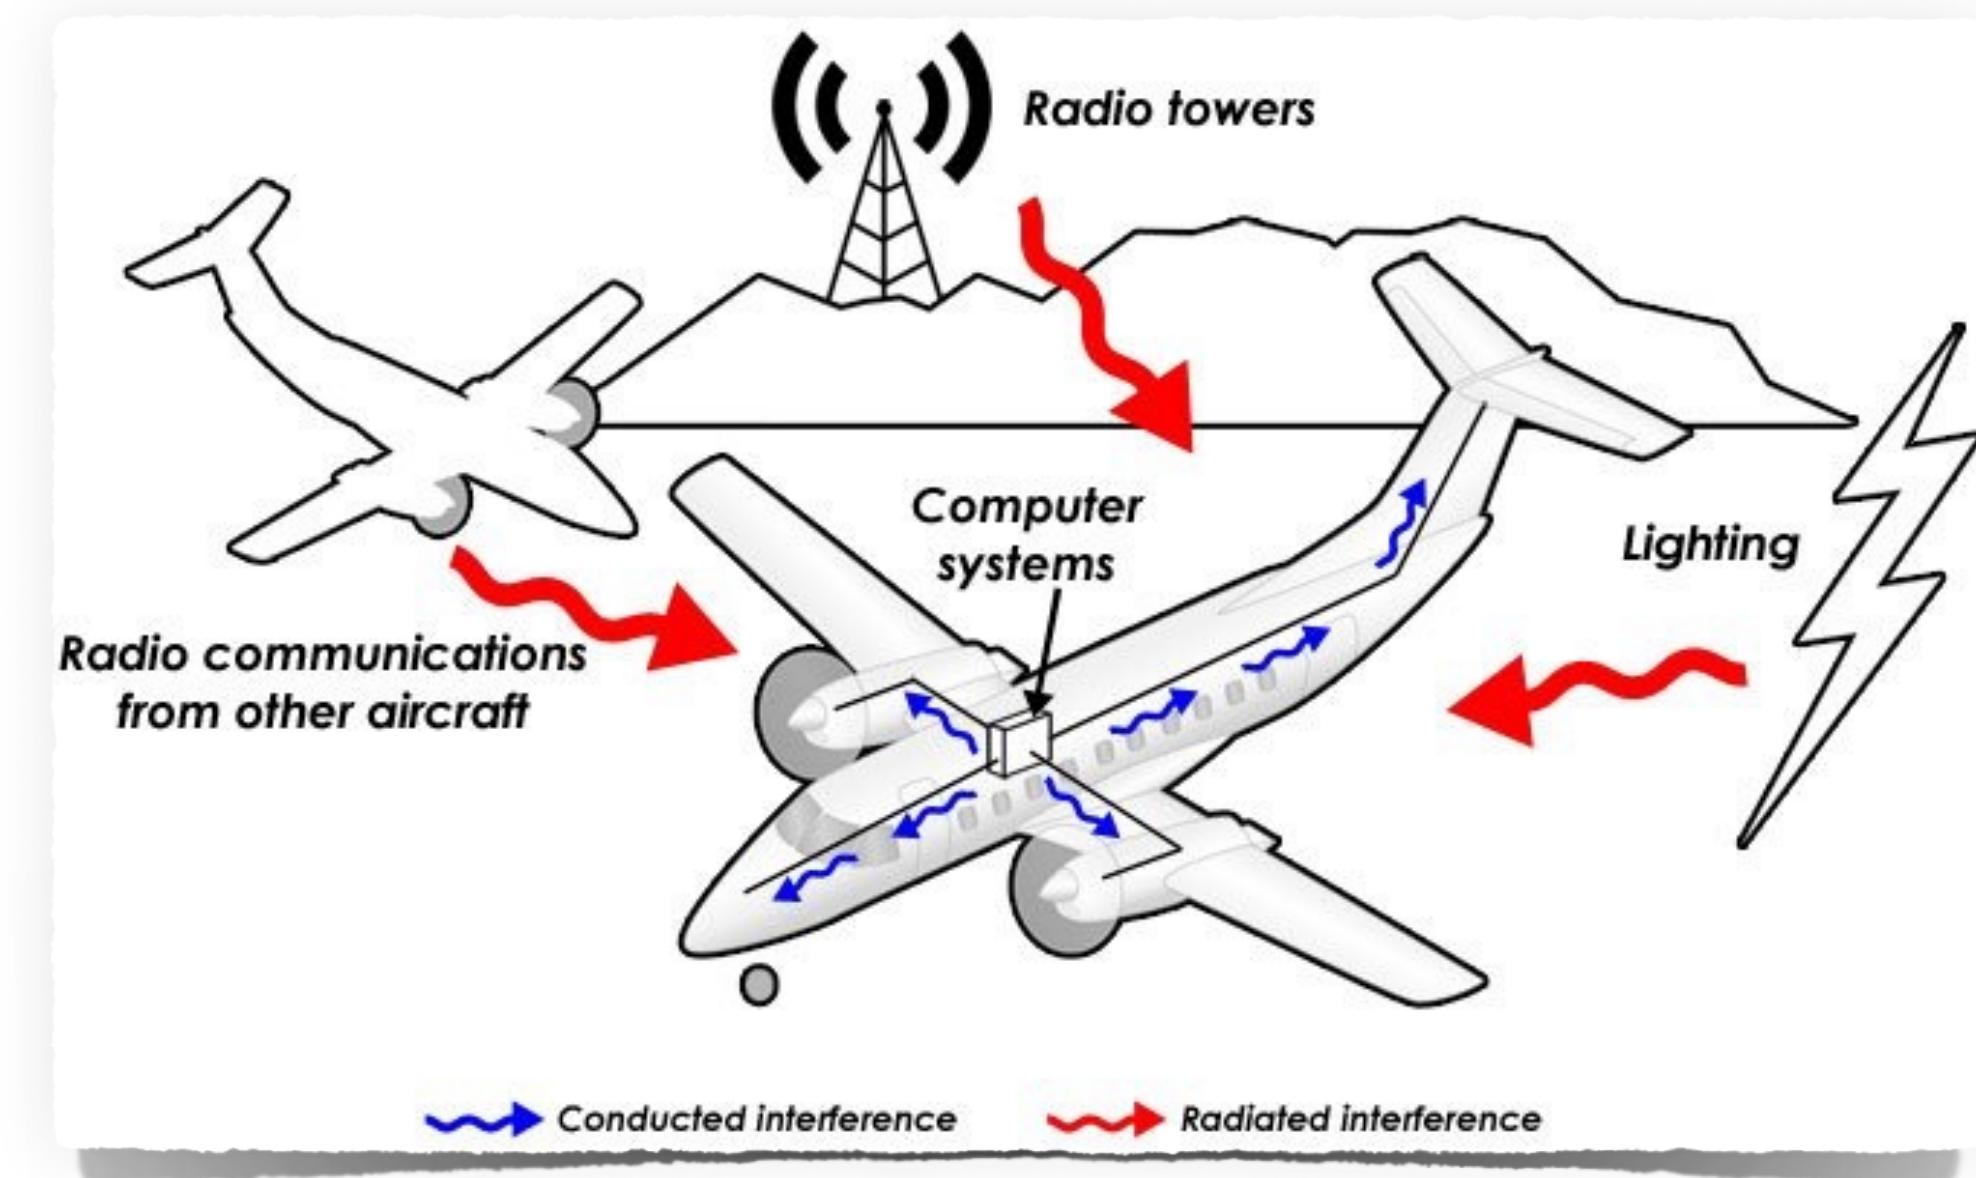
Environmentally-induced transient faults

- Harsh environments
 - ▶ Robots operating under **hard radiation**
 - ▶ Industrial systems near **high power machinery**
 - ▶ **Electric motors, spark plugs** inside automobiles



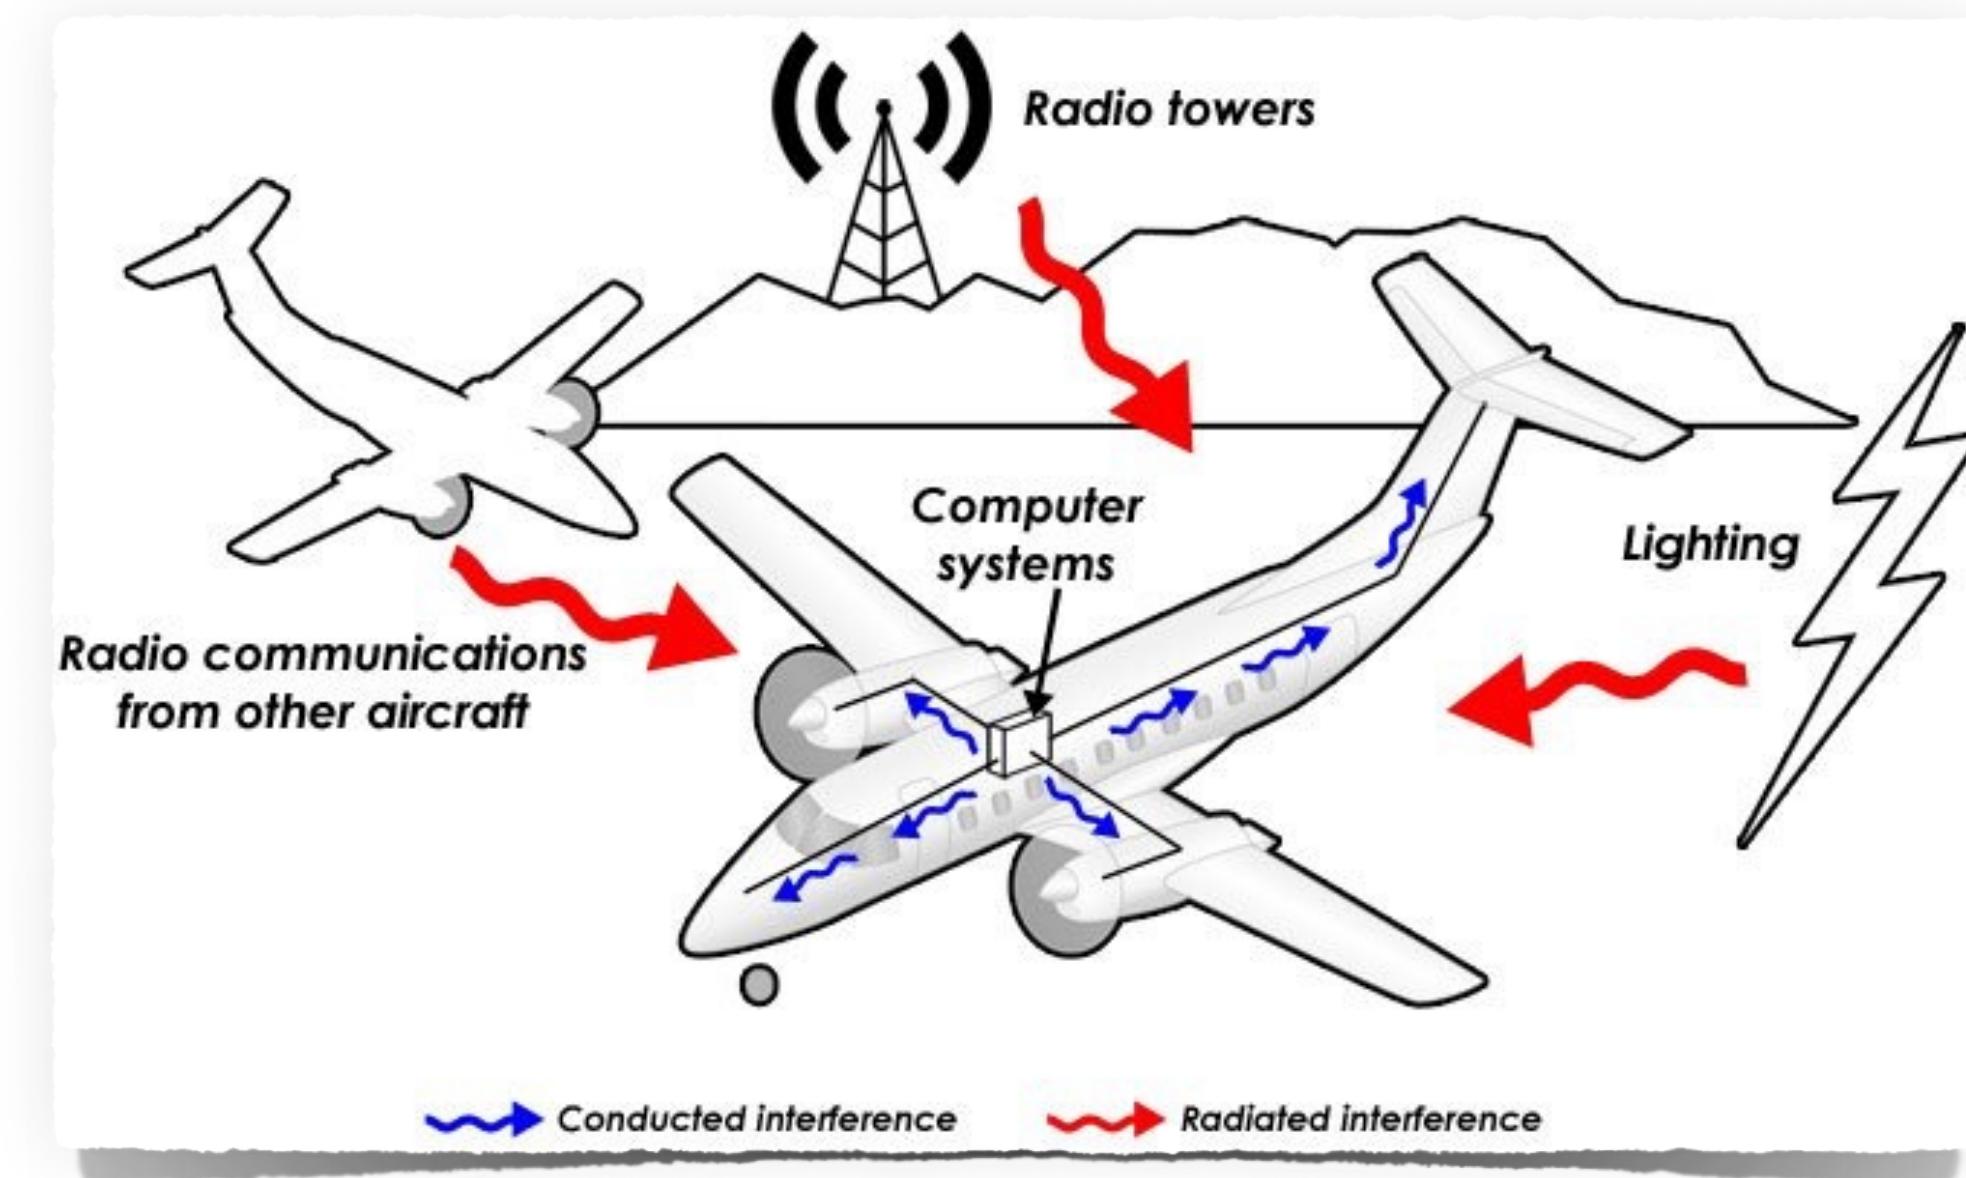
Environmentally-induced transient faults

- Harsh environments
 - ▶ Robots operating under **hard radiation**
 - ▶ Industrial systems near **high power machinery**
 - ▶ **Electric motors, spark plugs** inside automobiles
- **Bit-flips** in registers, buffers, networks



Environmentally-induced transient faults

- Harsh environments
 - ▶ Robots operating under **hard radiation**
 - ▶ Industrial systems near **high power machinery**
 - ▶ **Electric motors, spark plugs** inside automobiles
- **Bit-flips** in registers, buffers, networks



Example*

- ▶ One bit-flip in a 1 MB SRAM every 10^{12} hours of operation
- ▶ 0.5 billion cars with an average daily operation time of 5%
- ▶ **About 5000 cars are affected by a bit-flip every day**

* Mancuso. "Next-generation safety-critical systems on multi-core platforms." PhD thesis, UIUC (2017)

Errors and failures due to transient faults

Errors and failures due to transient faults

- Transmission errors
 - ▶ Faults on the network
- Omission errors
 - ▶ Fault-induced kernel panics, hangs
- Incorrect computation errors
 - ▶ Faults in memory buffers
- Inconsistent broadcast errors
 - ▶ Faults in networked systems

Errors and failures due to transient faults

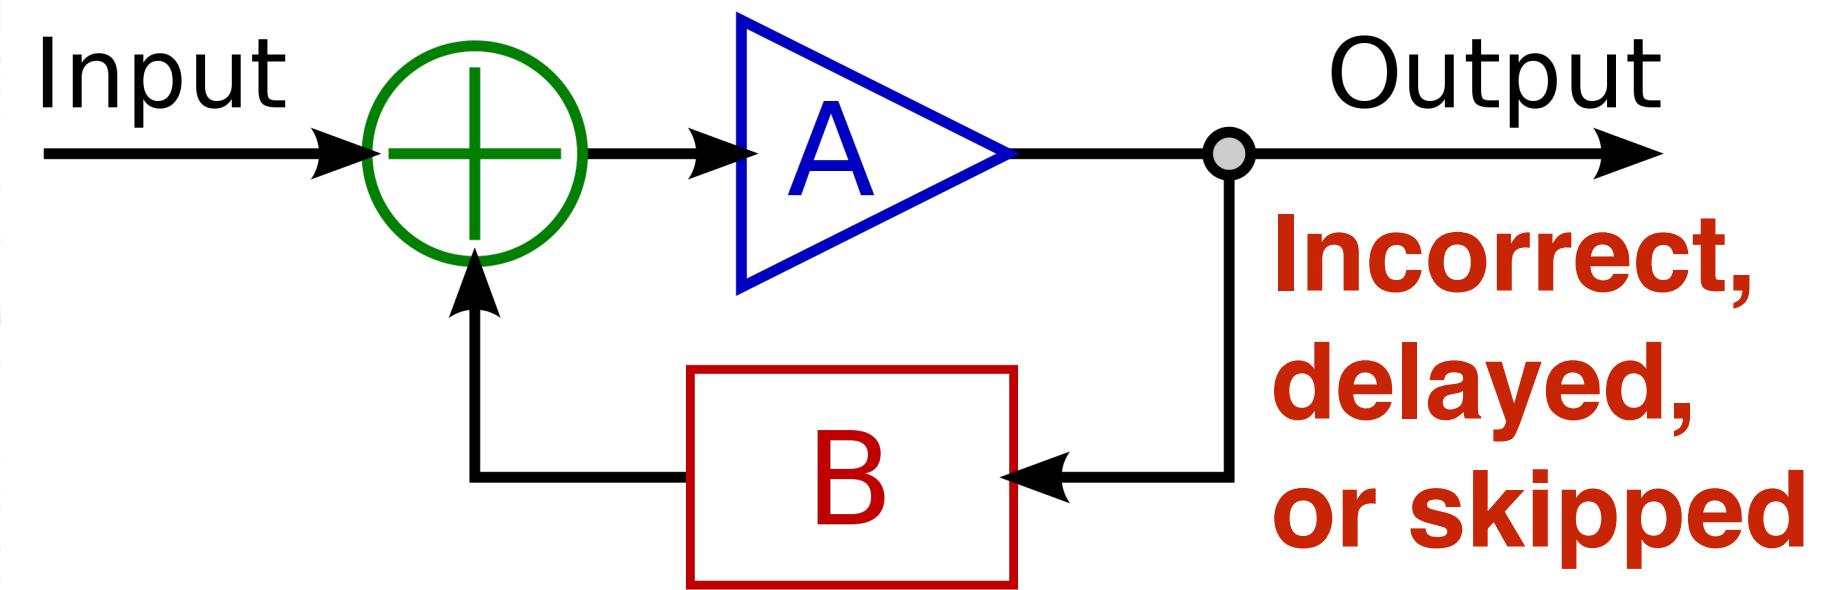
- Transmission errors
 - ▶ Faults on the network
- Omission errors
 - ▶ Fault-induced kernel panics, hangs
- Incorrect computation errors
 - ▶ Faults in memory buffers
- Inconsistent broadcast errors
 - ▶ Faults in networked systems



Failures in

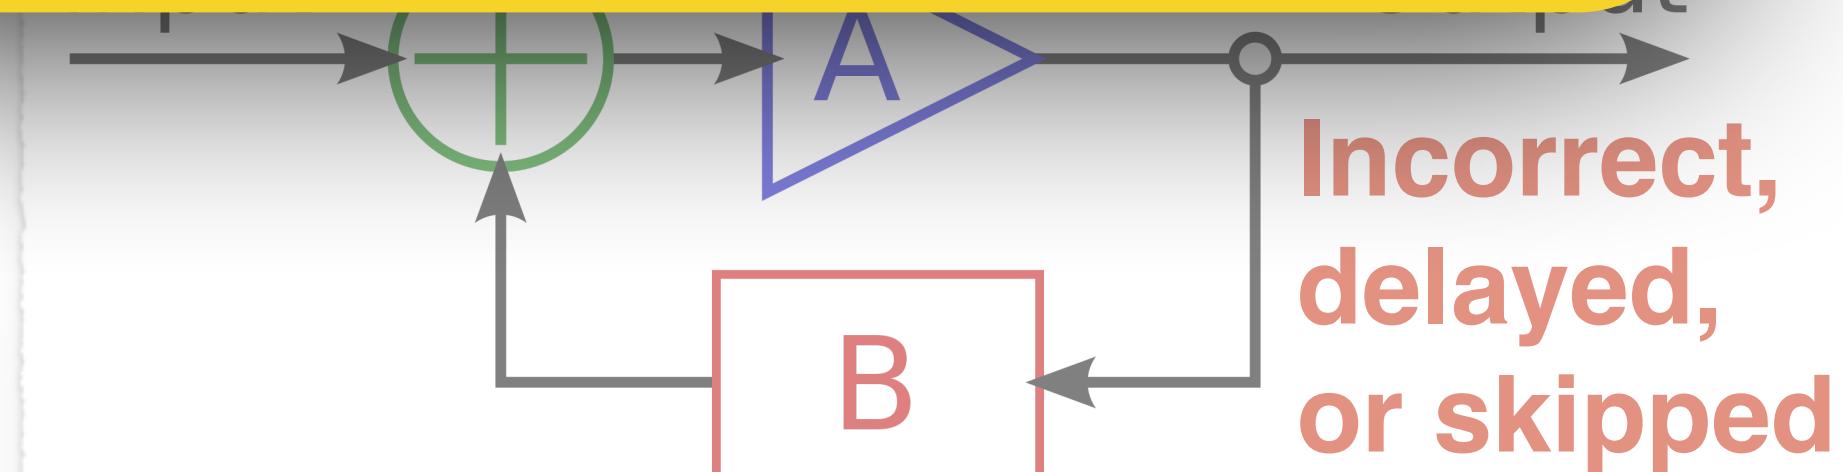
- ▶ **value domain (incorrect output)**
- ▶ **time domain (delayed output)**

E.g., safety-critical control system



Errors and failures due to transient faults

- Transmission errors
 - ▶ Faults on the network
 - Omission
 - ▶ Fault-induced omission
 - Incorrect
 - ▶ Faults in memory buffers
 - Inconsistent broadcast errors
 - ▶ Faults in networked systems
- Fault-induced errors are **random events**
- ▶ **Cannot be predicted** in advance
 - ▶ Must be tolerated at runtime using **fault tolerance mechanisms**
- Failures in
- ▶ value domain (incorrect output)
 - ▶ time domain (delayed output)
- Incorrect, delayed, or skipped



Errors and failures due to transient faults

- Transmission errors
 - ▶ Faults on the network
- Omission errors
 - ▶ Fault-induced kernel panics, hangs
- Incorrect computation errors
 - ▶ Faults in memory buffers
- Inconsistent broadcast errors
 - ▶ Faults in networked systems

Checksums and retransmissions

Dual Modular Redundancy (DMR)

**ECC Memory +
Triple Modular Redundancy (TMR)**

Byzantine Fault Tolerance (BFT)

Which of these mechanisms (or a combination thereof) should be **used** in practice?

- Transmission errors
 - ▶ Faults on the network
- Omission errors
 - ▶ Fault-induced kernel panics, hangs
- Incorrect computation errors
 - ▶ Faults in memory buffers
- Inconsistent broadcast errors
 - ▶ Faults in networked systems

Checksums and retransmissions

Dual Modular Redundancy (DMR)

ECC Memory + Triple Modular Redundancy (TMR)

Byzantine Fault Tolerance (BFT)

Which of these mechanisms (or a combination thereof) should be **used** in practice?

- Transmission errors
 - ▶ Faults on the network
- Omission errors
 - ▶ Fault-induced kernel panics, hangs
- Incorrect computation errors
 - ▶ Faults in memory buffers
- Inconsistent broadcast errors
 - ▶ Faults in networked systems

Industry:  **RULE**

Checksums and retransmissions

Dual Modular Redundancy (DMR)

ECC Memory + Triple Modular Redundancy (TMR)

Byzantine Fault Tolerance (BFT)

Which of these mechanisms (or a combination thereof) should be **used** in practice?

- Transmission errors
 - ▶ Faults on the network
- Omission errors
 - ▶ Fault-induced kernel panics
- Incorrect computations
 - ▶ Faults in memory buffers
- Inconsistent broadcast errors
 - ▶ Faults in networked systems

Industry:  **RULE**

SWaP-C
Size, Weight, and Power ...
plus Cost

Checksums and retransmissions

Dual Modular Redundancy (DMR)

ECC Memory + Triple Modular Redundancy (TMR)

Byzantine Fault Tolerance (BFT)

Which of these mechanisms (or a combination thereof) should be **used** in practice?

- Transmission errors

Real-time requirements

Industry:  **RULE**

SWaP-C

Size, Weight, and Power ...
plus Cost

- Incorrect computations
- ▶ Fault-induced kernel panics
- Faults in memory buffers

- Inconsistent broadcast errors
- ▶ Faults in networked systems

Checksums and retransmissions

Dual Modular Redundancy (DMR)

ECC Memory + Triple Modular Redundancy (TMR)

Byzantine Fault Tolerance (BFT)

Which of these mechanisms (or a combination thereof) should be **used** in practice?

- Transmission errors

Real-time requirements

Industry:  **RULE**

Safety certification

- ▶ Reliability thresholds
- ▶ $< 10^{-9}$ failures/hour

SWaP-C
Size, Weight, and Power ...
plus Cost

- Inconsistent broadcast errors
- ▶ Faults in networked systems

Checksums and retransmissions

Dual Modular Redundancy (DMR)

ECC Memory + Triple Modular Redundancy (TMR)

Byzantine Fault Tolerance (BFT)

Which of these mechanisms (or a combination thereof) should be **used** in practice?

- Transmission errors

Real-time requirements

Industry:  **RULE**

Safety certification

- ▶ Reliability thresholds
- ▶ $< 10^{-9}$ failures/hour

SWaP-C
Size, Weight, and Power ...
plus Cost

- Inconsistent broadcast errors

Reliability Analyses

Checksums and retransmissions

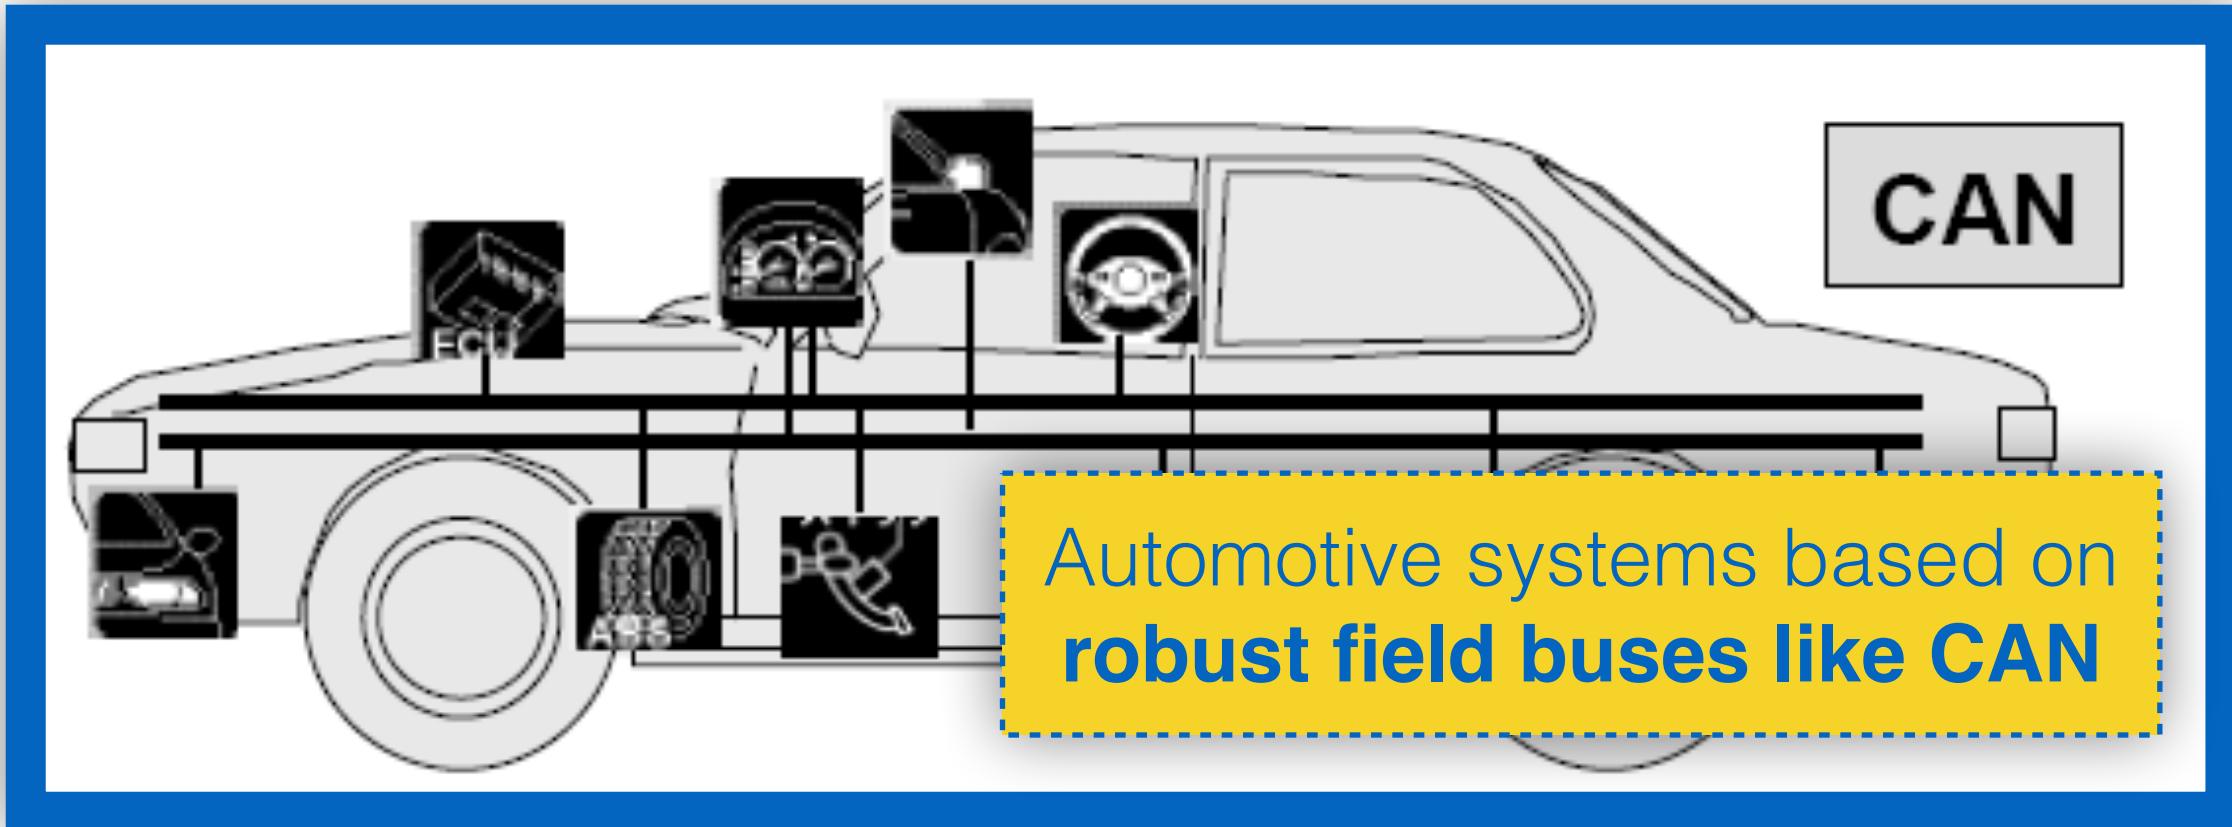
Dual Modular Redundancy (DMR)

**ECC Memory +
Triple Modular Redundancy (TMR)**

Byzantine Fault Tolerance (BFT)

Autonomous CPS landscape is changing!
New reliability analyses are necessary.

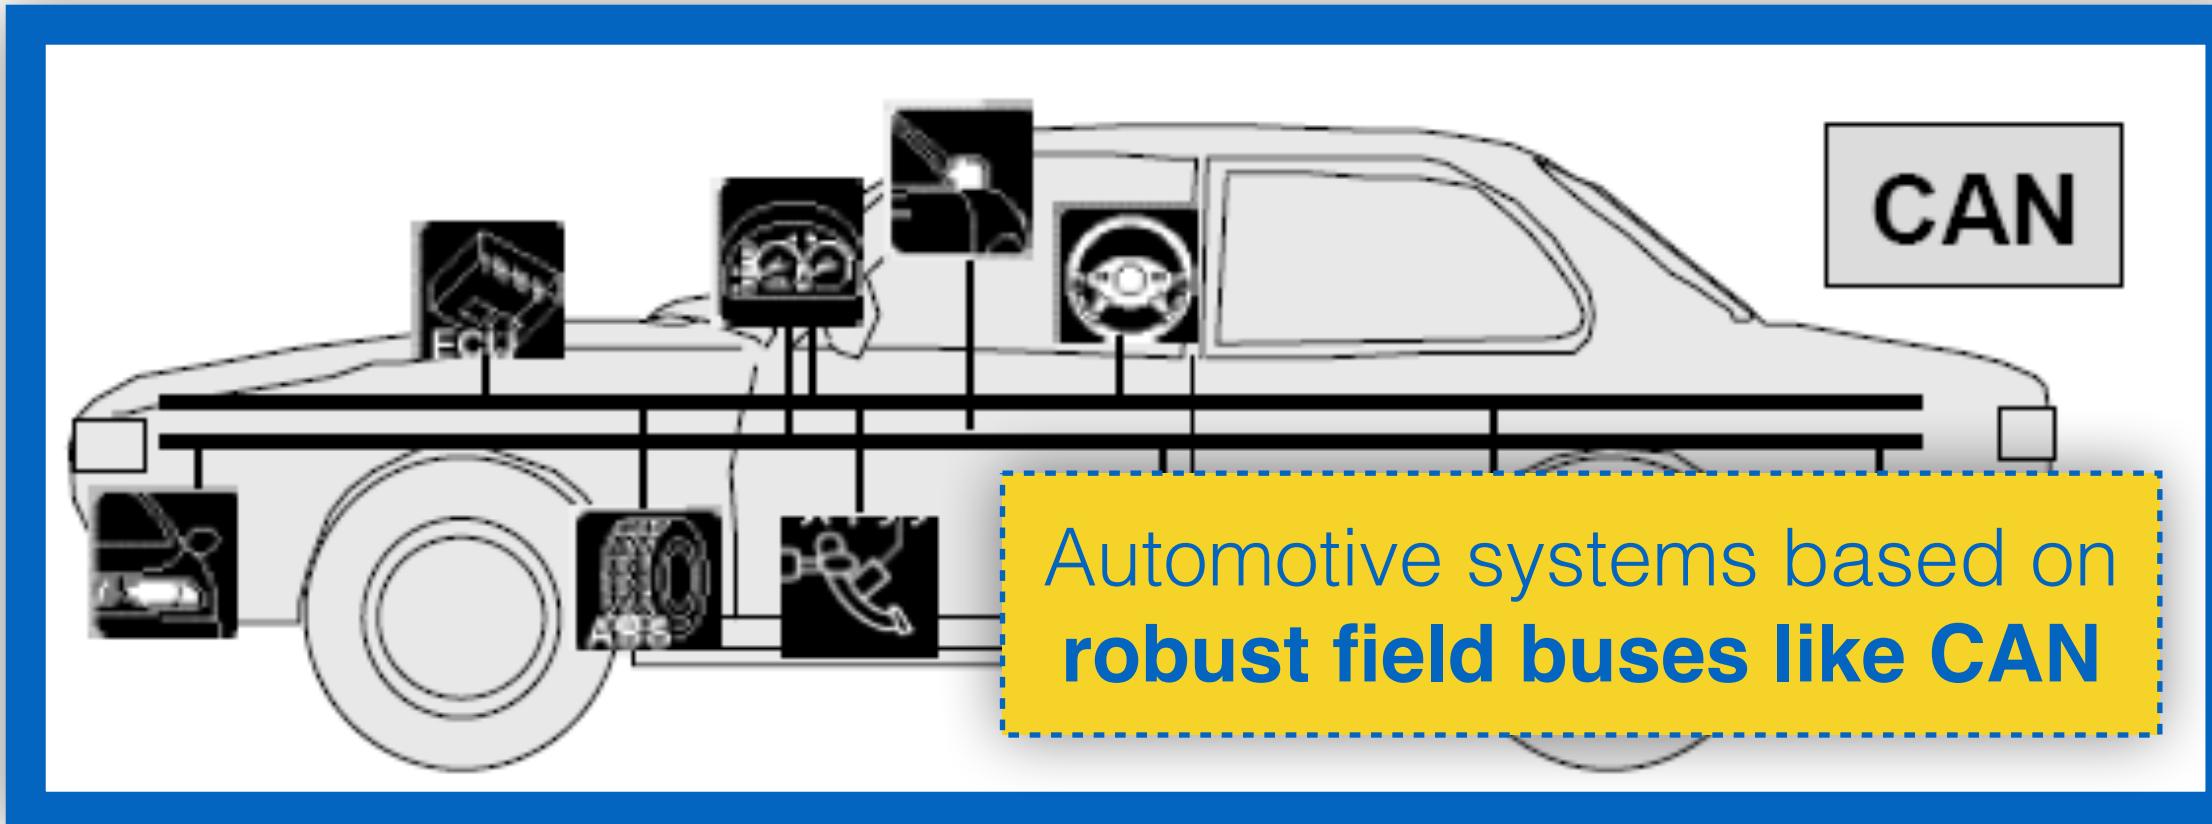
Motivating Trends



Ultra-reliability

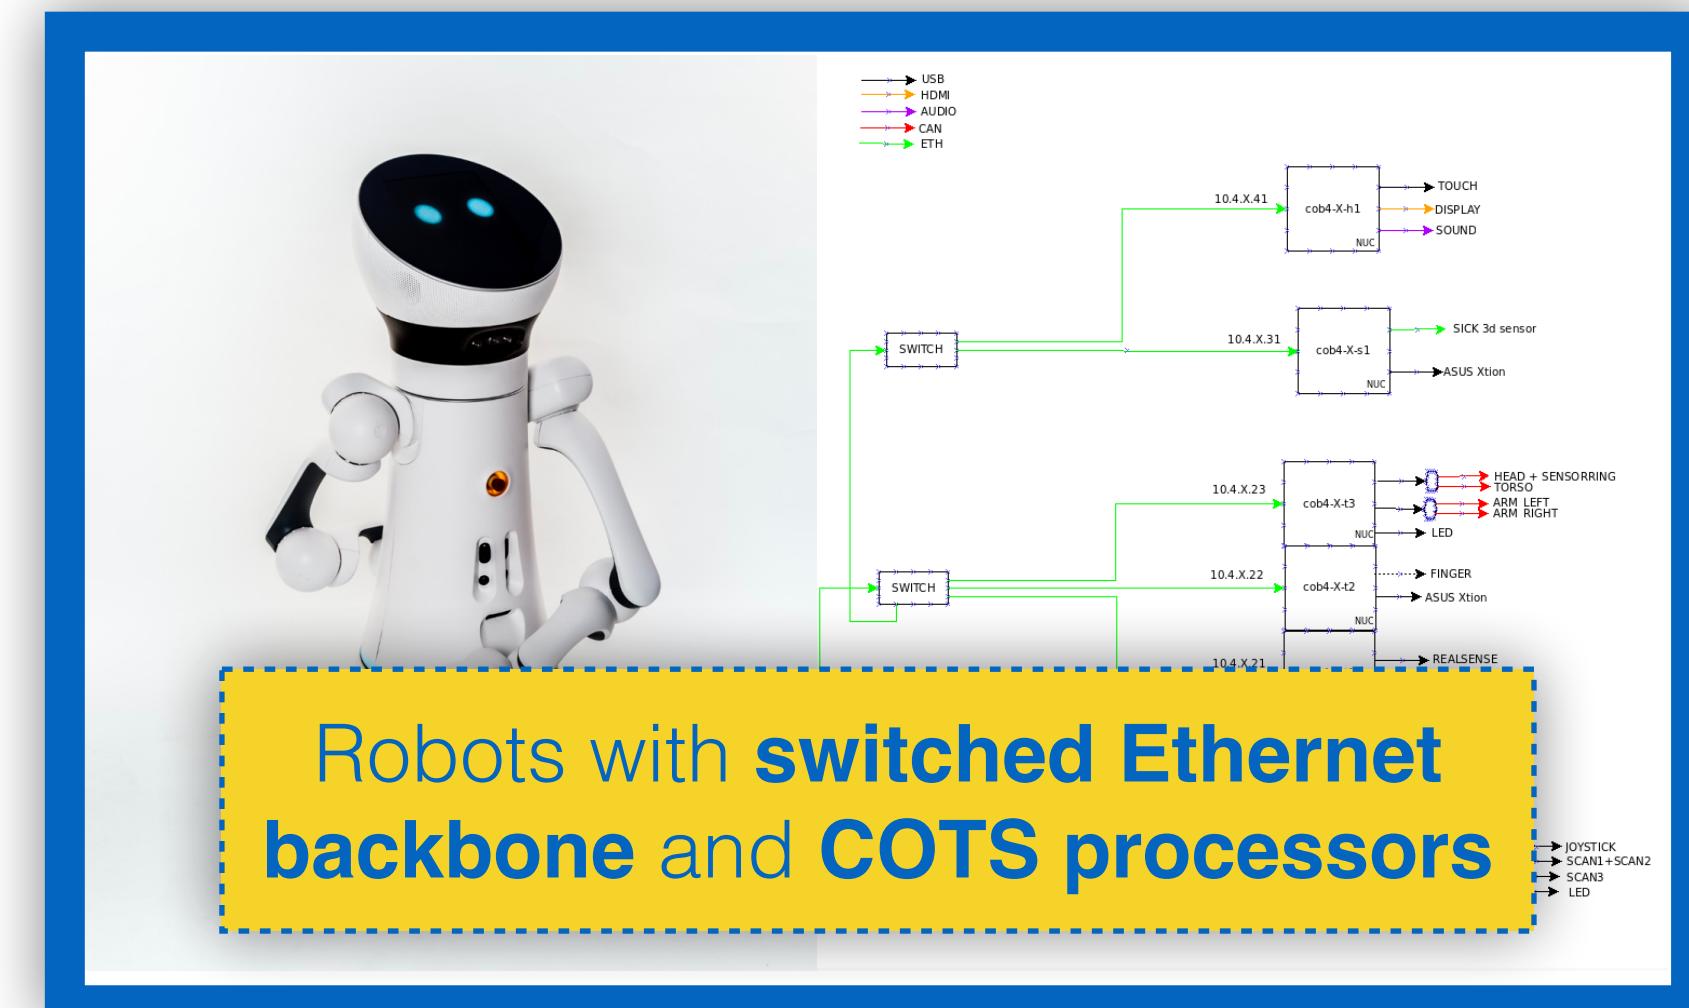
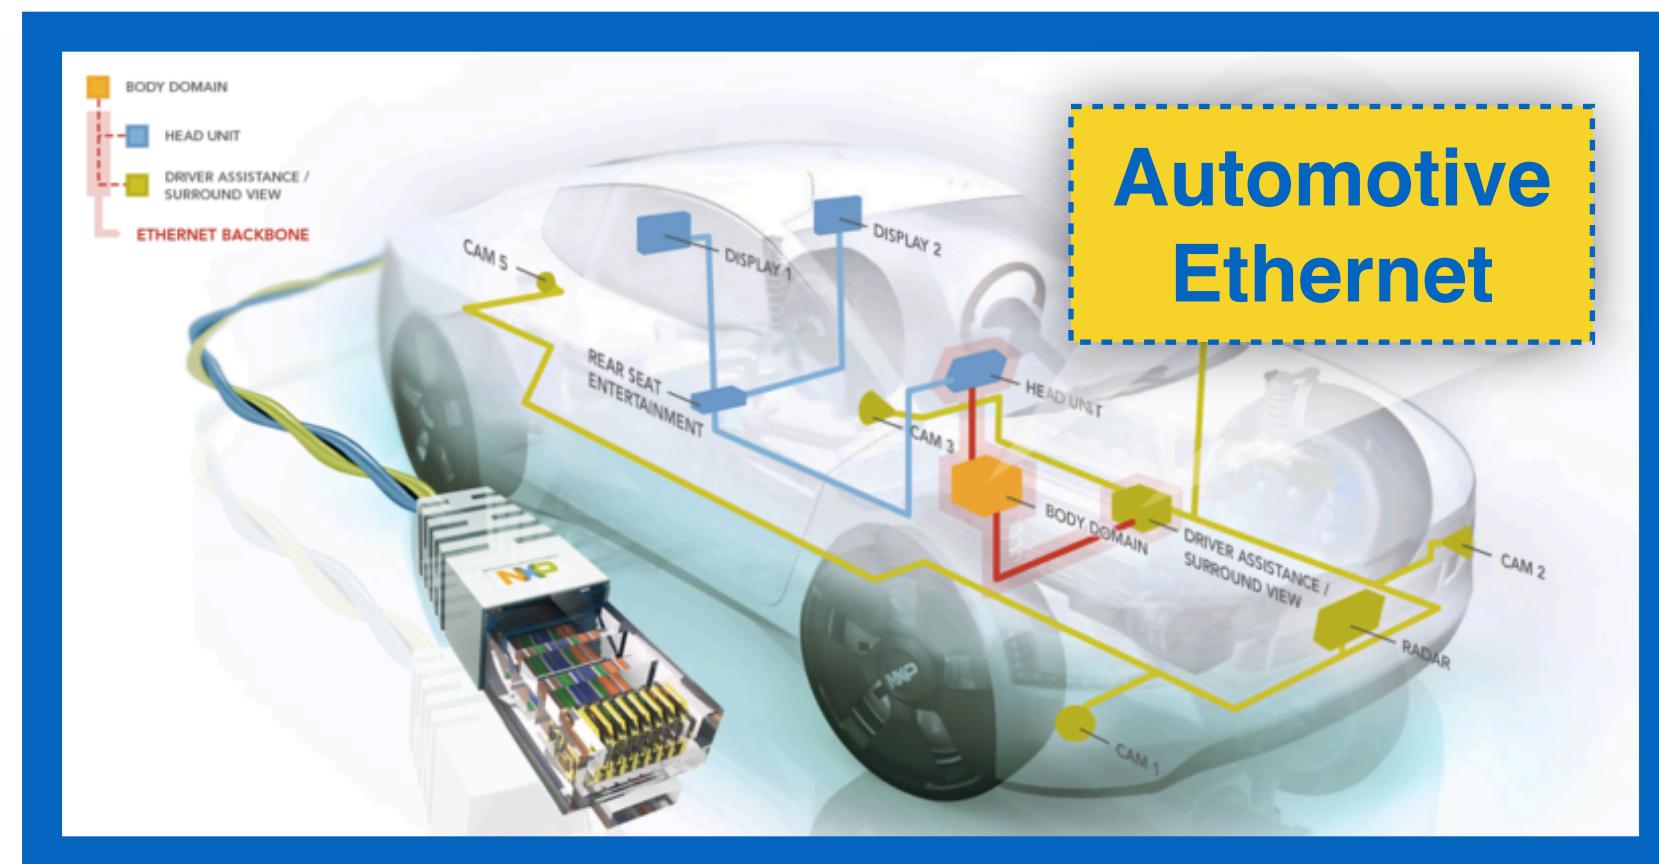
- ▶ Quantifiably negligible failure rates

Motivating Trends

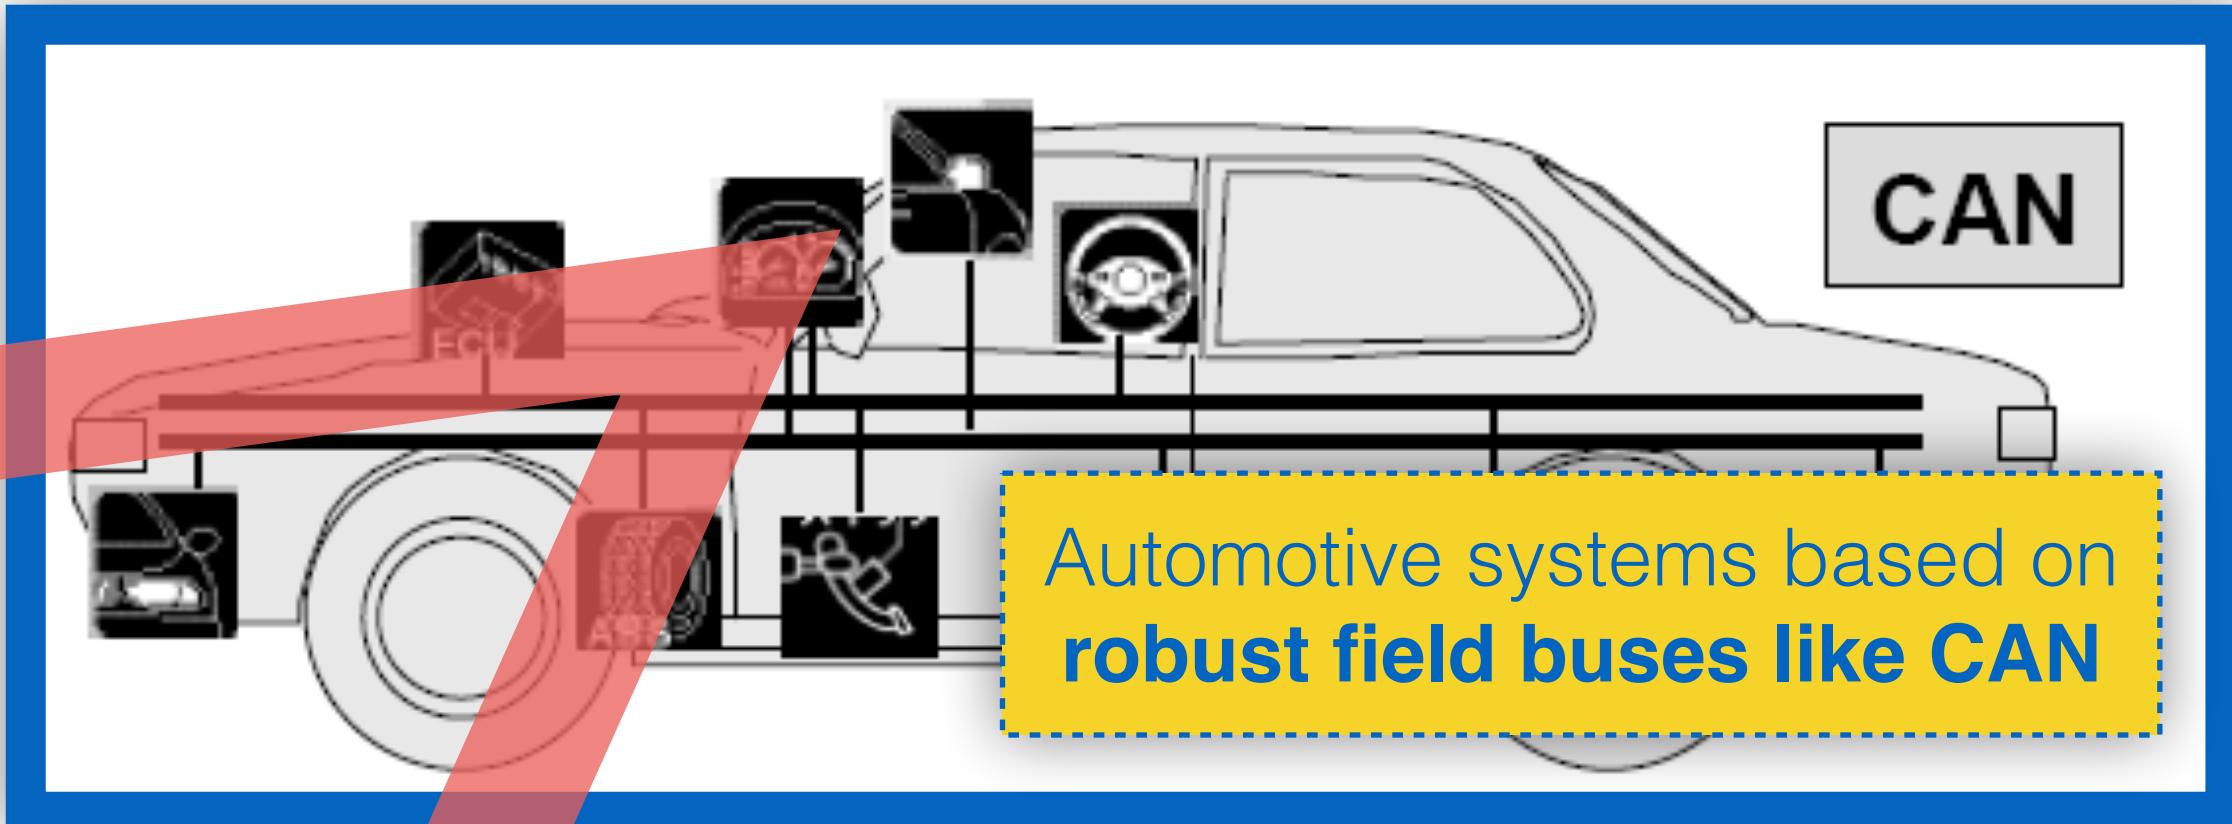


Ultra-reliability

- Quantifiably negligible failure rates



Motivating Trends

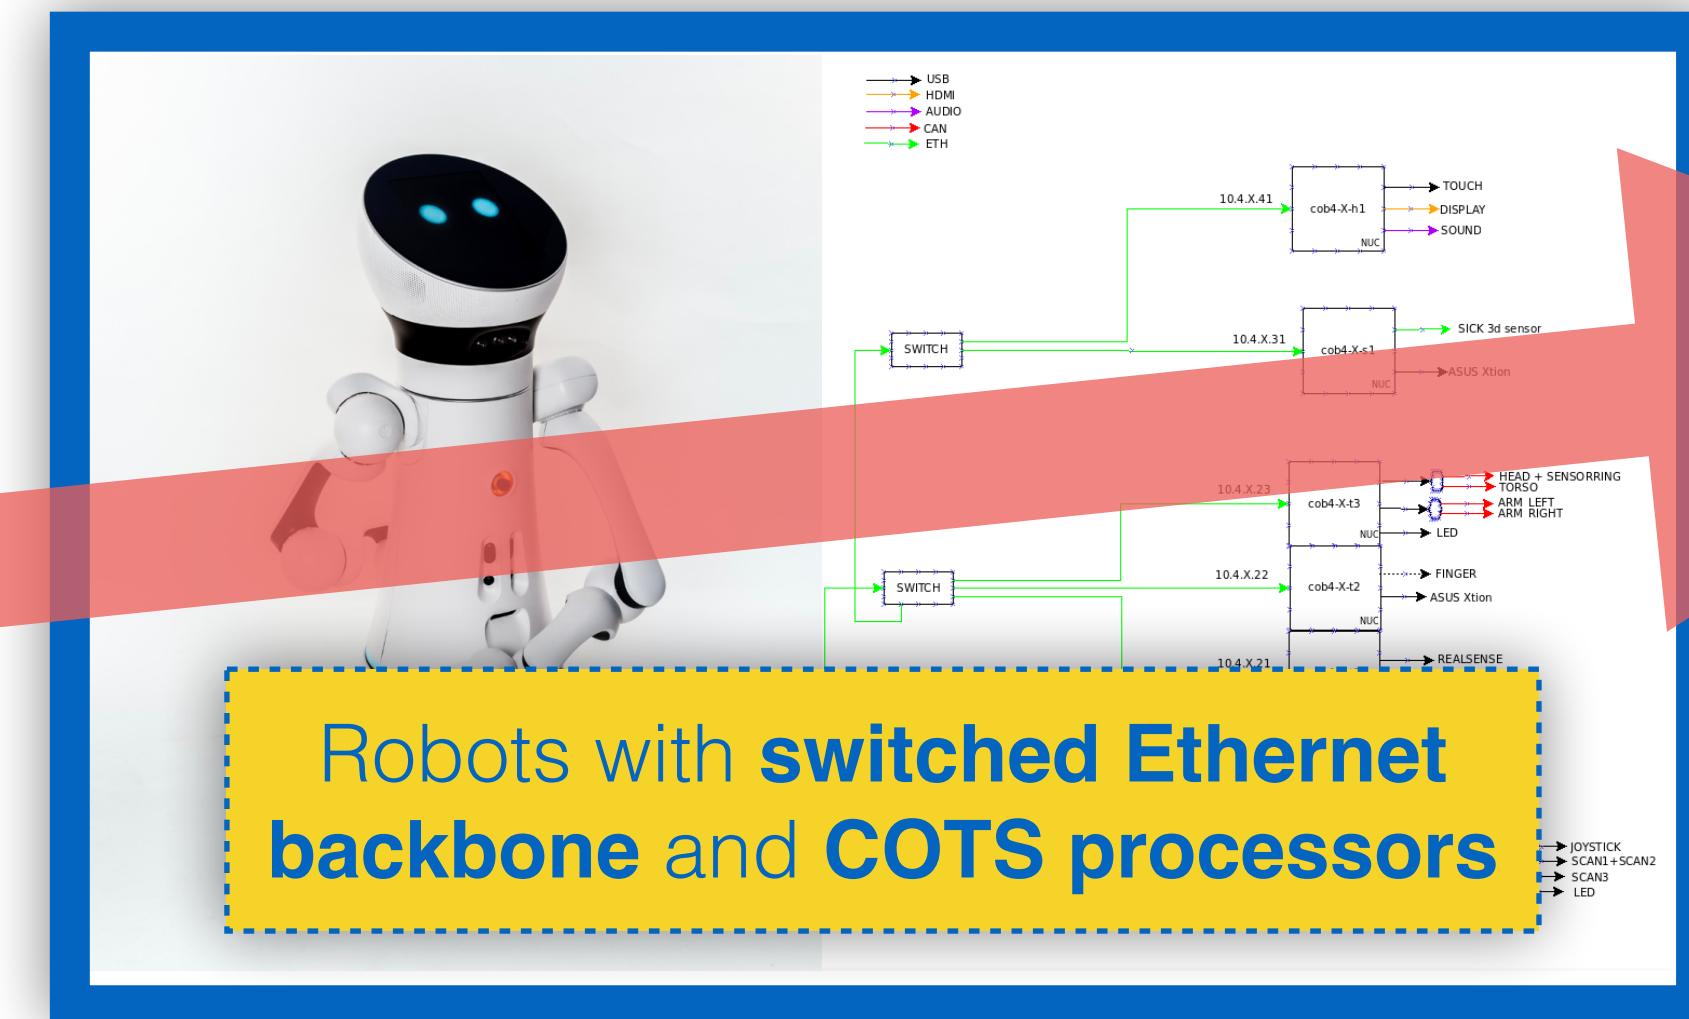


COTS hardware
(inexpensive, but unreliable)

Inadequate resources
(developer hours, computing power, component costs)

Time-to-market
pressures

Ultra-reliability
► Quantifiably negligible failure rates



Focus

Reliability analysis of Ethernet-based distributed real-time systems

PhD Thesis

Reliability Analysis of Distributed Real-Time Systems

Best Presentation Award

Controller Area Network [ECRTS '18]

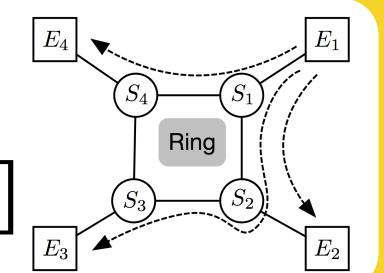


Periodic Weakly-Hard Systems [ECRTS '19]



Distinguished Paper Award

Replica Consistency over Ethernet [RTAS '20]

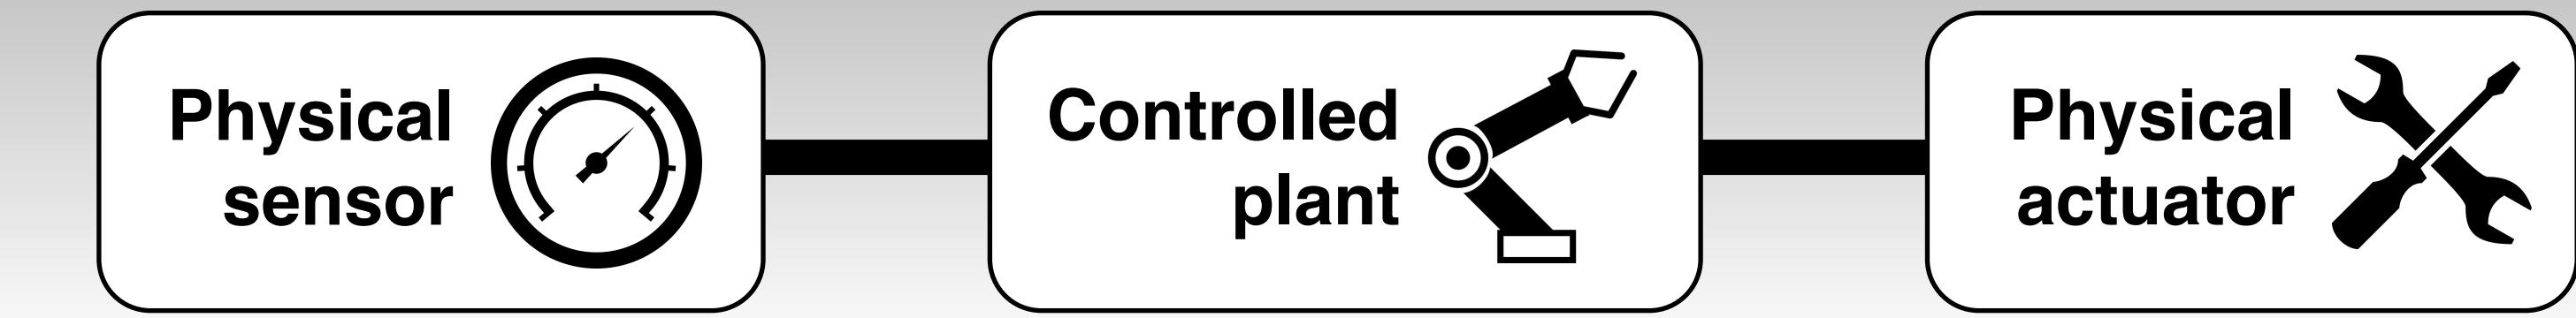


↑
This Talk

↑
This Talk

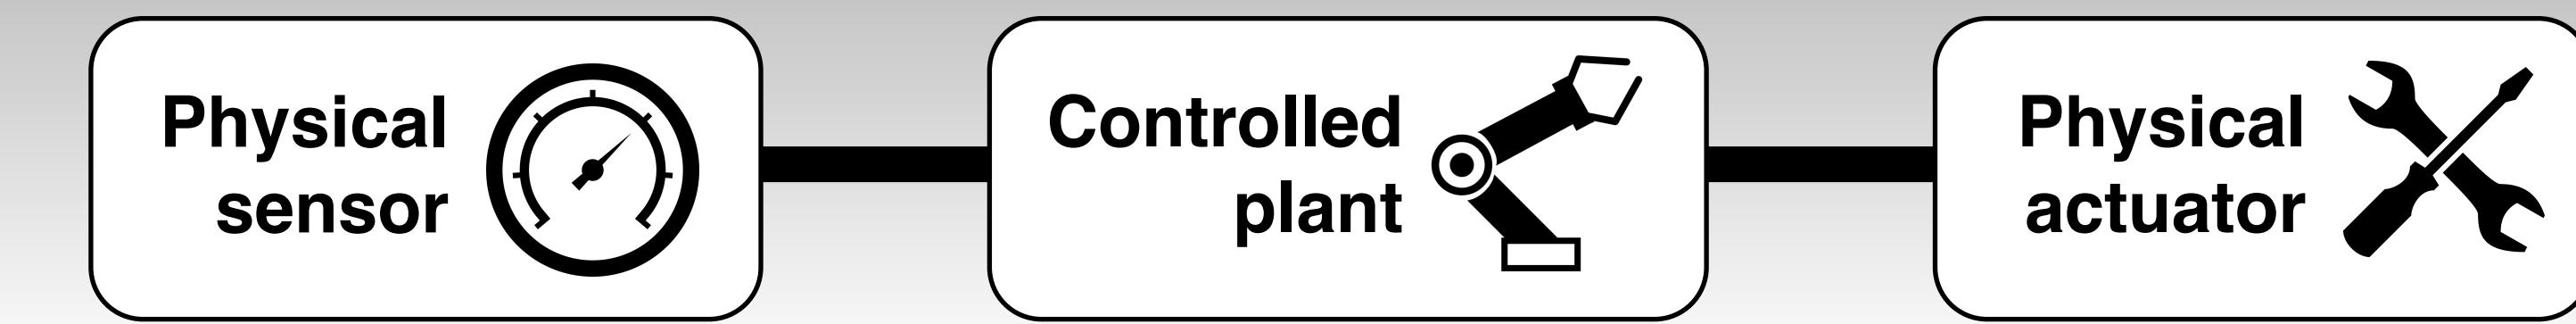
System Model

Network control system

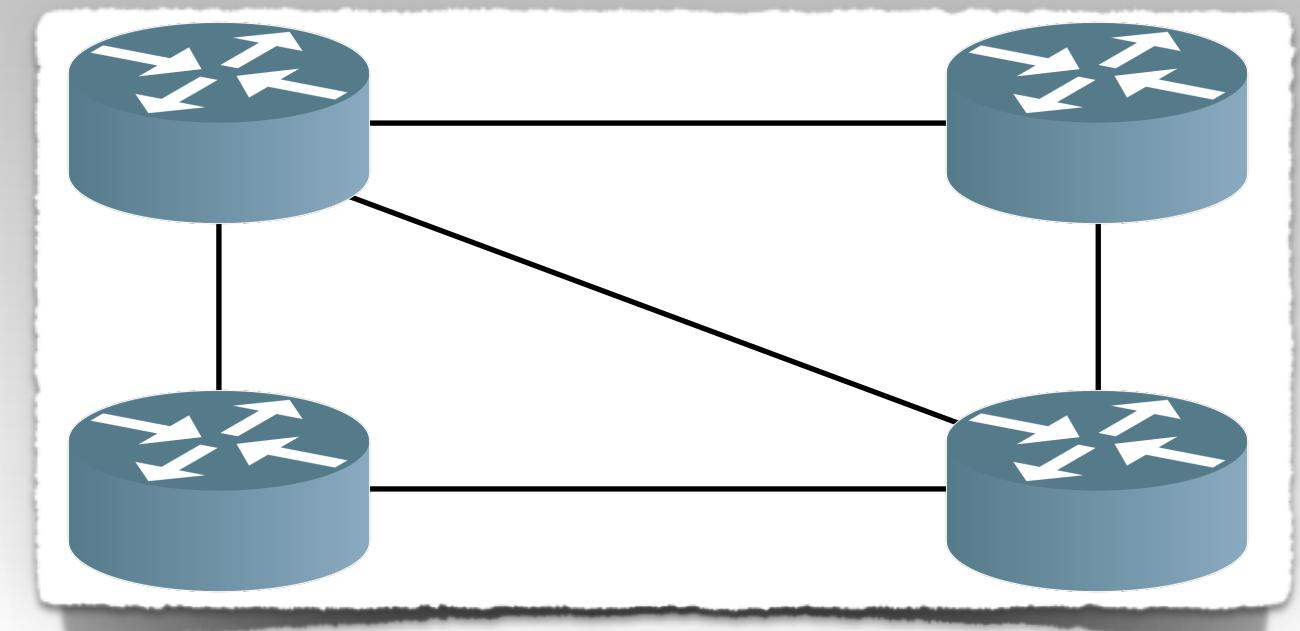


System Model

Network control system

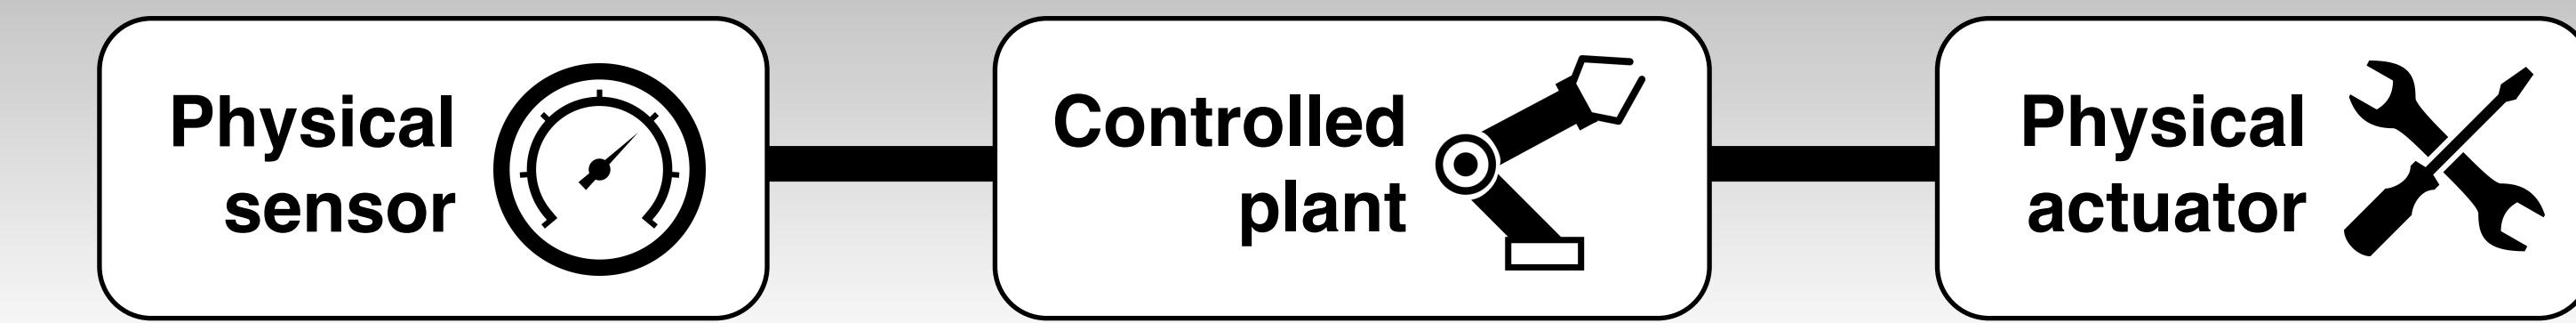


Ethernet Time-Sensitive Networking (TSN)



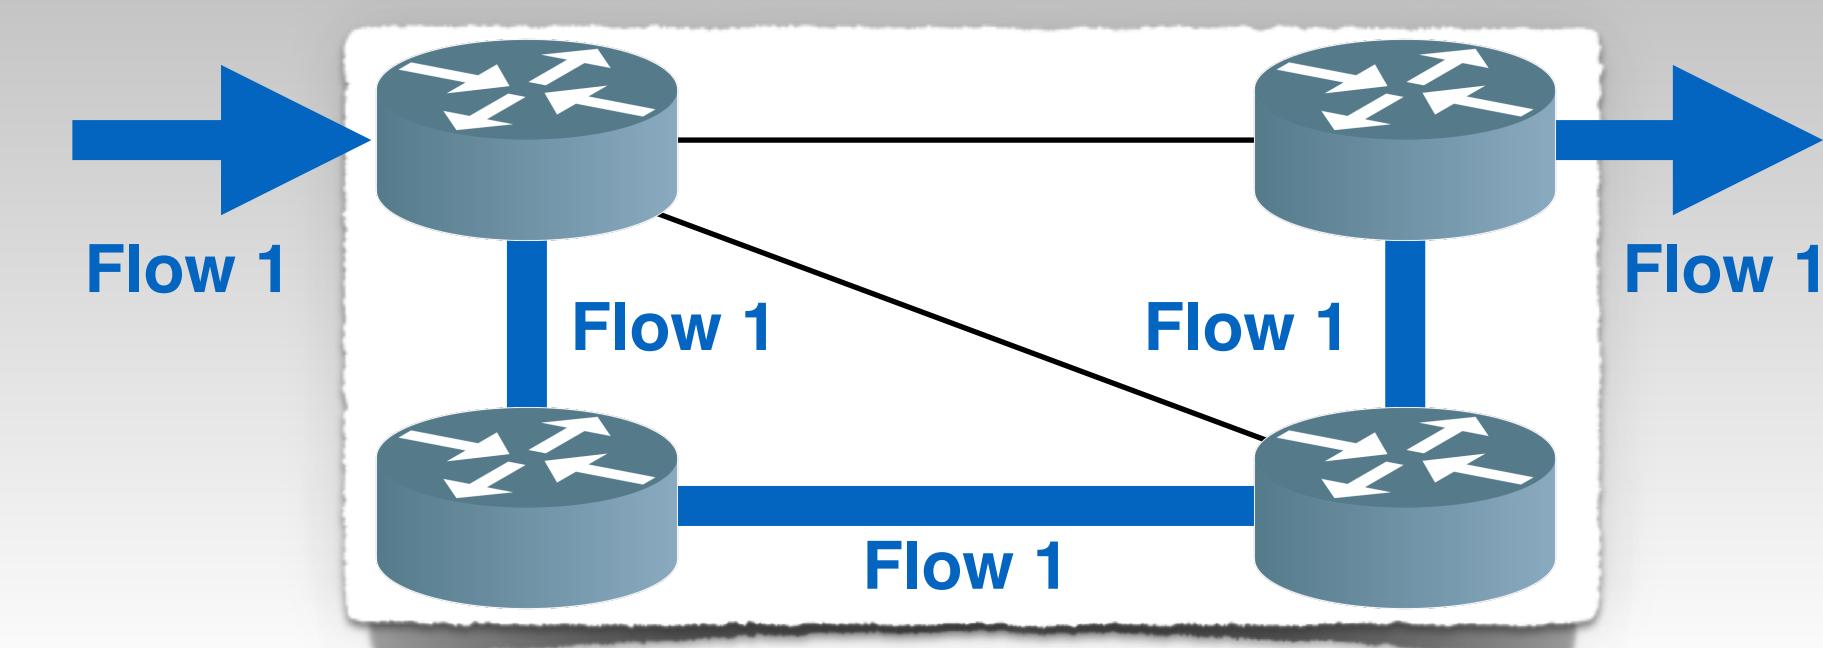
System Model

Network control system



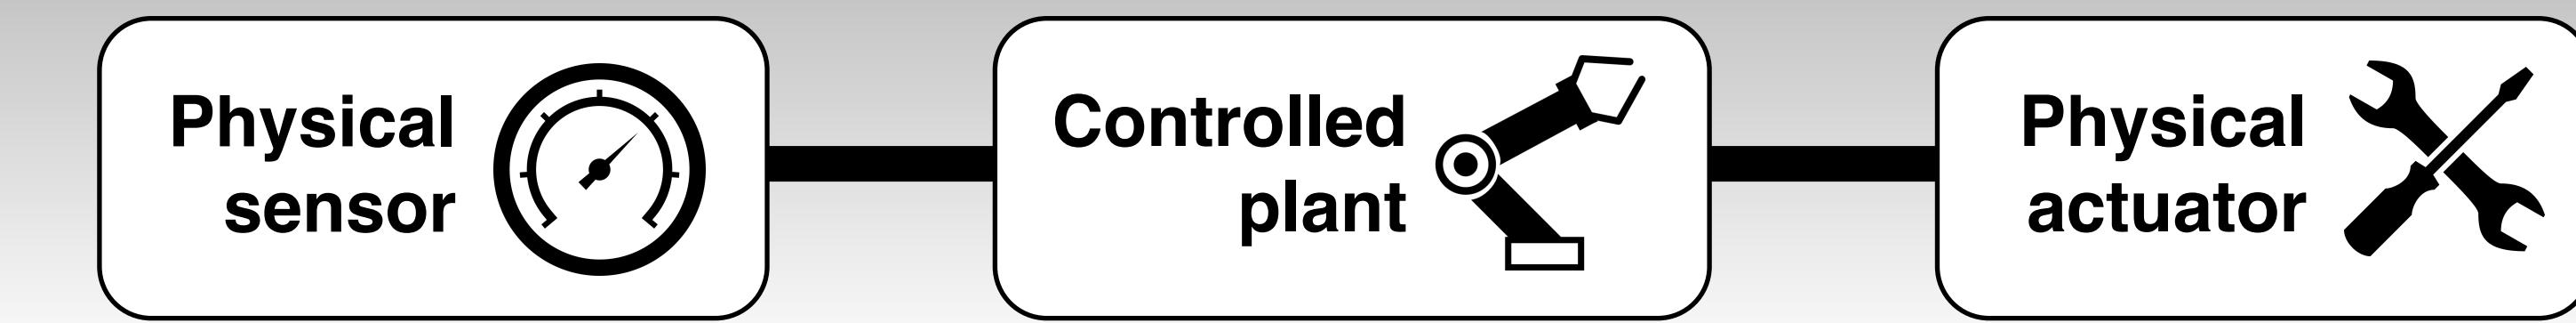
Ethernet Time-Sensitive Networking (TSN)

Statically reserved routes



System Model

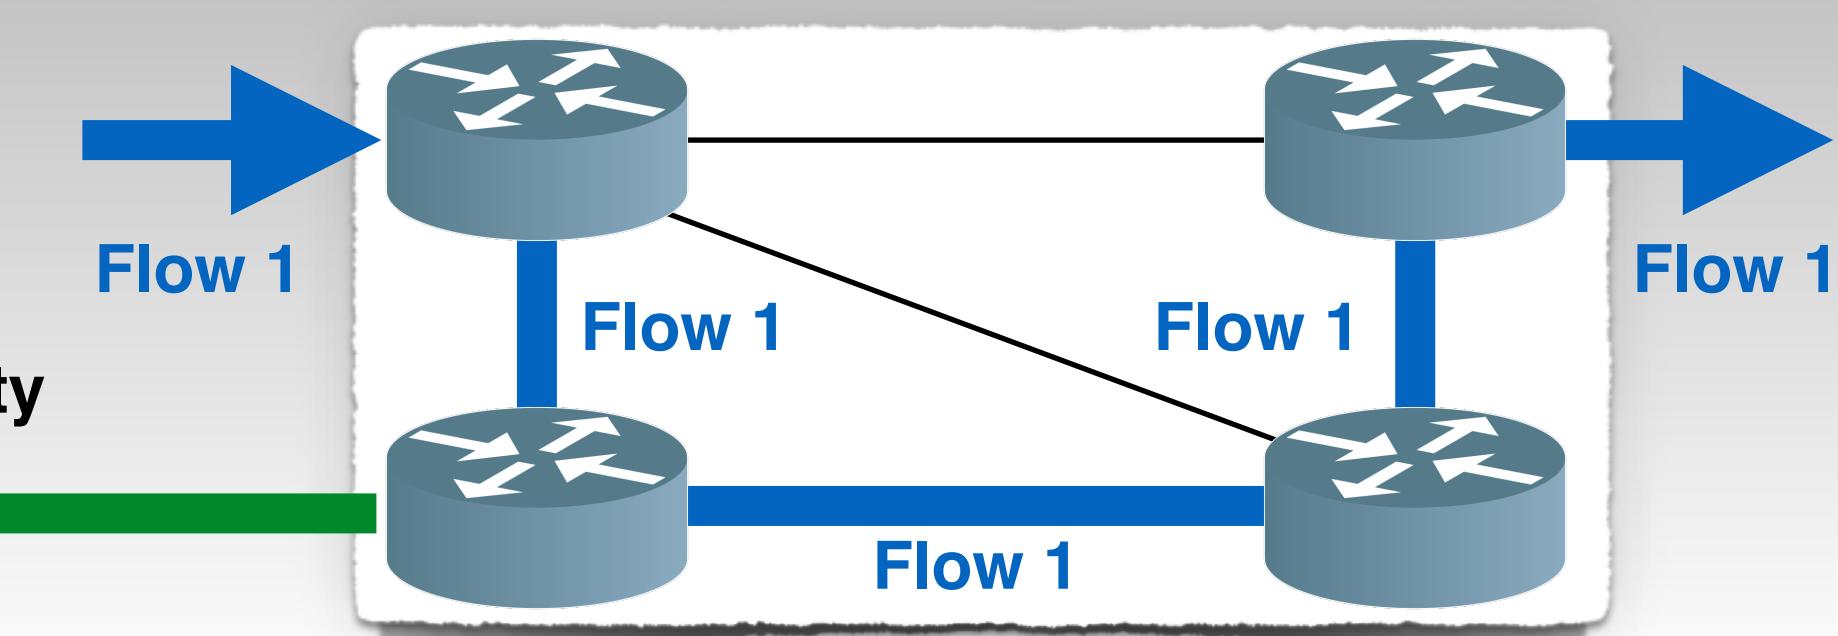
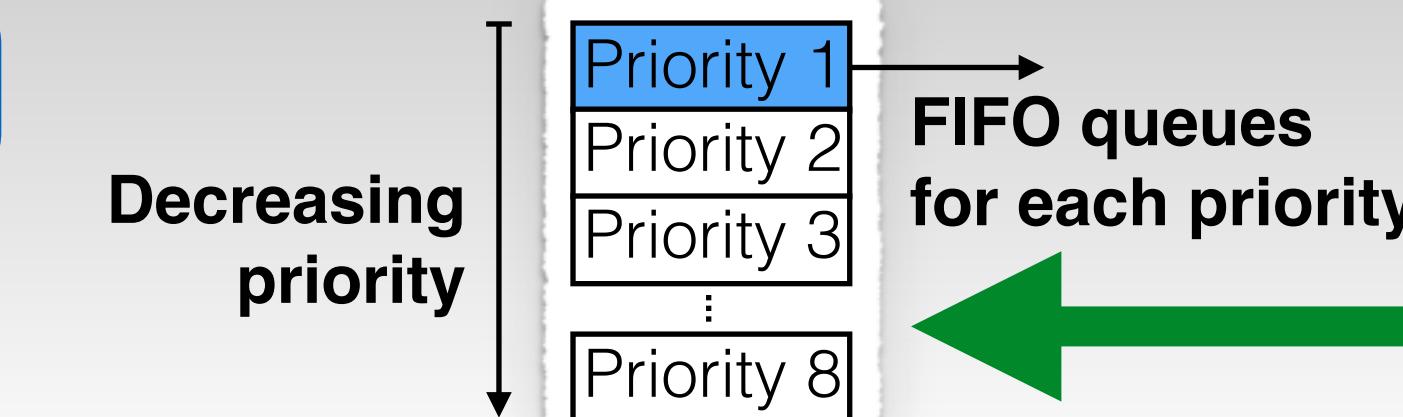
Network control system



Ethernet Time-Sensitive Networking (TSN)

Statically reserved routes

Priority classes



System Model

Network control system

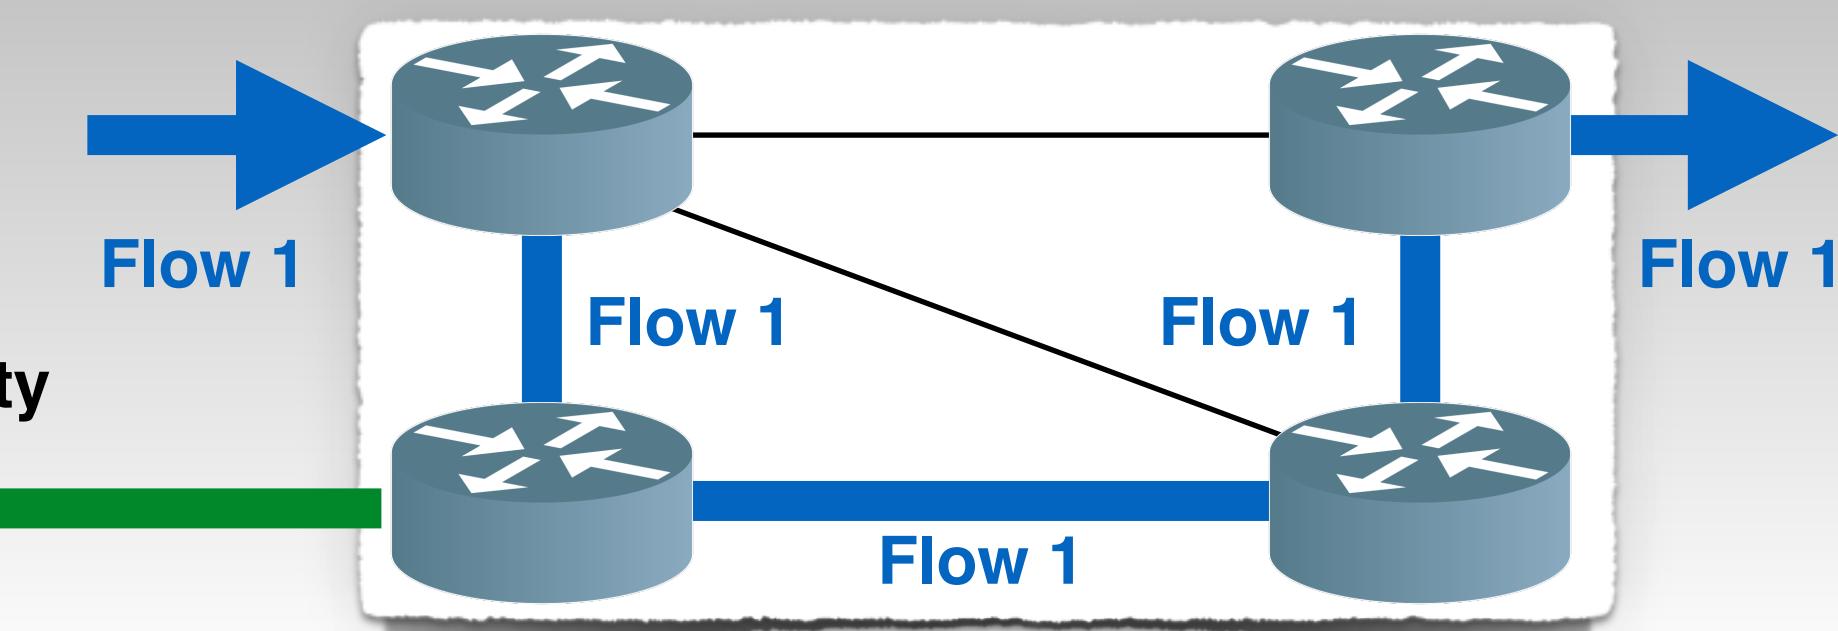
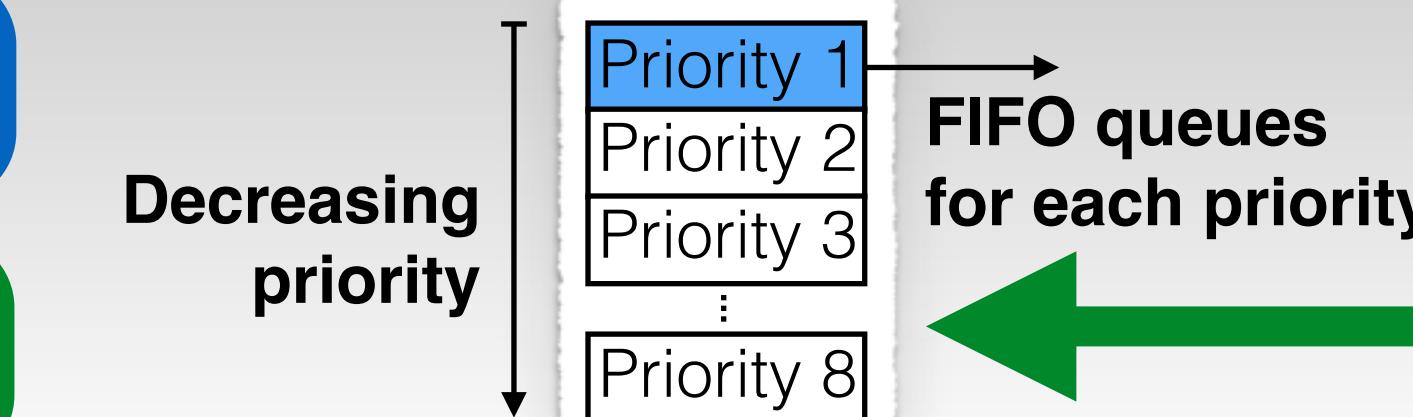


Atomic Broadcast

Ethernet Time-Sensitive Networking (TSN)

Statically reserved routes

Priority classes



System Model

Network control system

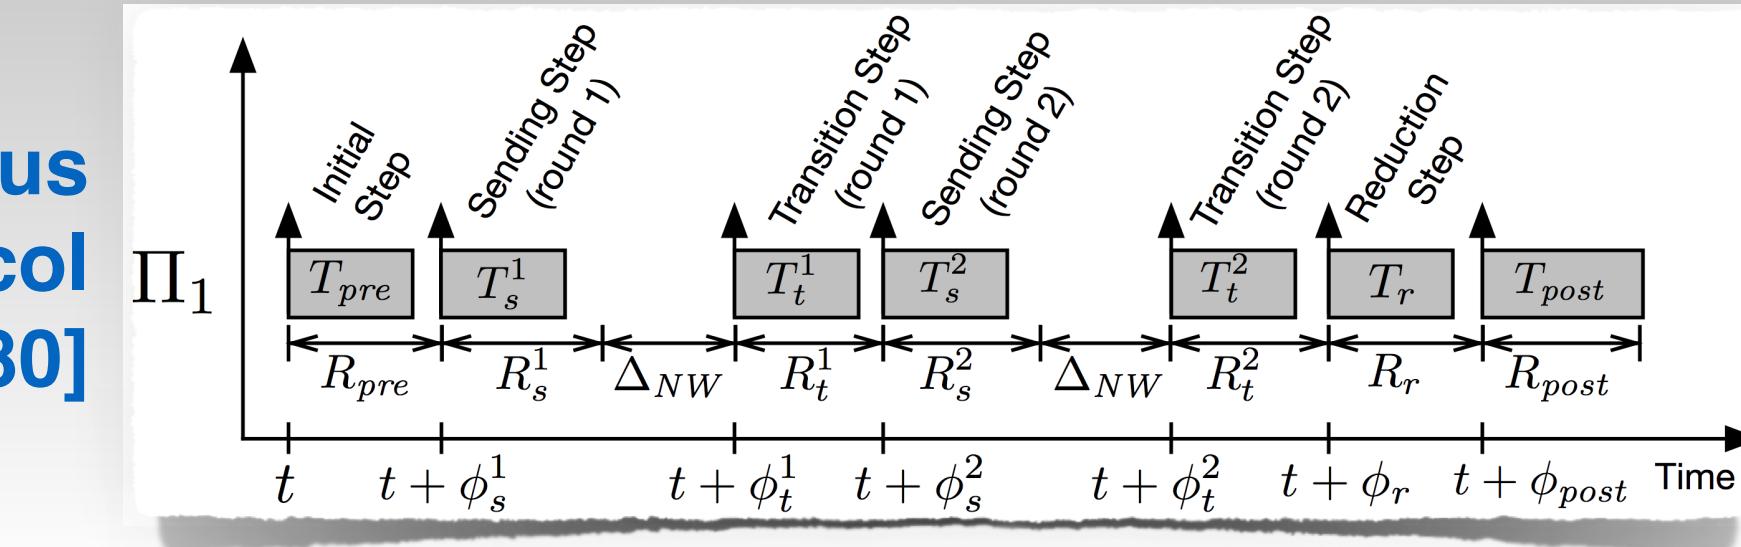


Atomic Broadcast

Statically-checked
hard real-time protocol

Synchronous
BFT protocol
[Pease et al., 1980]

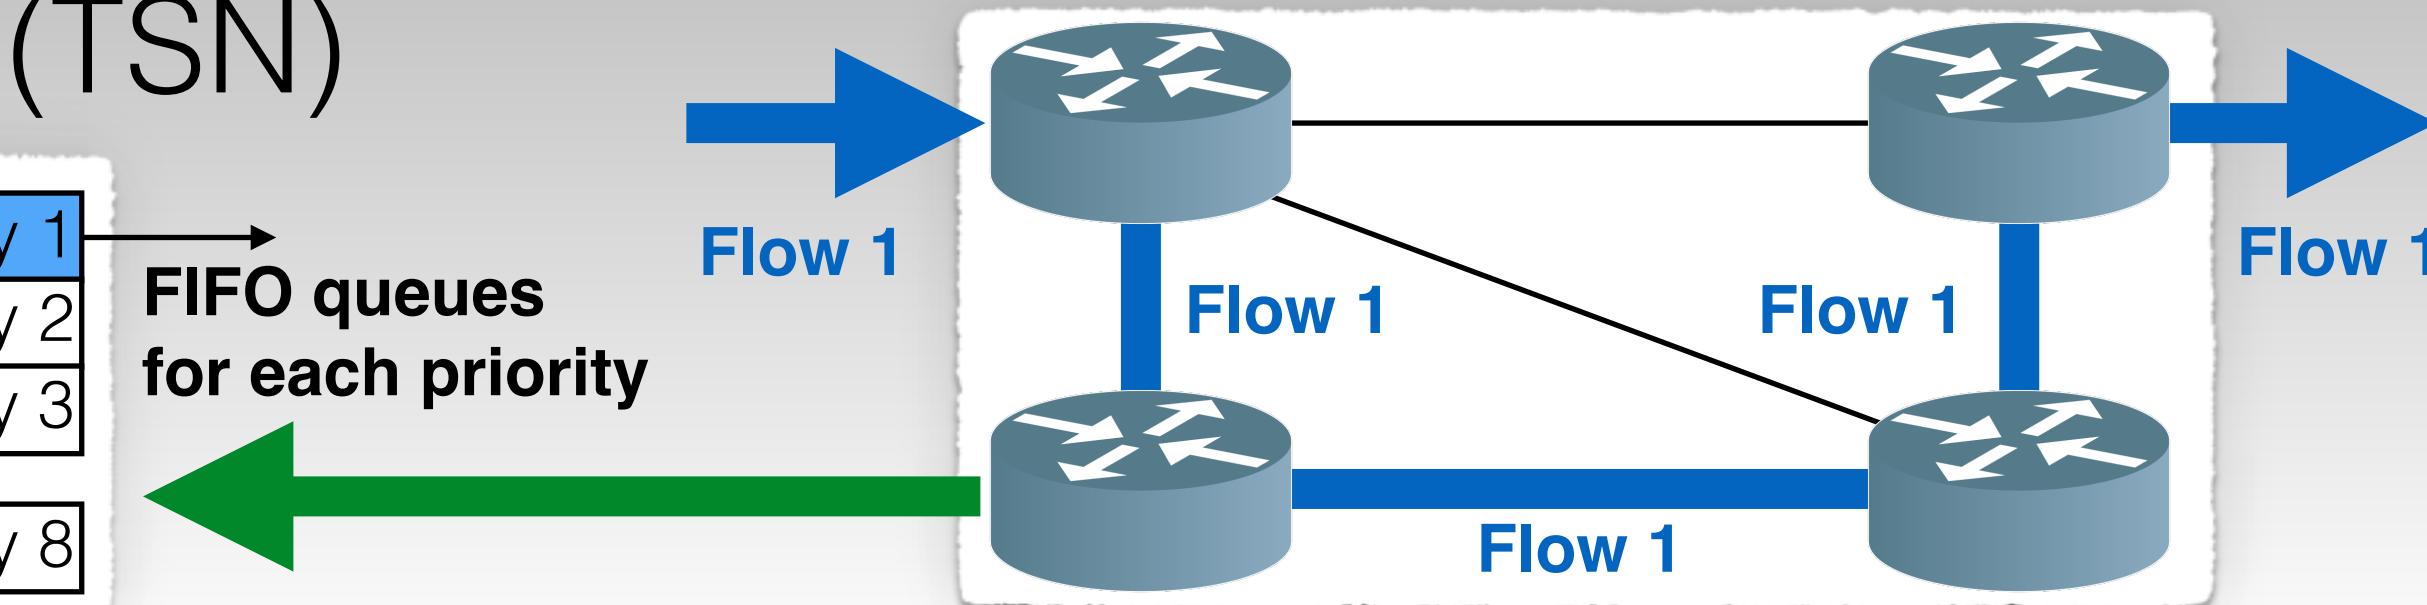
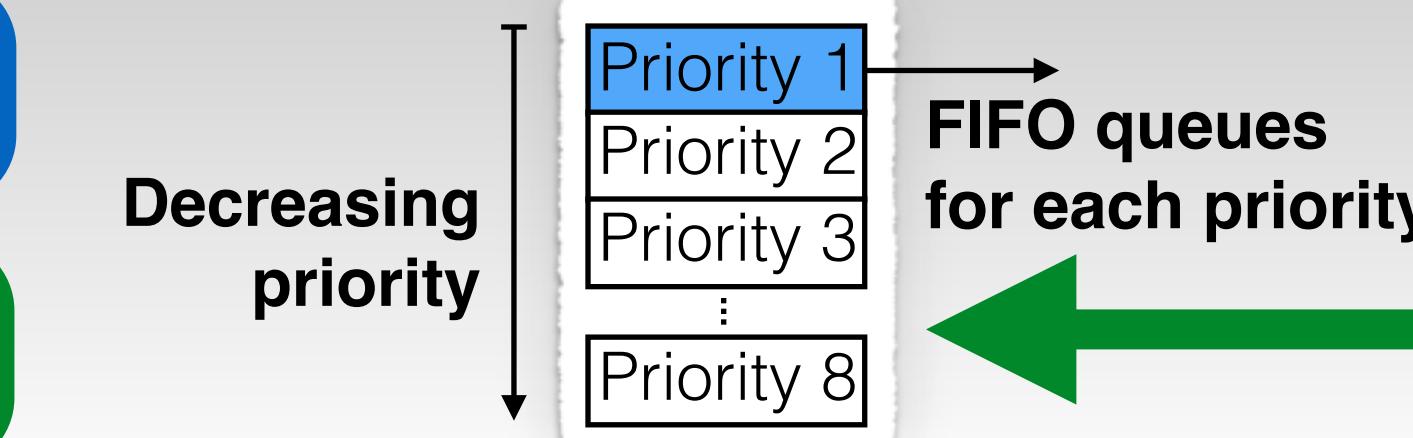
Periodic tasks and messages



Ethernet Time-Sensitive Networking (TSN)

Statically reserved routes

Priority classes



System Model

Network control system

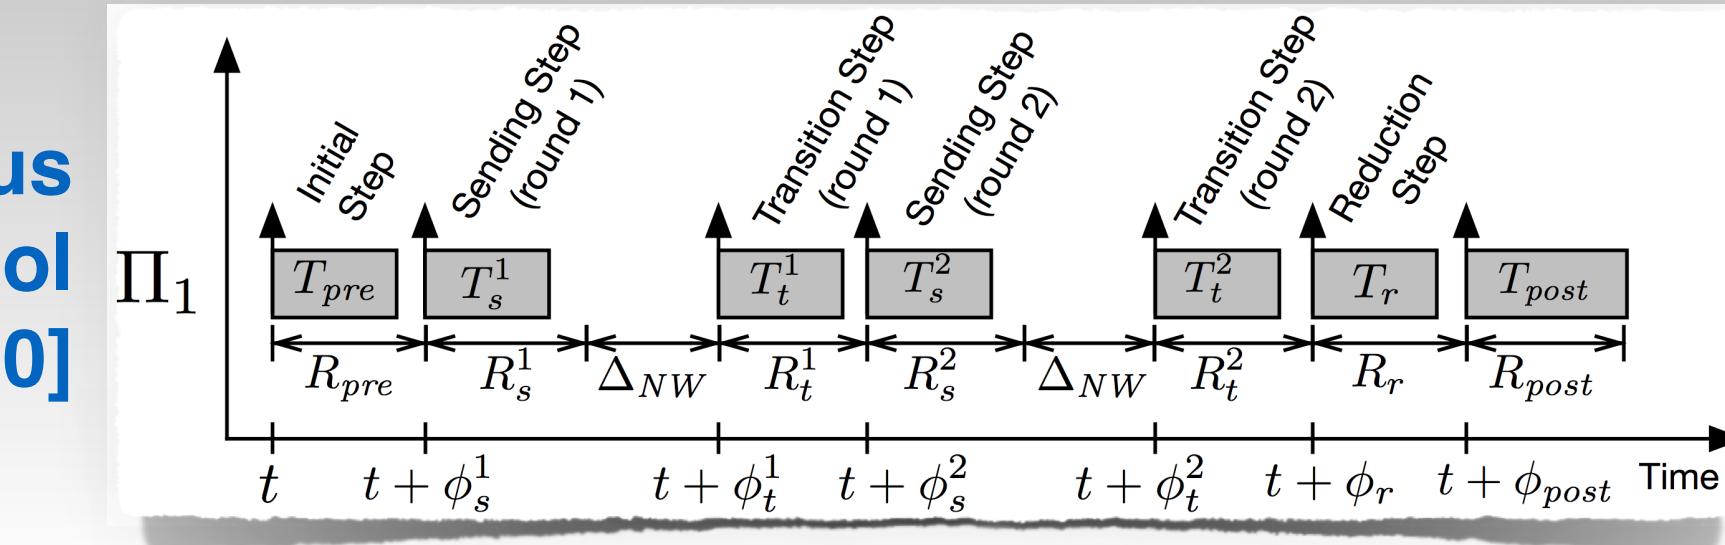


Atomic Broadcast

Statically-checked hard real-time protocol

Synchronous BFT protocol
[Pease et al., 1980]

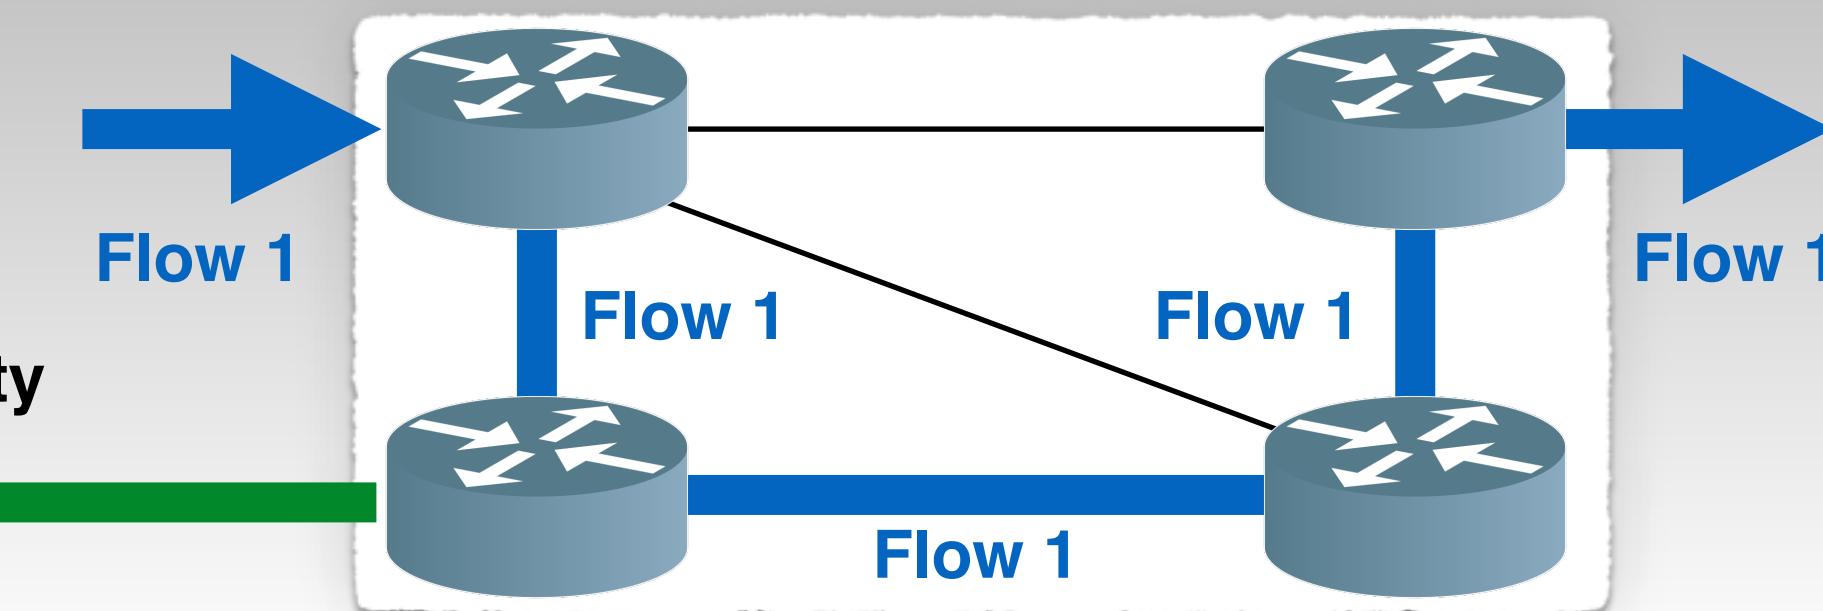
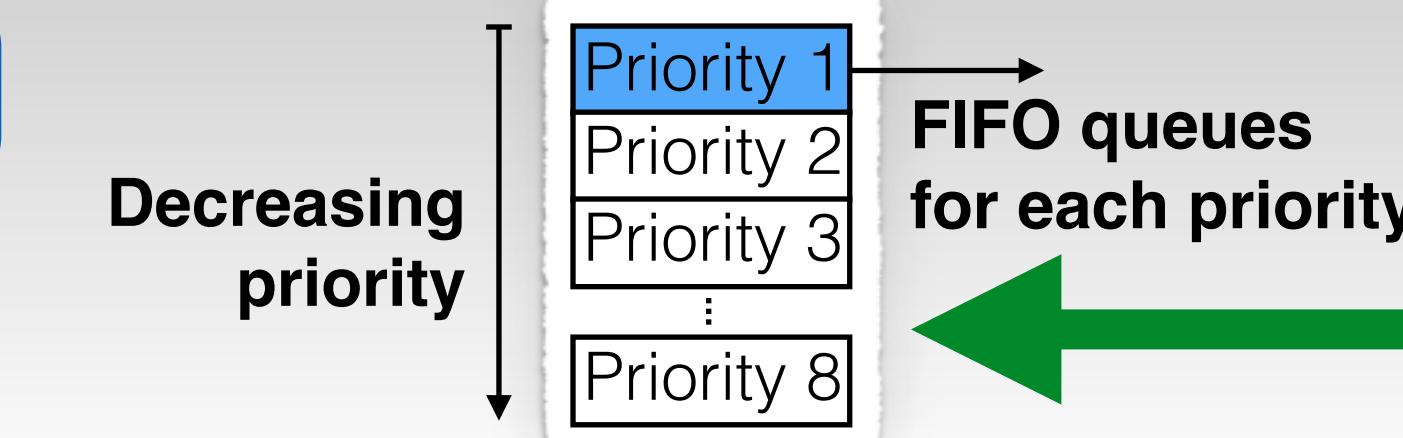
Periodic tasks and messages



Ethernet Time-Sensitive Networking (TSN)

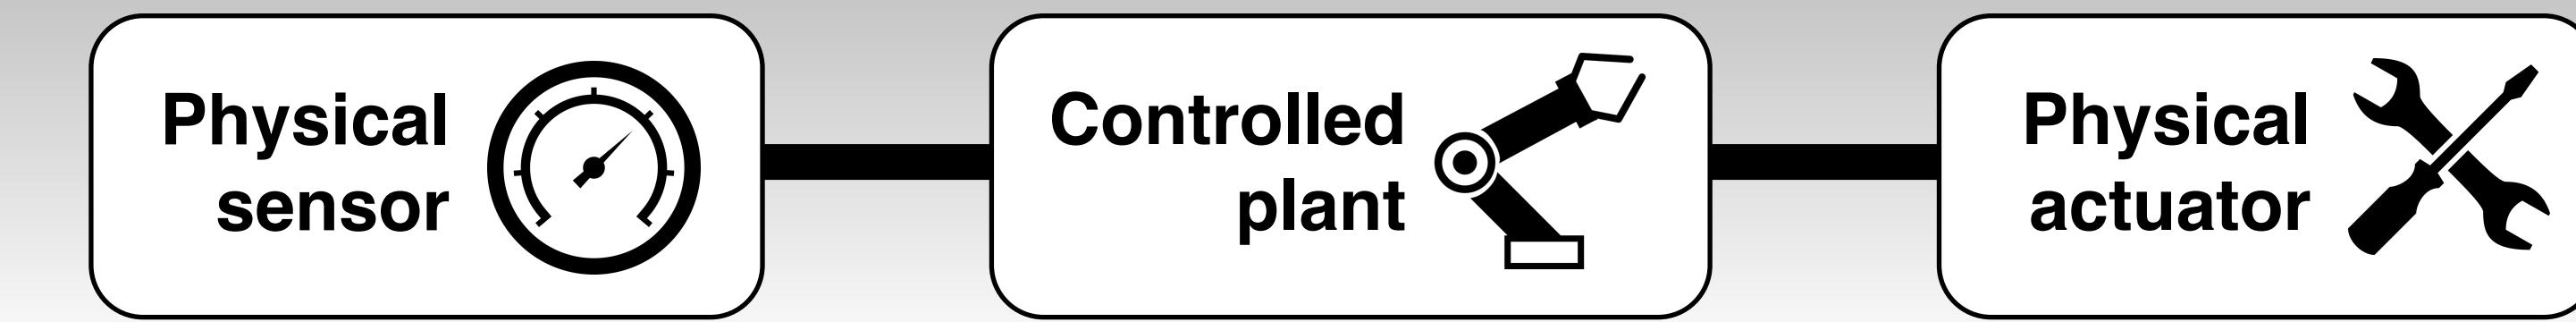
Statically reserved routes

Priority classes



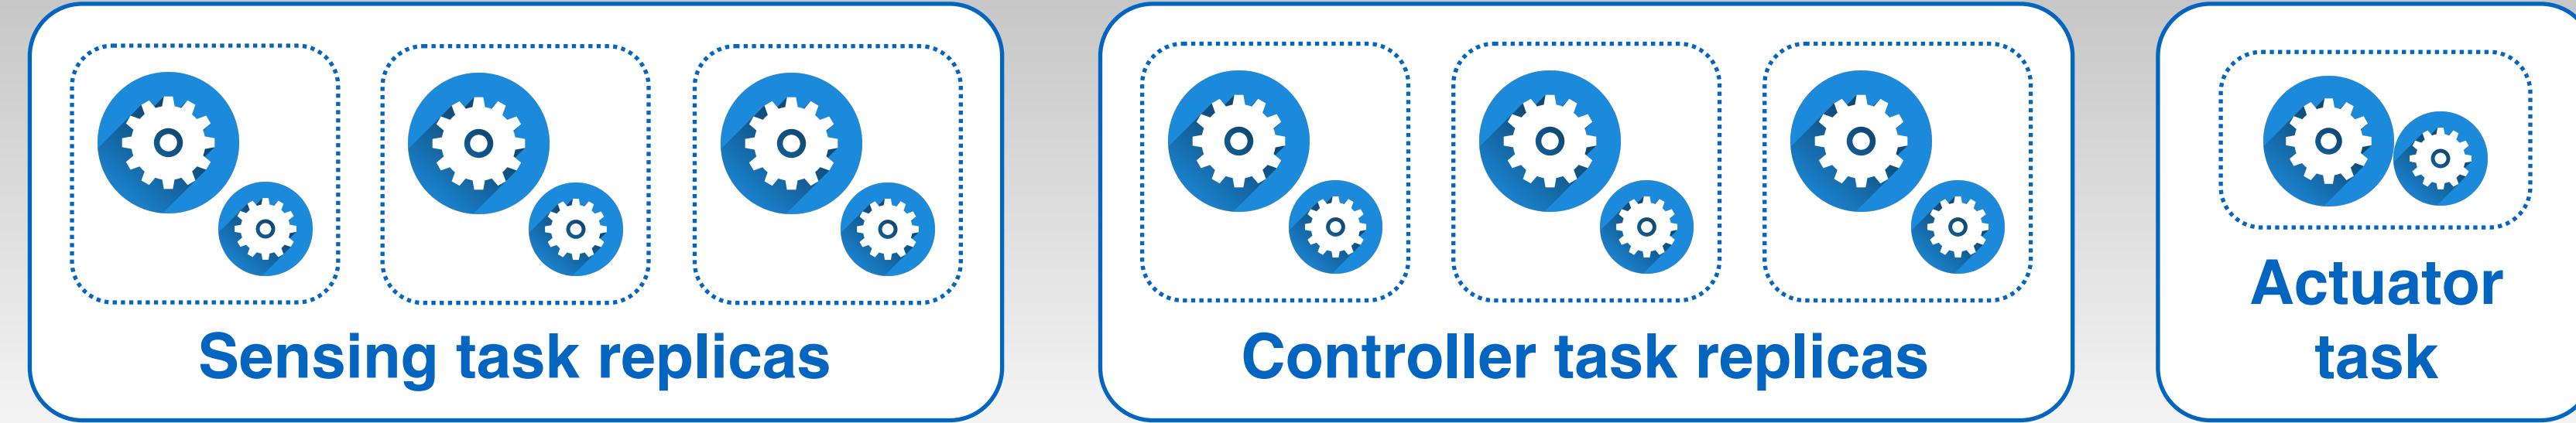
System Model

Network control system



Active Replication

DMR / TMR / Hybrid

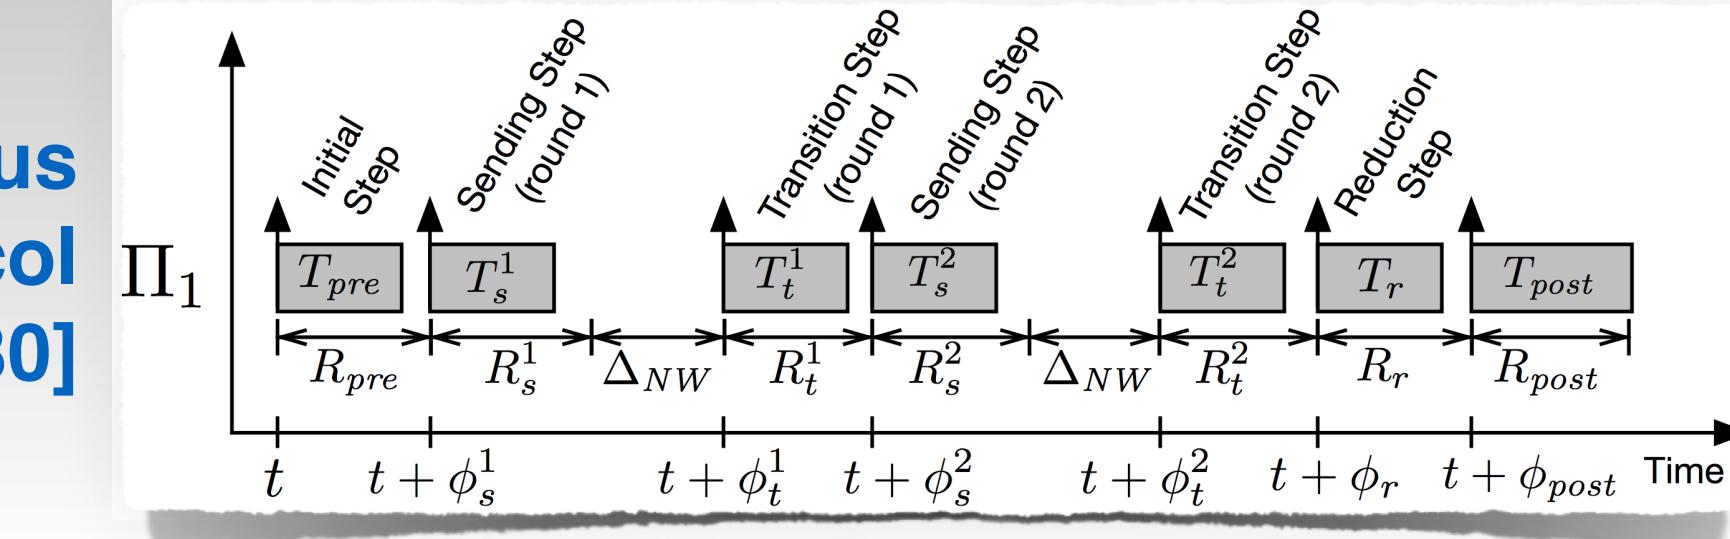


Atomic Broadcast

Statically-checked hard real-time protocol

Synchronous BFT protocol
[Pease et al., 1980]

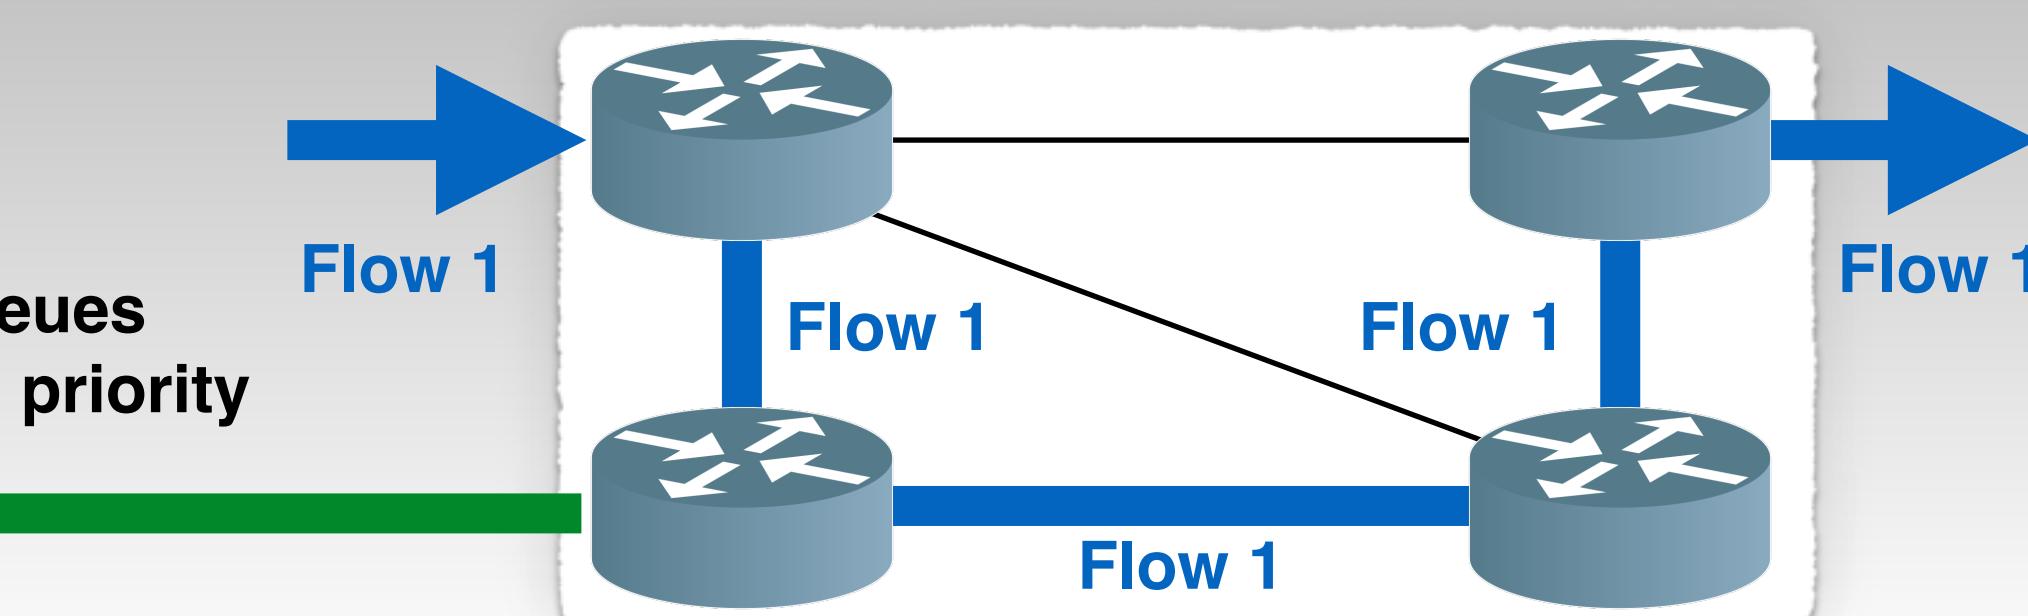
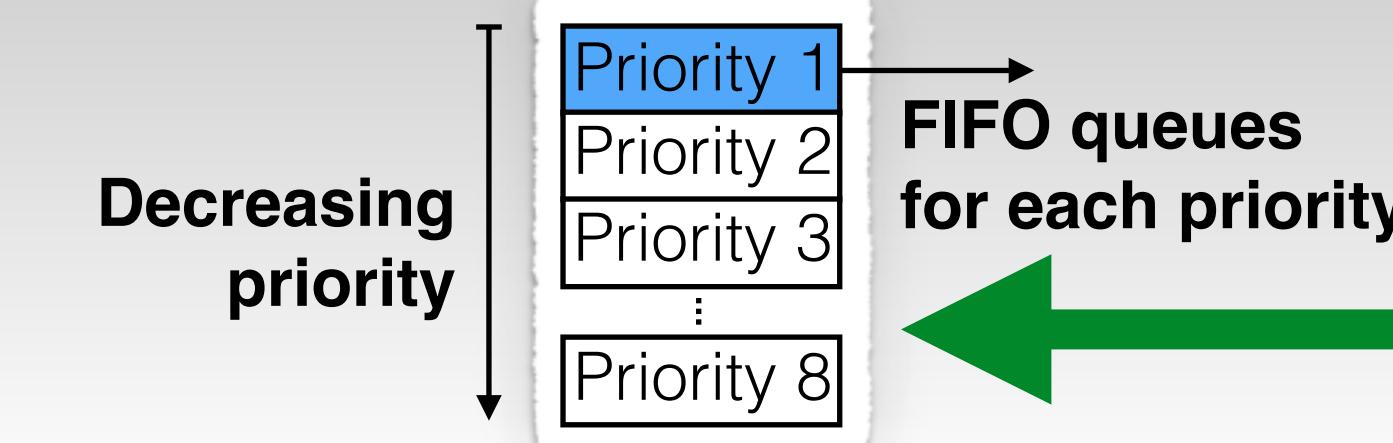
Periodic tasks and messages



Ethernet Time-Sensitive Networking (TSN)

Statically reserved routes

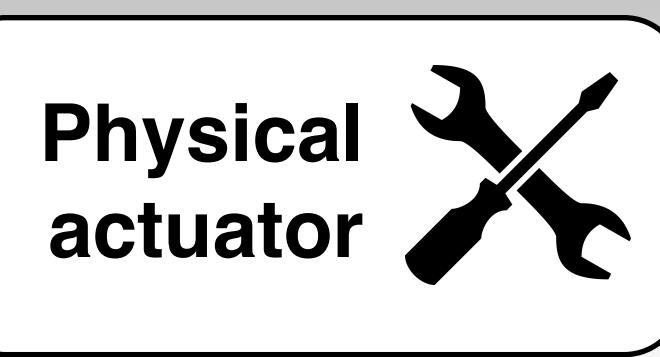
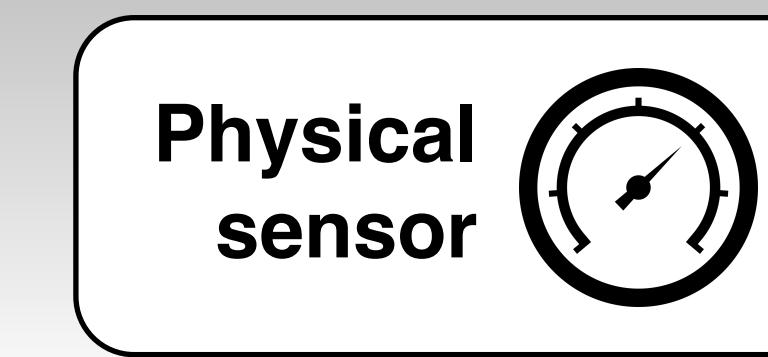
Priority classes



System Model

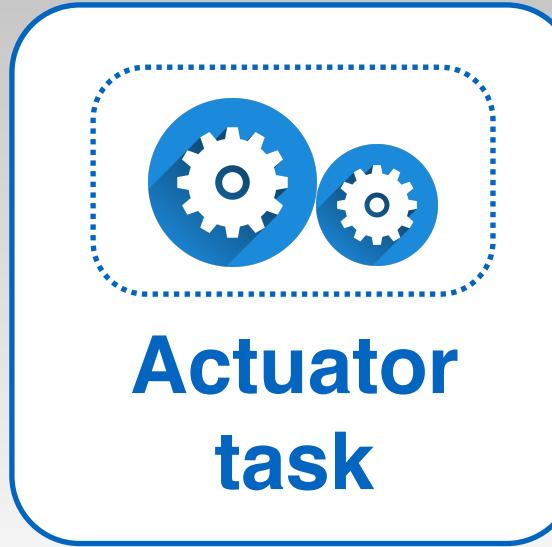
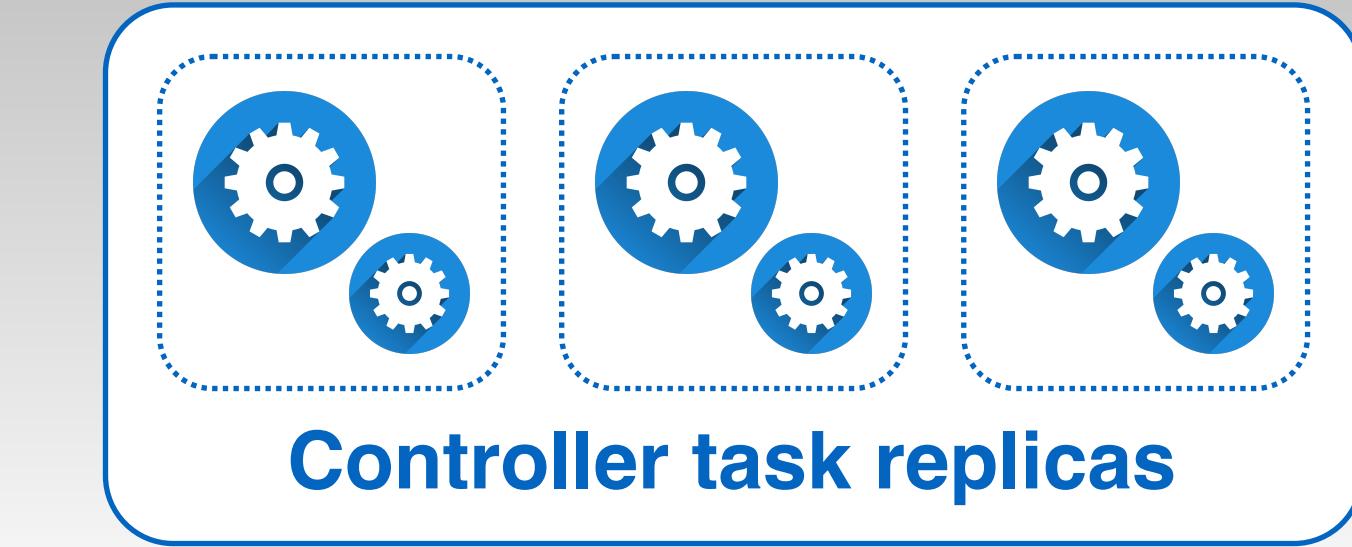
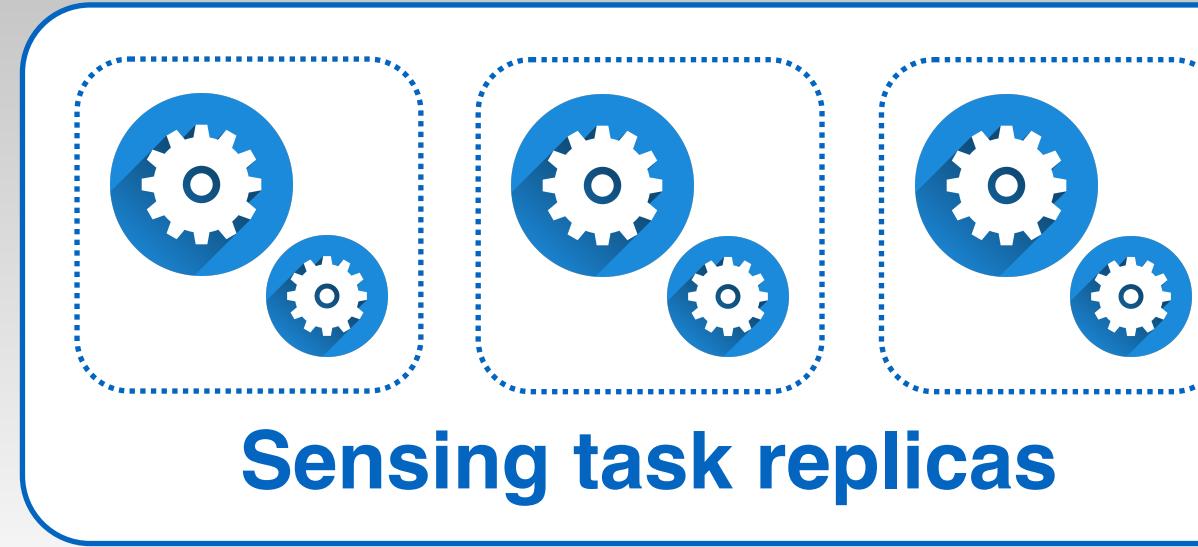
Network control system

Physical plant reliable



Active Replication

DMR / TMR / Hybrid

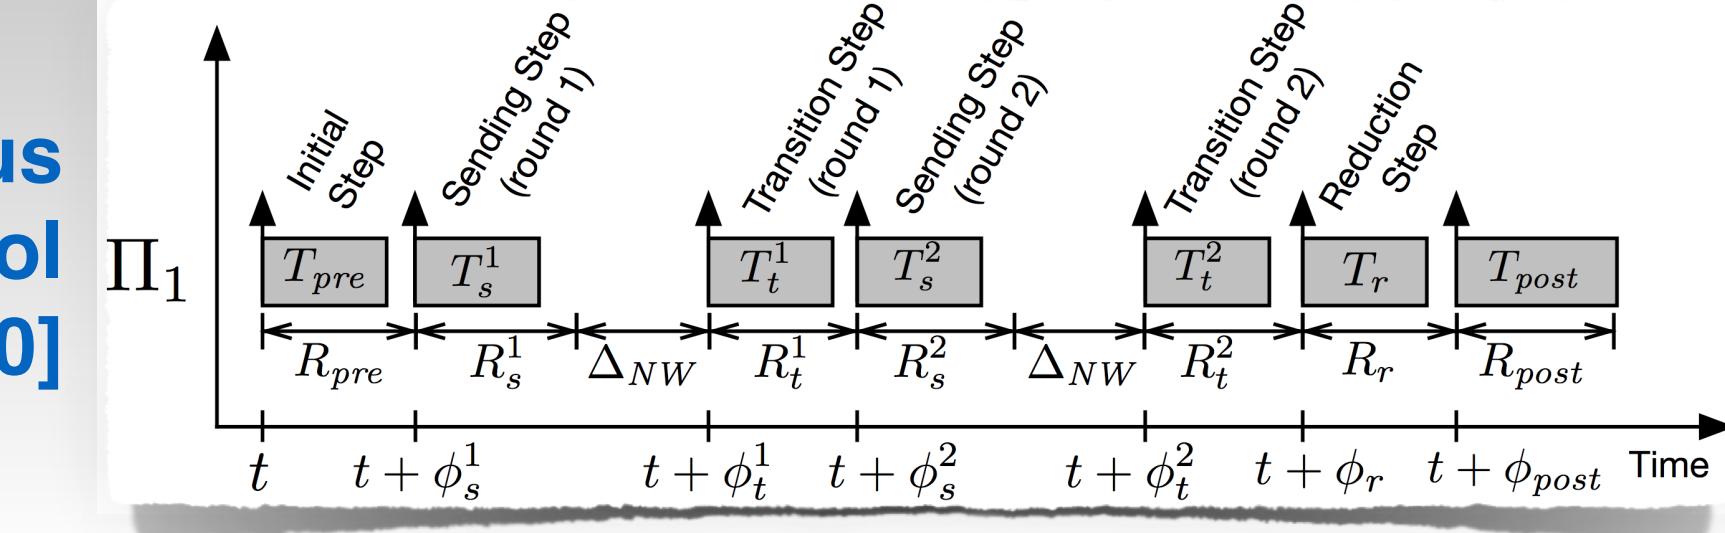


Atomic Broadcast

Statically-checked hard real-time protocol

Synchronous BFT protocol
[Pease et al., 1980]

Periodic tasks and messages

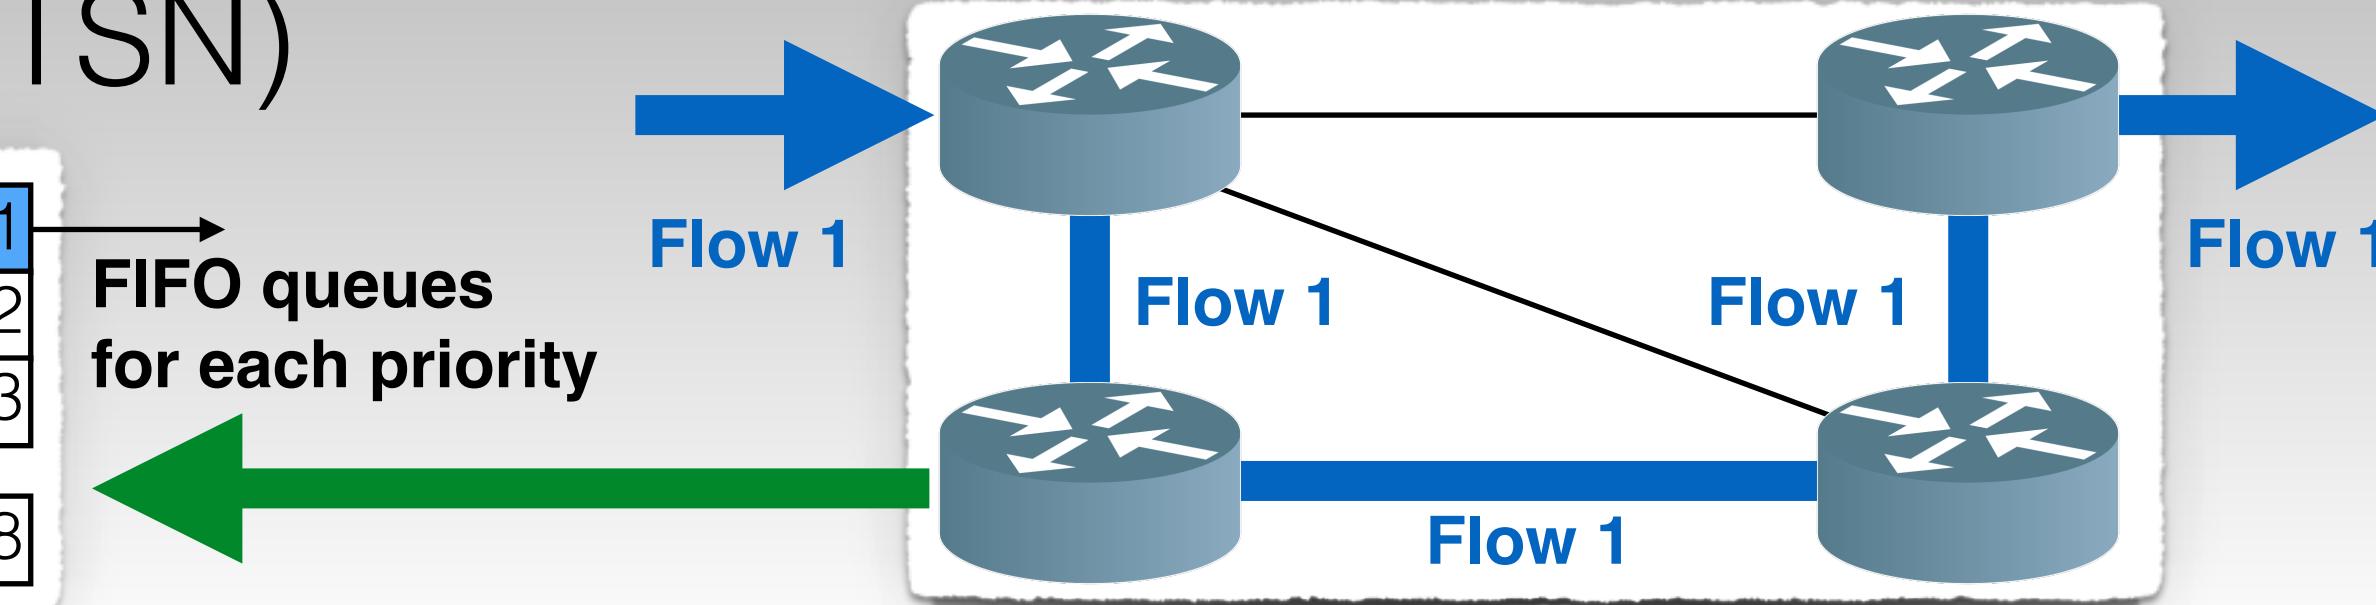
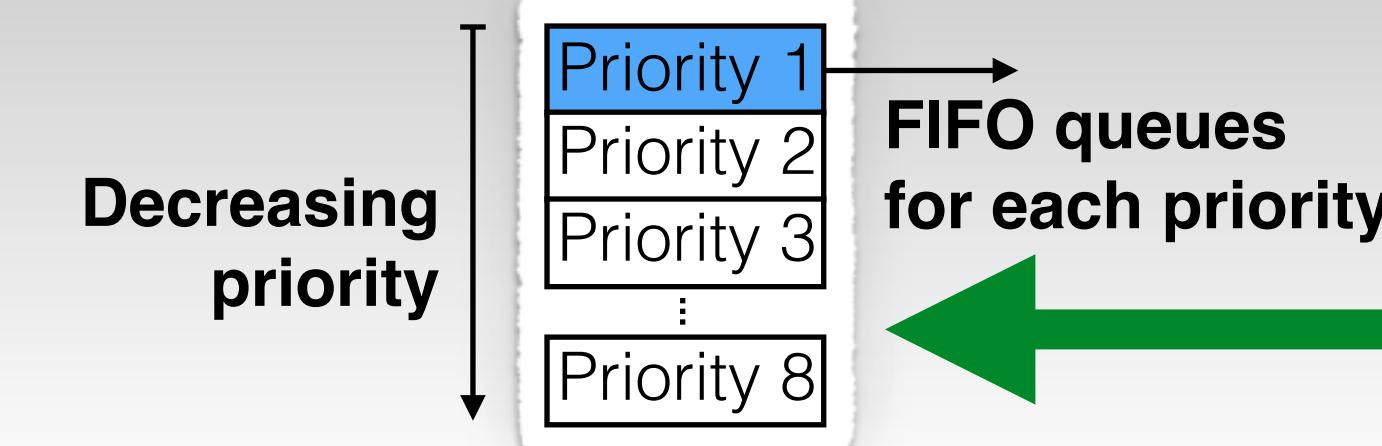


Clock synchronization

Ethernet Time-Sensitive Networking (TSN)

Statically reserved routes

Priority classes



System Model

Network control system

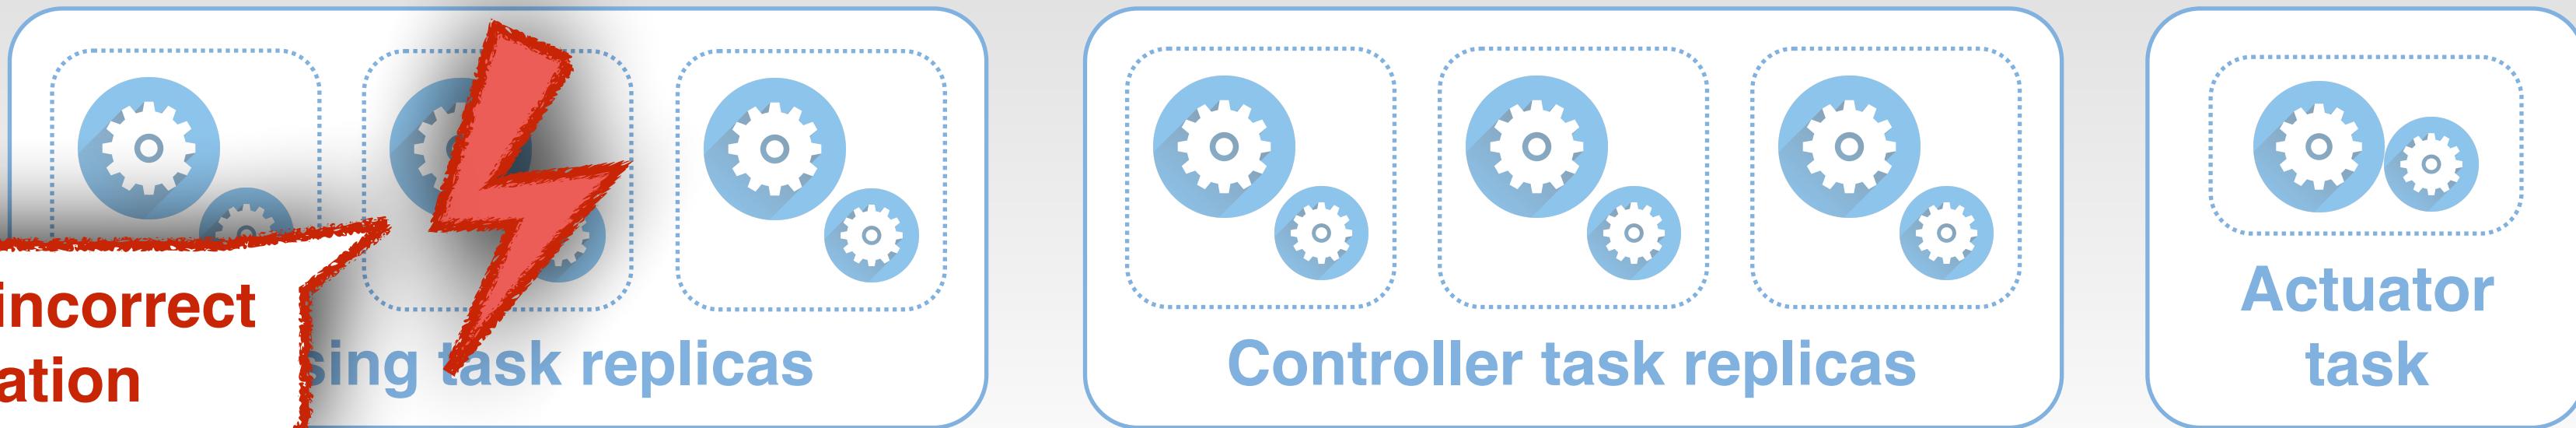
Physical plant reliable



Active Replication

DMR / TMR / HMR

Omission / incorrect computation



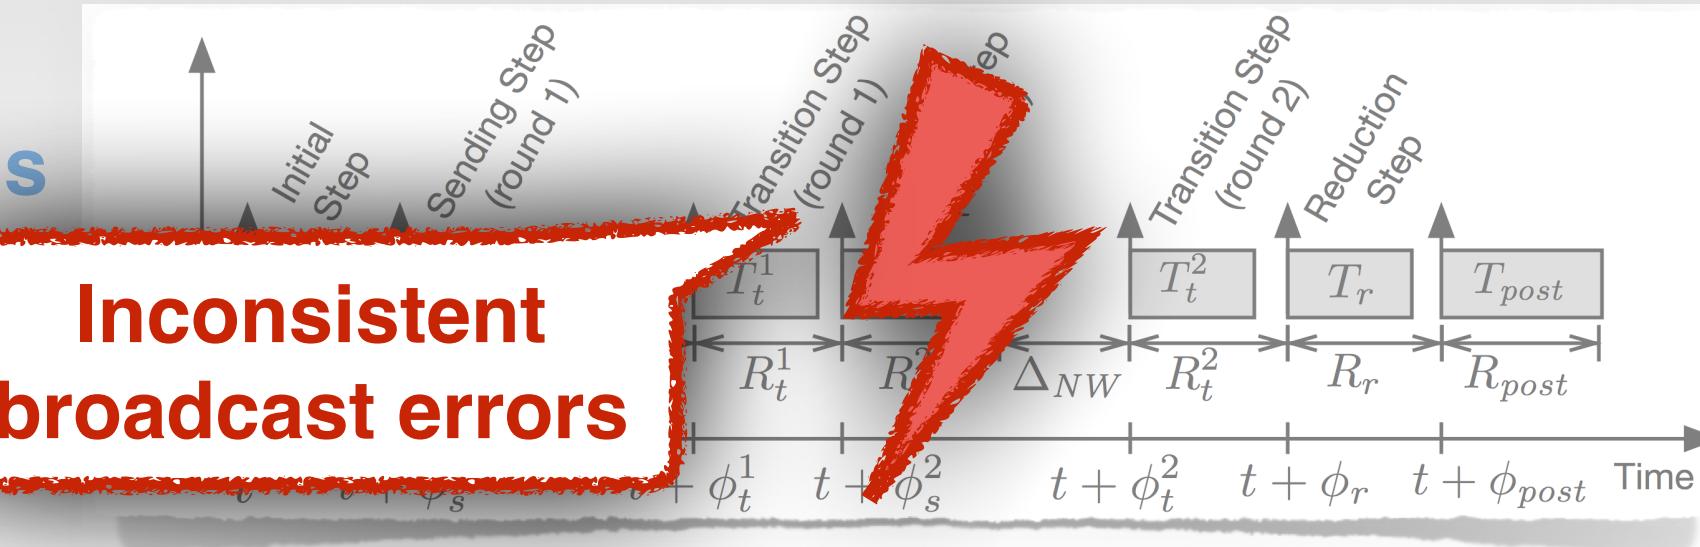
Atomic Broadcast

Statically-checked hard real-time protocol

Synchronous BFT protocol
[Pease et al., 1980]

Inconsistent broadcast errors

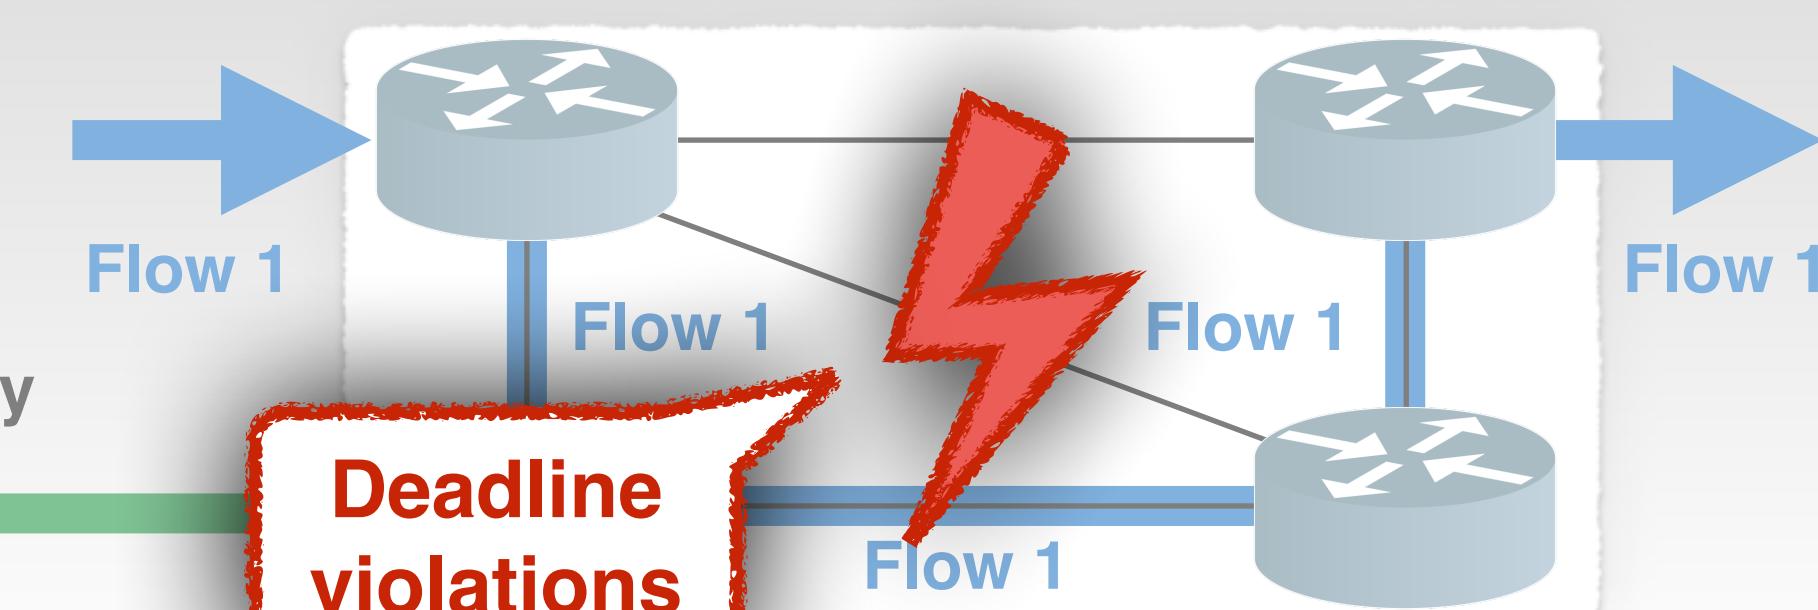
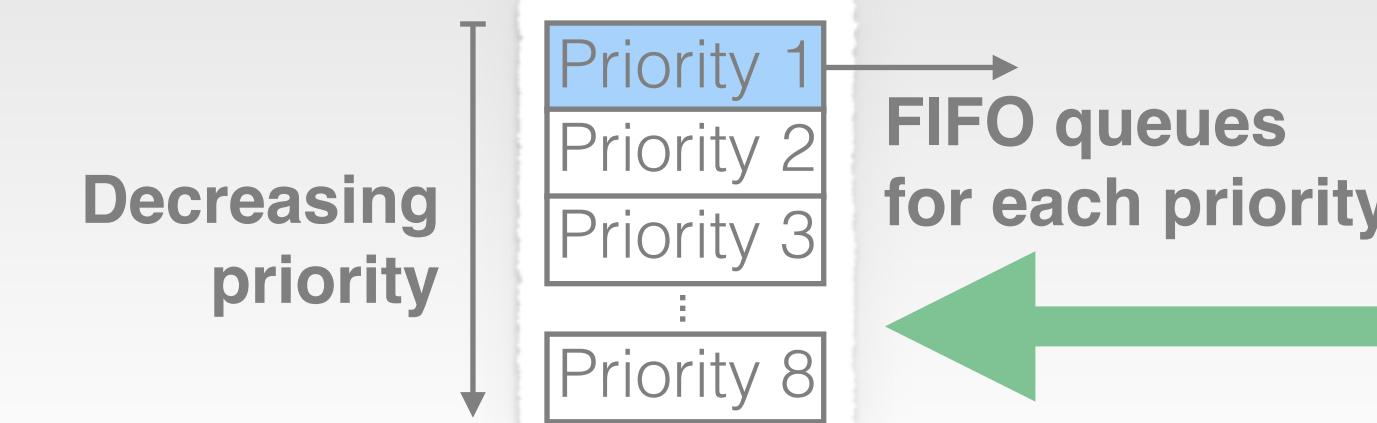
Periodic tasks and messages



Ethernet Time Sensitive Networking (TSN)

Statically reserved routes

Priority classes



System Model

Network control system

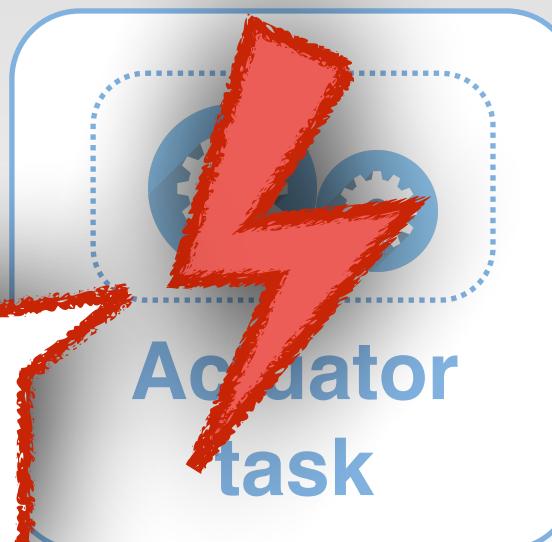
Physical plant reliable



Active Replication

DMR / TMR / HMR

Omission / incorrect computation



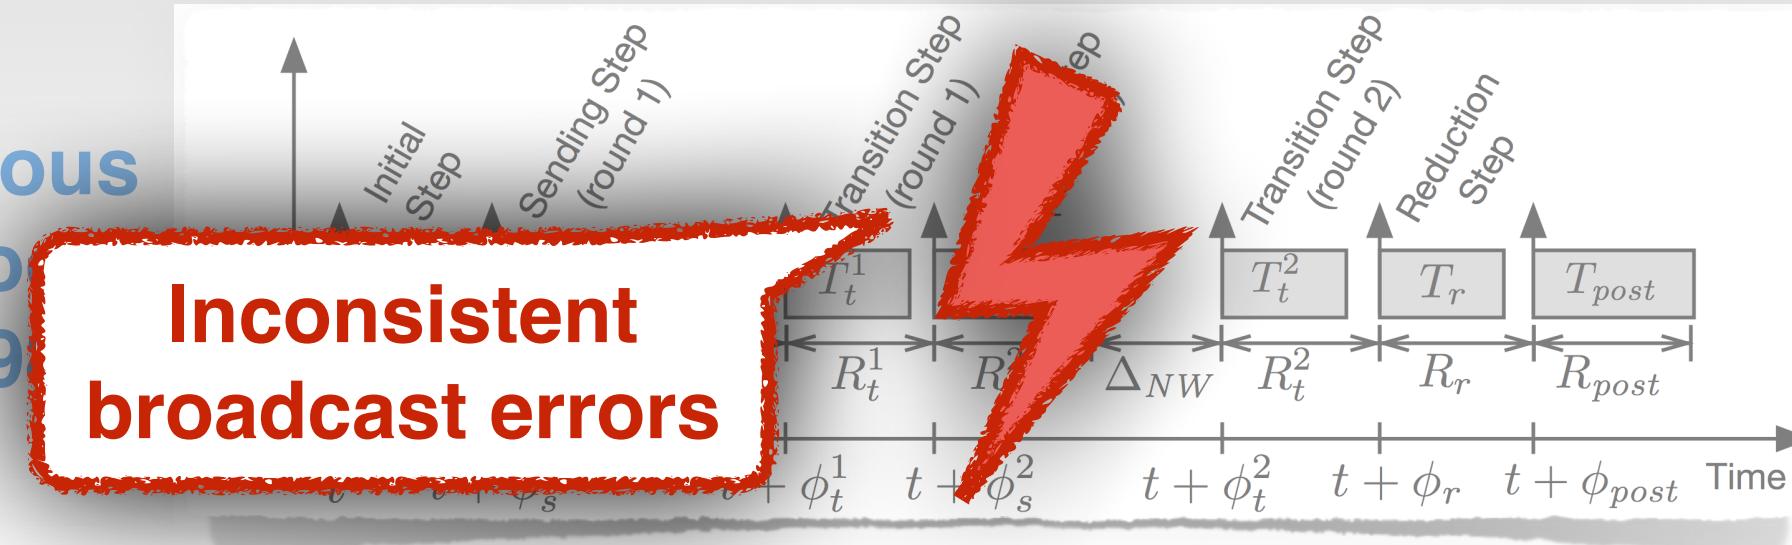
Redundant, but correlated & erroneous, inputs

Atomic Broadcast

Statically-checked hard real-time protocol

Synchronous BFT protocol
[Pease et al., 1980]

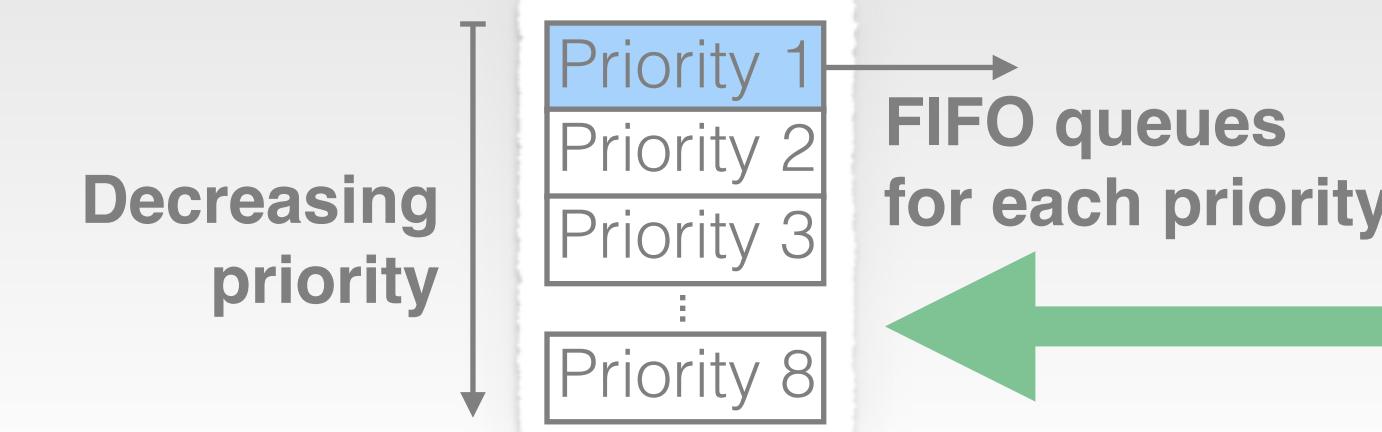
Periodic tasks and messages



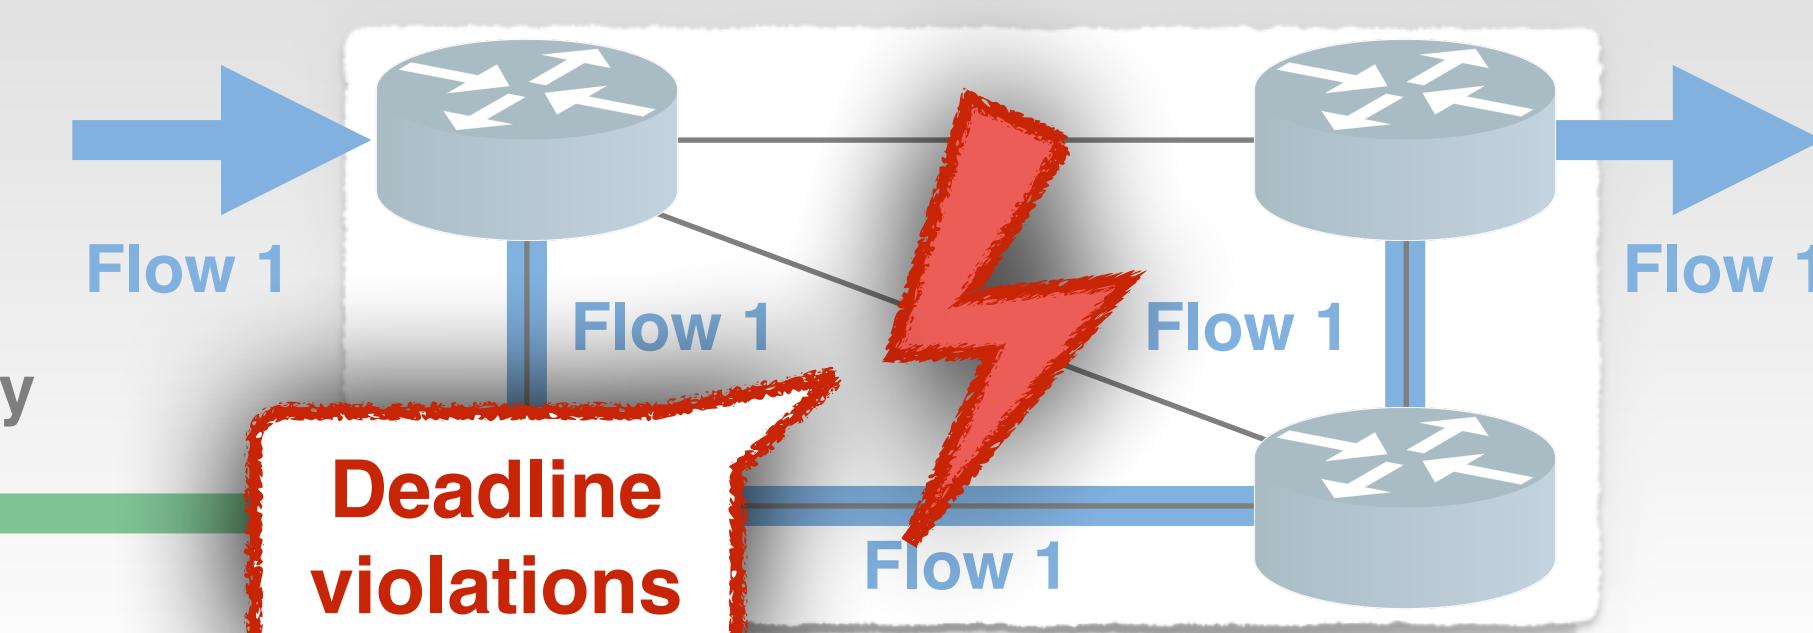
Ethernet Time Sensitive Networking (TSN)

Statically reserved routes

Priority classes



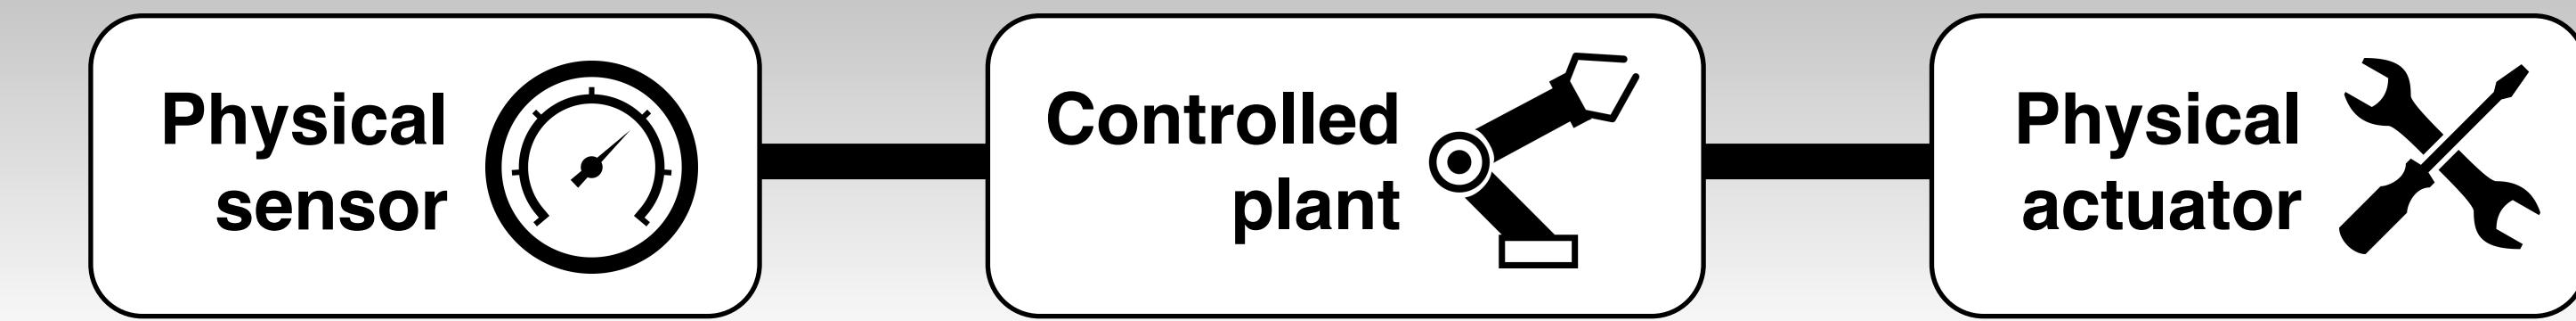
Deadline violations



System Model

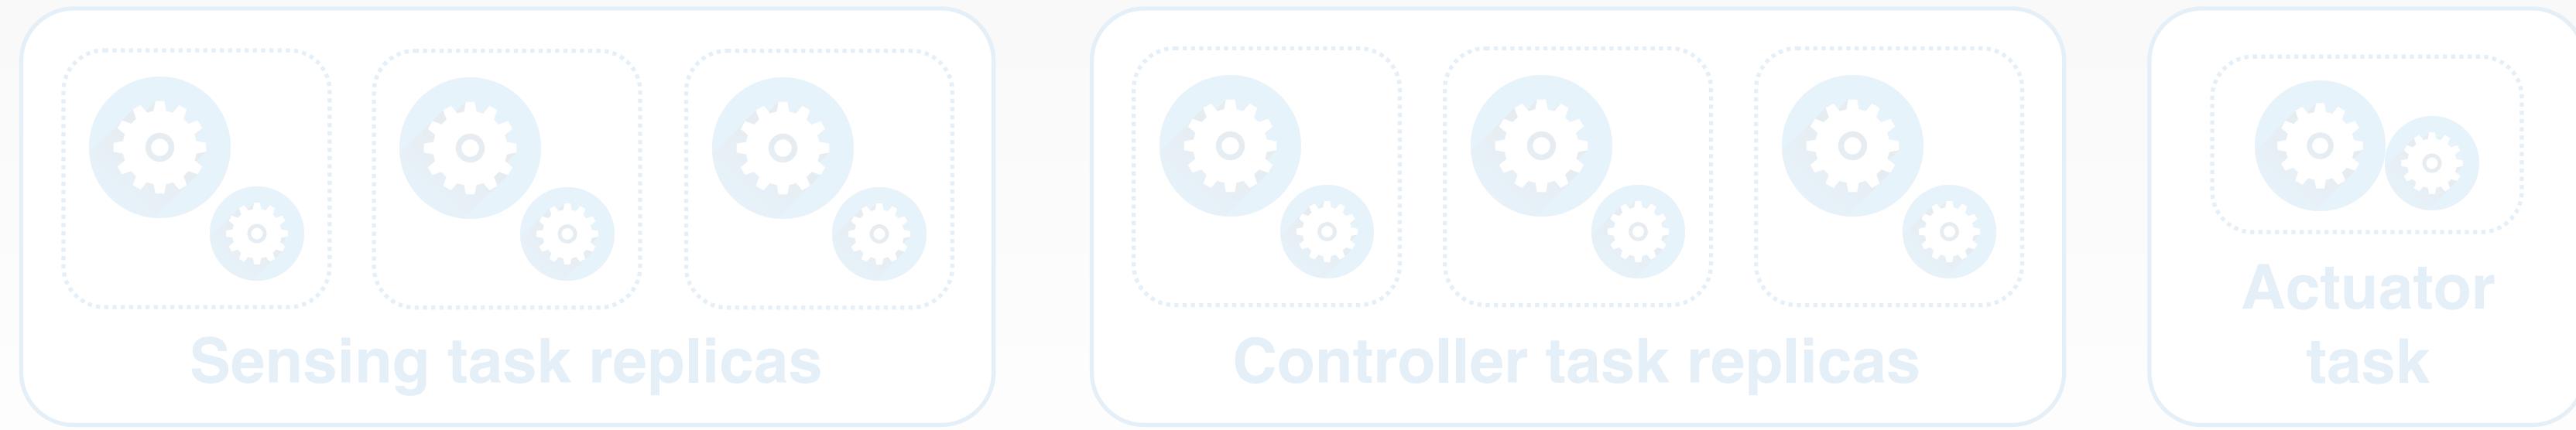
Network control system

Physical plant reliable



Active Replication

DMR / TMR / Hybrid

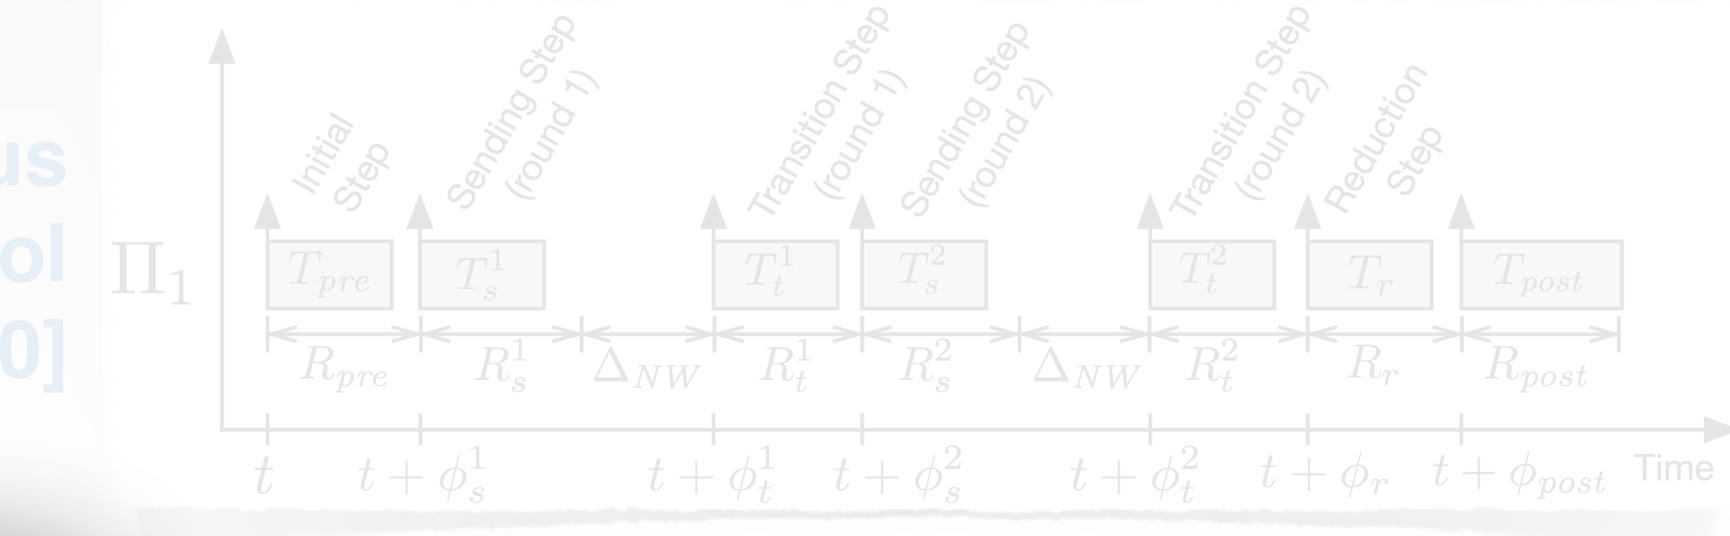


Atomic Broadcast

Statically-checked
hard real-time protocol

Synchronous
BFT protocol
[Pease et al., 1980]

Periodic tasks and messages



Ethernet Time

Statically reserved

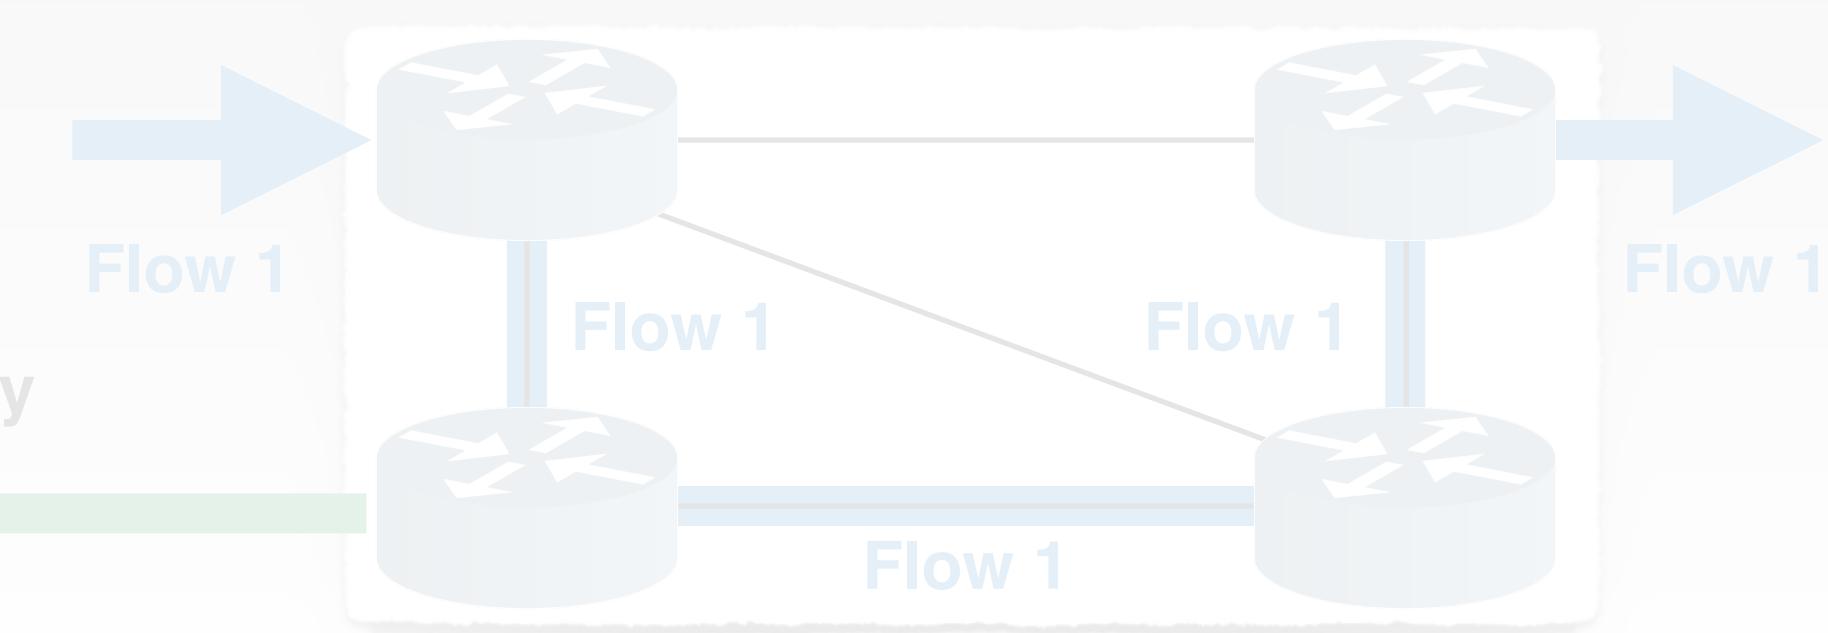
Priority queues

Transient
fault-induced
errors



(TSN)

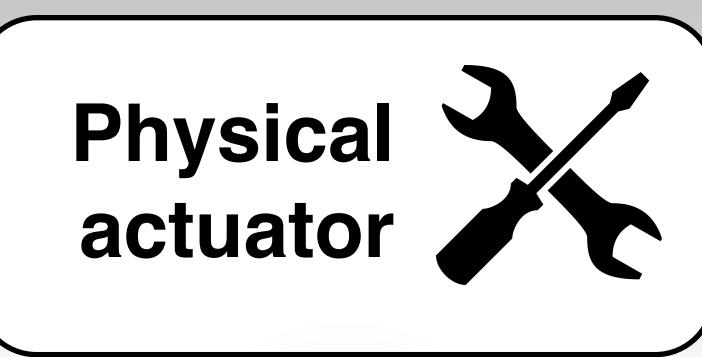
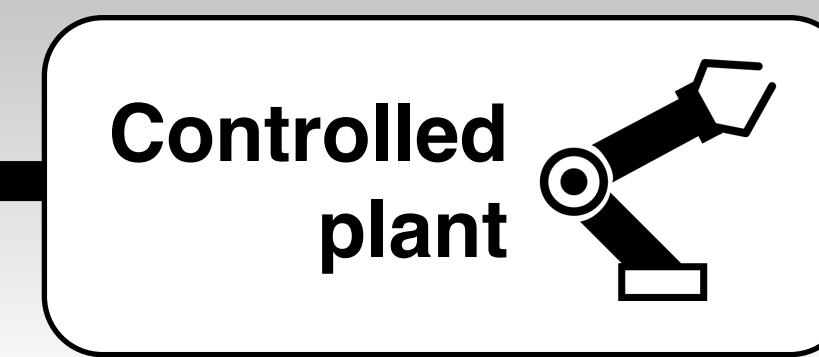
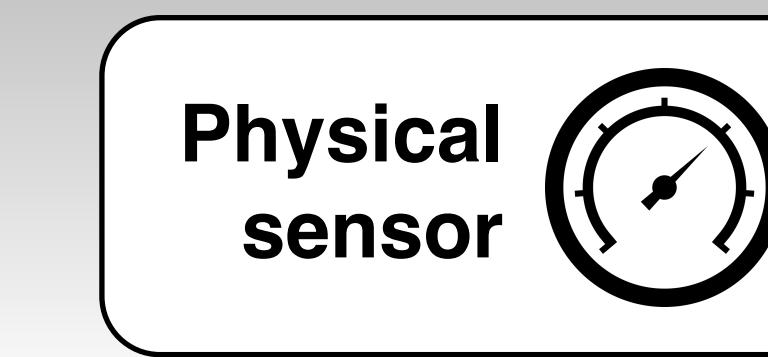
FIFO queues
for each priority



System Model

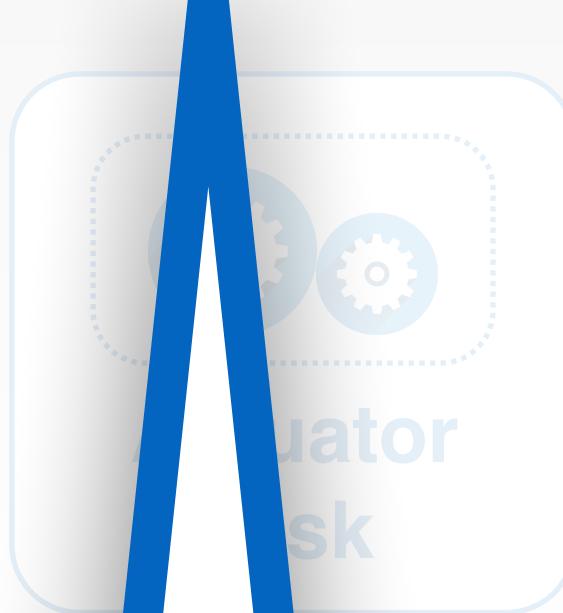
Network control system

Physical plant reliable



Active Replication

DMR / TMR / Hybrid



Atomic Broadcast

Statically-checked hard real-time protocol

Synchronous BFT protocol
[Pease et al., 1980]

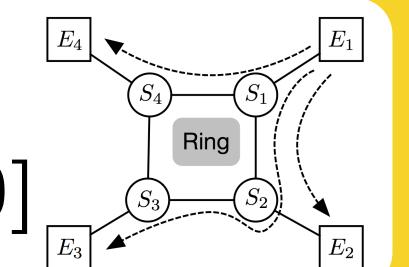
Periodic tasks and messages

Transient fault-induced errors



Q1. How often does the **final actuation deviate** from an error-free scenario (iteration failure)?

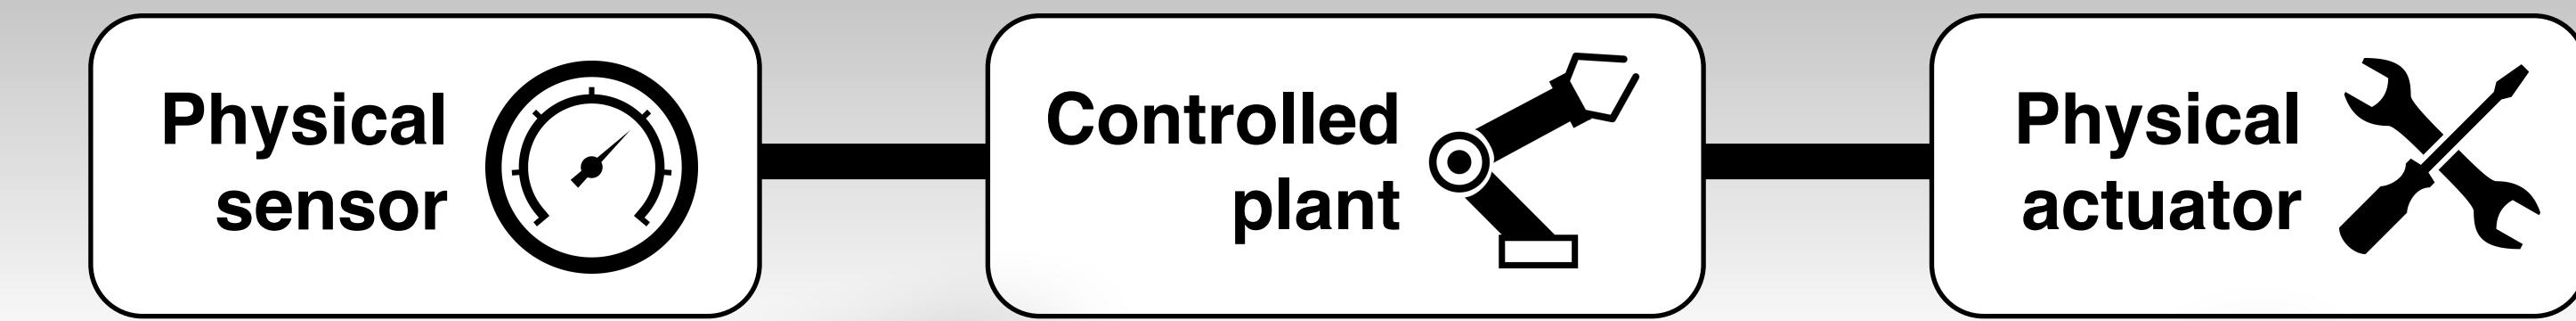
Replica Consistency over Ethernet [RTAS '20]



System Model

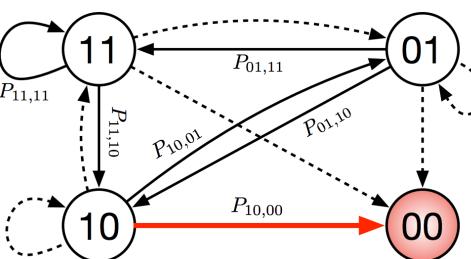
Network control system

Physical plant reliable

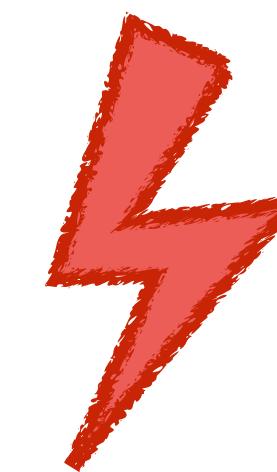


Q2. What is the likelihood of a **control failure**?

Periodic Weakly-Hard Systems [ECRTS '19]

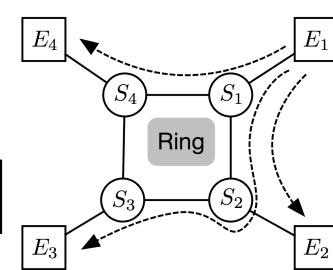


Transient
fault-induced
errors



Q1. How often does the **final actuation deviate** from an error-free scenario (iteration failure)?

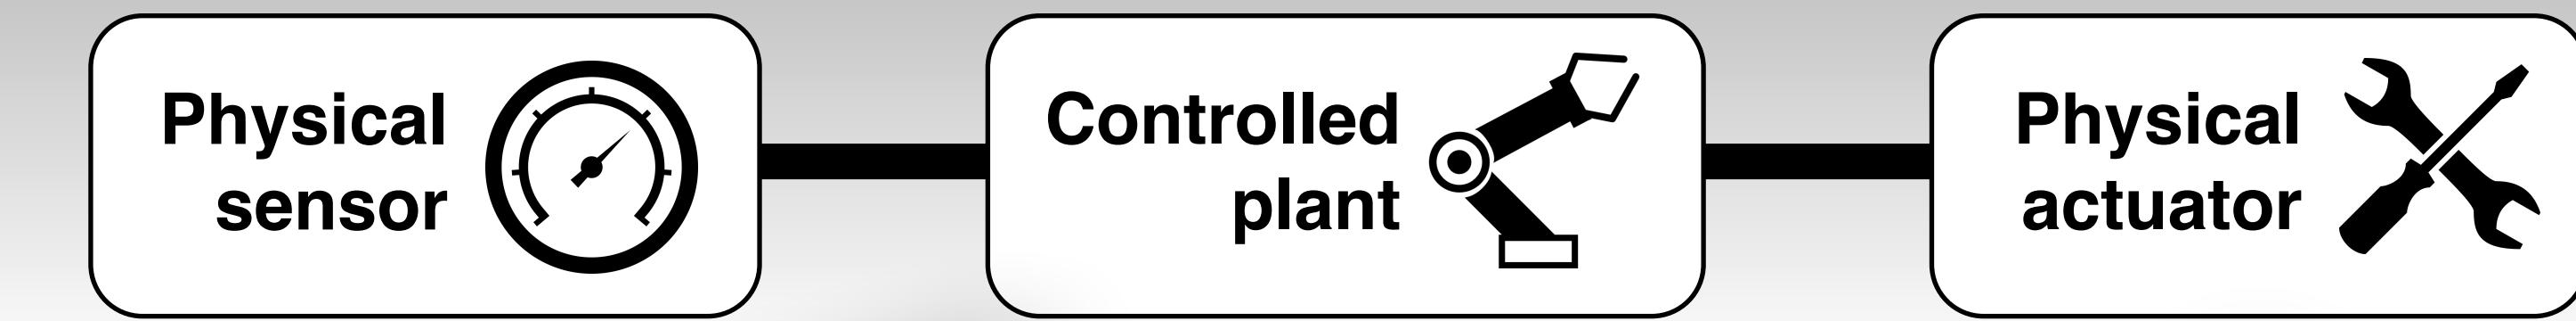
Replica Consistency over Ethernet [RTAS '20]



System Model

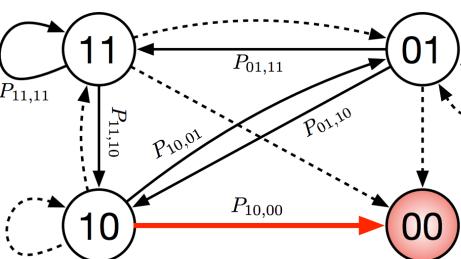
Network control system

Physical plant reliable



Q2. What is the likelihood of a **control failure**?

Periodic Weakly-Hard Systems [ECRTS '19]

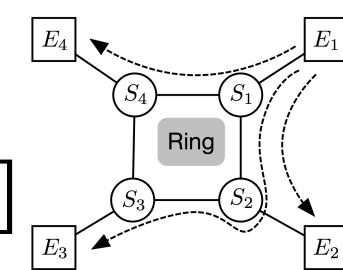


Transient
fault-induced
errors



Q1. How often does the **final actuation deviate** from an error-free scenario (iteration failure)?

Replica Consistency over Ethernet [RTAS '20]



Stochastically modeled basic errors

Basic errors due to transient faults are random, independent events

- E.g., node crashes, link corruption

Stochastically modeled basic errors

Basic errors due to transient faults are random, independent events

- E.g., node crashes, link corruption

Poisson distribution using peak rates
from maximum interference periods

Stochastically modeled basic errors

Basic errors due to transient faults are random, independent events

- E.g., node crashes, link corruption

Poisson distribution using peak rates
from maximum interference periods

For processors
and switches

Poisson(n , δ , λ_{crash})
 $= \Pr(n \text{ crashes in an interval of length } \delta \mid \text{crash rate } \lambda_{\text{crash}})$

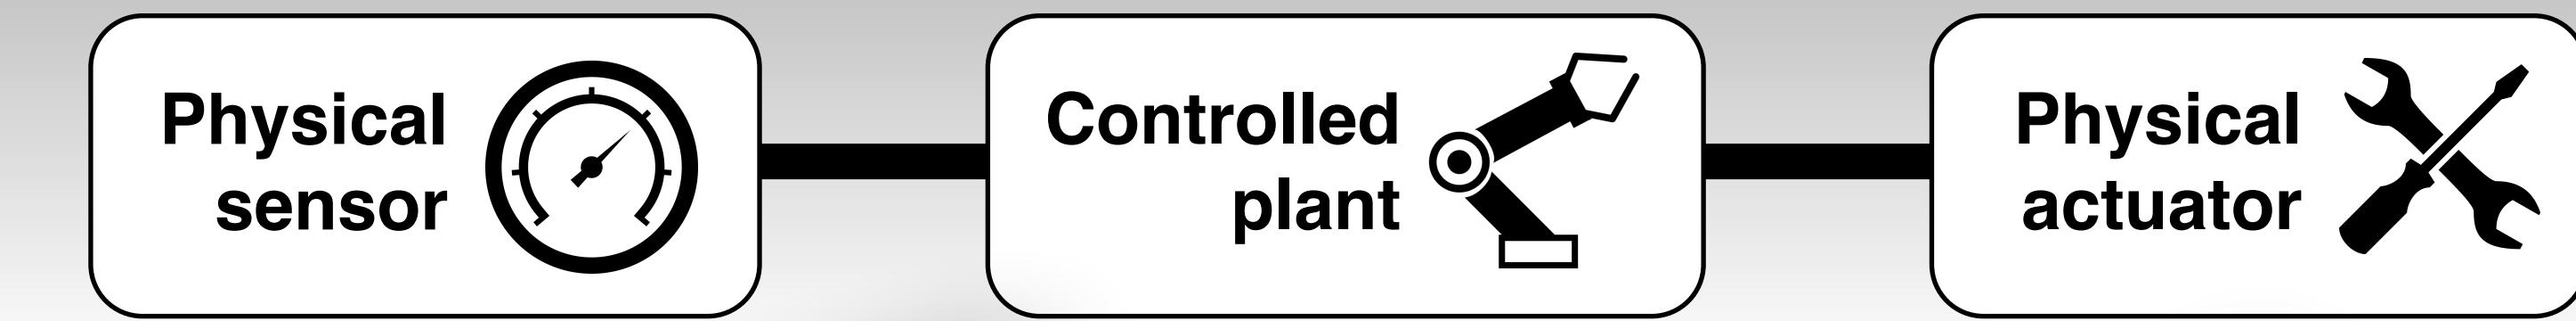
For processors, switches,
and network links

Poisson(n , δ , $\lambda_{\text{corruption}}$)
 $= \Pr(n \text{ corruptions in an interval of length } \delta \mid \text{corruption rate } \lambda_{\text{corruption}})$

System Model

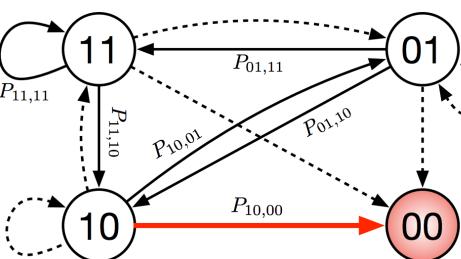
Network control system

Physical plant reliable

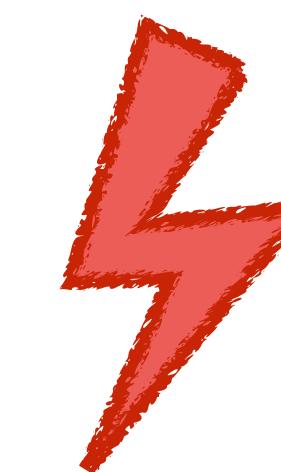


Q2. What is the likelihood of a **control failure**?

Periodic Weakly-Hard Systems [ECRTS '19]

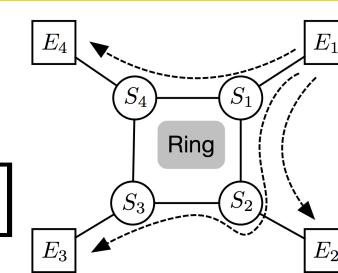


Transient
fault-induced
errors

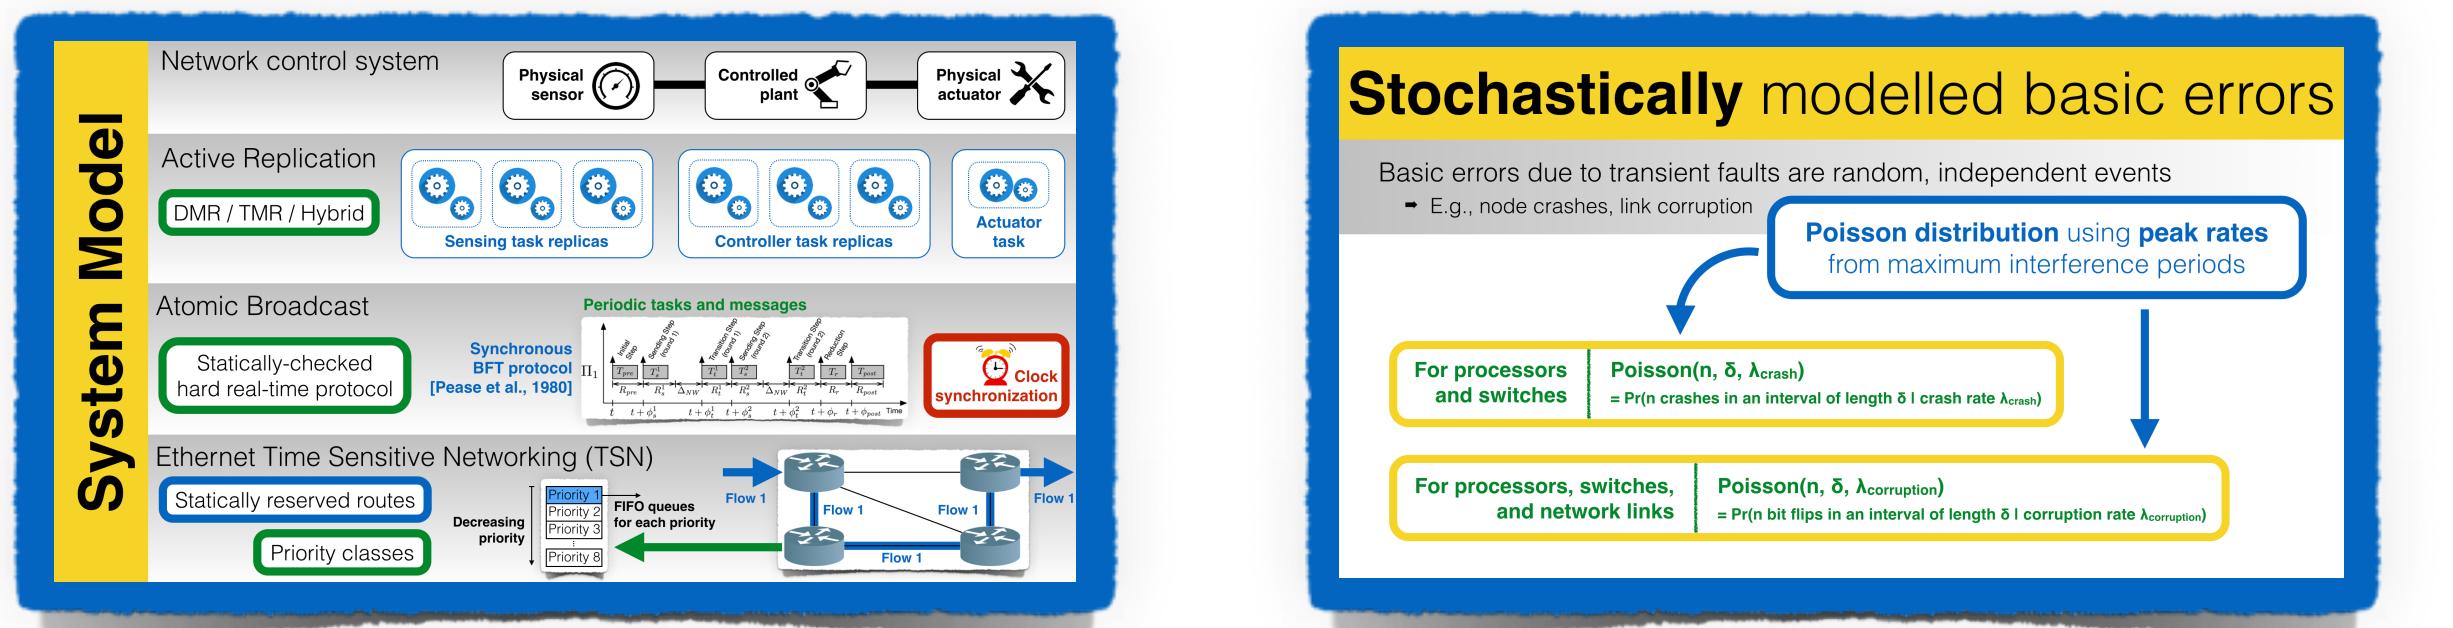


Q1. How often does the **final actuation deviate** from an error-free scenario (iteration failure)?

Replica Consistency over Ethernet [RTAS '20]



Straw-man Solutions



Stochastically modelled basic errors

Basic errors due to transient faults are random, independent events

- E.g., node crashes, link corruption

Poisson distribution using peak rates
from maximum interference periods

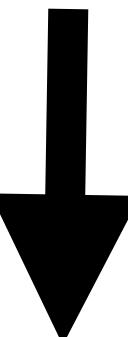
For processors and switches

Poisson(n , δ , λ_{crash})
= Pr(n crashes in an interval of length δ | crash rate λ_{crash})

For processors, switches, and network links

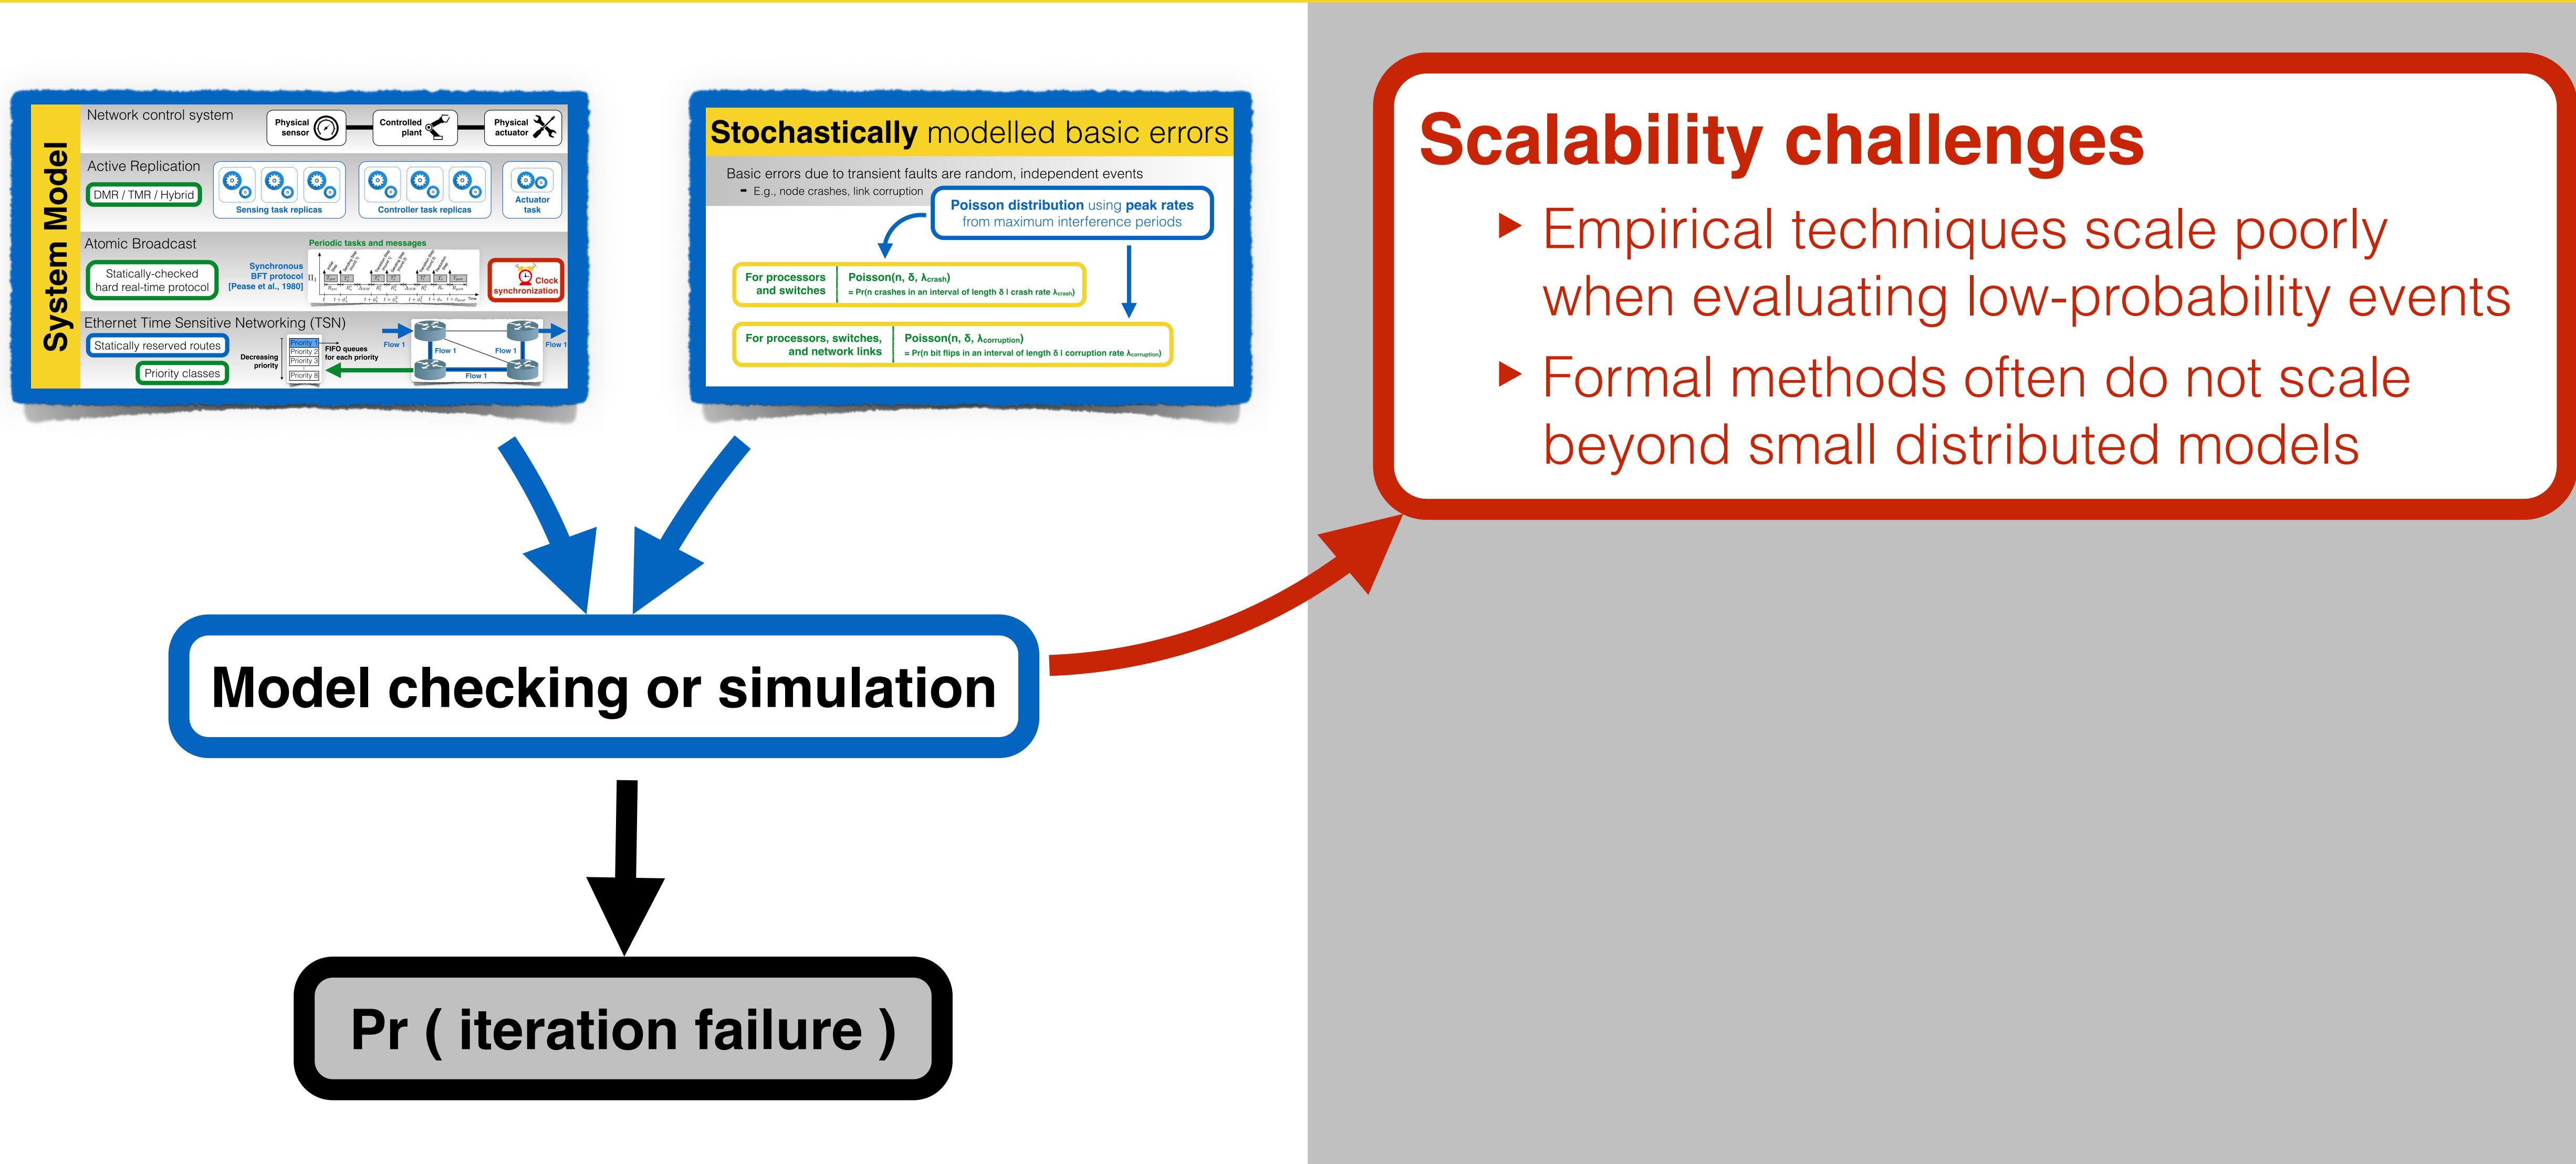
Poisson(n , δ , $\lambda_{\text{corruption}}$)
= Pr(n bit flips in an interval of length δ | corruption rate $\lambda_{\text{corruption}}$)

Model checking or simulation

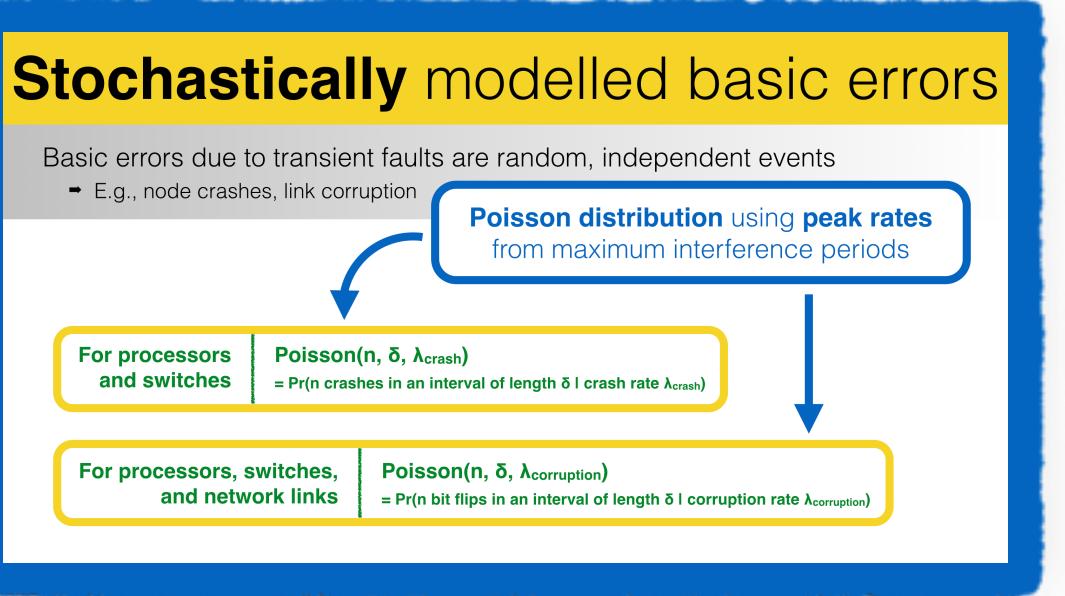
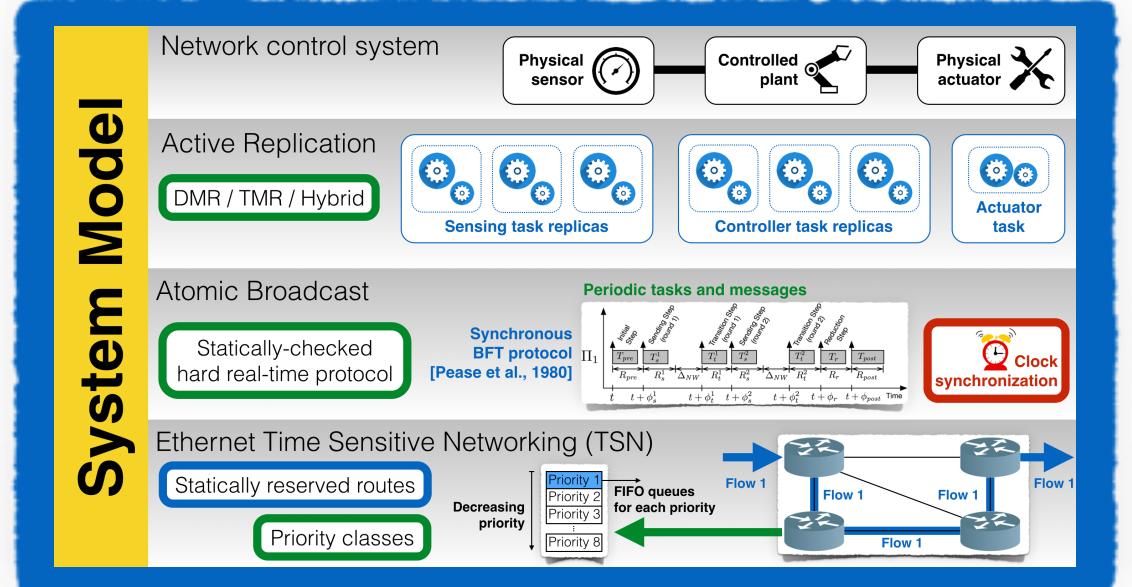


Pr (iteration failure)

Straw-man Solutions



Straw-man Solutions



Model checking or simulation

$\Pr(\text{ iteration failure})$

Scalability challenges

- Empirical techniques scale poorly when evaluating low-probability events
- Formal methods often do not scale beyond small distributed models

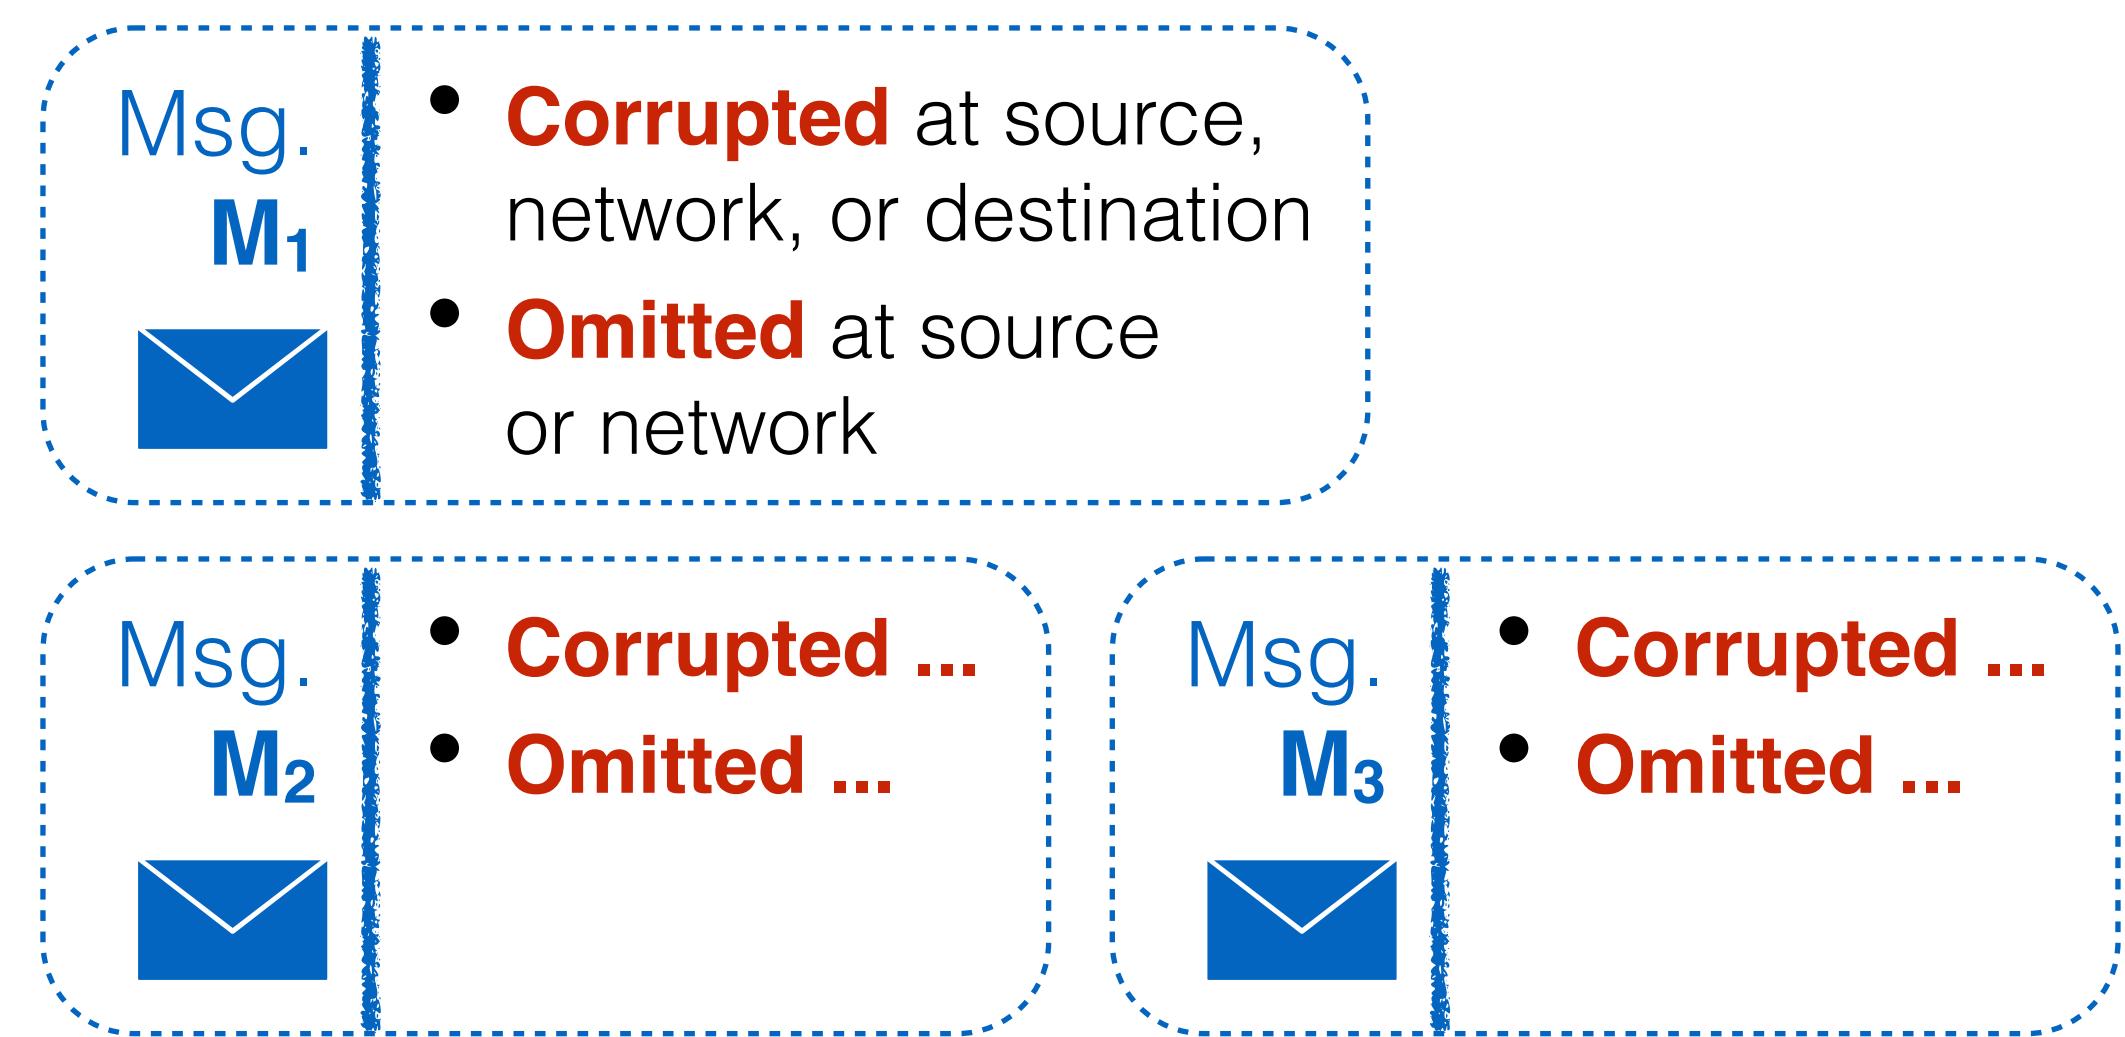
Reliability anomalies

- In practice, the iteration failure probability may **significantly exceed** the estimated $\Pr(\text{ iteration failure})$

Key idea 1: Scalability through **abstraction** and **pruning**

Goal: $P_{UB} > Pr(\text{iteration failure})$

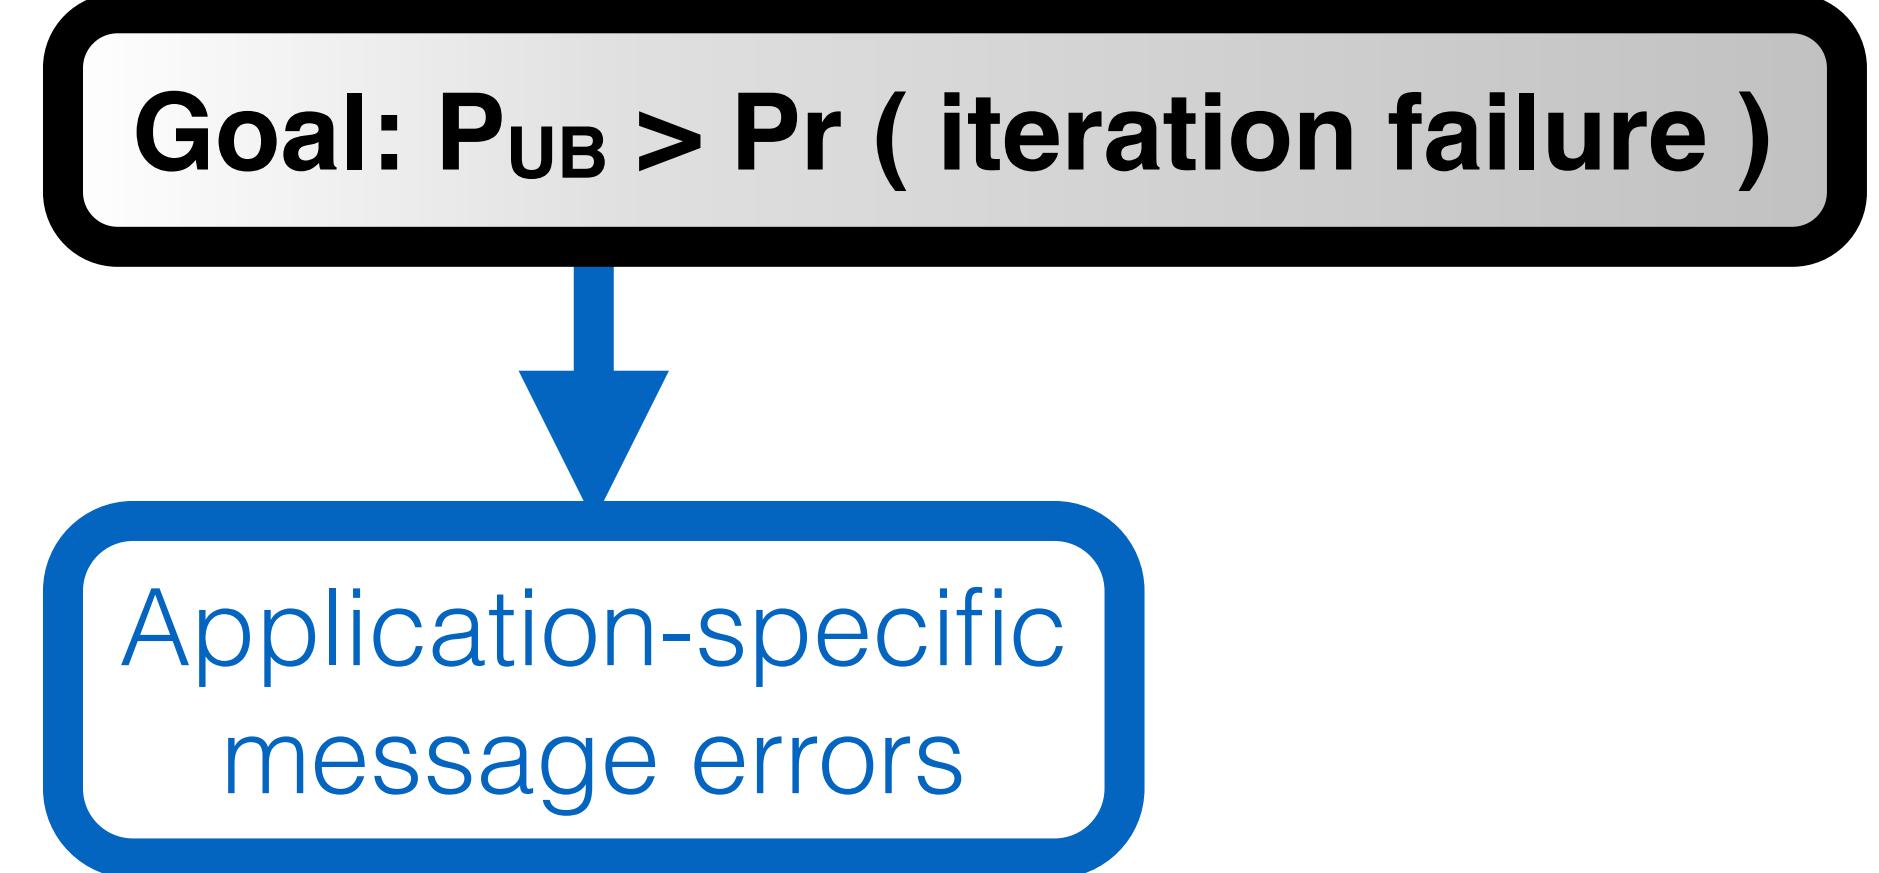
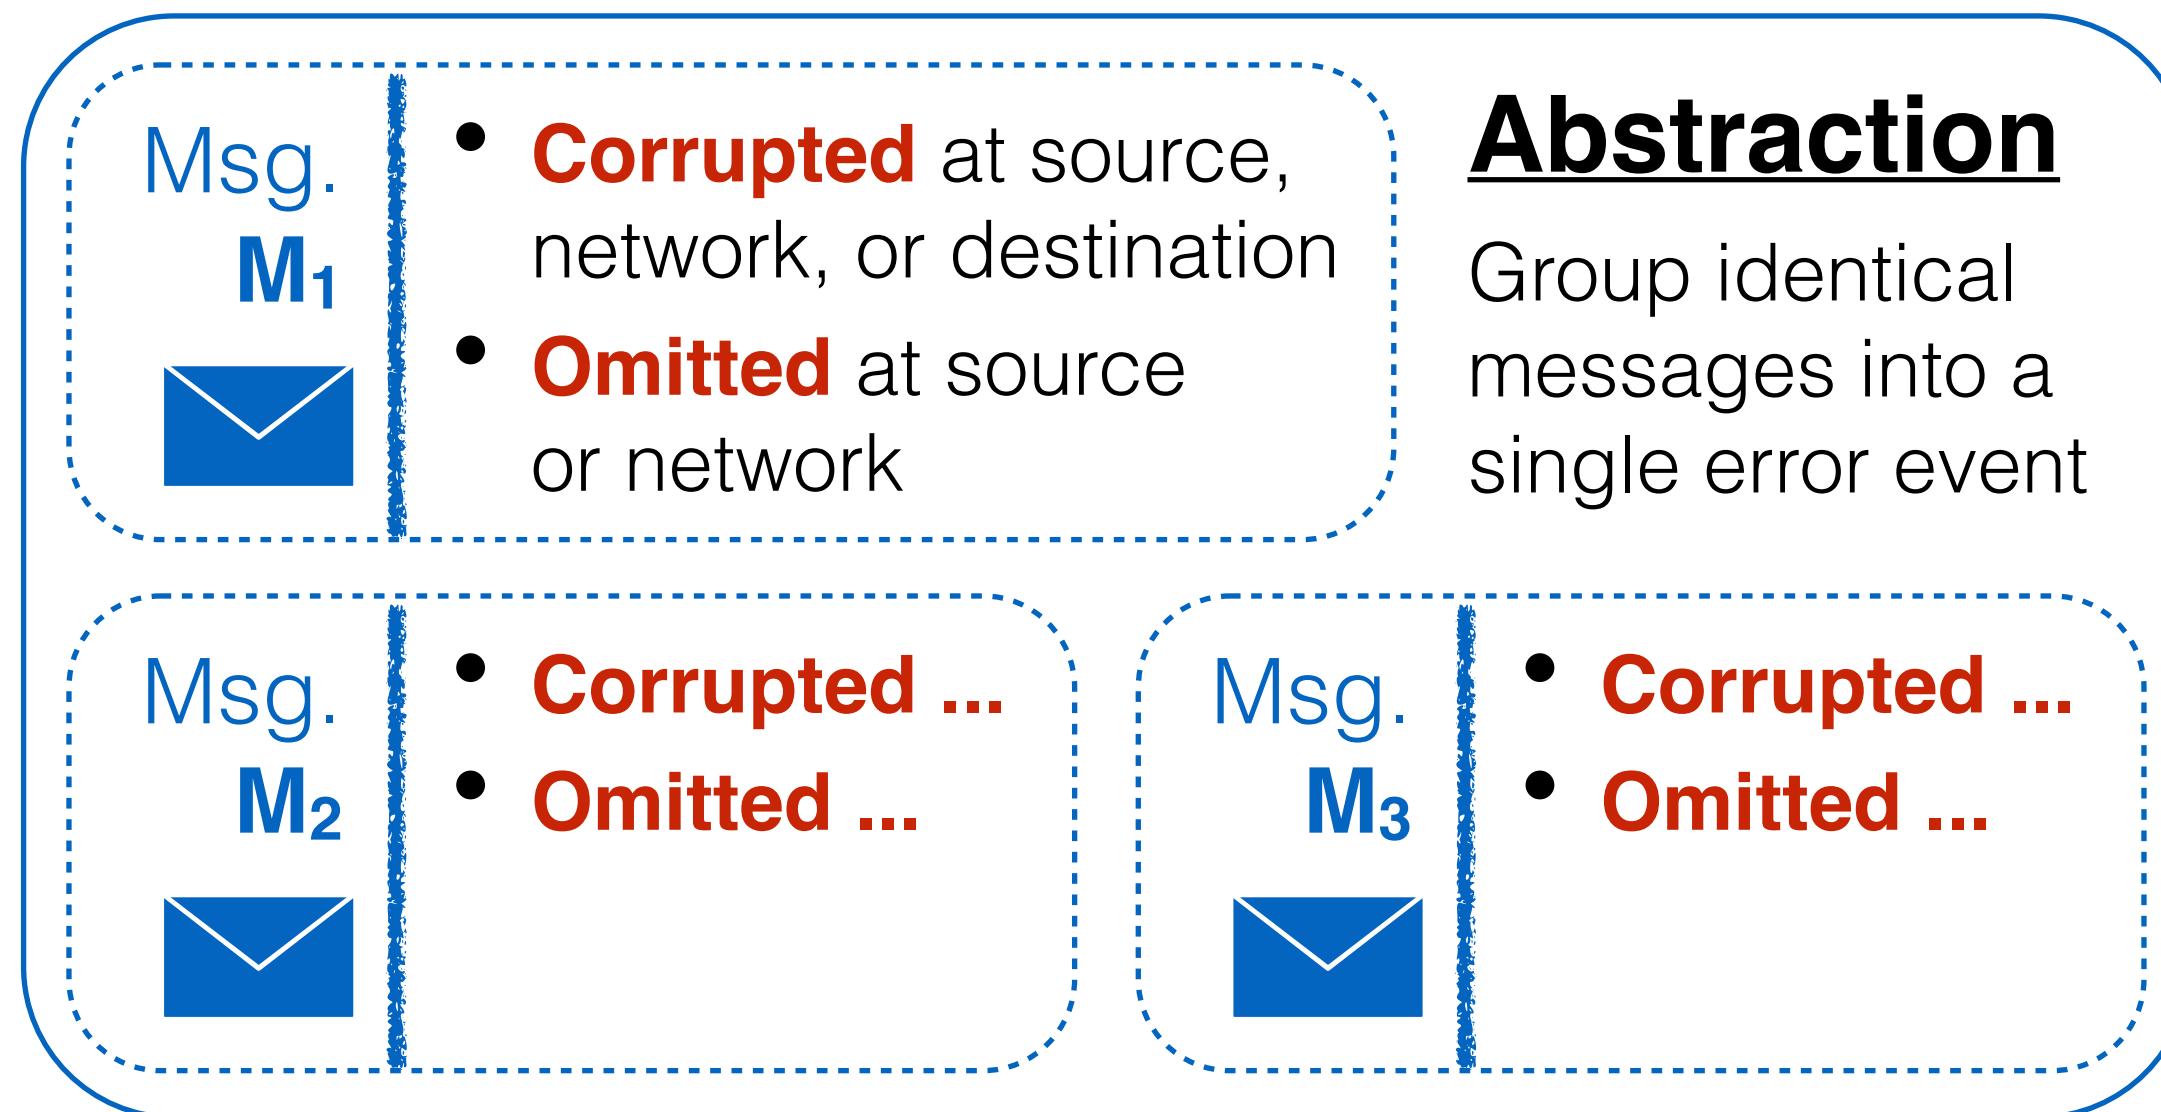
Key idea 1: Scalability through **abstraction** and **pruning**



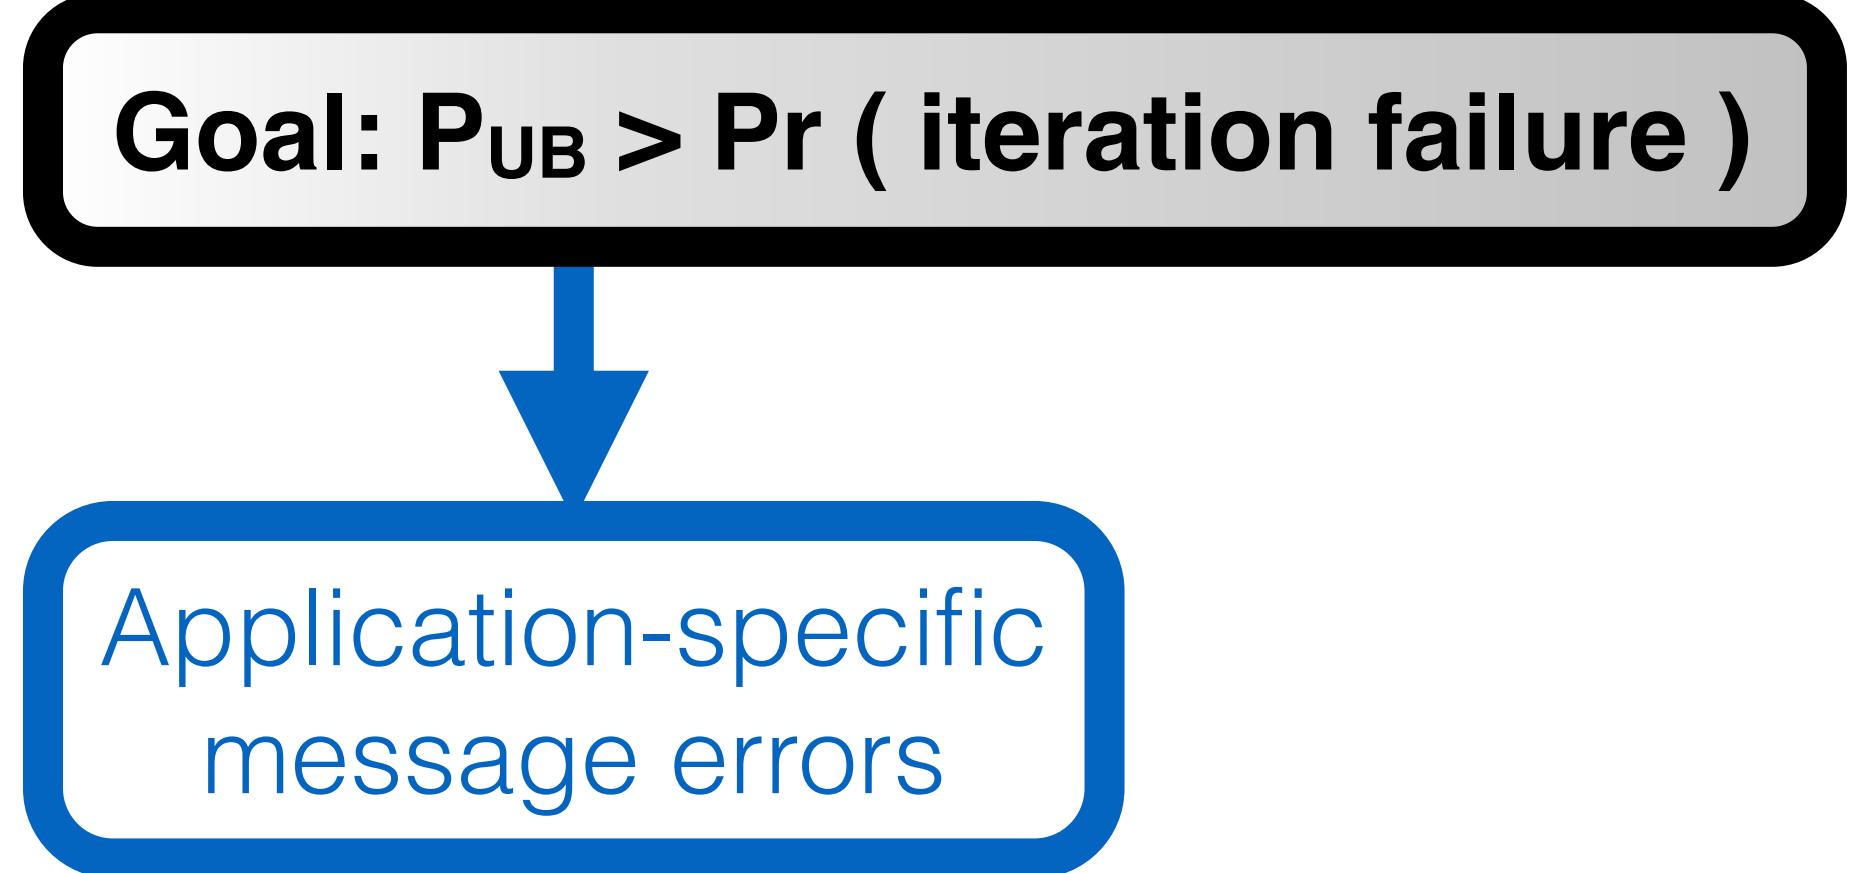
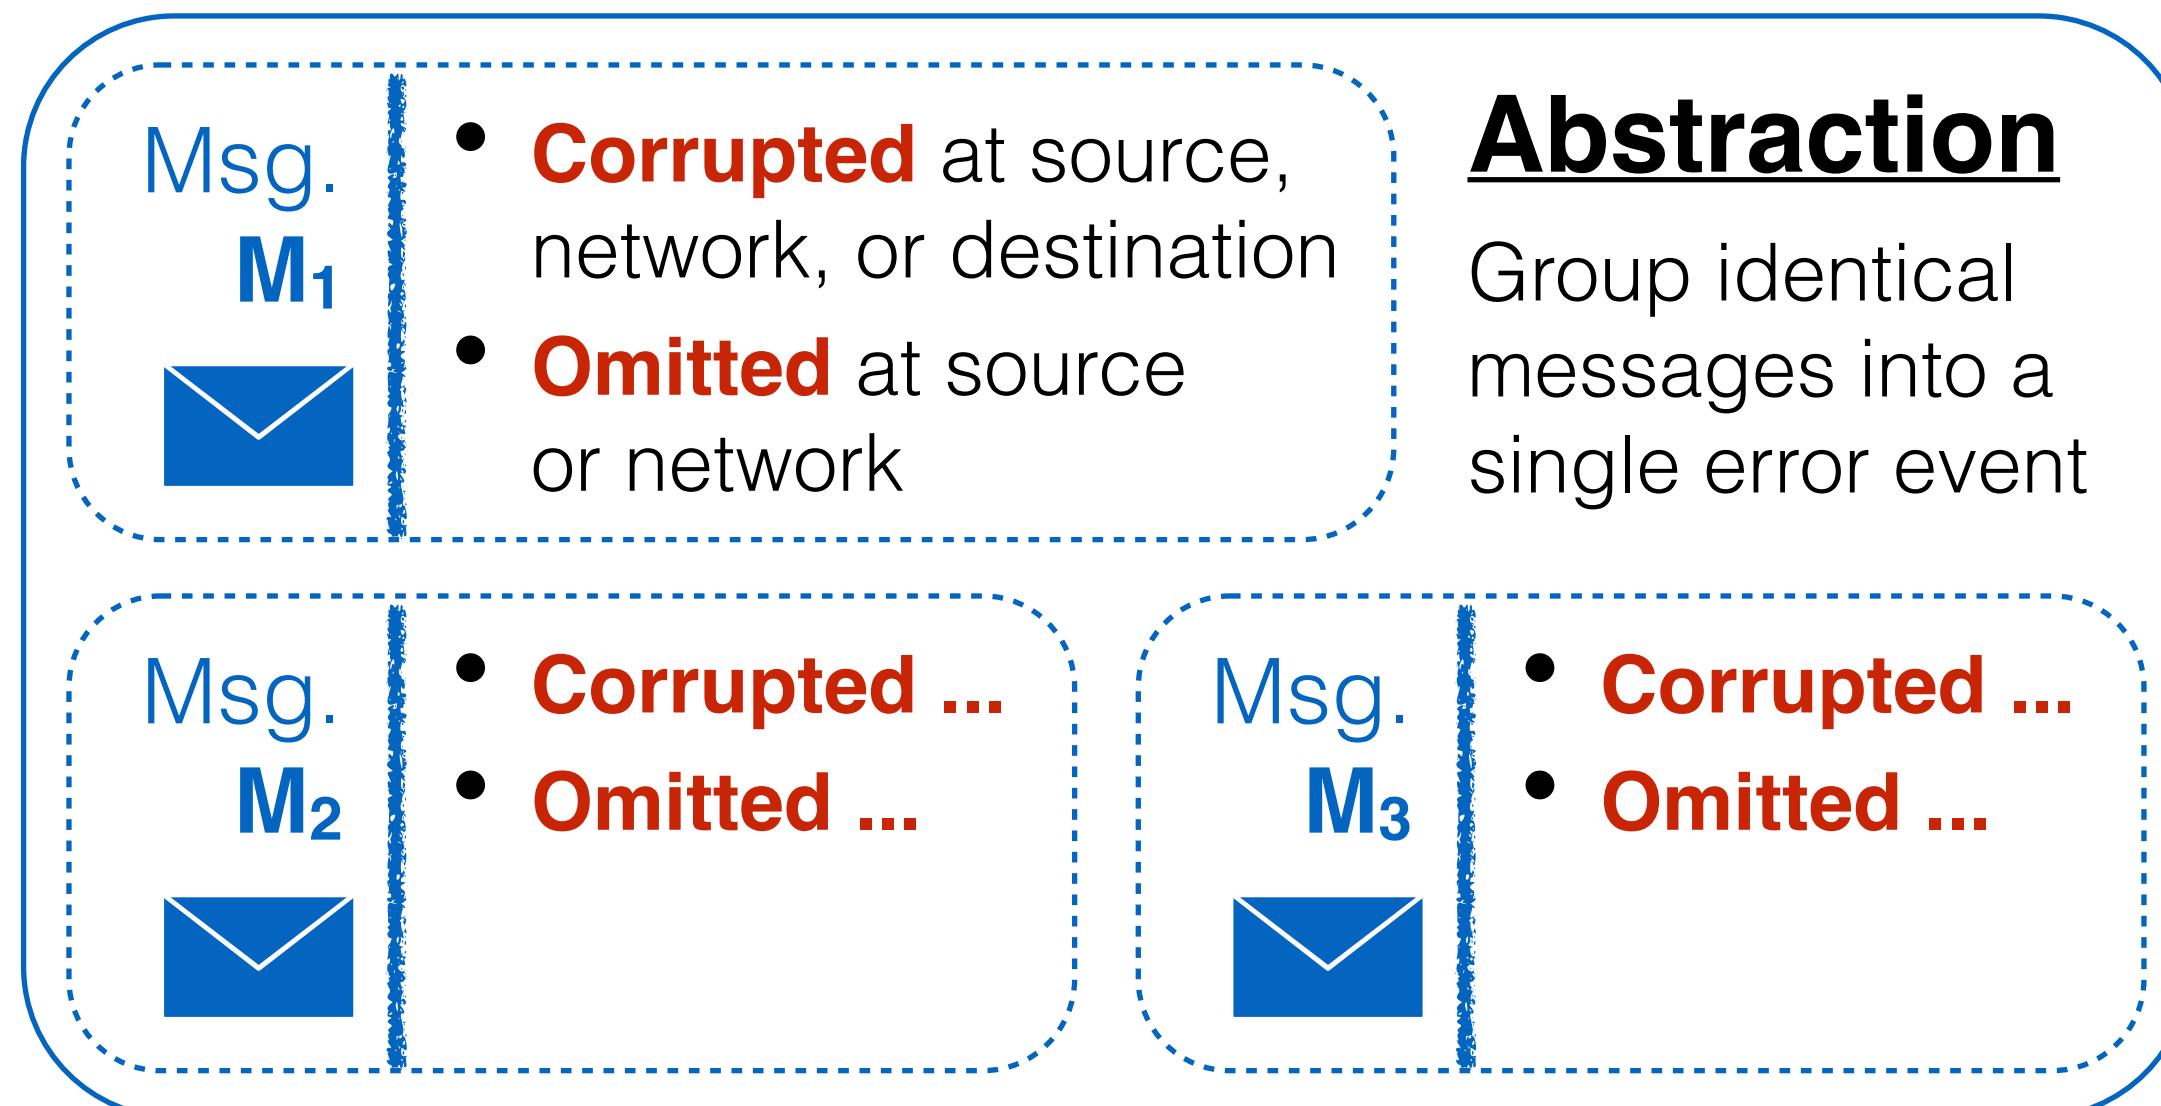
Goal: $P_{UB} > P_r$ (iteration failure)

Application-specific
message errors

Key idea 1: Scalability through **abstraction** and **pruning**

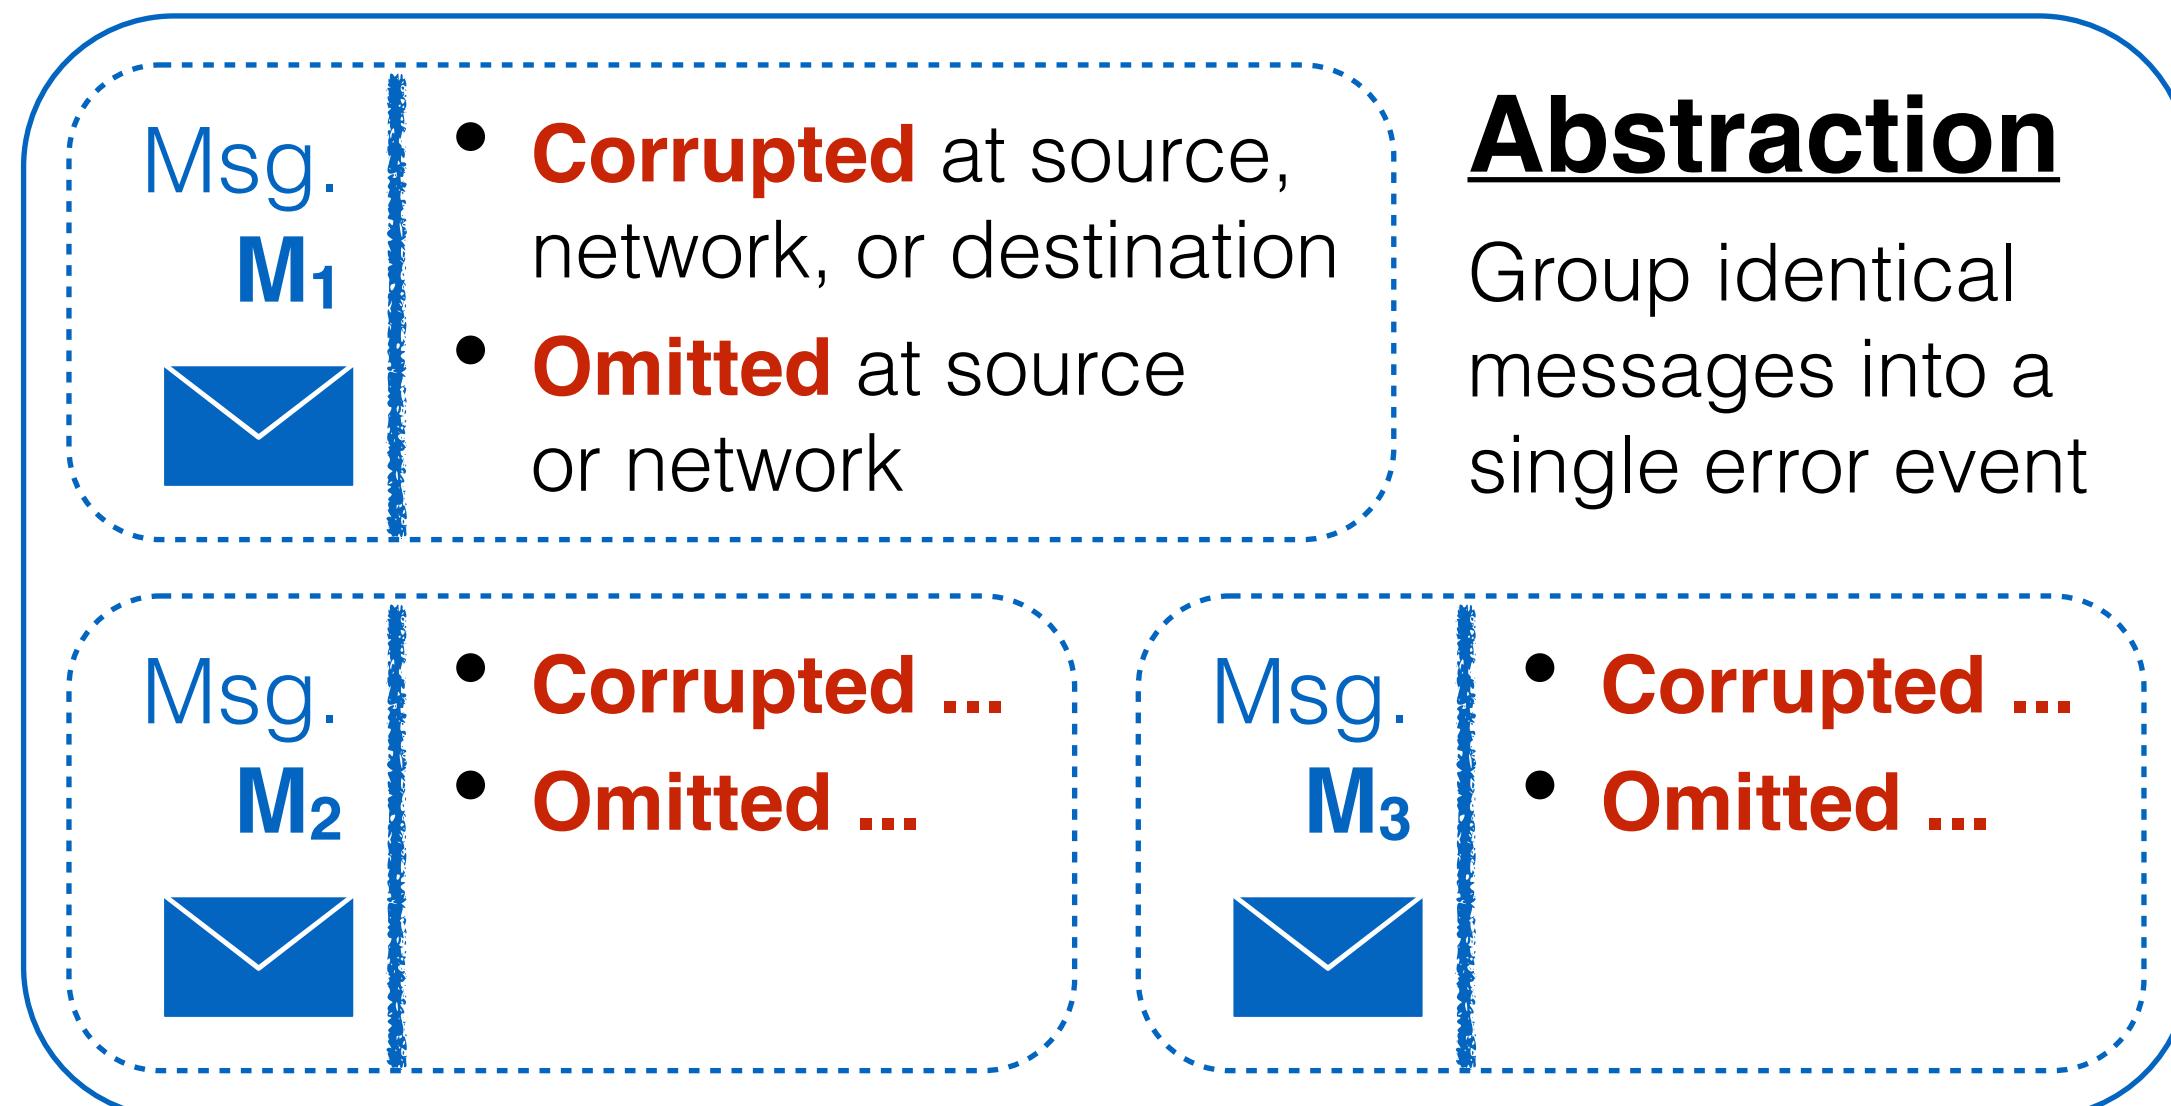


Key idea 1: Scalability through **abstraction** and **pruning**



Error event E_1
Round 1 messages sent by Π_1 omitted at source

Key idea 1: Scalability through **abstraction** and **pruning**



Goal: $P_{UB} > Pr(\text{iteration failure})$

Application-specific message errors

Error event E₁

Round 1 messages sent by Π_1 omitted at source

Error event E₂

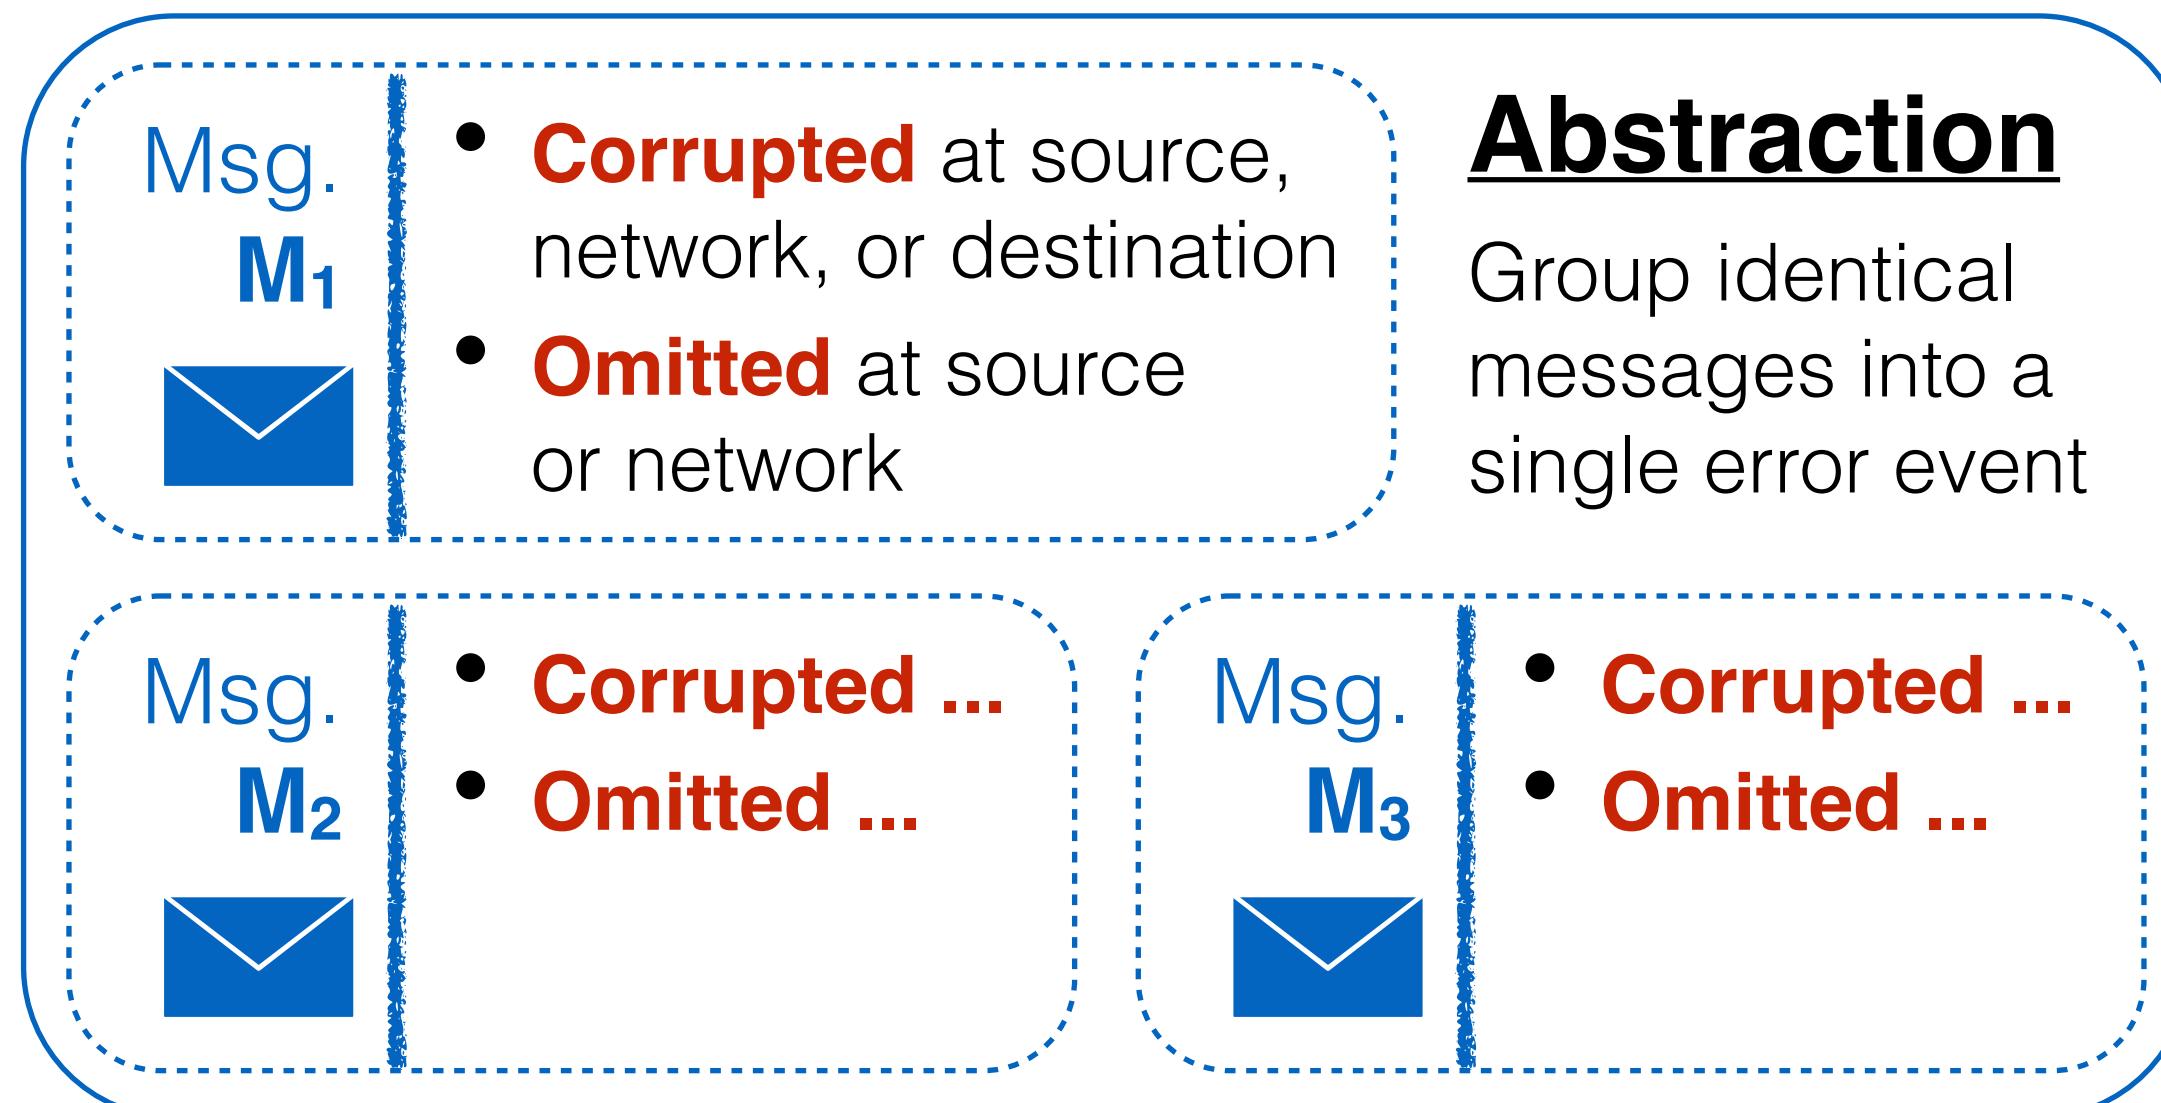
Round 1 messages sent by Π_1 corrupted at source

Network error event E₃

Frame carrying round 1 messages from Π_1 to Π_2 corrupted by the network

...

Key idea 1: Scalability through **abstraction** and **pruning**

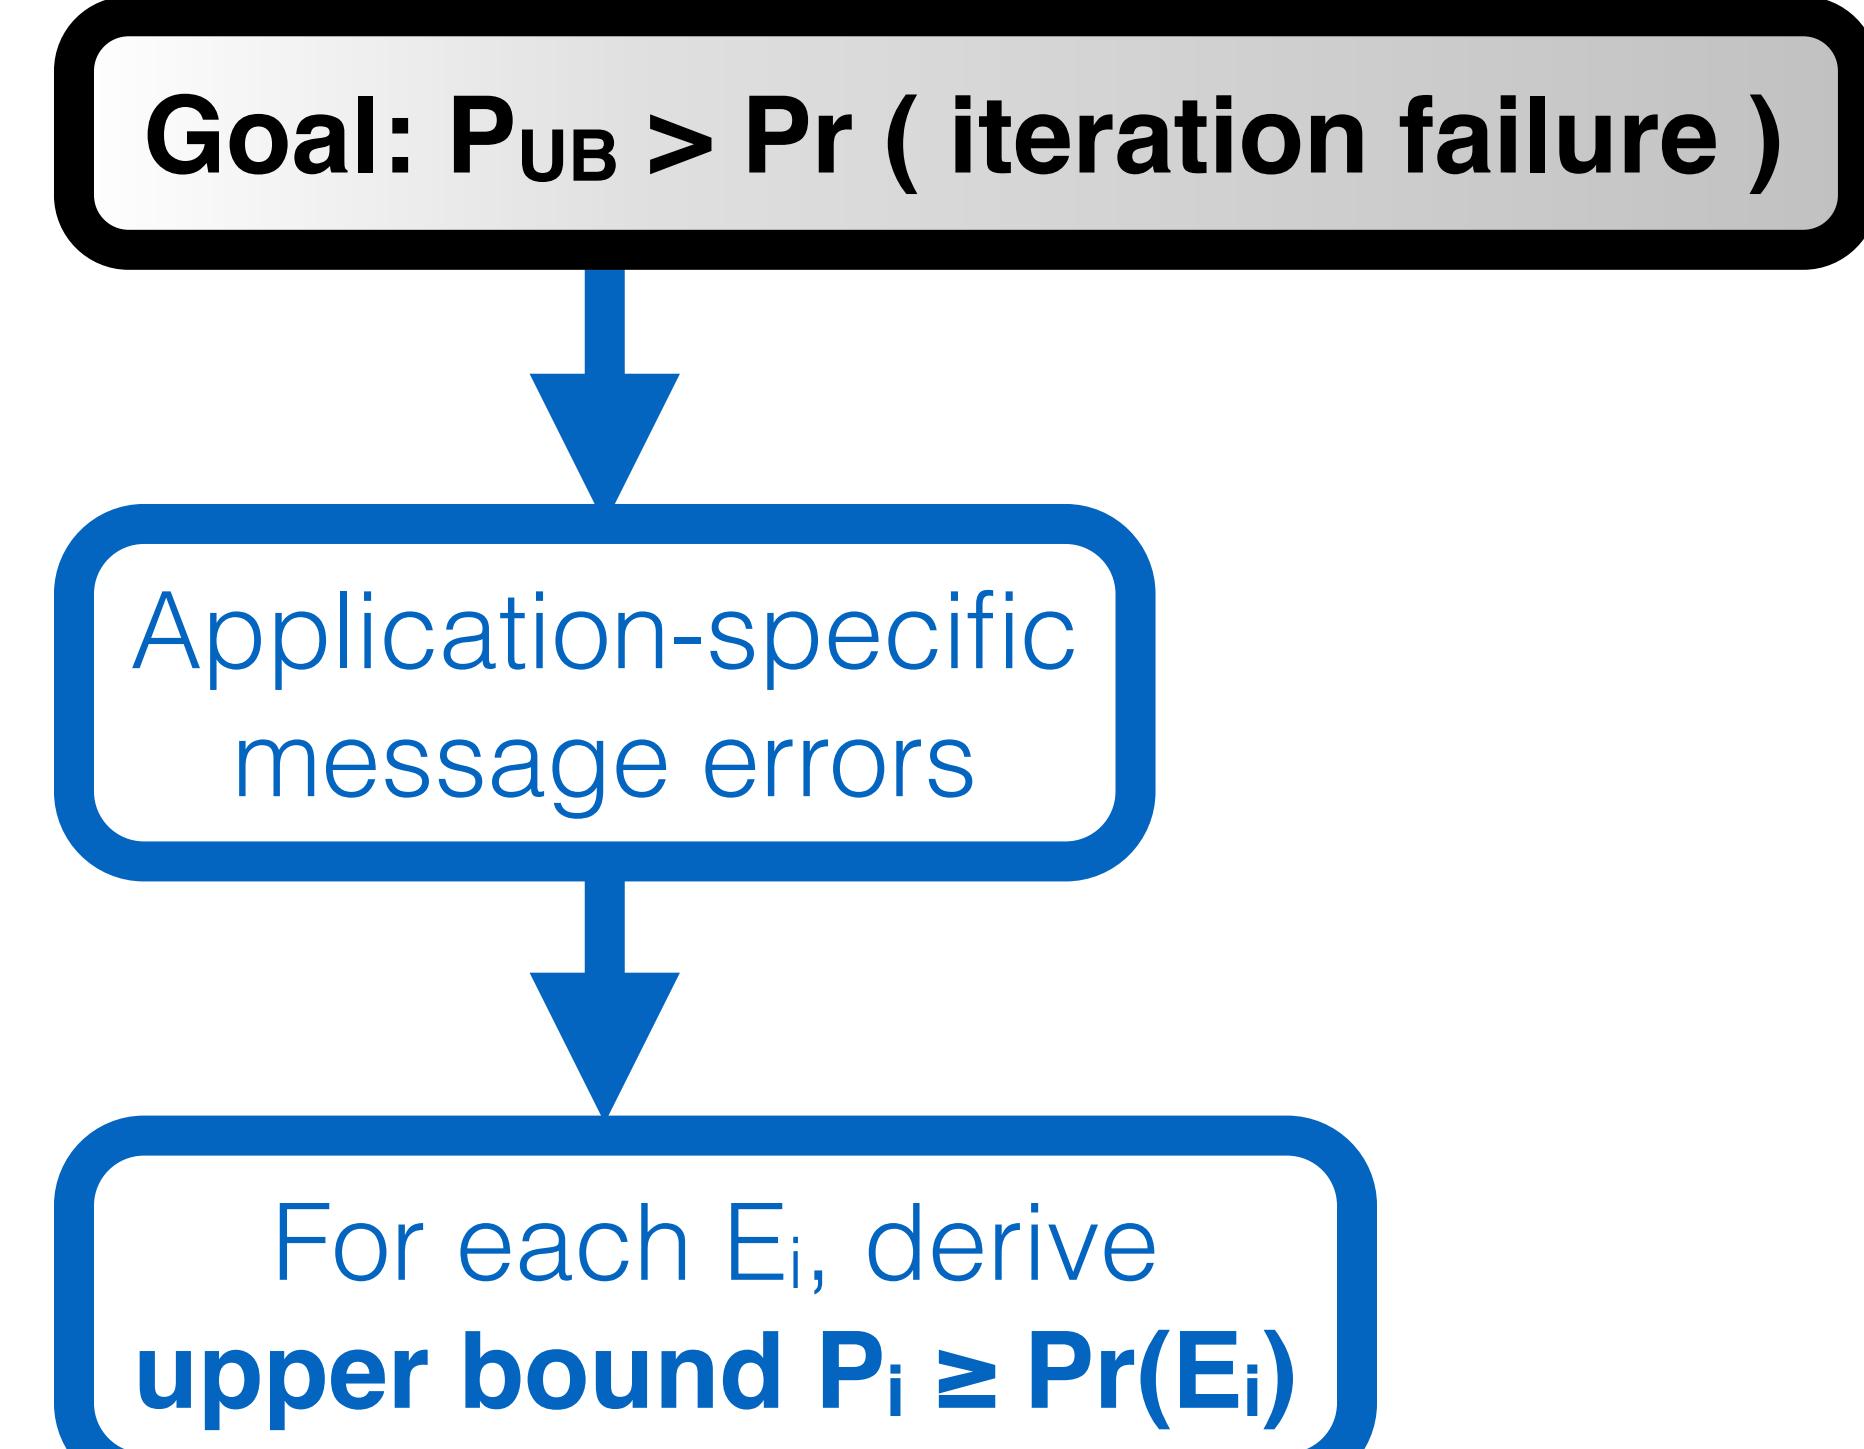


Error event E₁
Round 1 messages sent by Π_1 omitted at source

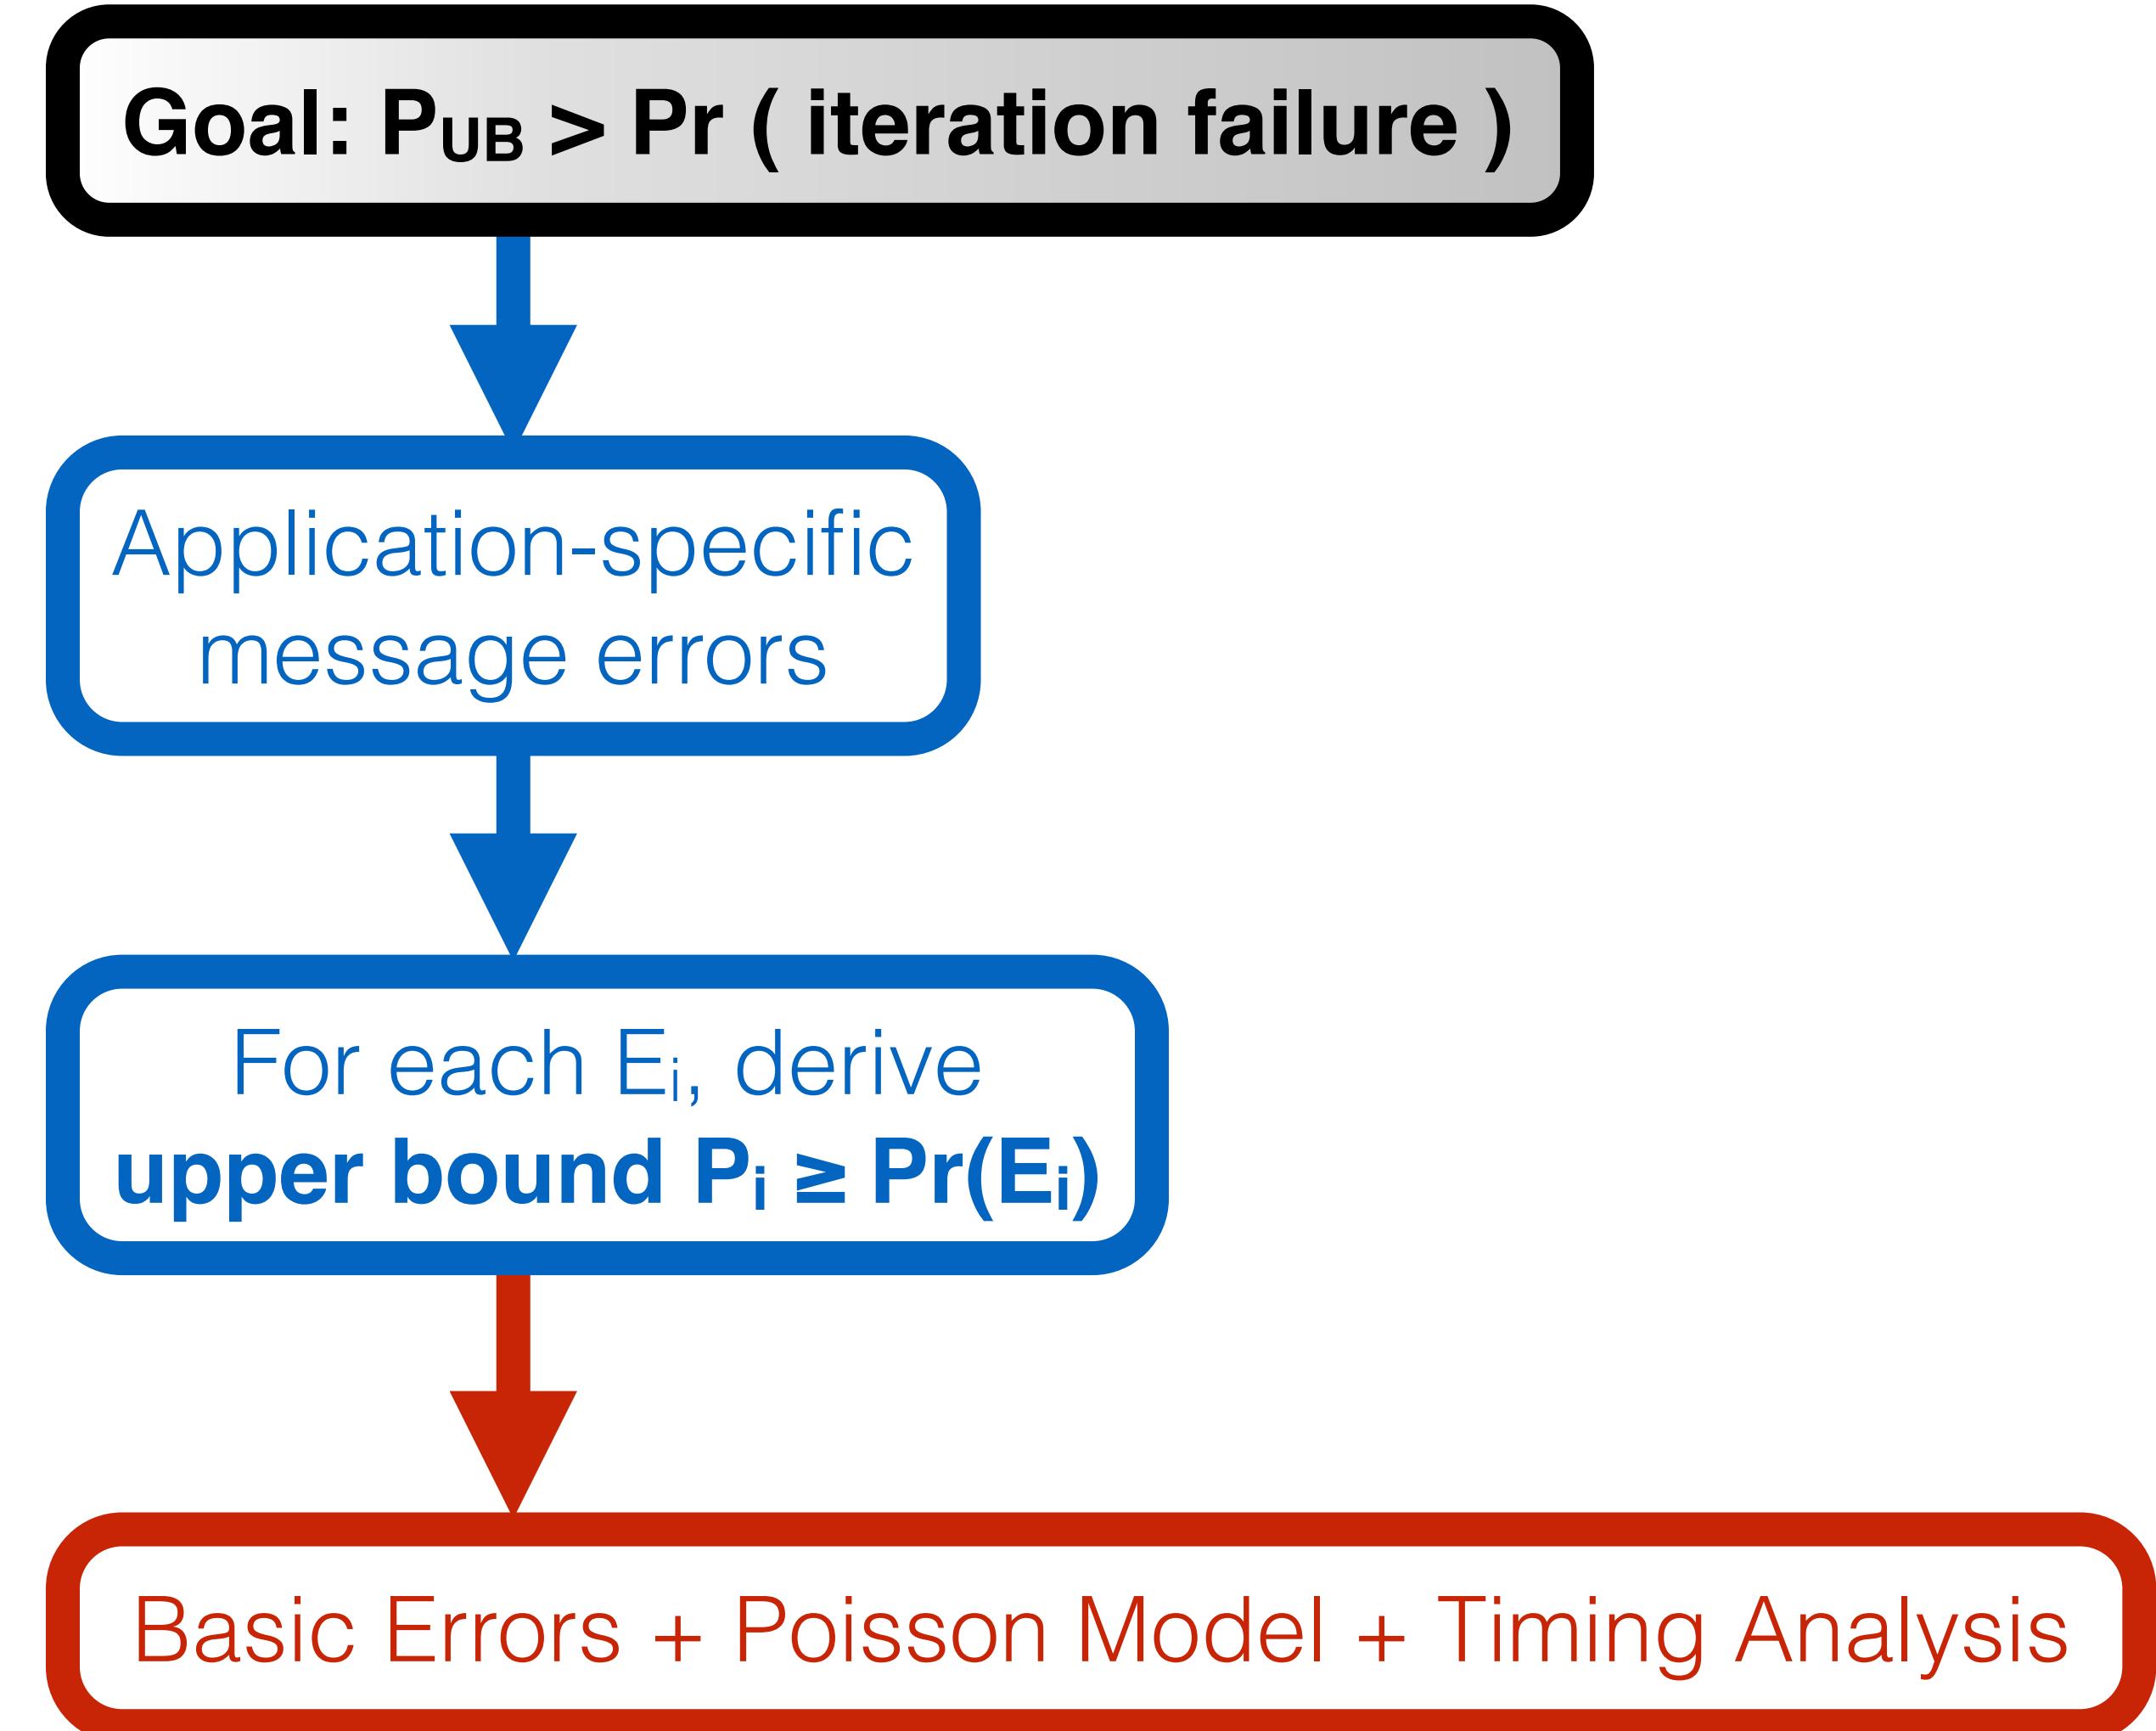
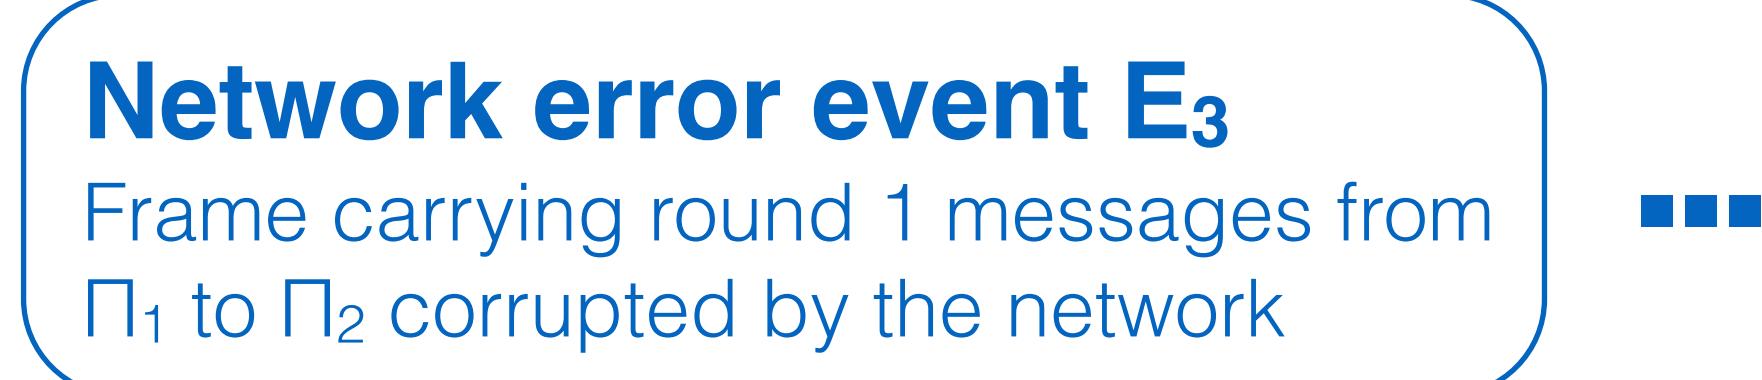
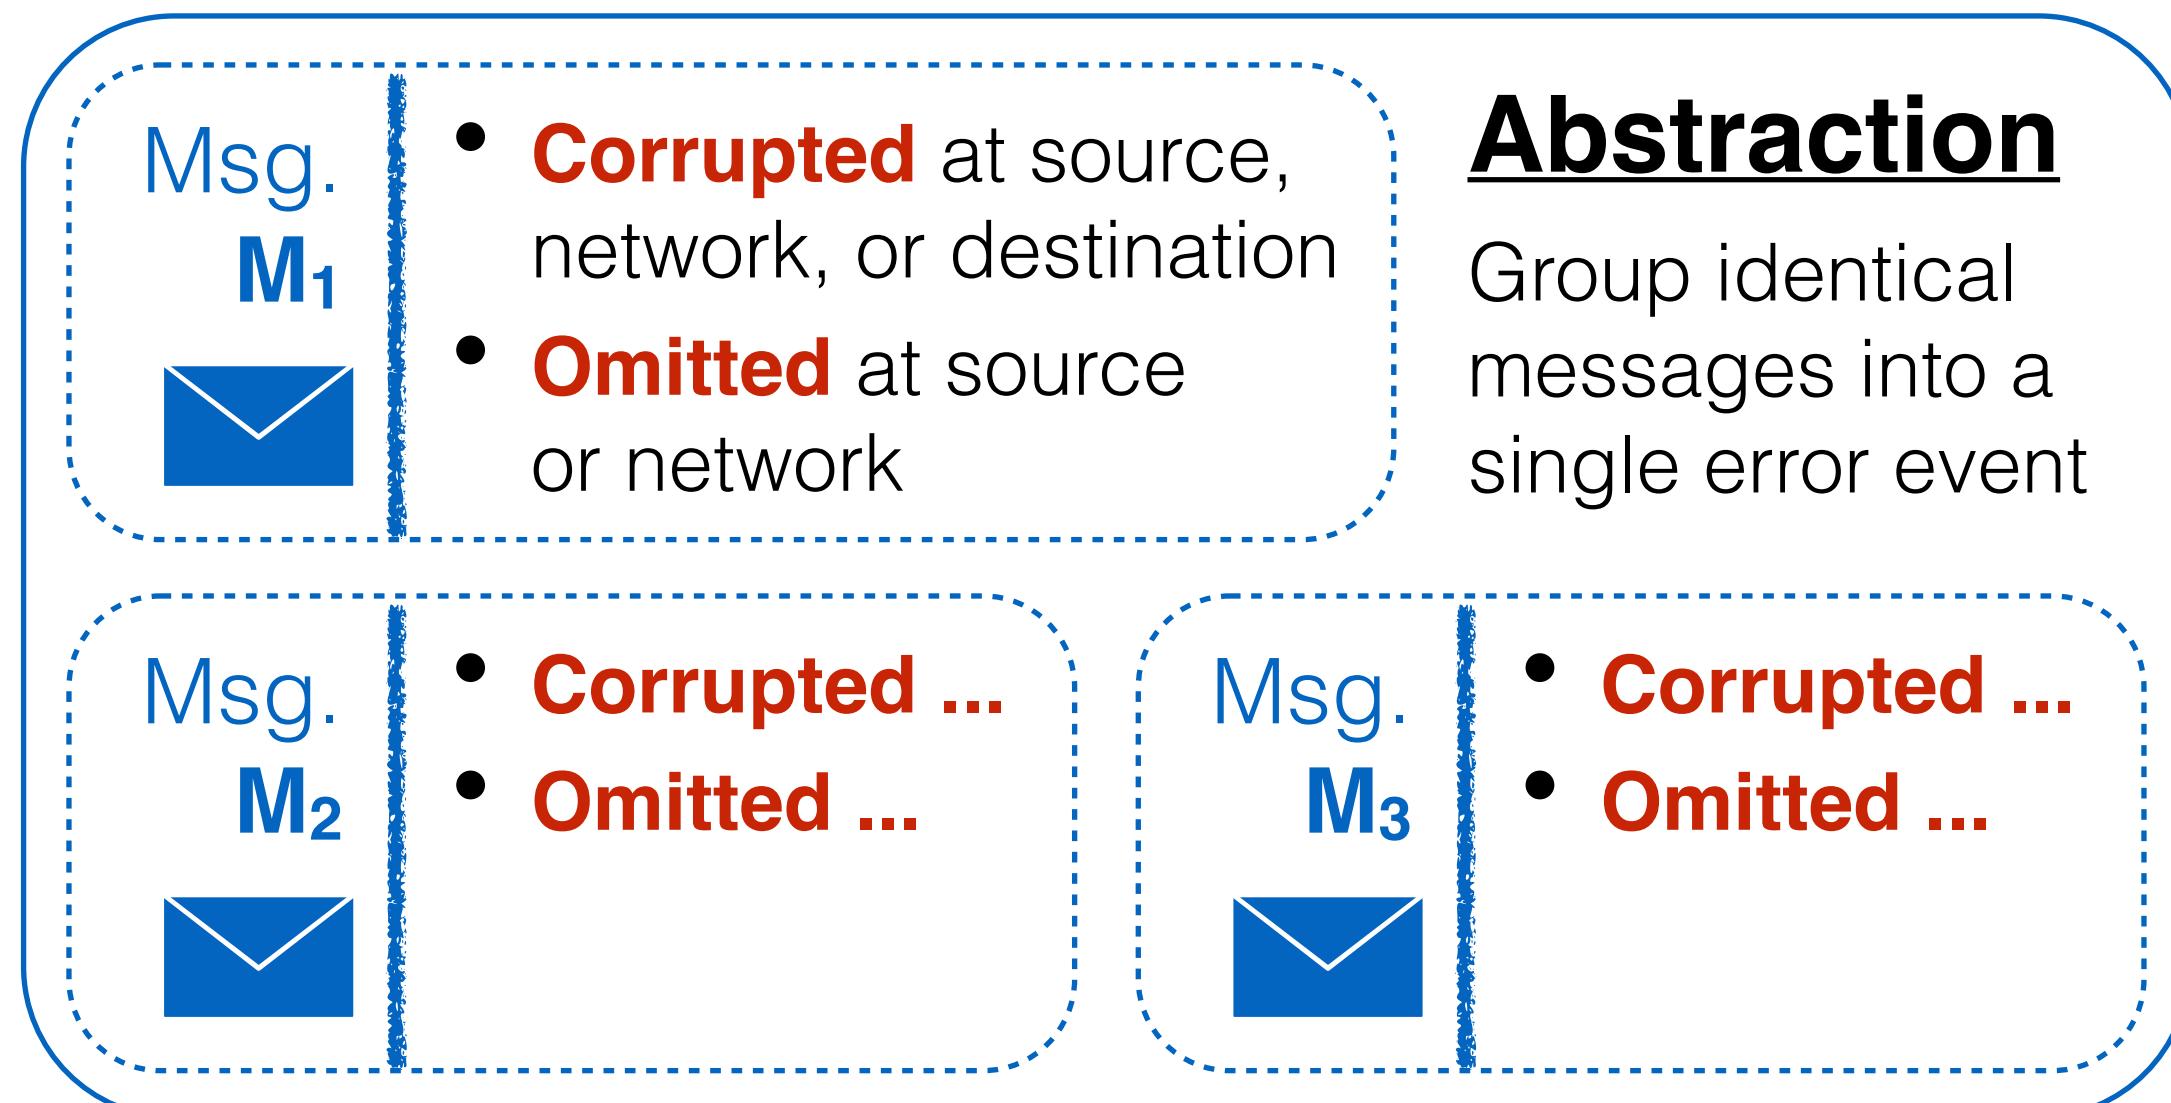
Error event E₂
Round 1 messages sent by Π_1 corrupted at source

Network error event E₃
Frame carrying round 1 messages from Π_1 to Π_2 corrupted by the network

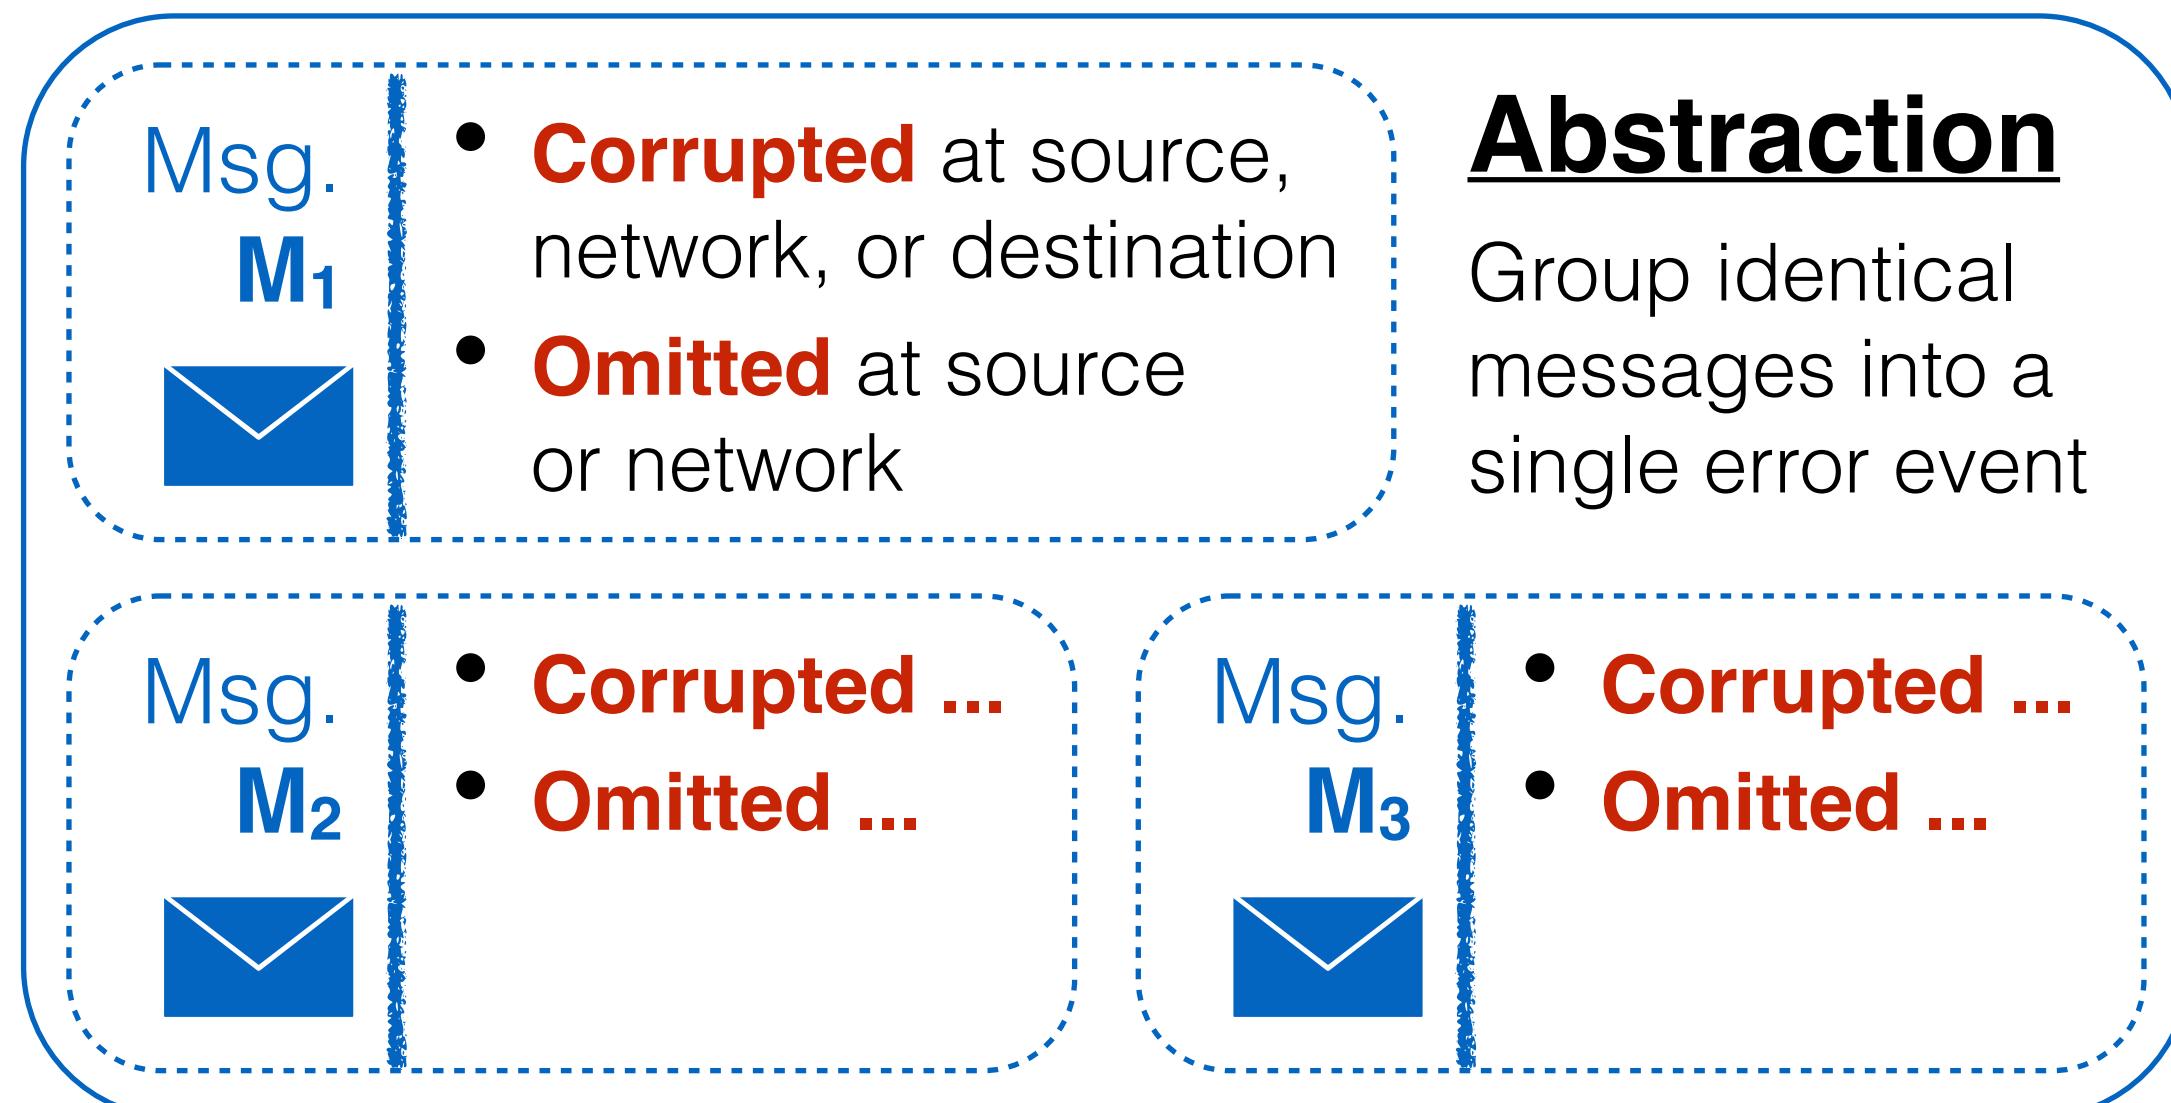
...



Key idea 1: Scalability through **abstraction** and **pruning**



Key idea 1: Scalability through **abstraction** and **pruning**

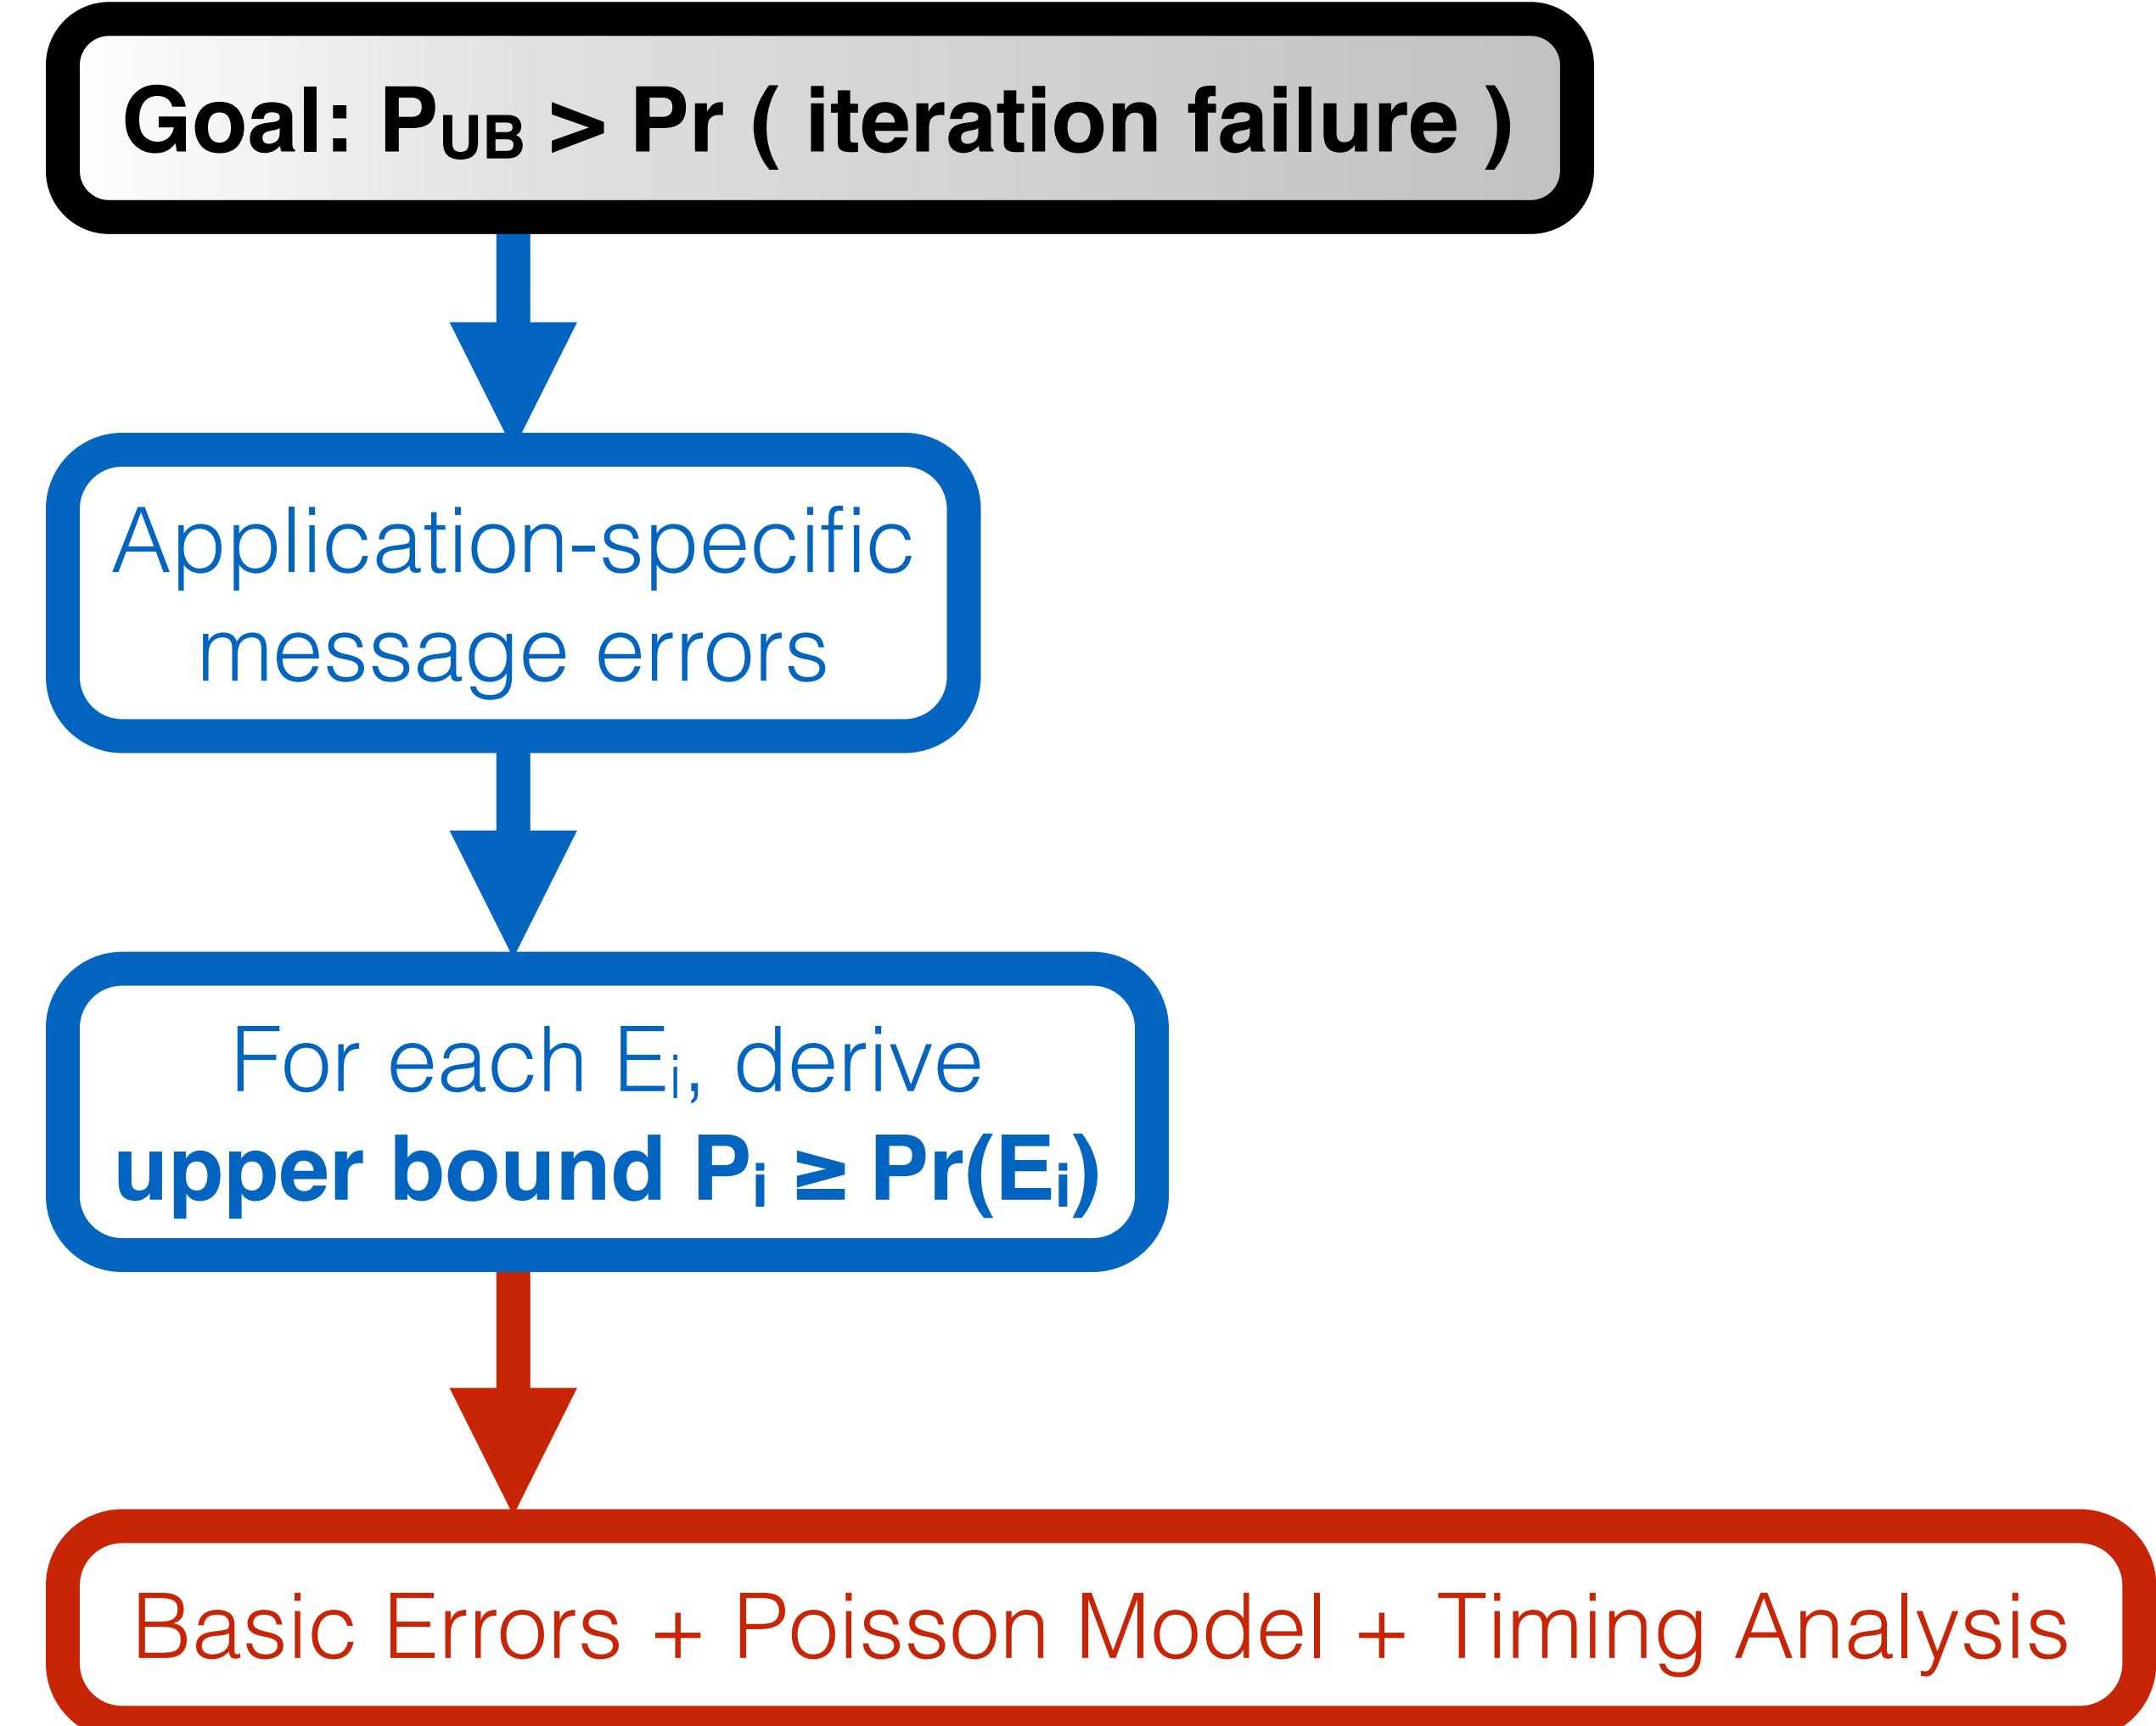


Error event E₁
Round 1 messages sent by Π_1 omitted at source

Error event E₂
Round 1 messages sent by Π_1 corrupted at source

Network error event E₃
Frame carrying round 1 messages from Π_1 to Π_2 corrupted by the network

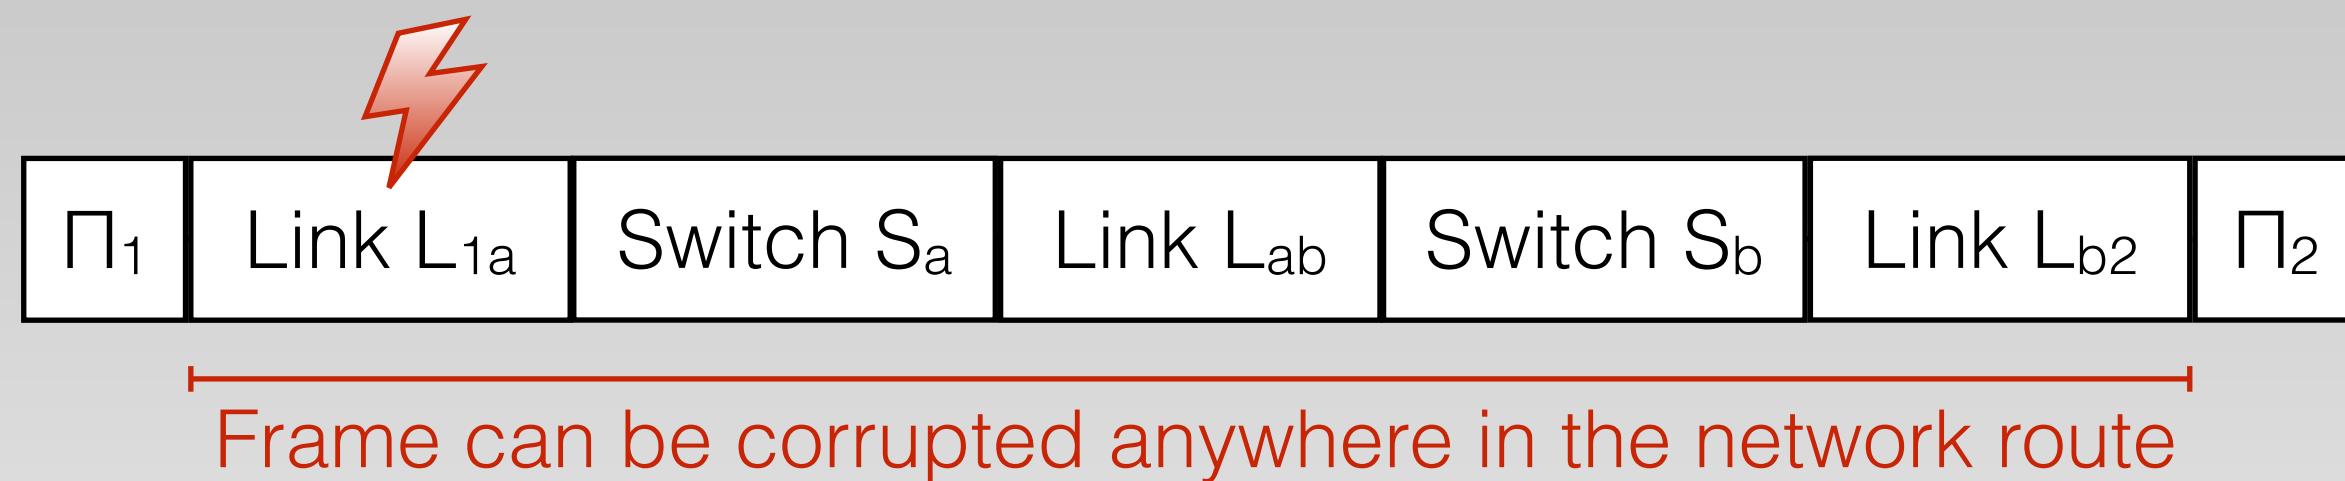
■■■
Example!



Example: E_3 = Frame carrying Round 1 messages from Π_1 to Π_2 corrupted by the network



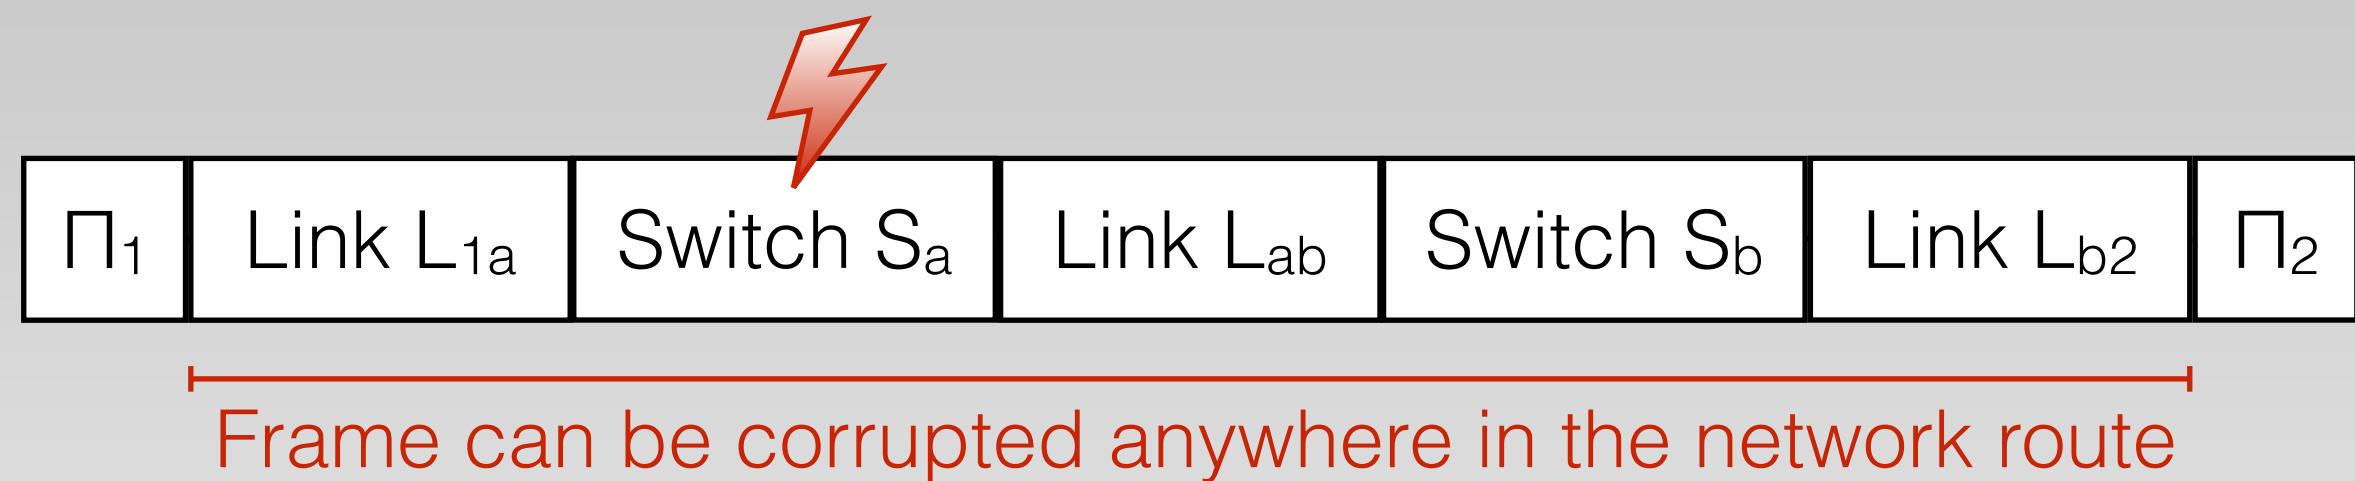
Example: E_3 = Frame carrying Round 1 messages from Π_1 to Π_2 corrupted by the network



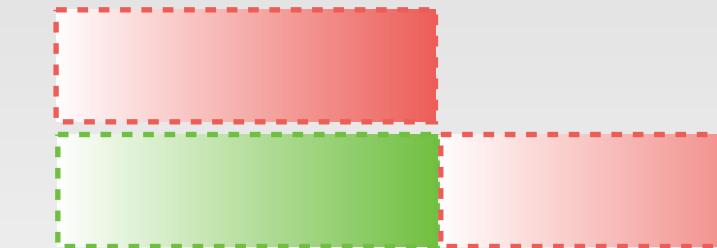
Scenario 1



Example: E_3 = Frame carrying Round 1 messages from Π_1 to Π_2 corrupted by the network



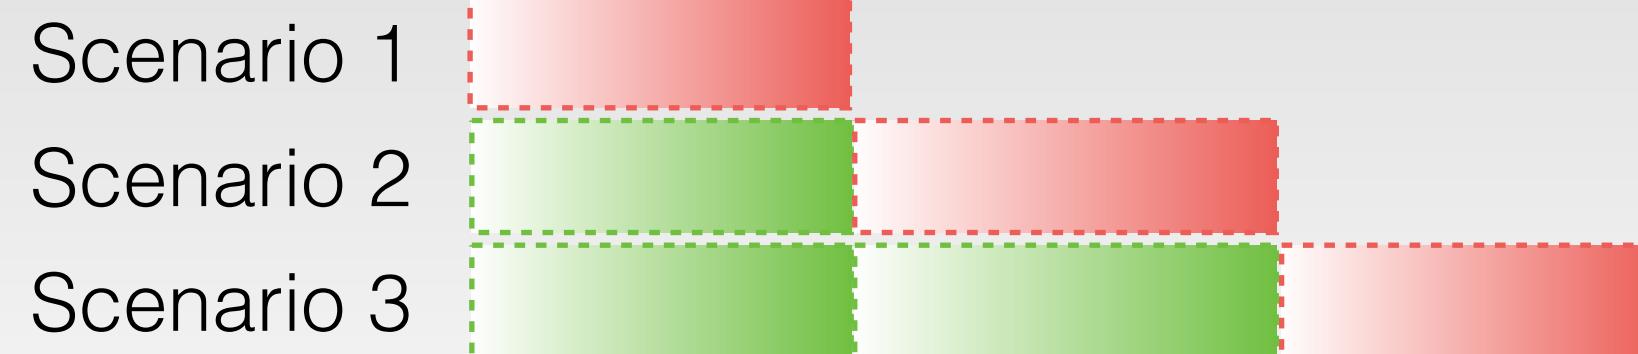
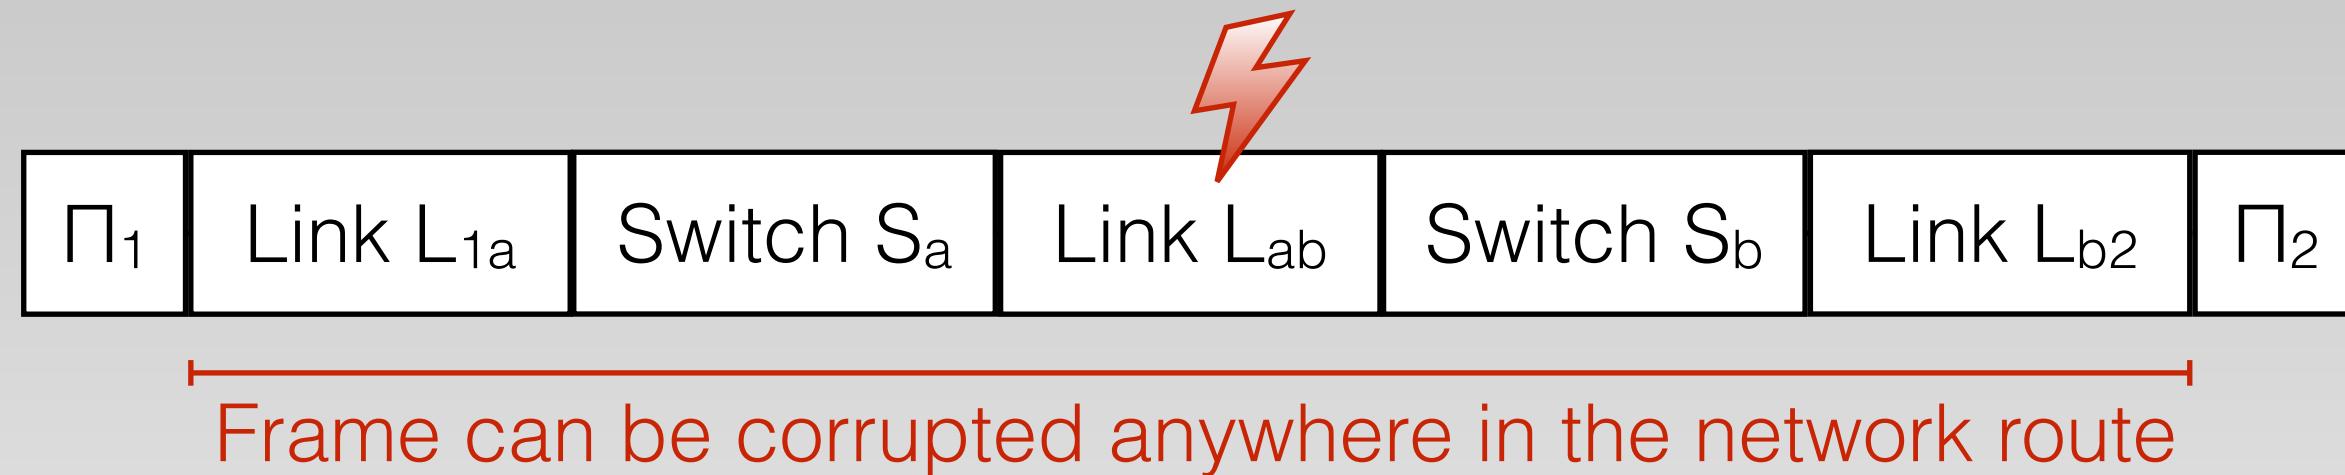
Scenario 1



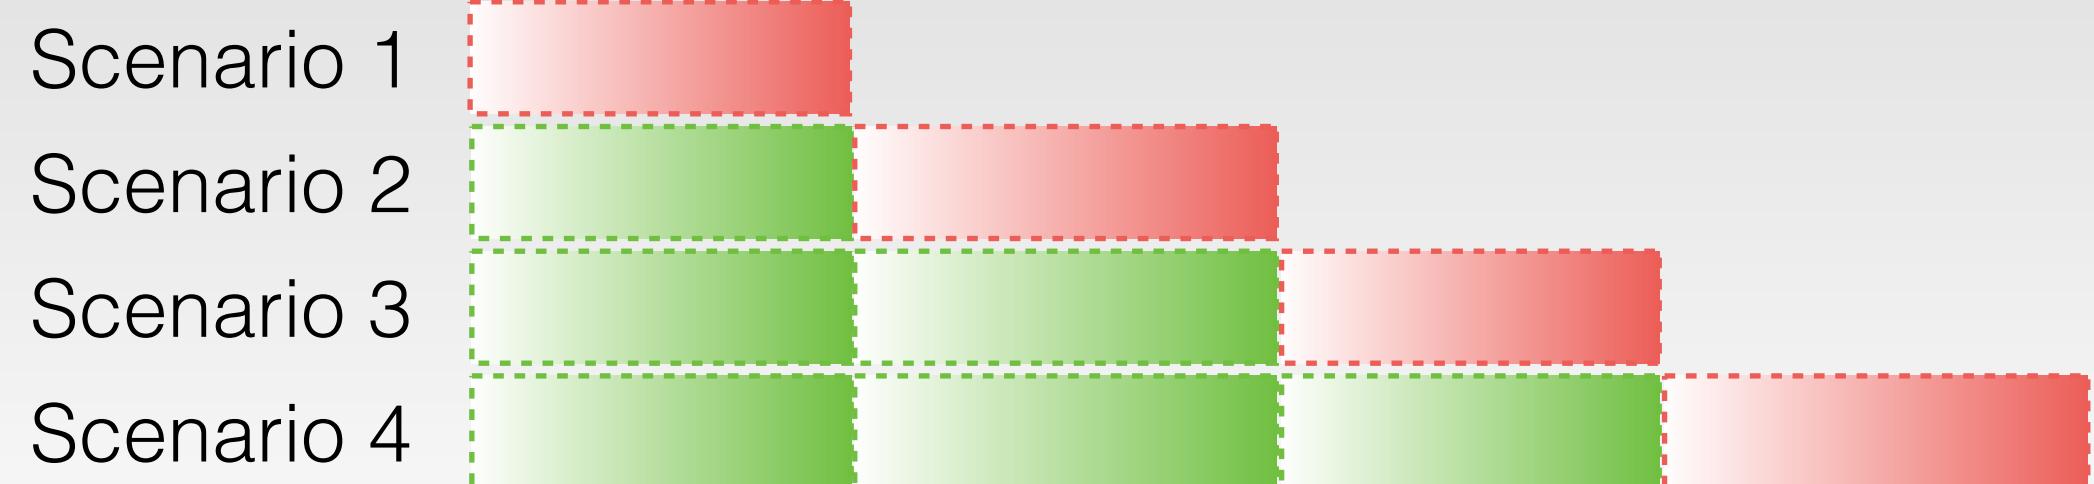
Scenario 2



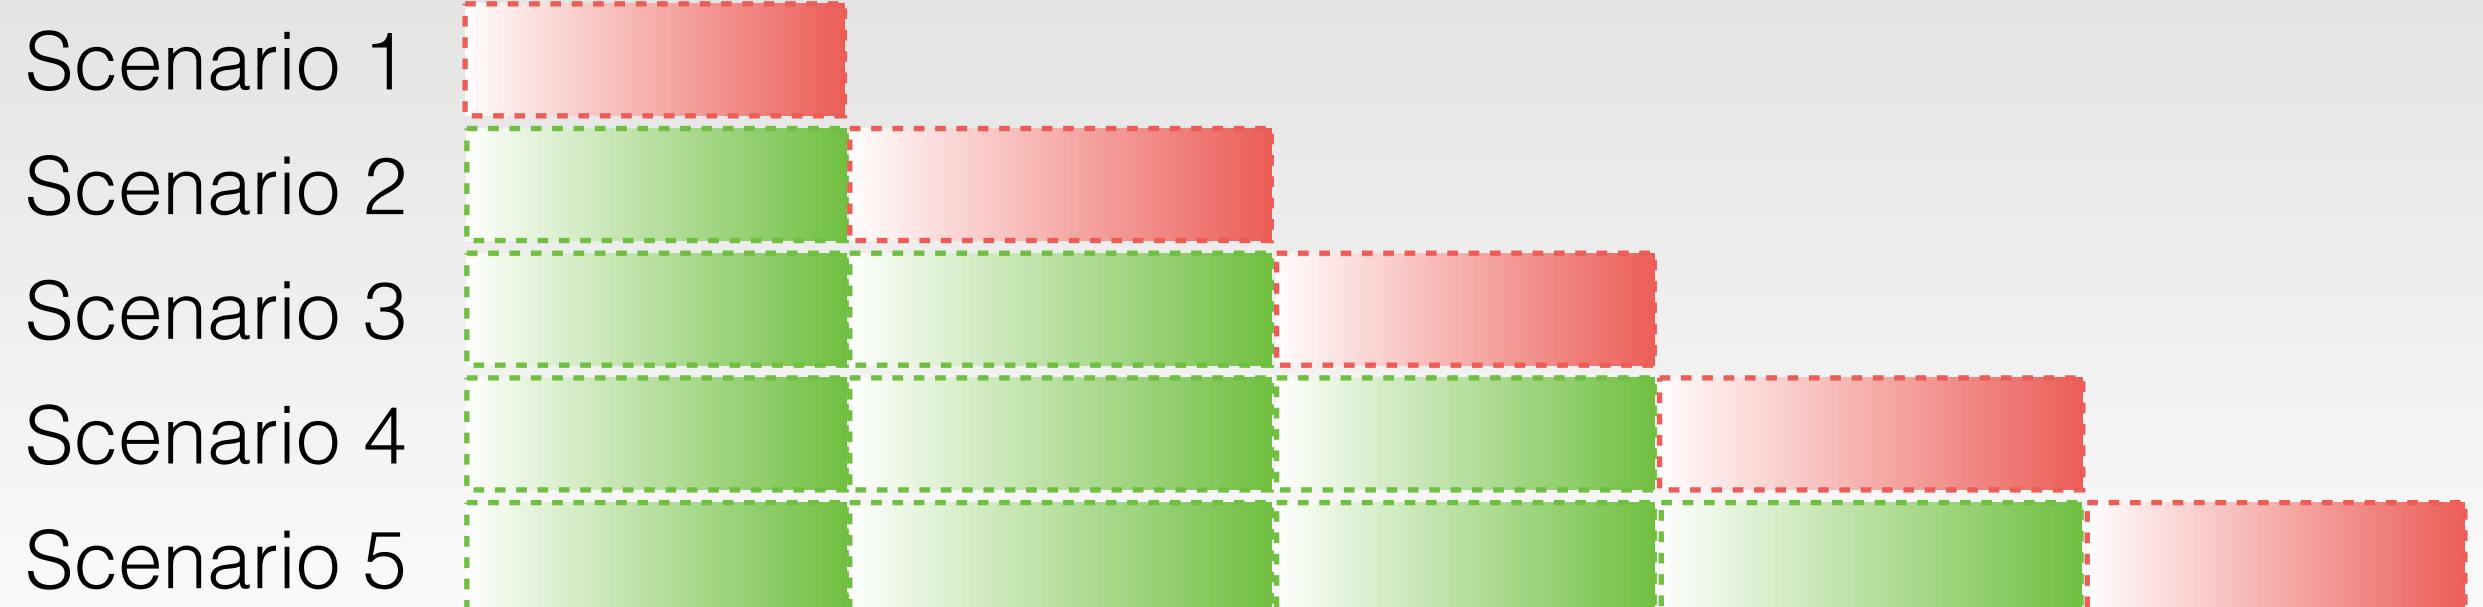
Example: E_3 = Frame carrying Round 1 messages from Π_1 to Π_2 corrupted by the network



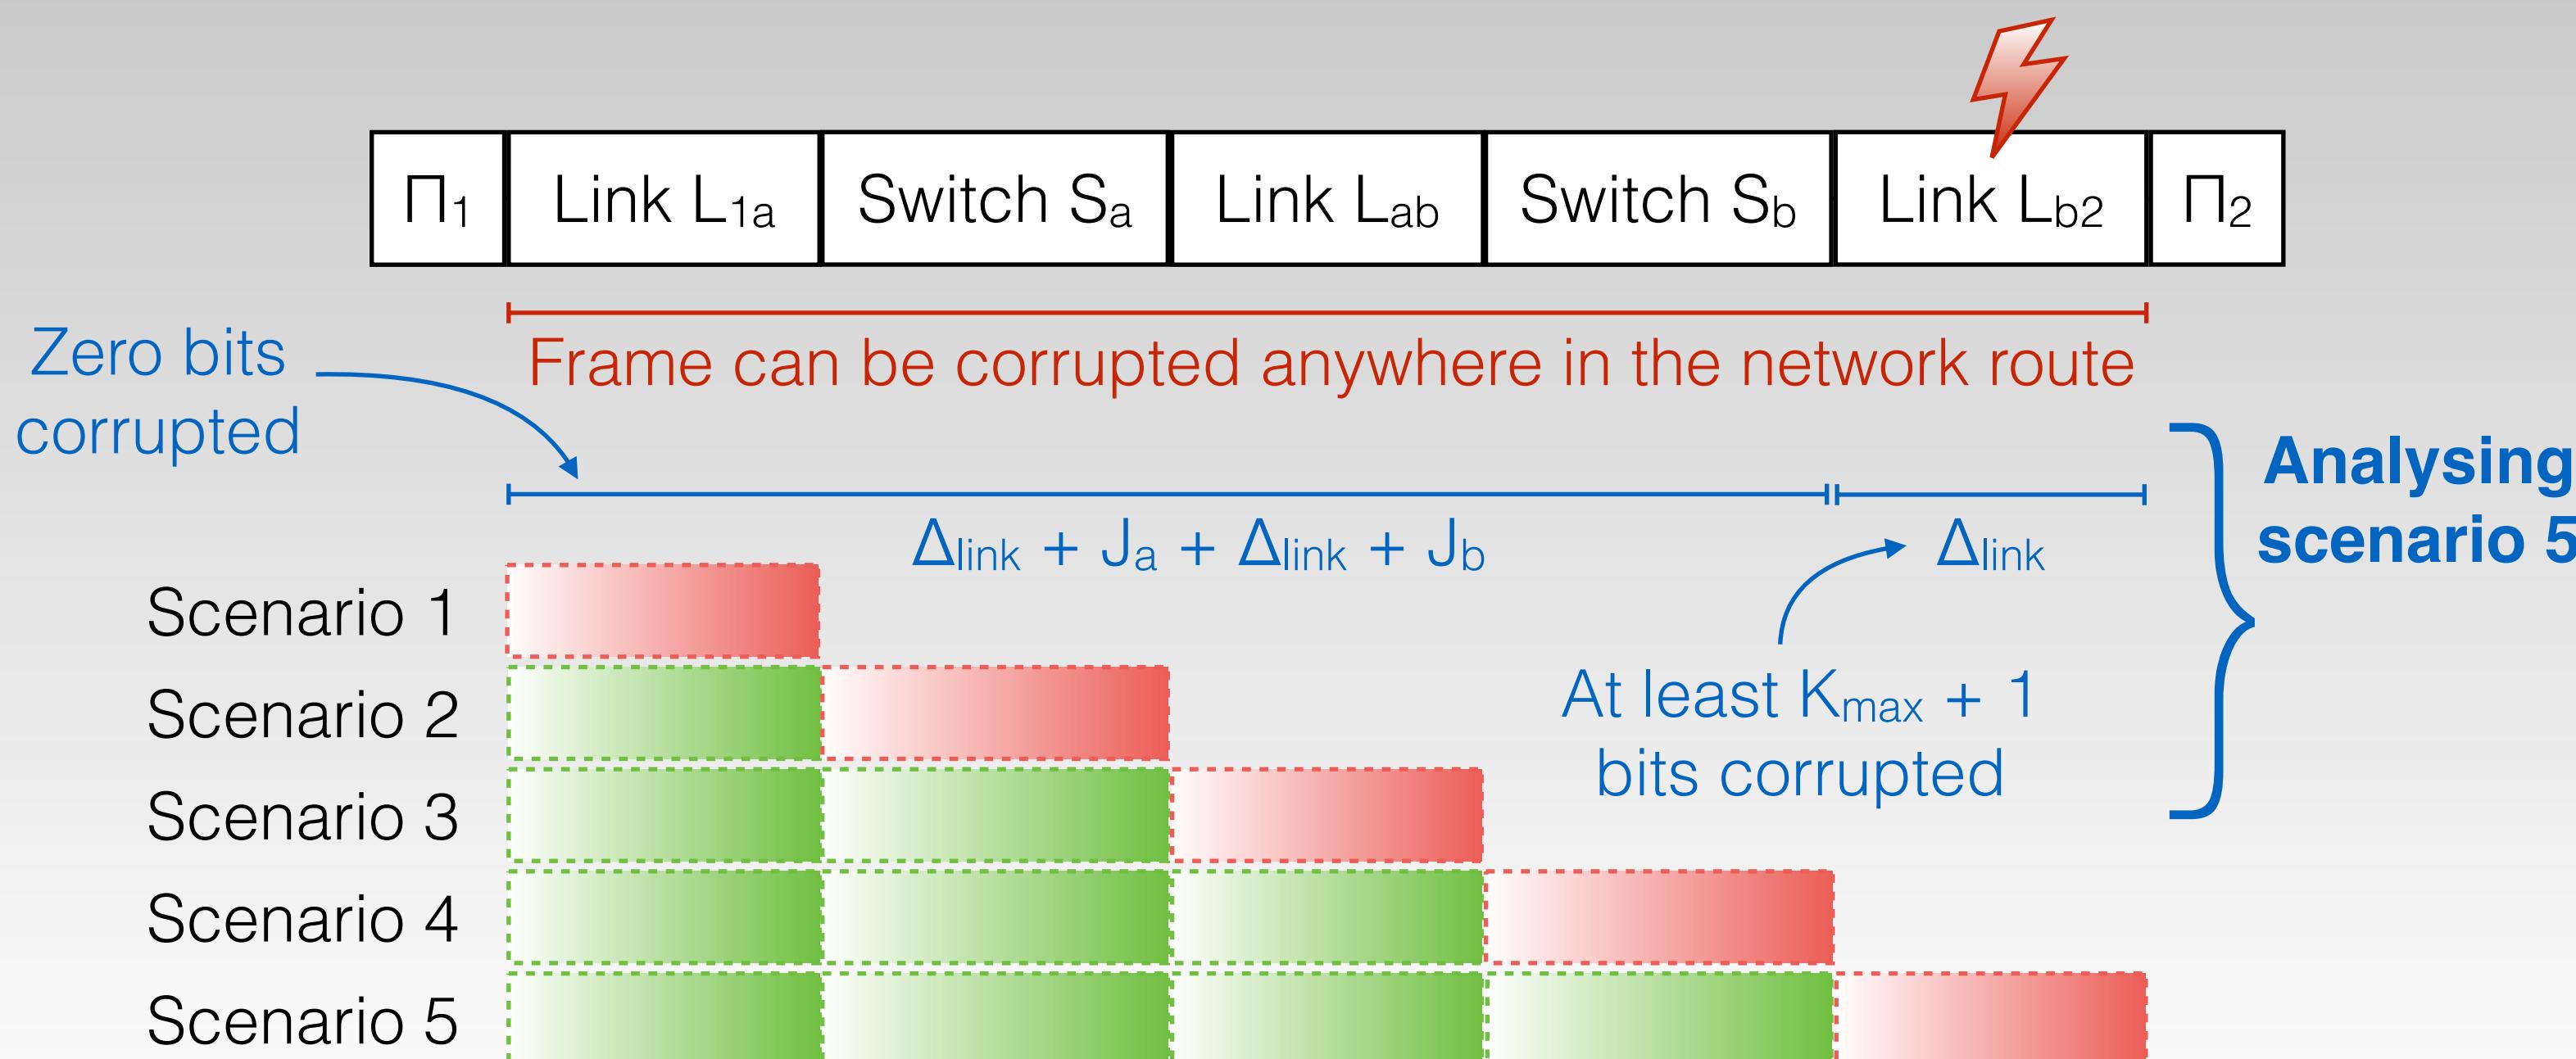
Example: E_3 = Frame carrying Round 1 messages from Π_1 to Π_2 corrupted by the network



Example: E_3 = Frame carrying Round 1 messages from Π_1 to Π_2 corrupted by the network



Example: E_3 = Frame carrying Round 1 messages from Π_1 to Π_2 corrupted by the network



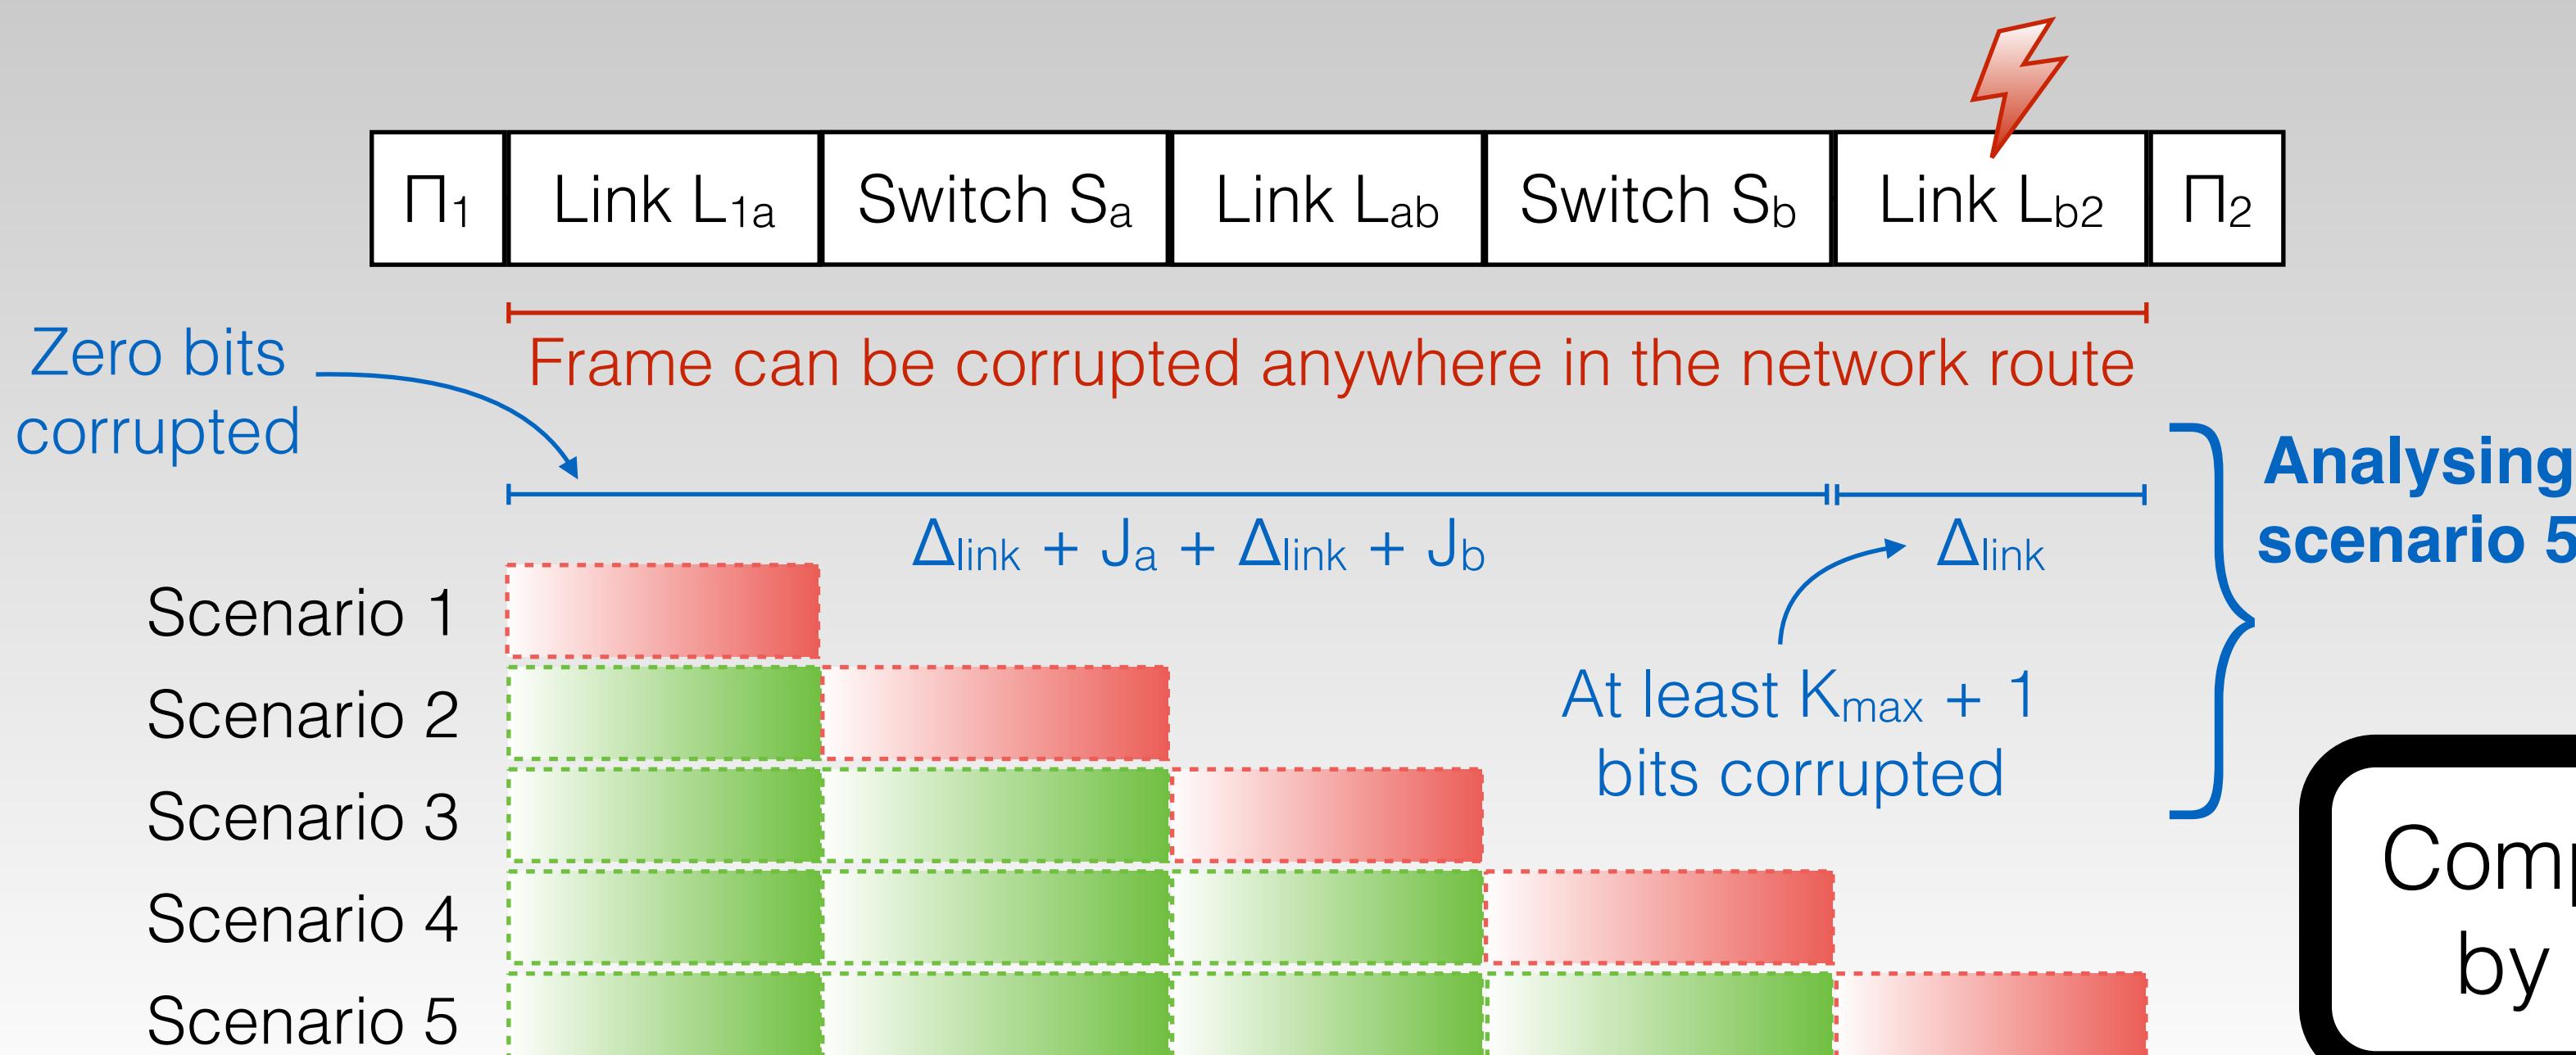
Notation

Δ_{link} = Link transfer time

J_a = Maximum scheduling jitter at Switch S_a

K_{max} = Maximum bit flips detected by the CRC

Example: E_3 = Frame carrying Round 1 messages from Π_1 to Π_2 corrupted by the network



Notation

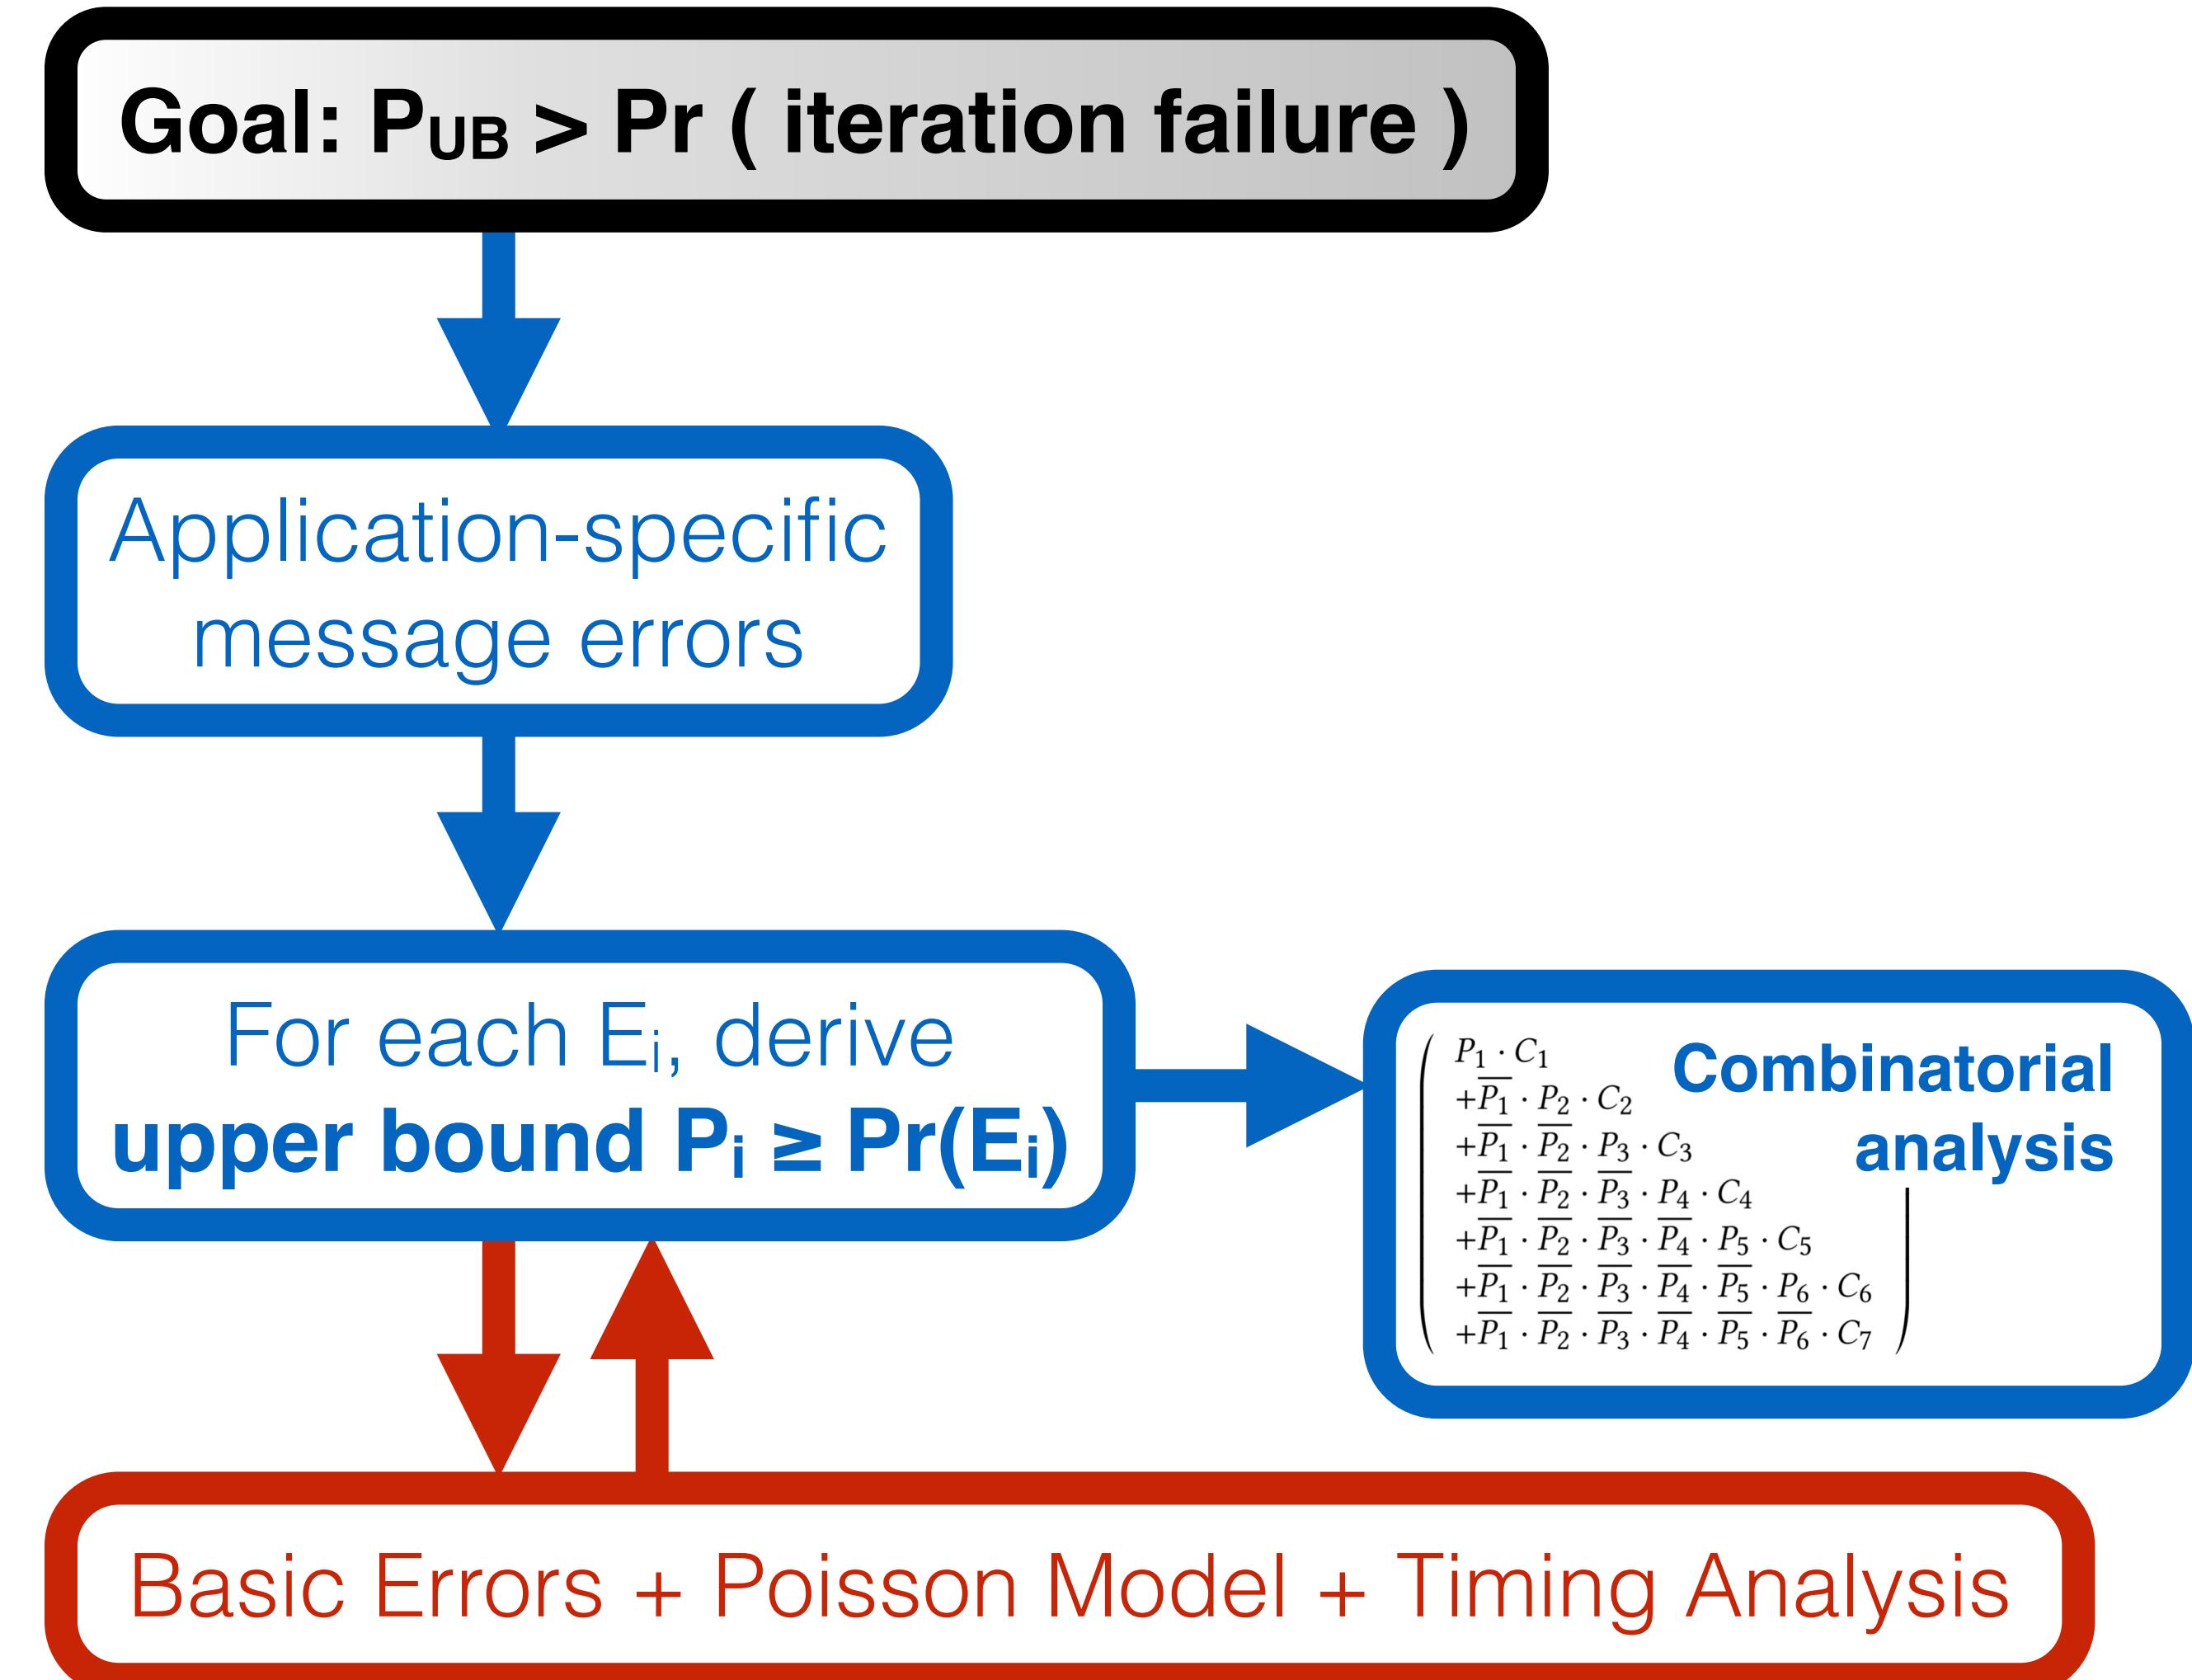
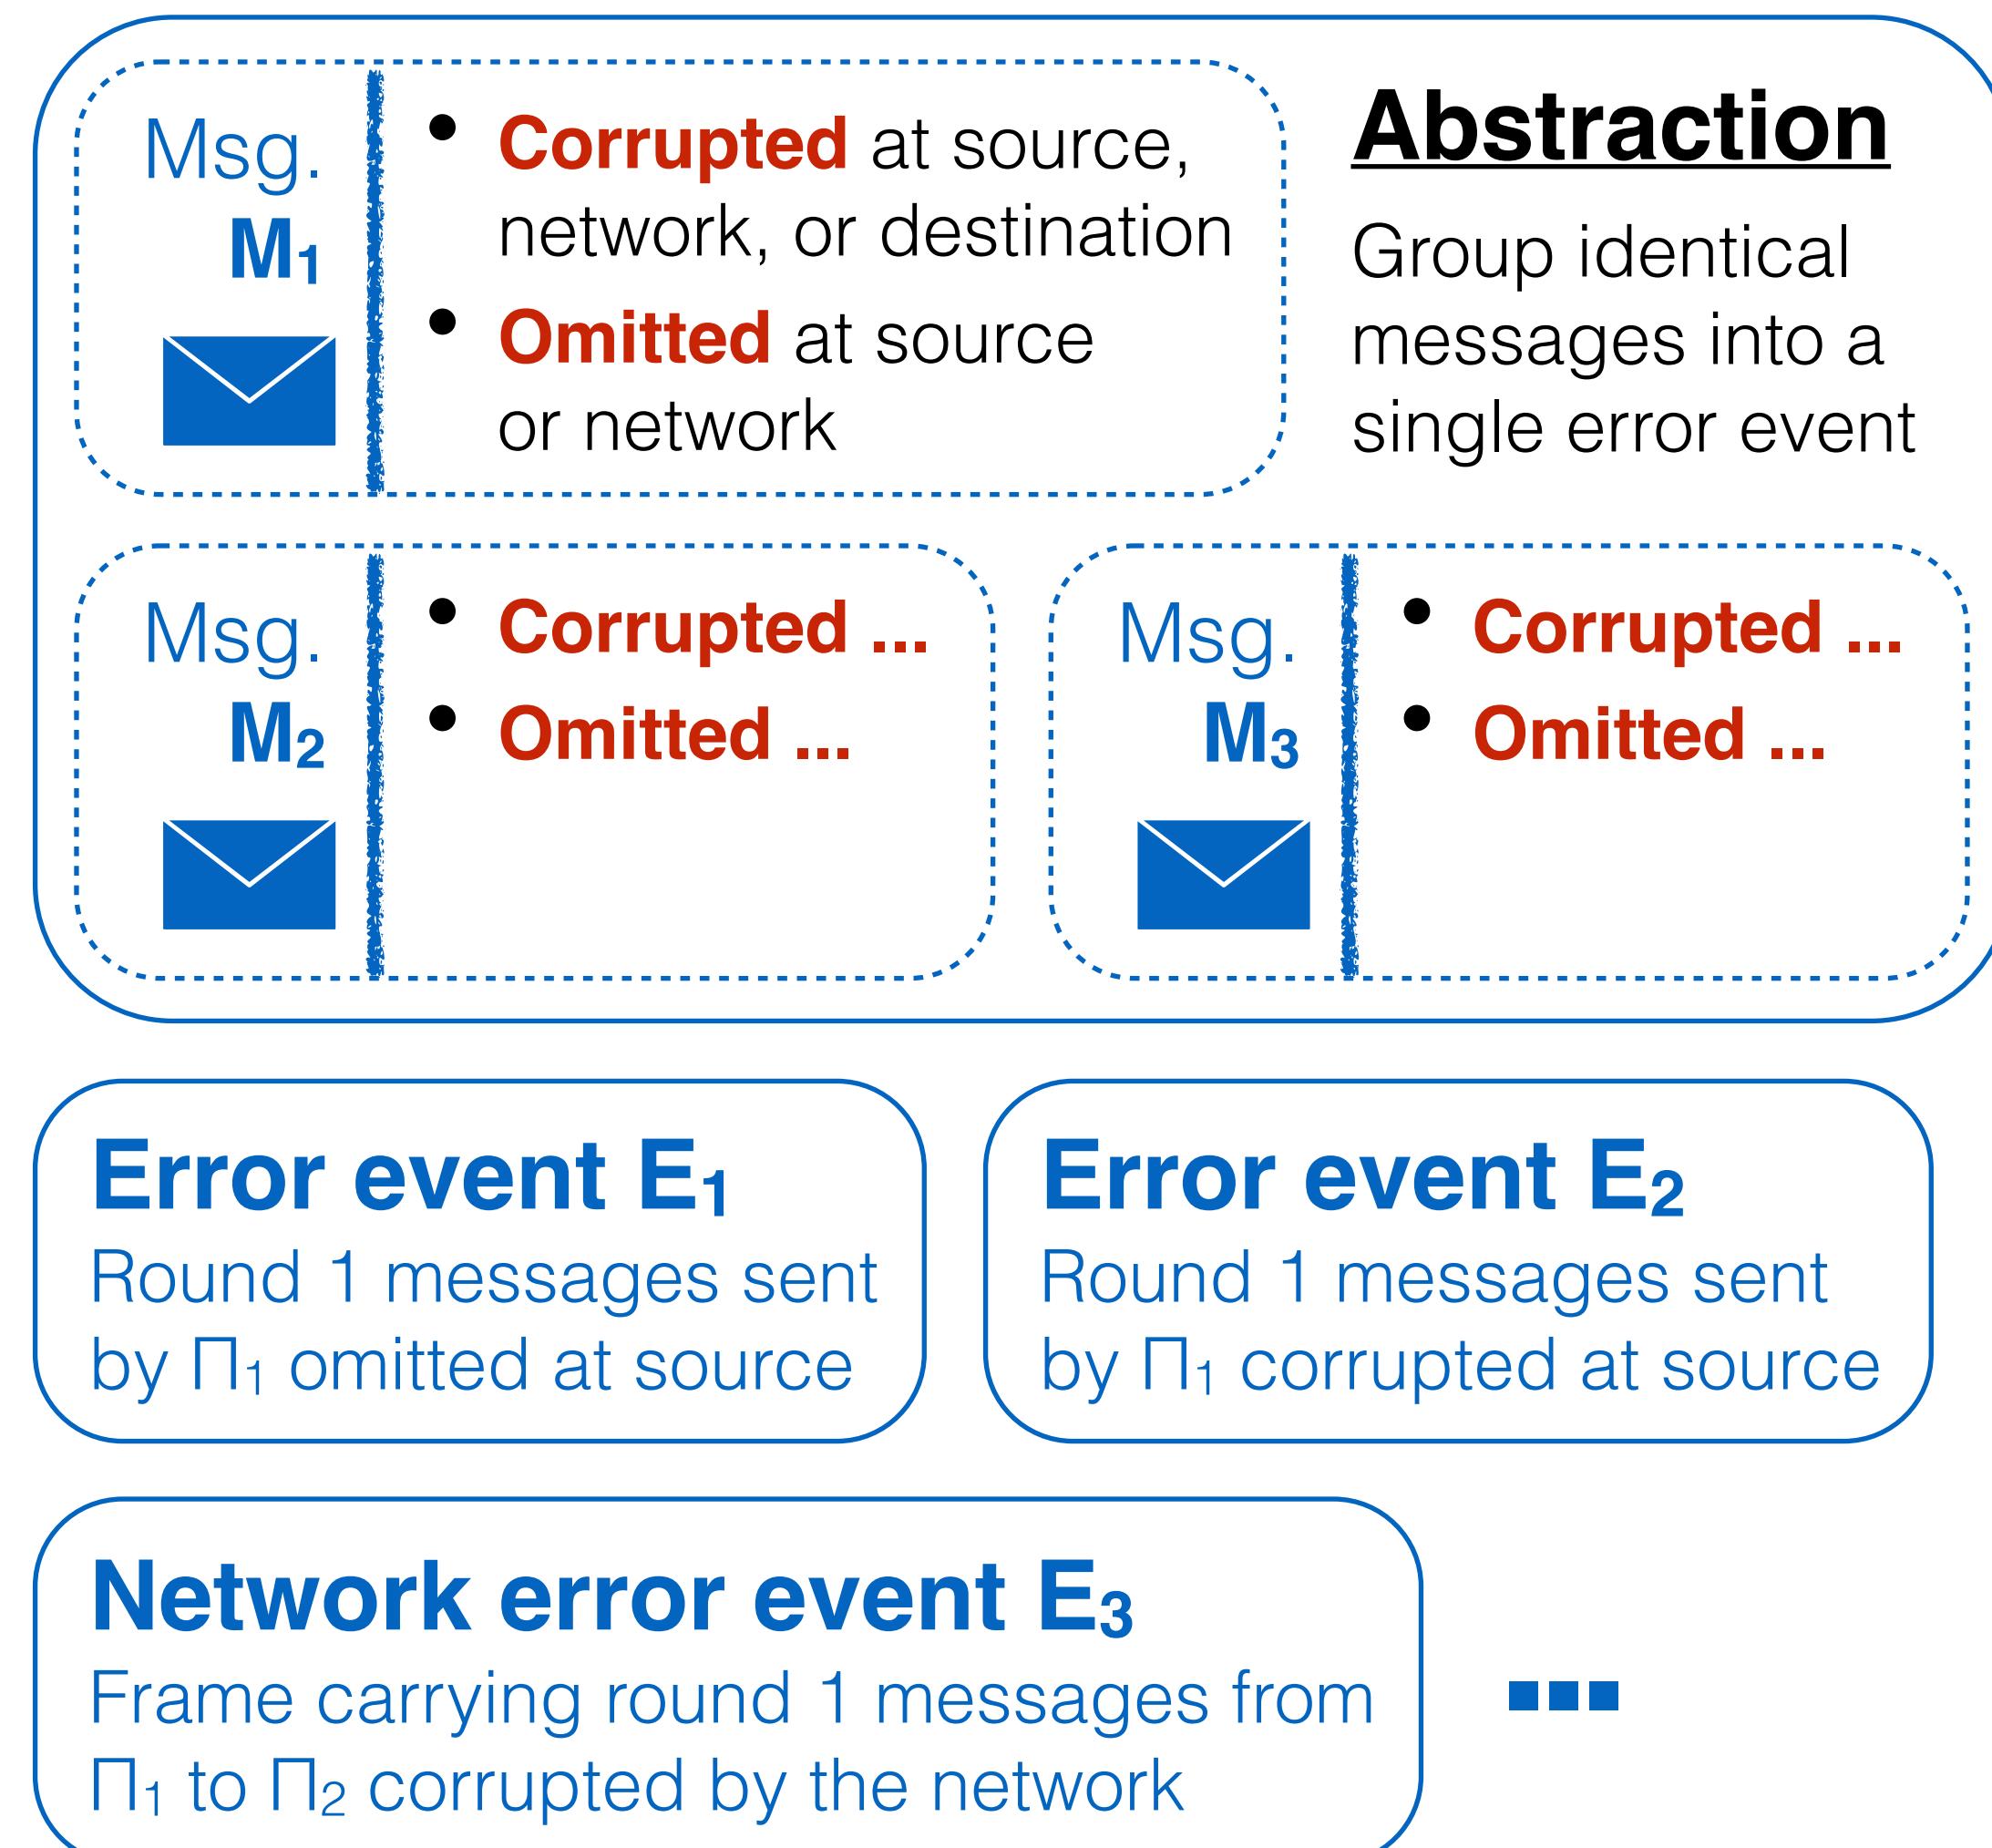
Δ_{link} = Link transfer time

J_a = Maximum scheduling jitter at Switch S_a

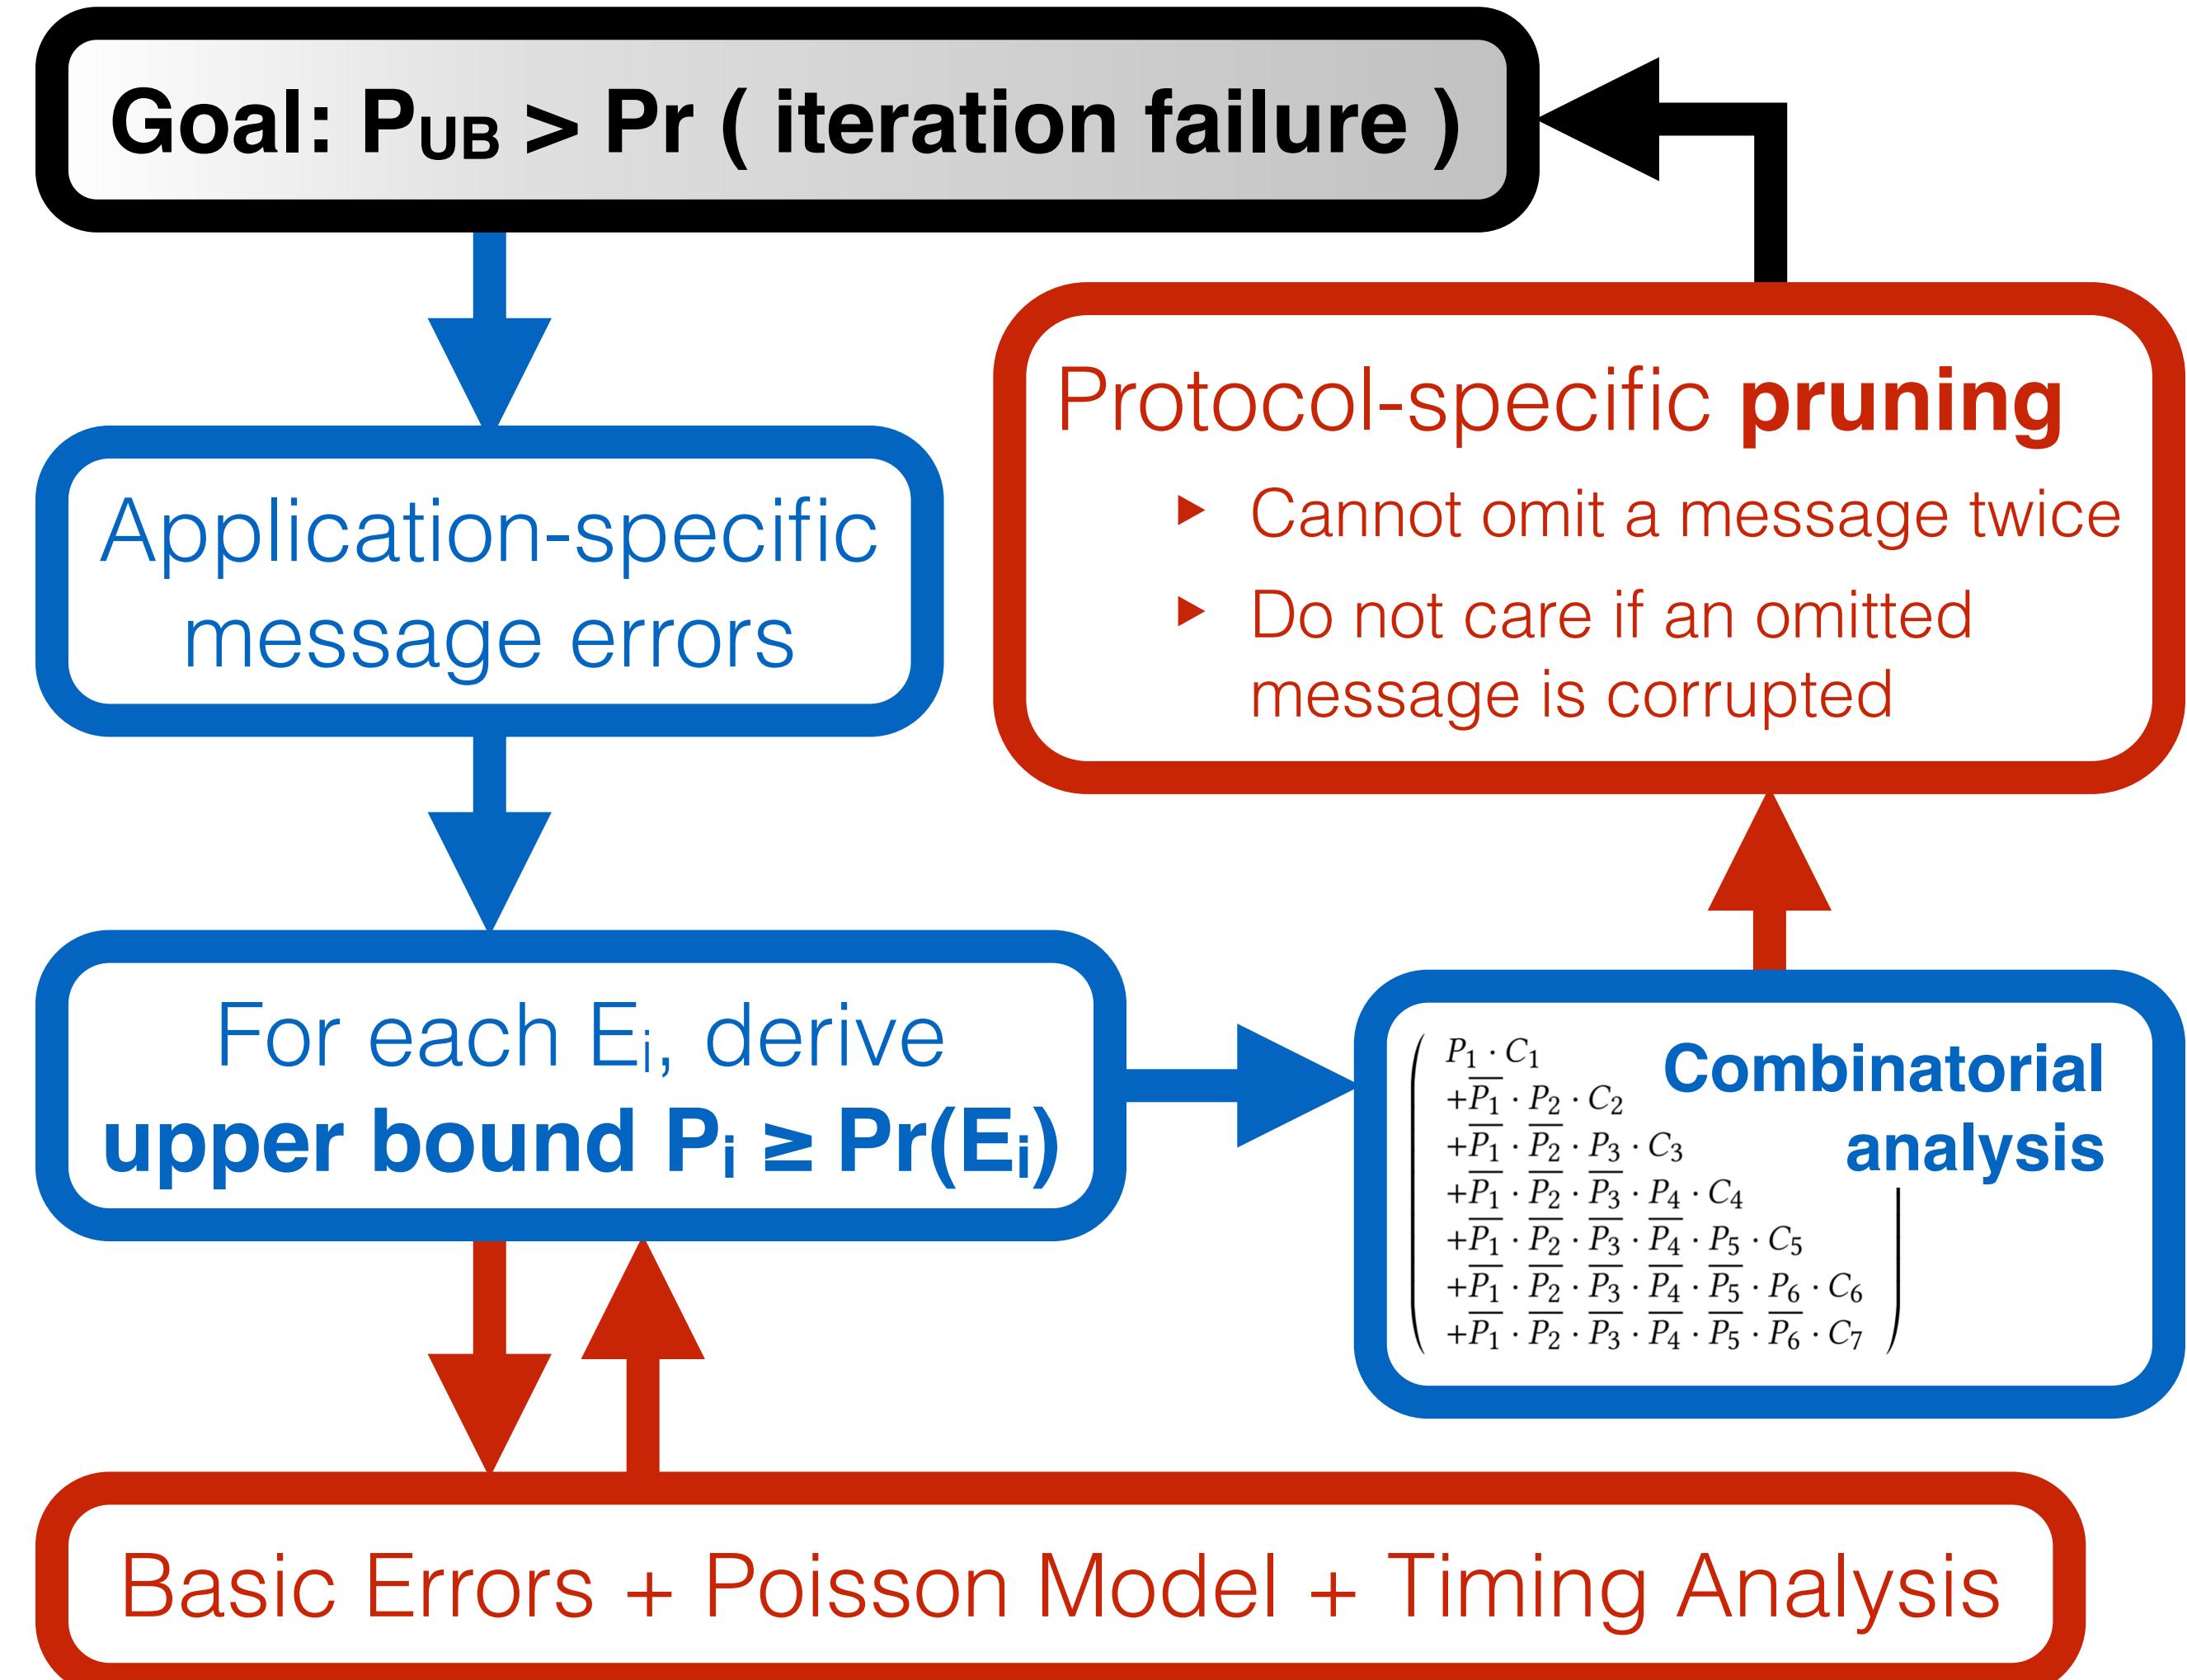
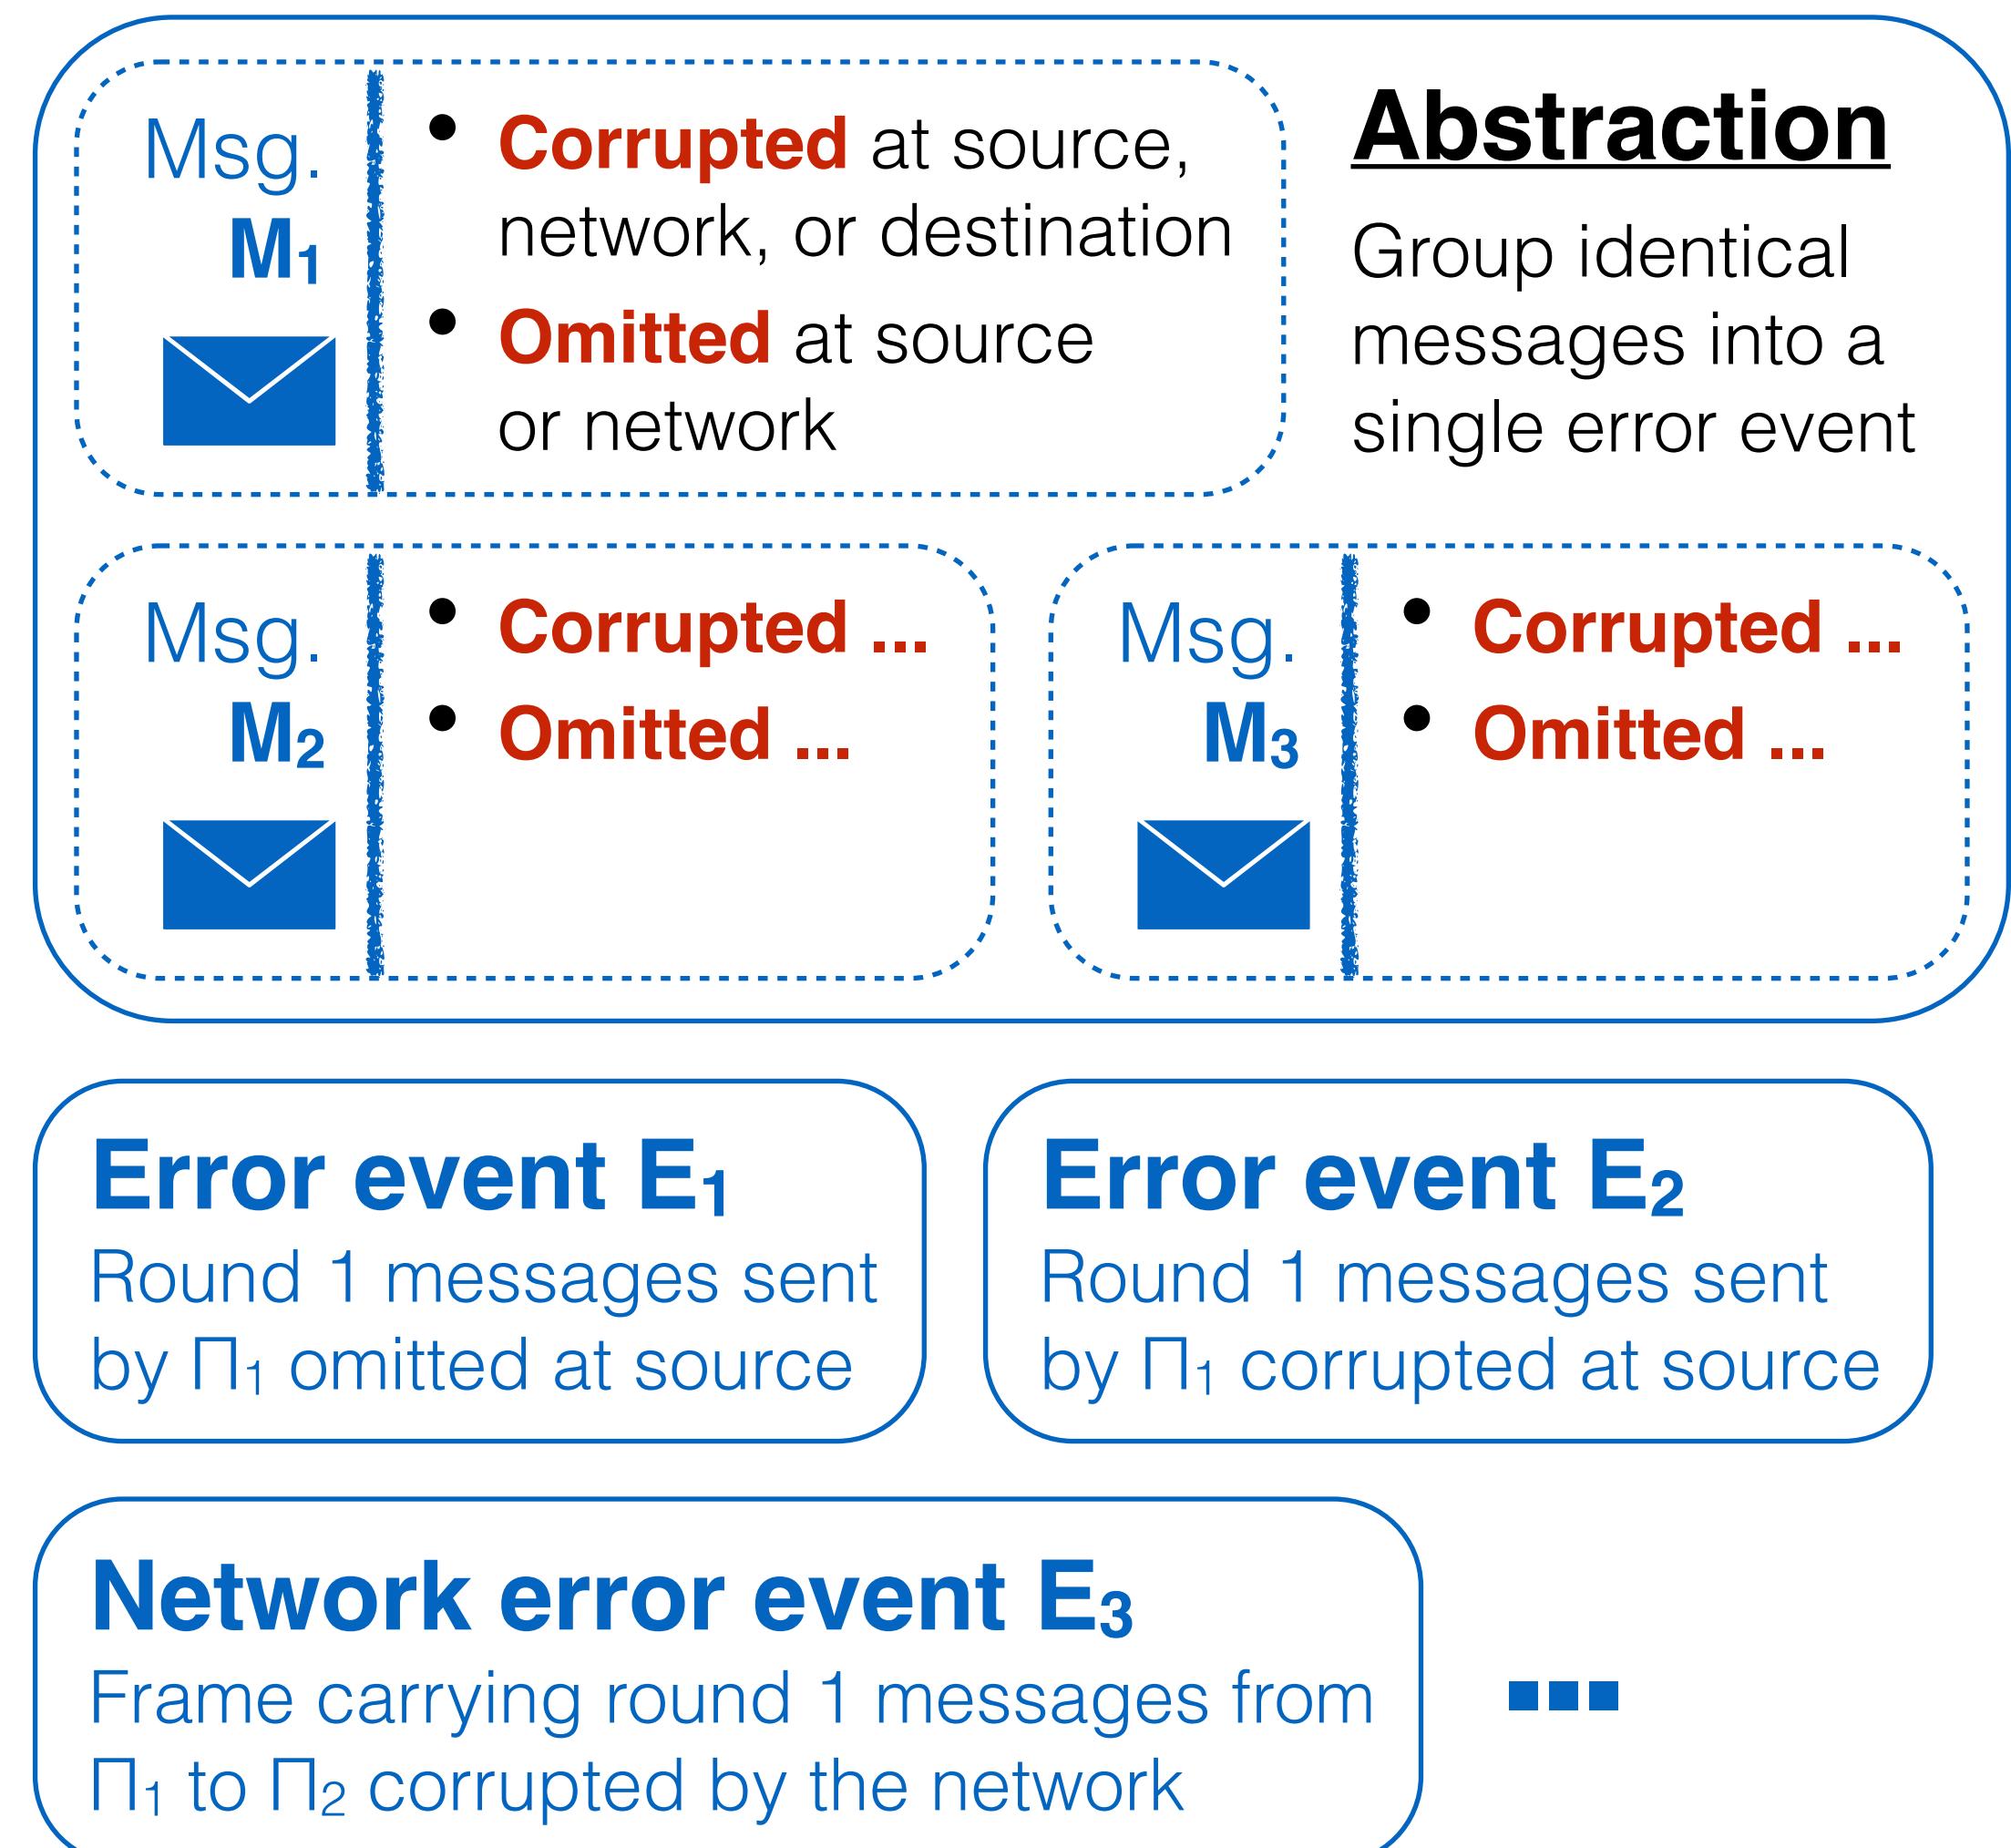
K_{\max} = Maximum bit flips detected by the CRC

Compute upper bound $P_3 \geq \Pr(E_3)$
by accounting for all scenarios

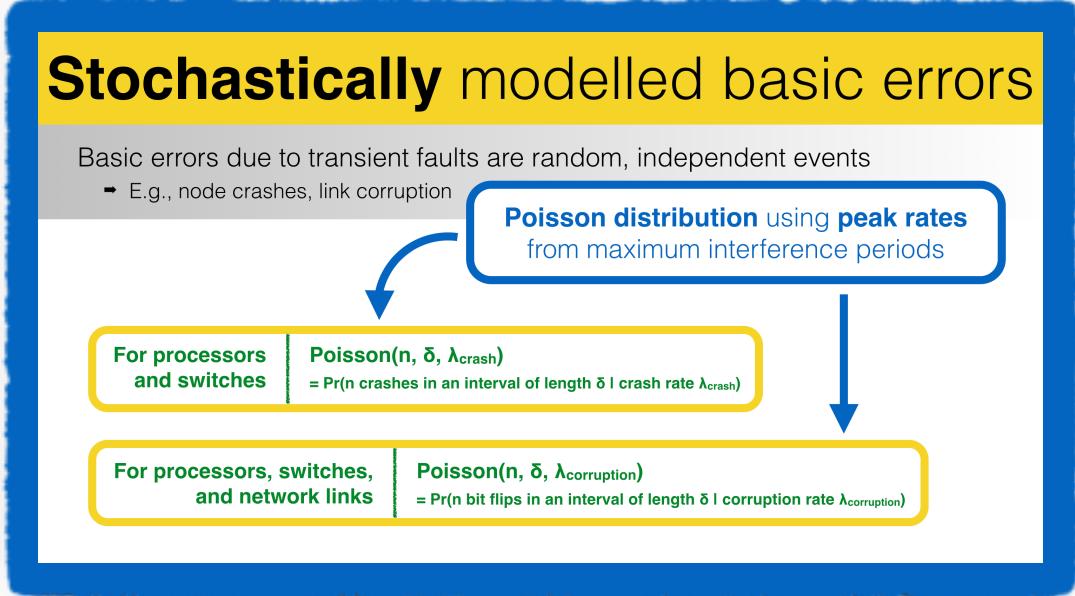
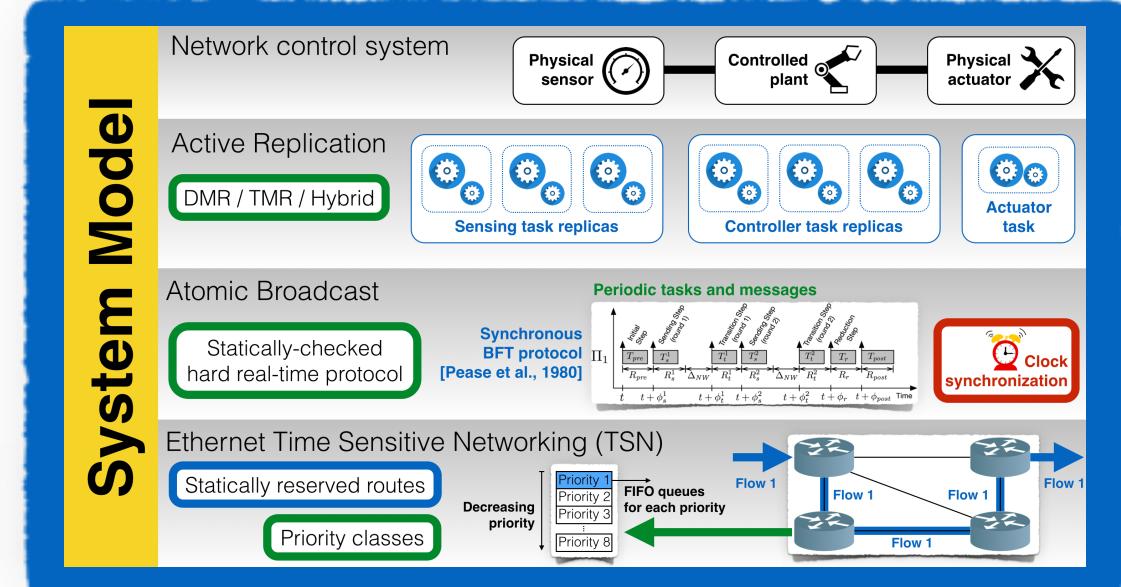
Key idea 1: Scalability through **abstraction** and **pruning**



Key idea 1: Scalability through **abstraction** and **pruning**



Key idea 1: Scalability through abstraction and pruning



Scalability challenges

Key idea 1: Tackle scalability through abstraction and pruning

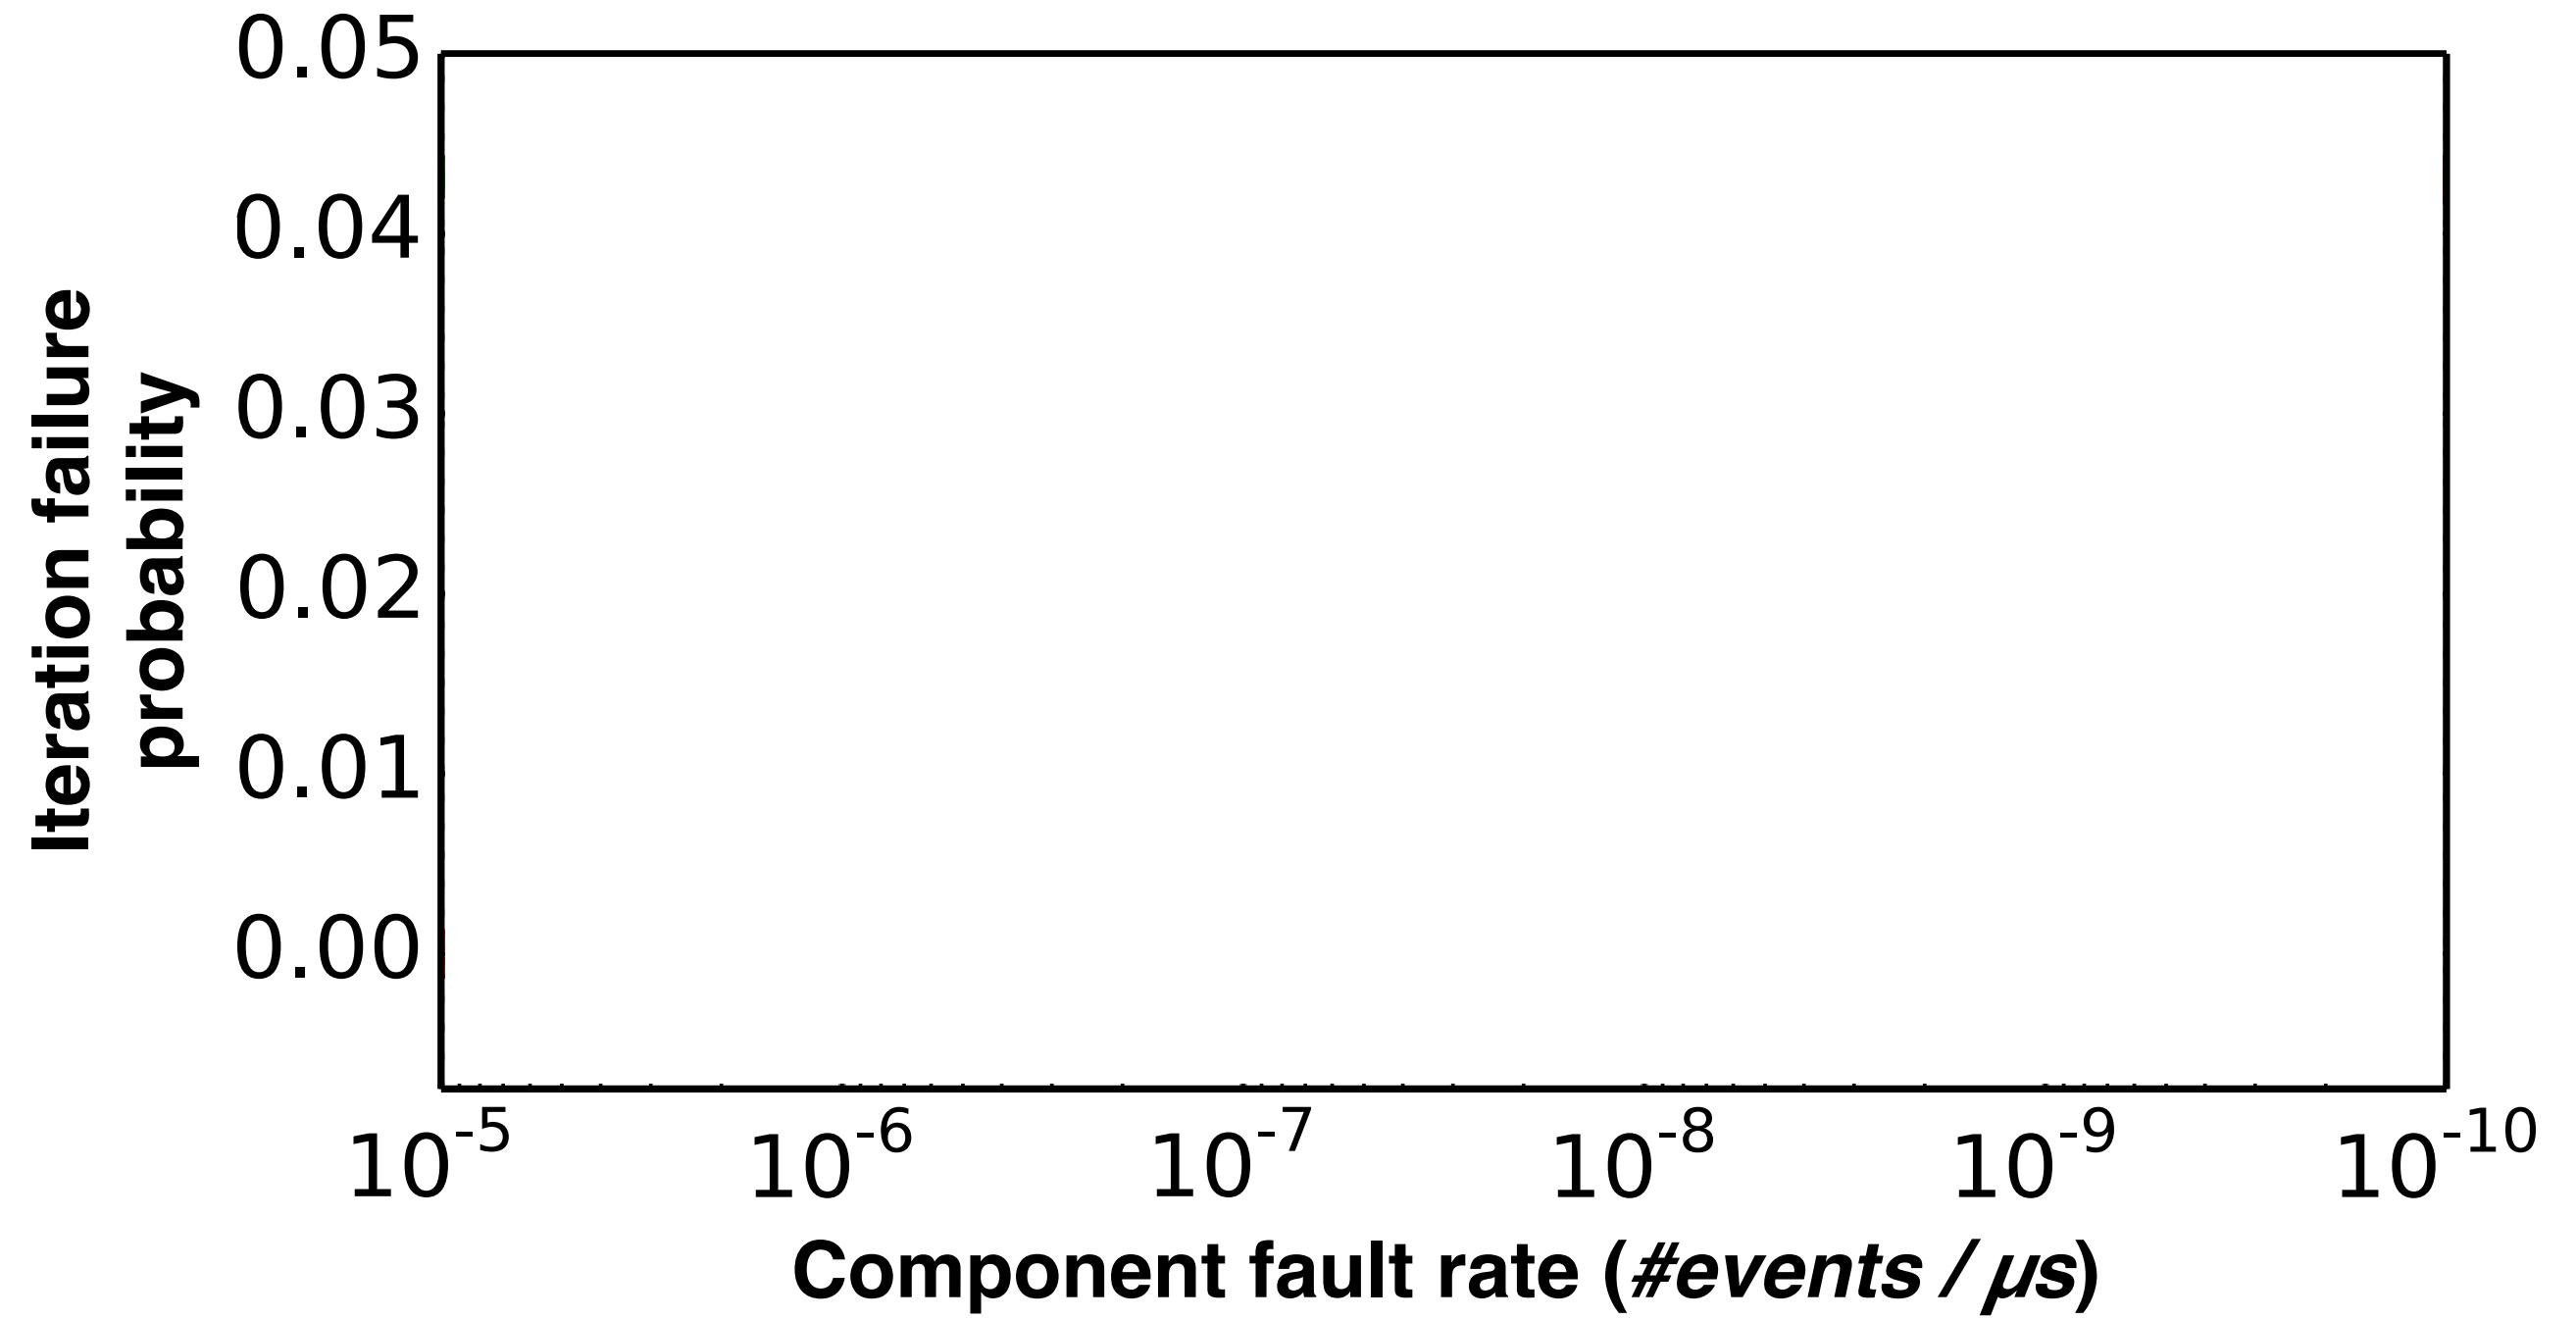
Model checking or simulation

$\Pr(\text{ iteration failure})$

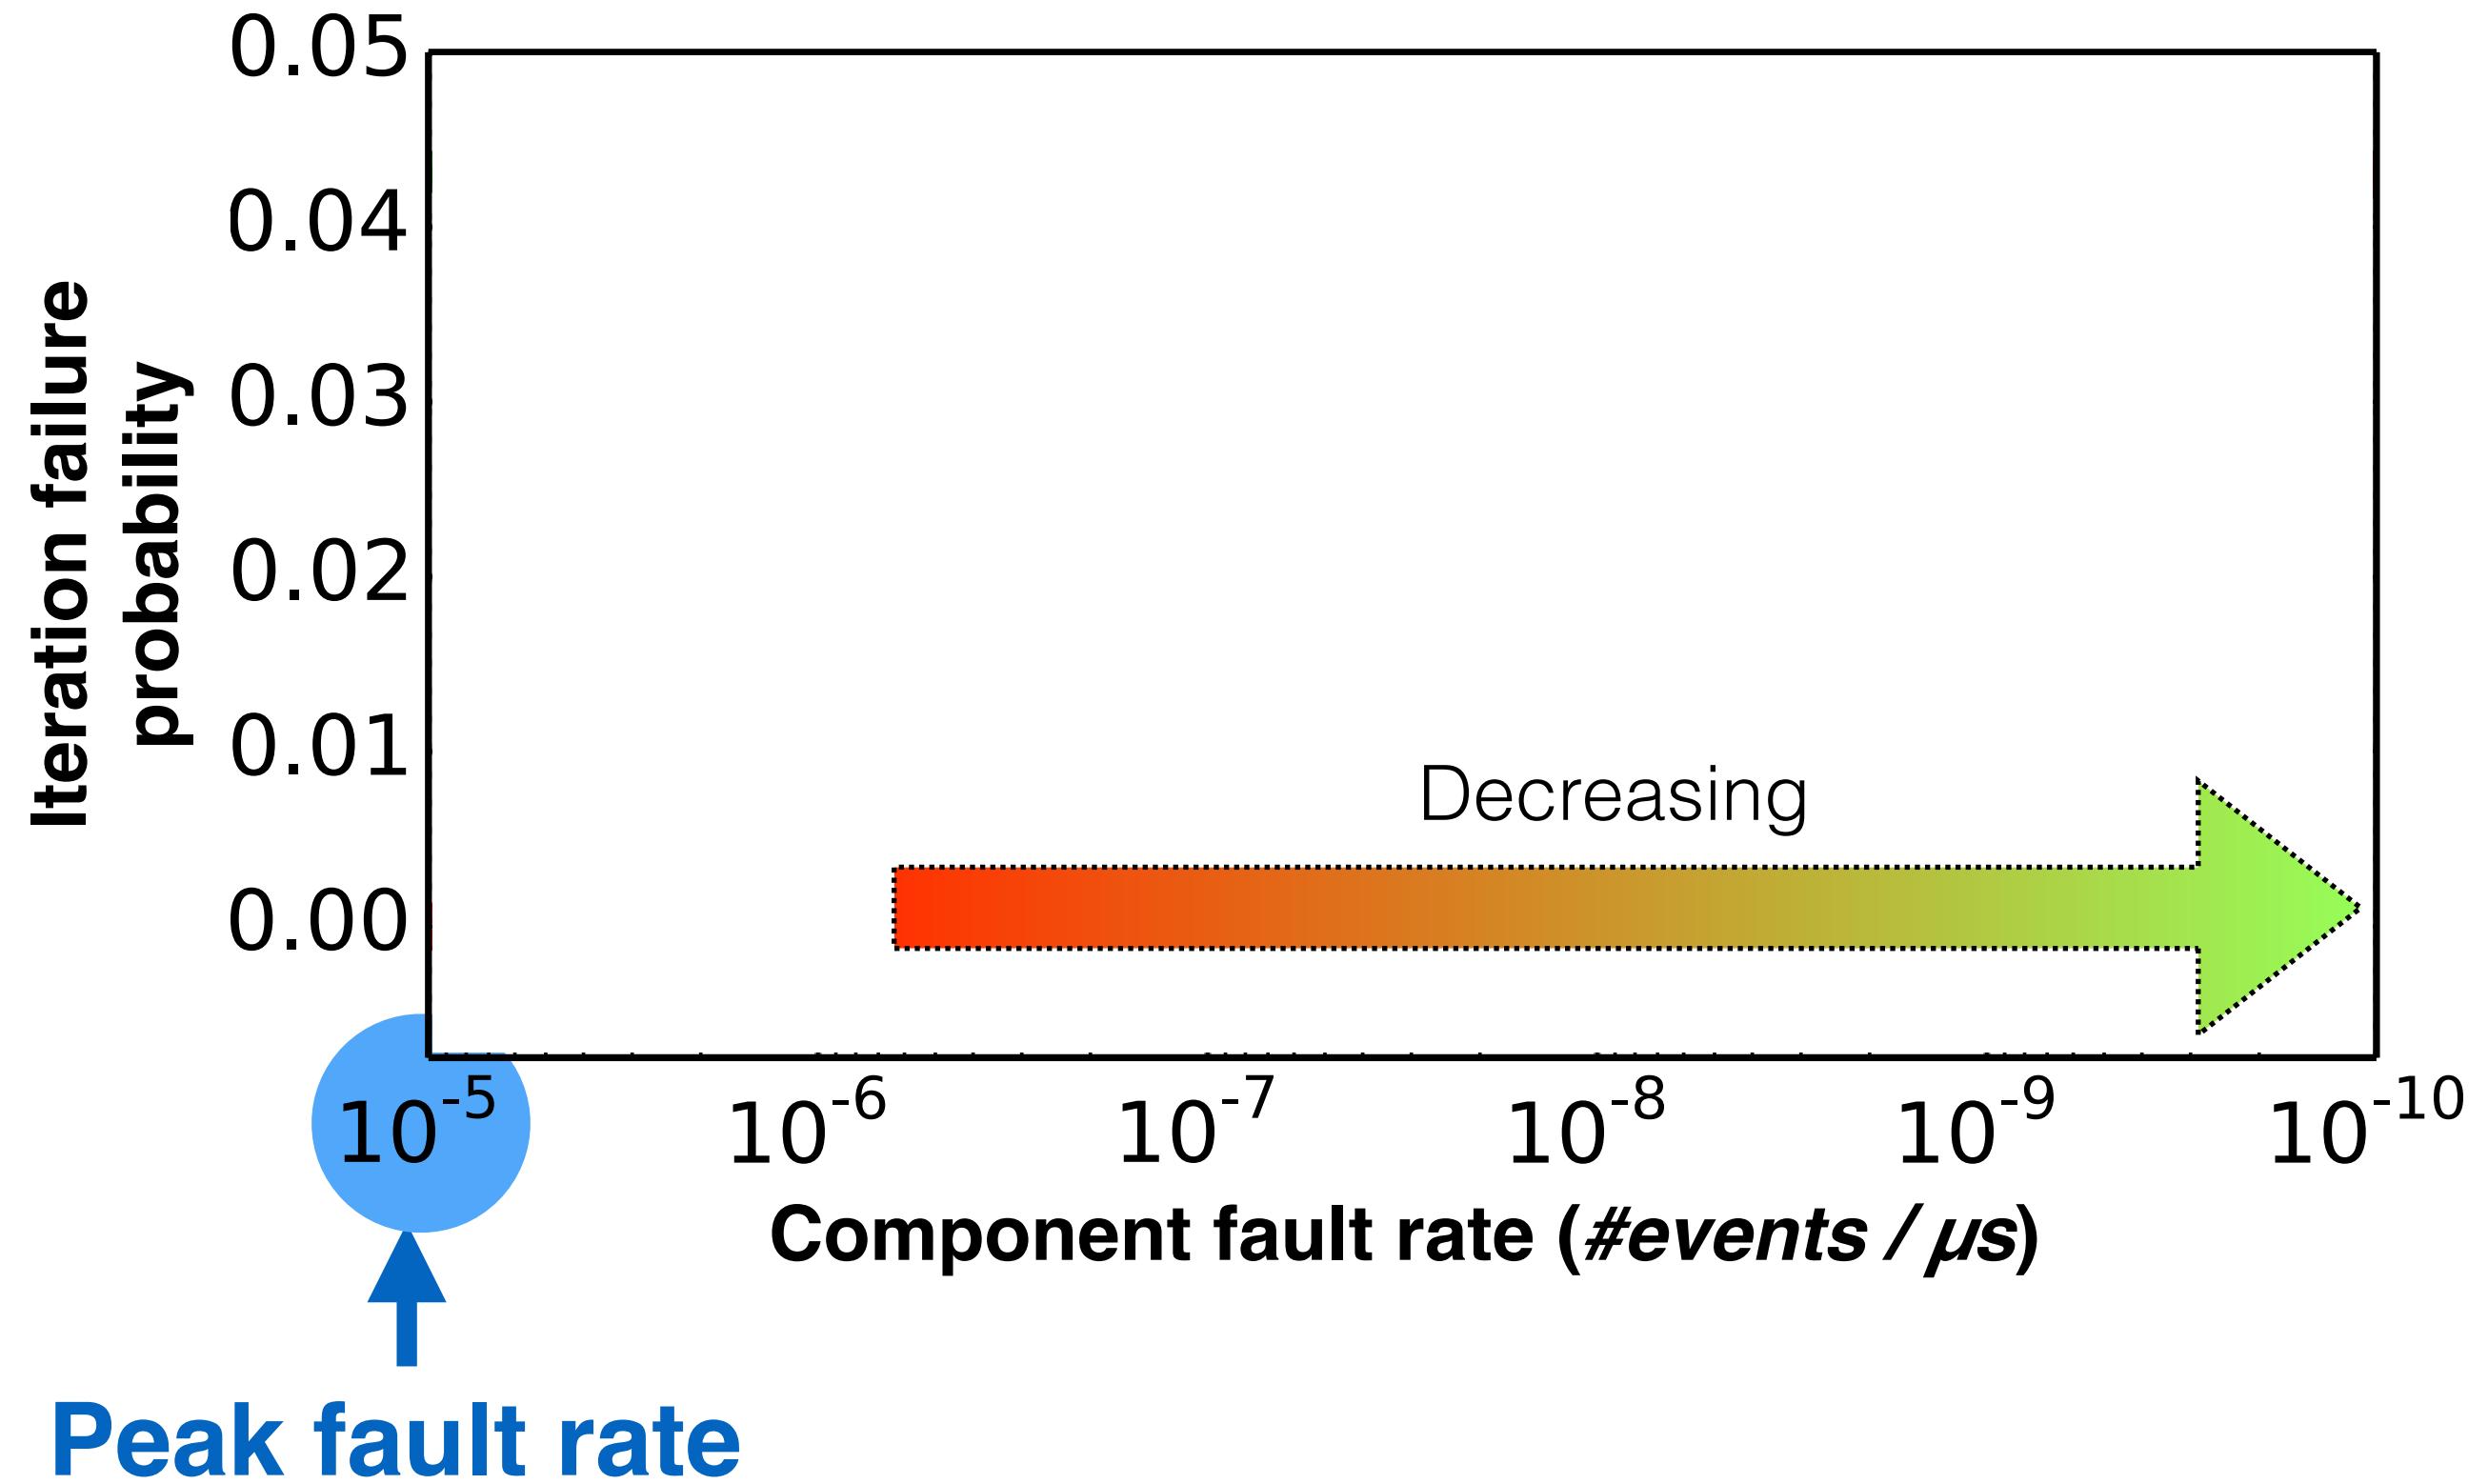
Reliability anomalies

- In practice, the iteration failure probability may **significantly exceed** the estimated $\Pr(\text{ iteration failure})$

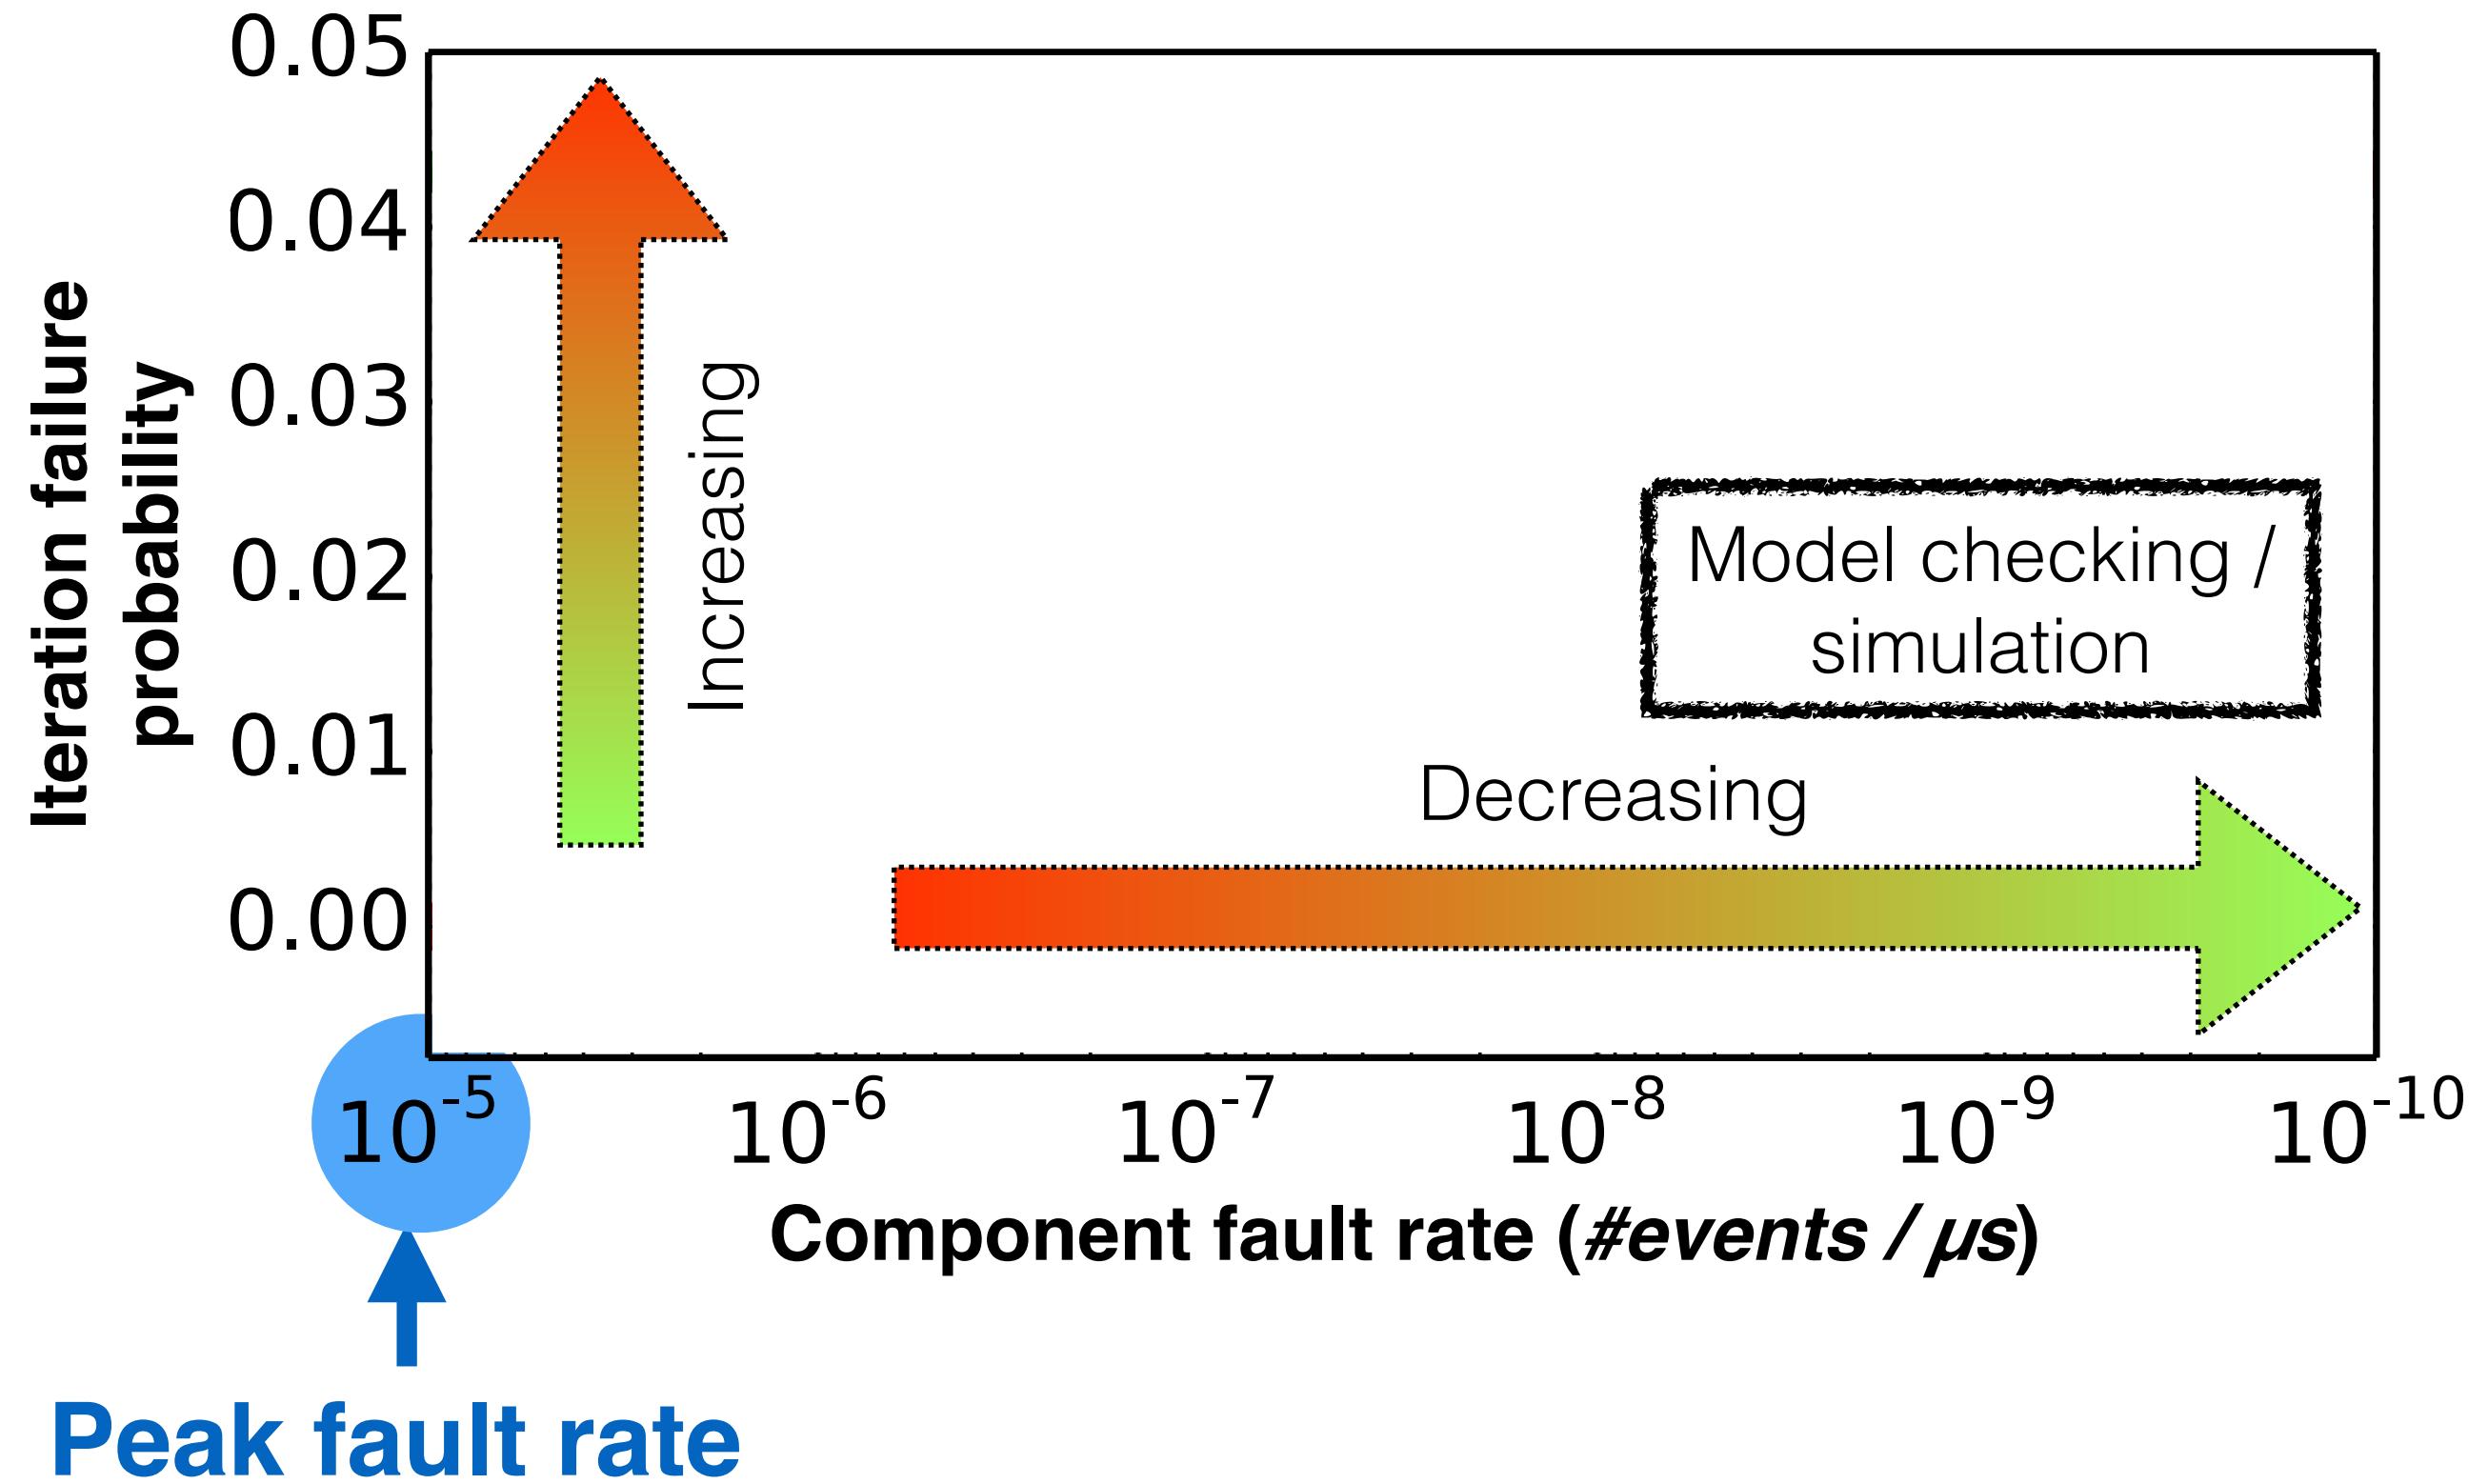
The problem of reliability anomalies



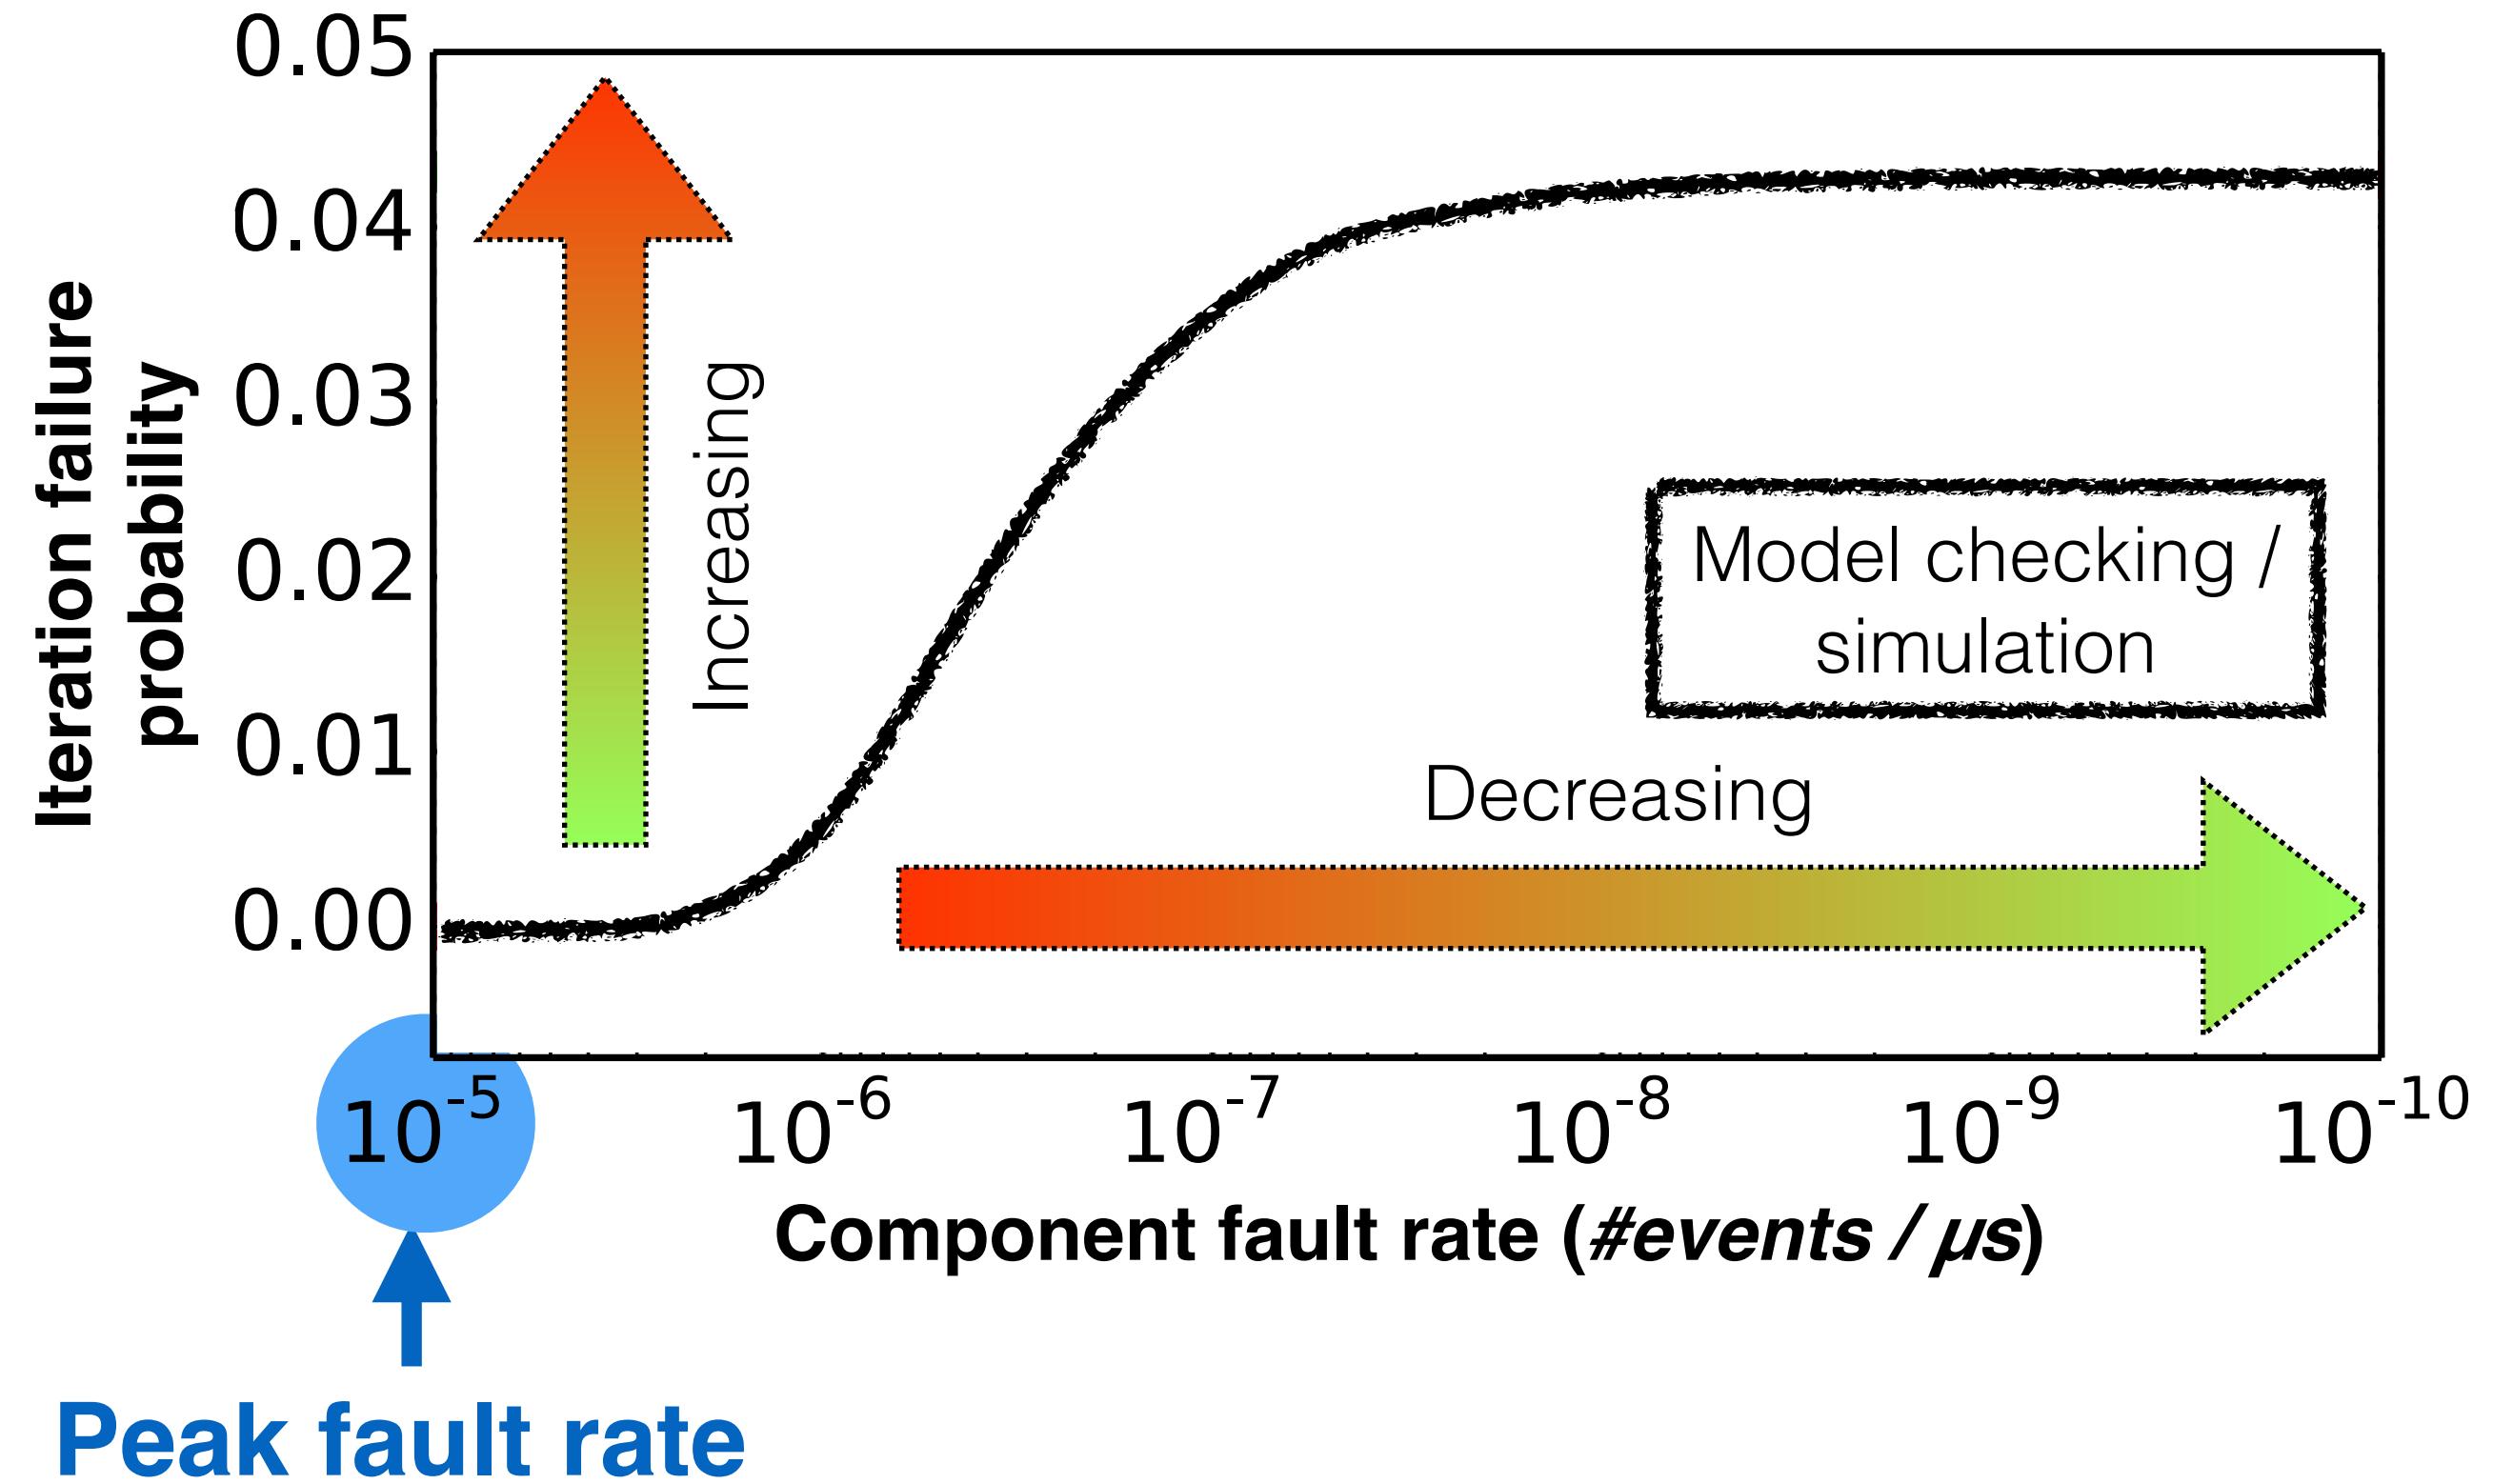
The problem of reliability anomalies



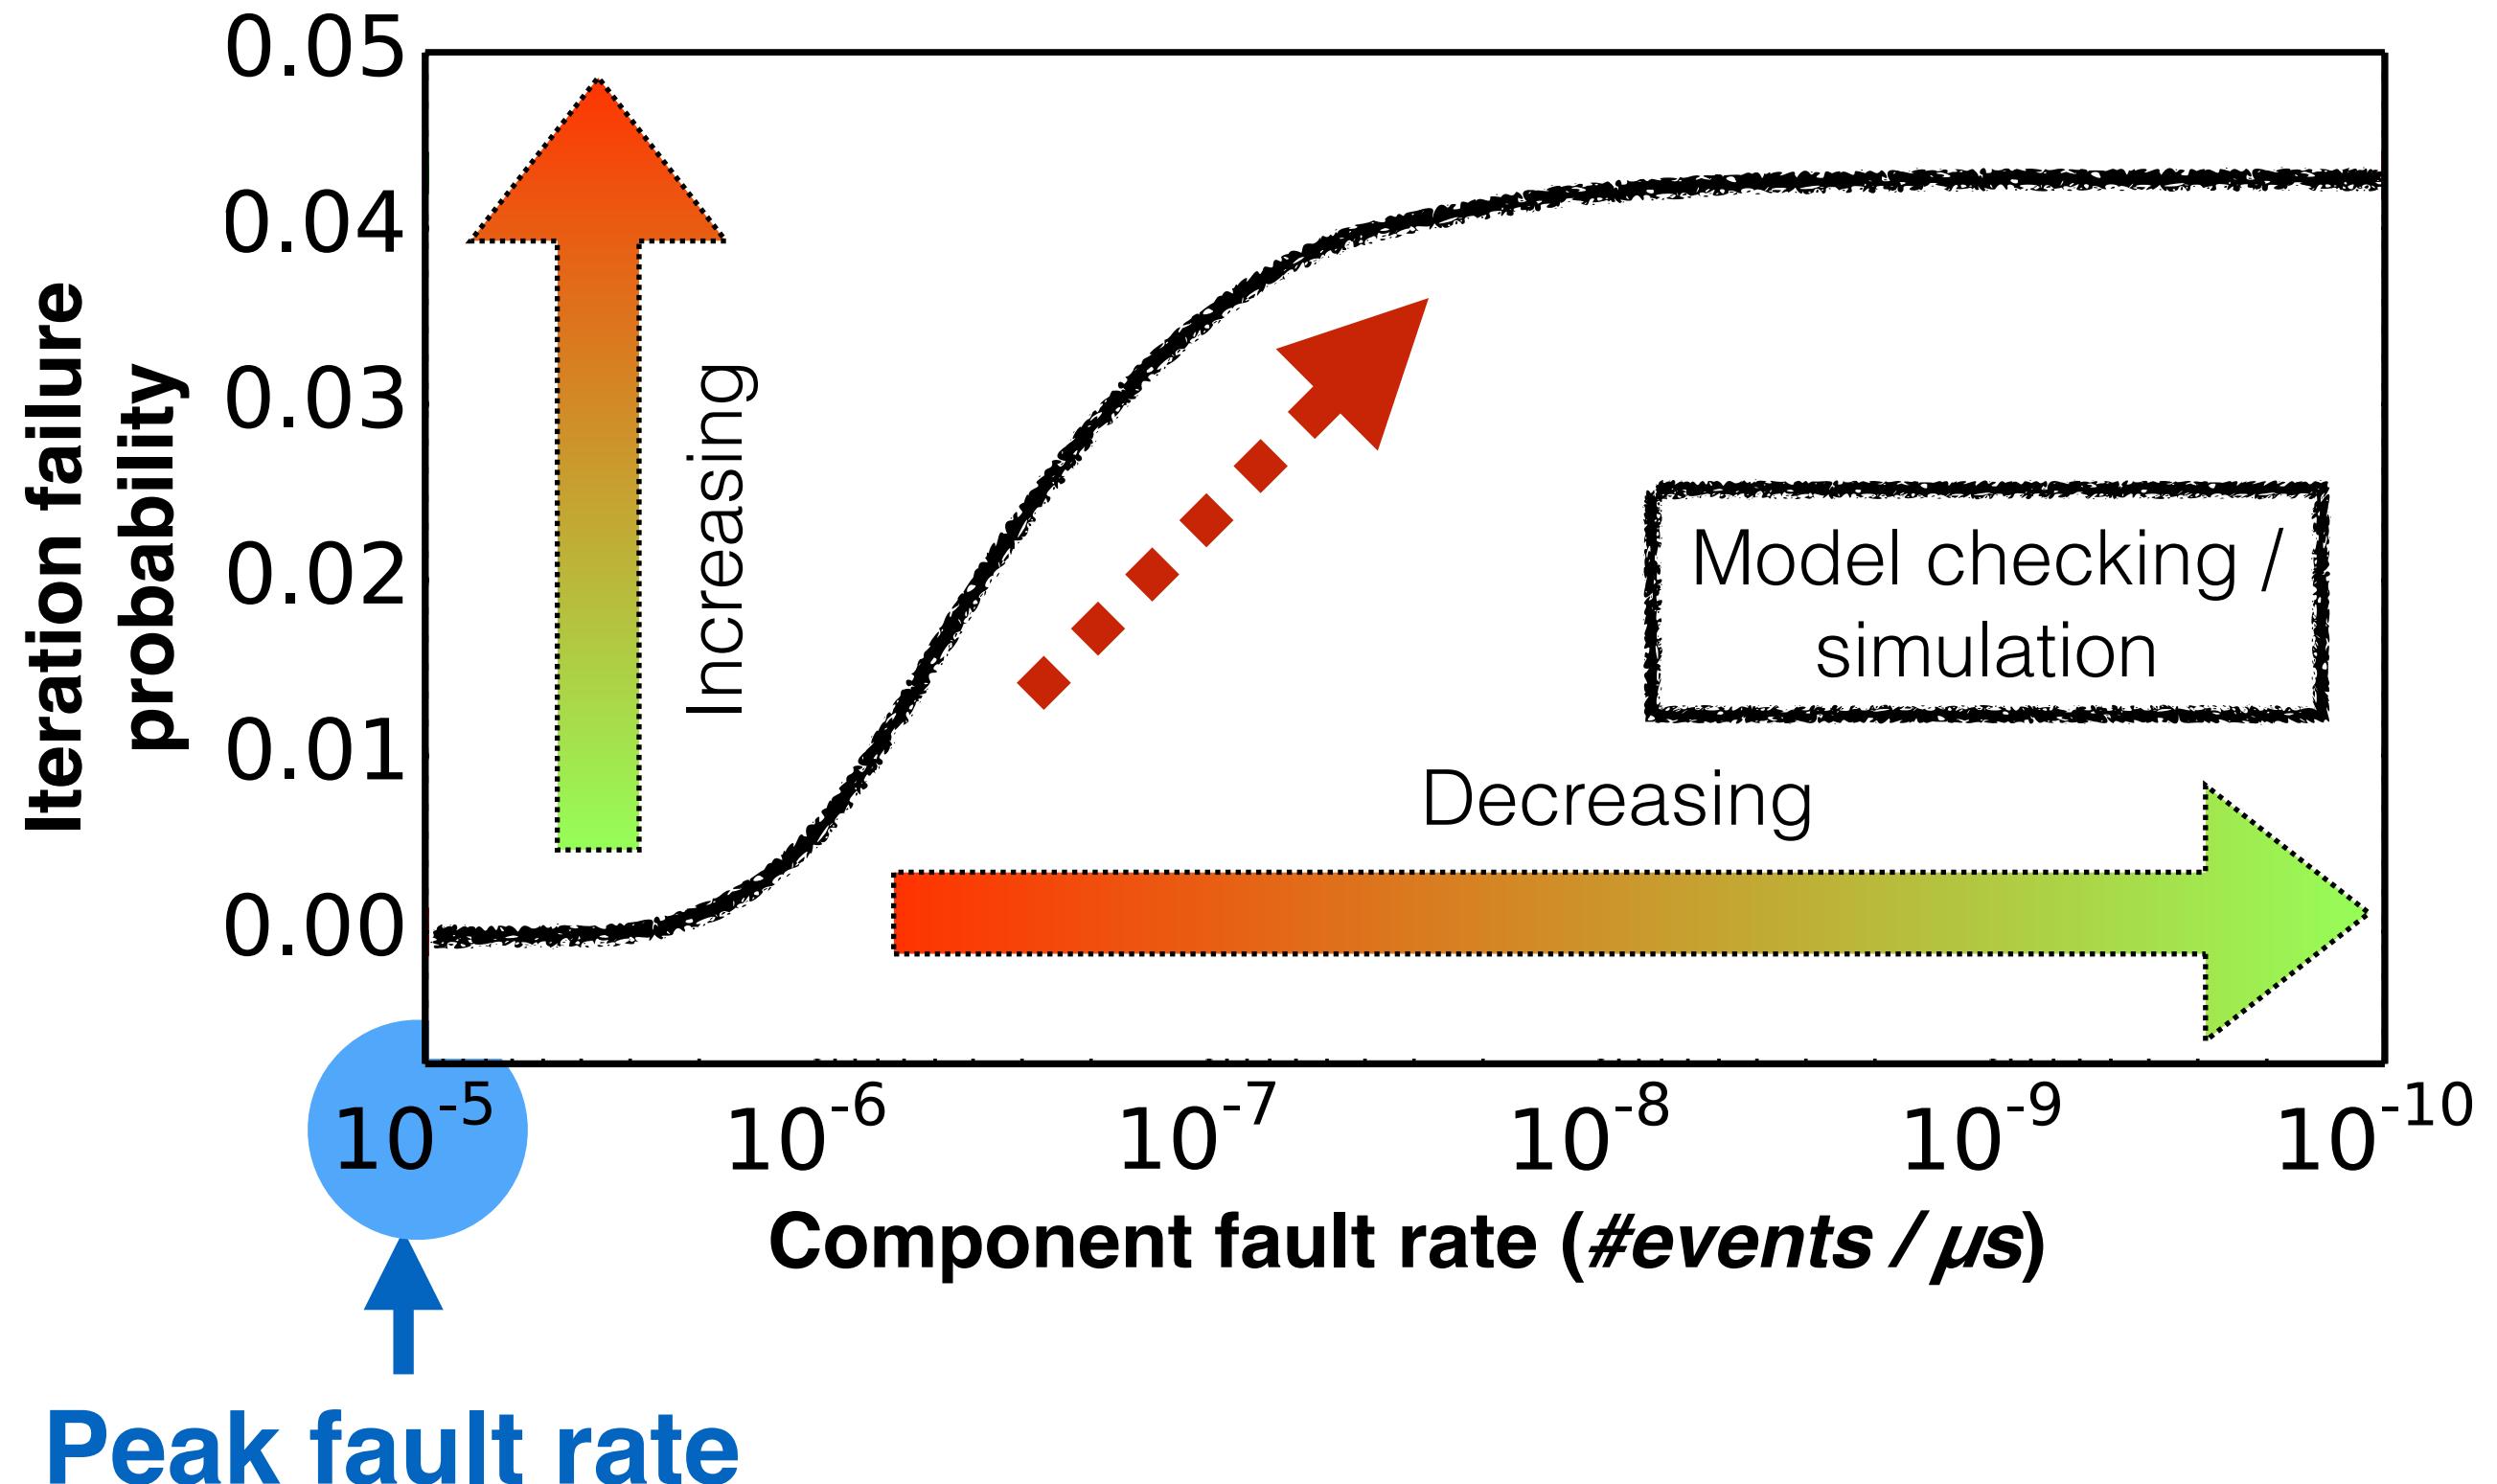
The problem of reliability anomalies



The problem of reliability anomalies



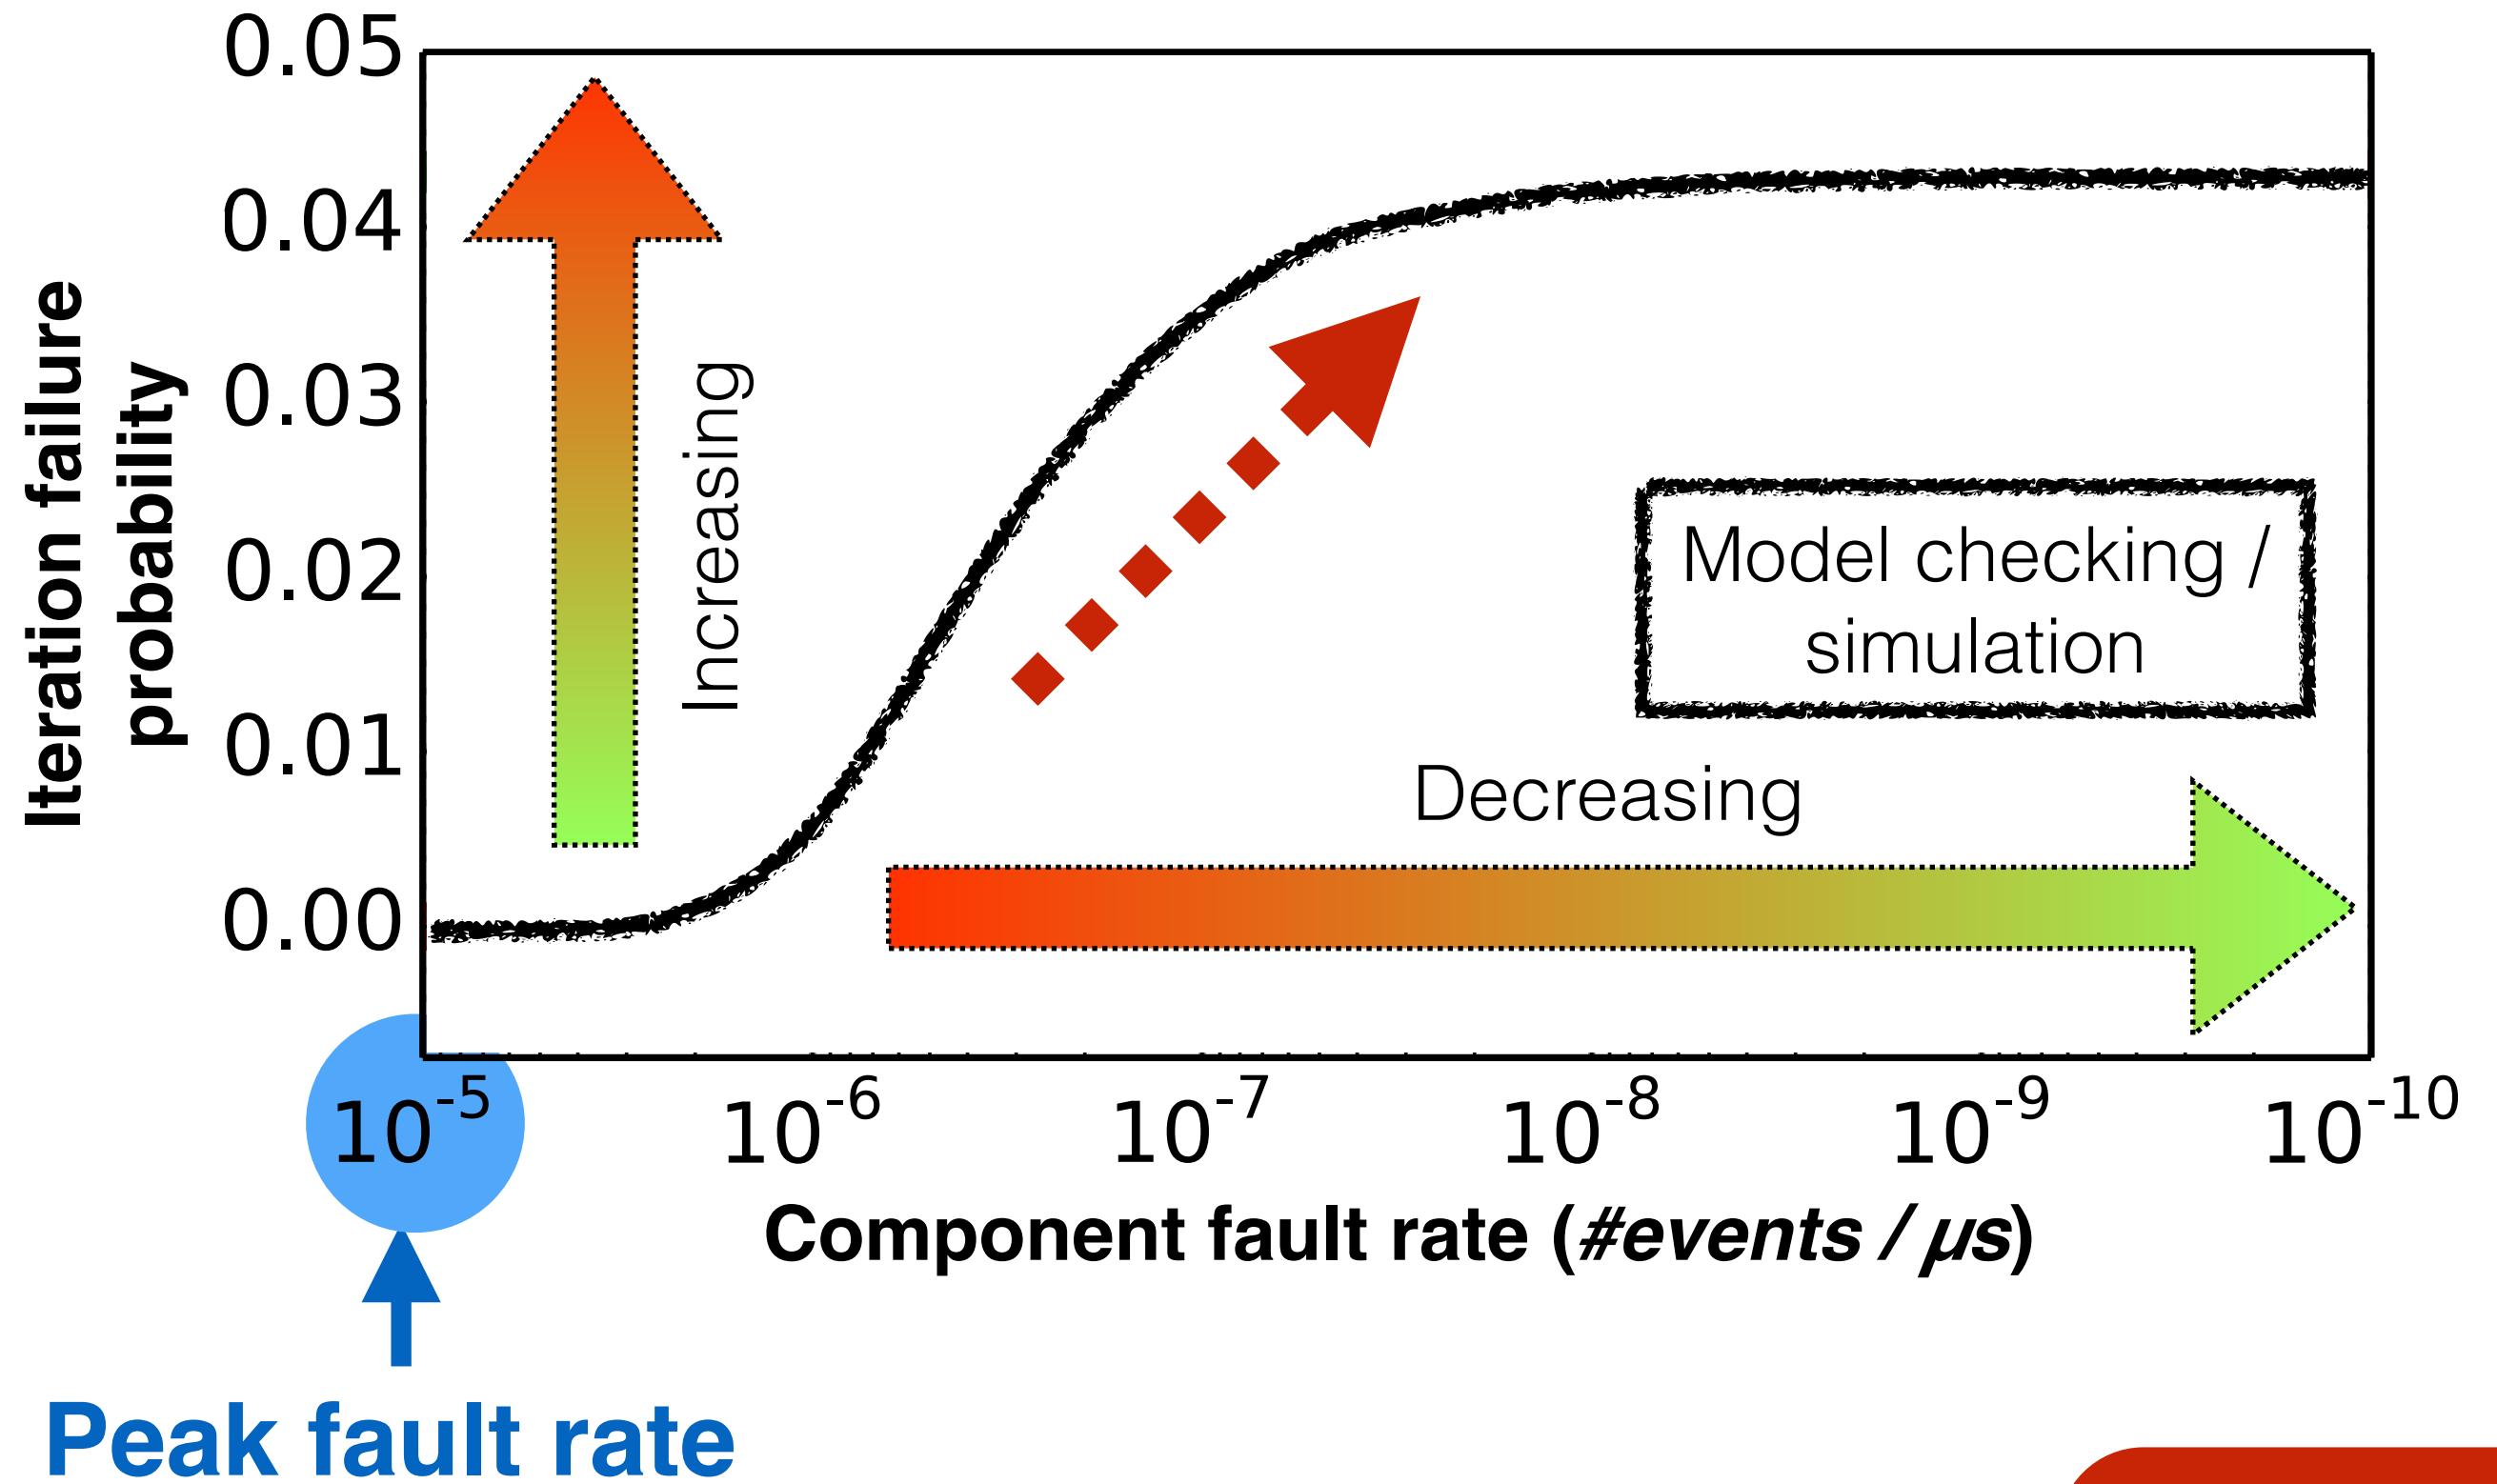
The problem of reliability anomalies



$\Pr(\text{iteration failure})$ increases despite decreasing component fault rate

Intuition: Sometimes, a node crash is good for the overall system, because it may reduce the probability of confusing a majority voting protocol in another part of the system!

The problem of reliability anomalies



$\Pr(\text{iteration failure})$ increases despite decreasing component fault rate

Intuition: Sometimes, a node crash is good for the overall system, because it may reduce the probability of confusing a majority voting protocol in another part of the system!



For soundness, need to estimate failure probabilities for the **entire search space** $[0, 10^{-5}]$

Key idea 2: Ensure **monotonicity** to eliminate anomalies

Combinatorial analysis

$$P_{UB} = \left(\begin{array}{l} P_1 \cdot C_1 \\ + \overline{P_1} \cdot P_2 \cdot C_2 \\ + \overline{P_1} \cdot \overline{P_2} \cdot P_3 \cdot C_3 \\ + \overline{P_1} \cdot \overline{P_2} \cdot \overline{P_3} \cdot P_4 \cdot C_4 \\ + \overline{P_1} \cdot \overline{P_2} \cdot \overline{P_3} \cdot \overline{P_4} \cdot P_5 \cdot C_5 \\ + \overline{P_1} \cdot \overline{P_2} \cdot \overline{P_3} \cdot \overline{P_4} \cdot \overline{P_5} \cdot P_6 \cdot C_6 \\ + \overline{P_1} \cdot \overline{P_2} \cdot \overline{P_3} \cdot \overline{P_4} \cdot \overline{P_5} \cdot \overline{P_6} \cdot C_7 \end{array} \right)$$

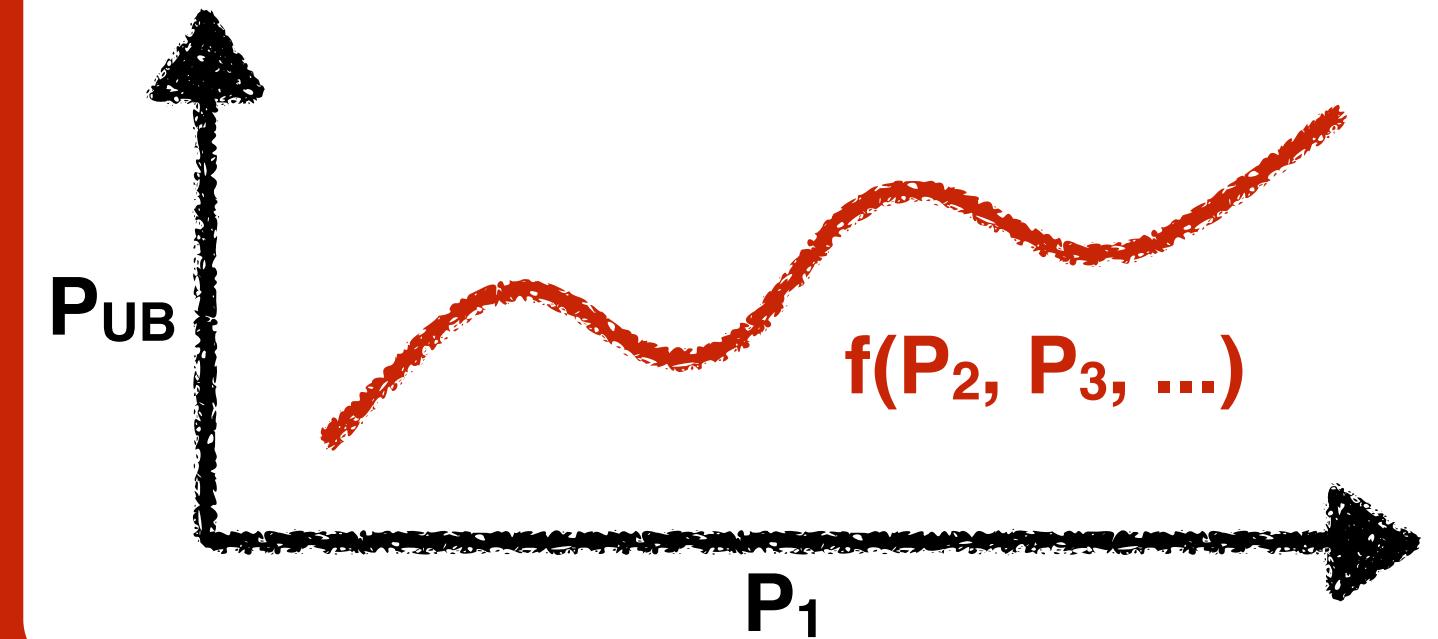
Key idea 2: Ensure **monotonicity** to eliminate anomalies

Combinatorial analysis

$$P_{UB} = \left(P_1 \cdot C_1 + \overline{P_1} \cdot P_2 \cdot C_2 + \overline{P_1} \cdot \overline{P_2} \cdot P_3 \cdot C_3 + \overline{P_1} \cdot \overline{P_2} \cdot \overline{P_3} \cdot P_4 \cdot C_4 + \overline{P_1} \cdot \overline{P_2} \cdot \overline{P_3} \cdot \overline{P_4} \cdot P_5 \cdot C_5 + \overline{P_1} \cdot \overline{P_2} \cdot \overline{P_3} \cdot \overline{P_4} \cdot \overline{P_5} \cdot P_6 \cdot C_6 + \overline{P_1} \cdot \overline{P_2} \cdot \overline{P_3} \cdot \overline{P_4} \cdot \overline{P_5} \cdot \overline{P_6} \cdot C_7 \right)$$

Root cause of reliability anomalies

P_{UB} may not monotonically increase with P₁, P₂, P₃, ...



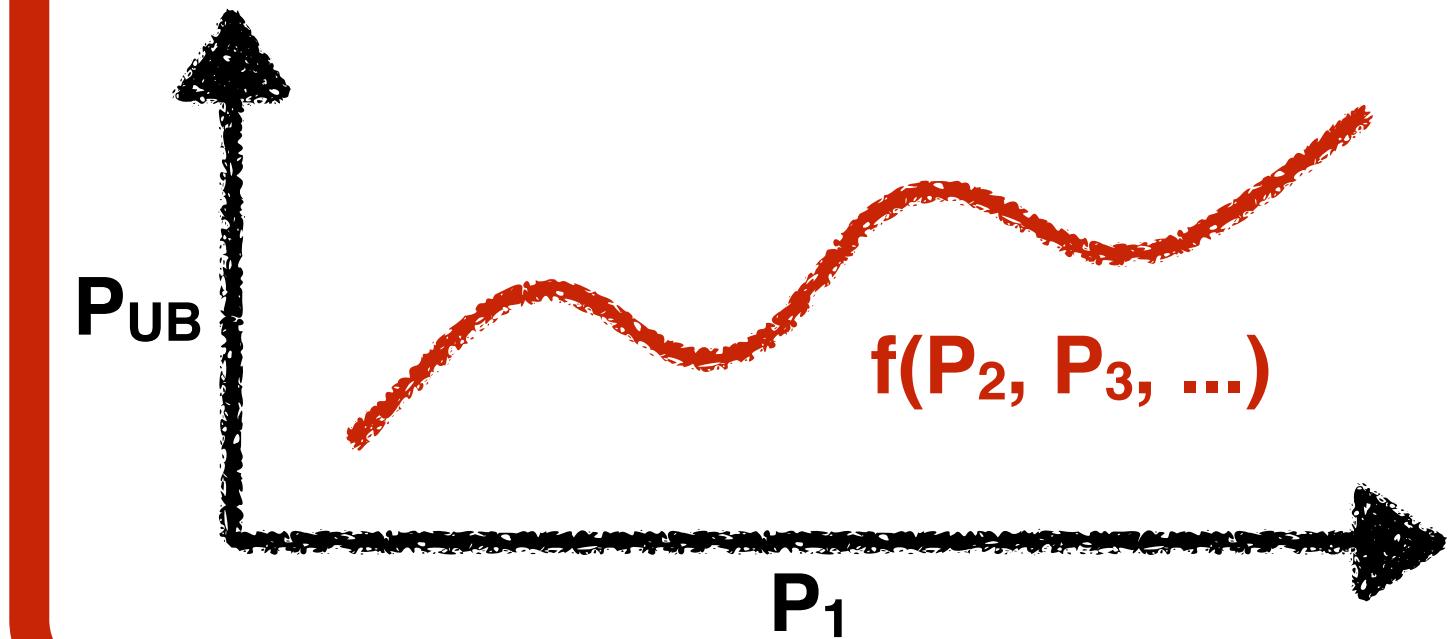
Key idea 2: Ensure **monotonicity** to eliminate anomalies

Combinatorial analysis

$$P_{UB} = \left(P_1 \cdot C_1 + \overline{P_1} \cdot P_2 \cdot C_2 + \overline{P_1} \cdot \overline{P_2} \cdot P_3 \cdot C_3 + \overline{P_1} \cdot \overline{P_2} \cdot \overline{P_3} \cdot P_4 \cdot C_4 + \overline{P_1} \cdot \overline{P_2} \cdot \overline{P_3} \cdot \overline{P_4} \cdot P_5 \cdot C_5 + \overline{P_1} \cdot \overline{P_2} \cdot \overline{P_3} \cdot \overline{P_4} \cdot \overline{P_5} \cdot P_6 \cdot C_6 + \overline{P_1} \cdot \overline{P_2} \cdot \overline{P_3} \cdot \overline{P_4} \cdot \overline{P_5} \cdot \overline{P_6} \cdot C_7 \right)$$

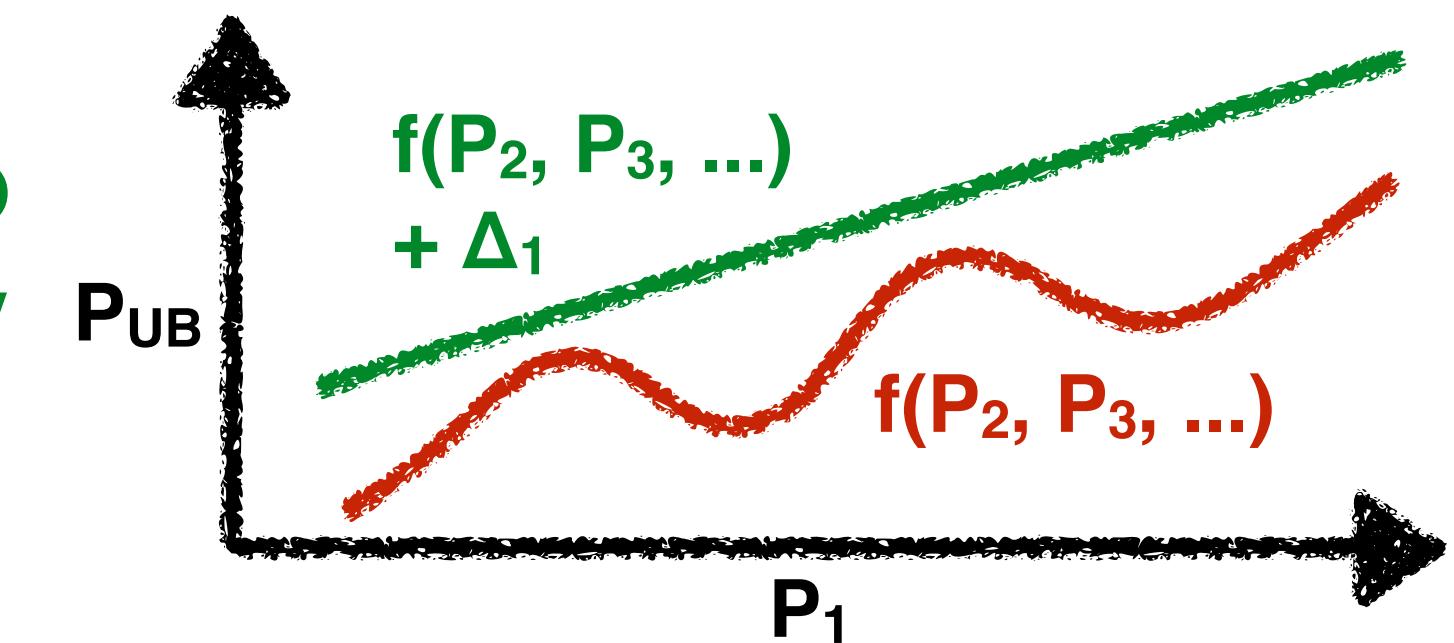
Root cause of reliability anomalies

P_{UB} may not monotonically increase with P₁, P₂, P₃, ...

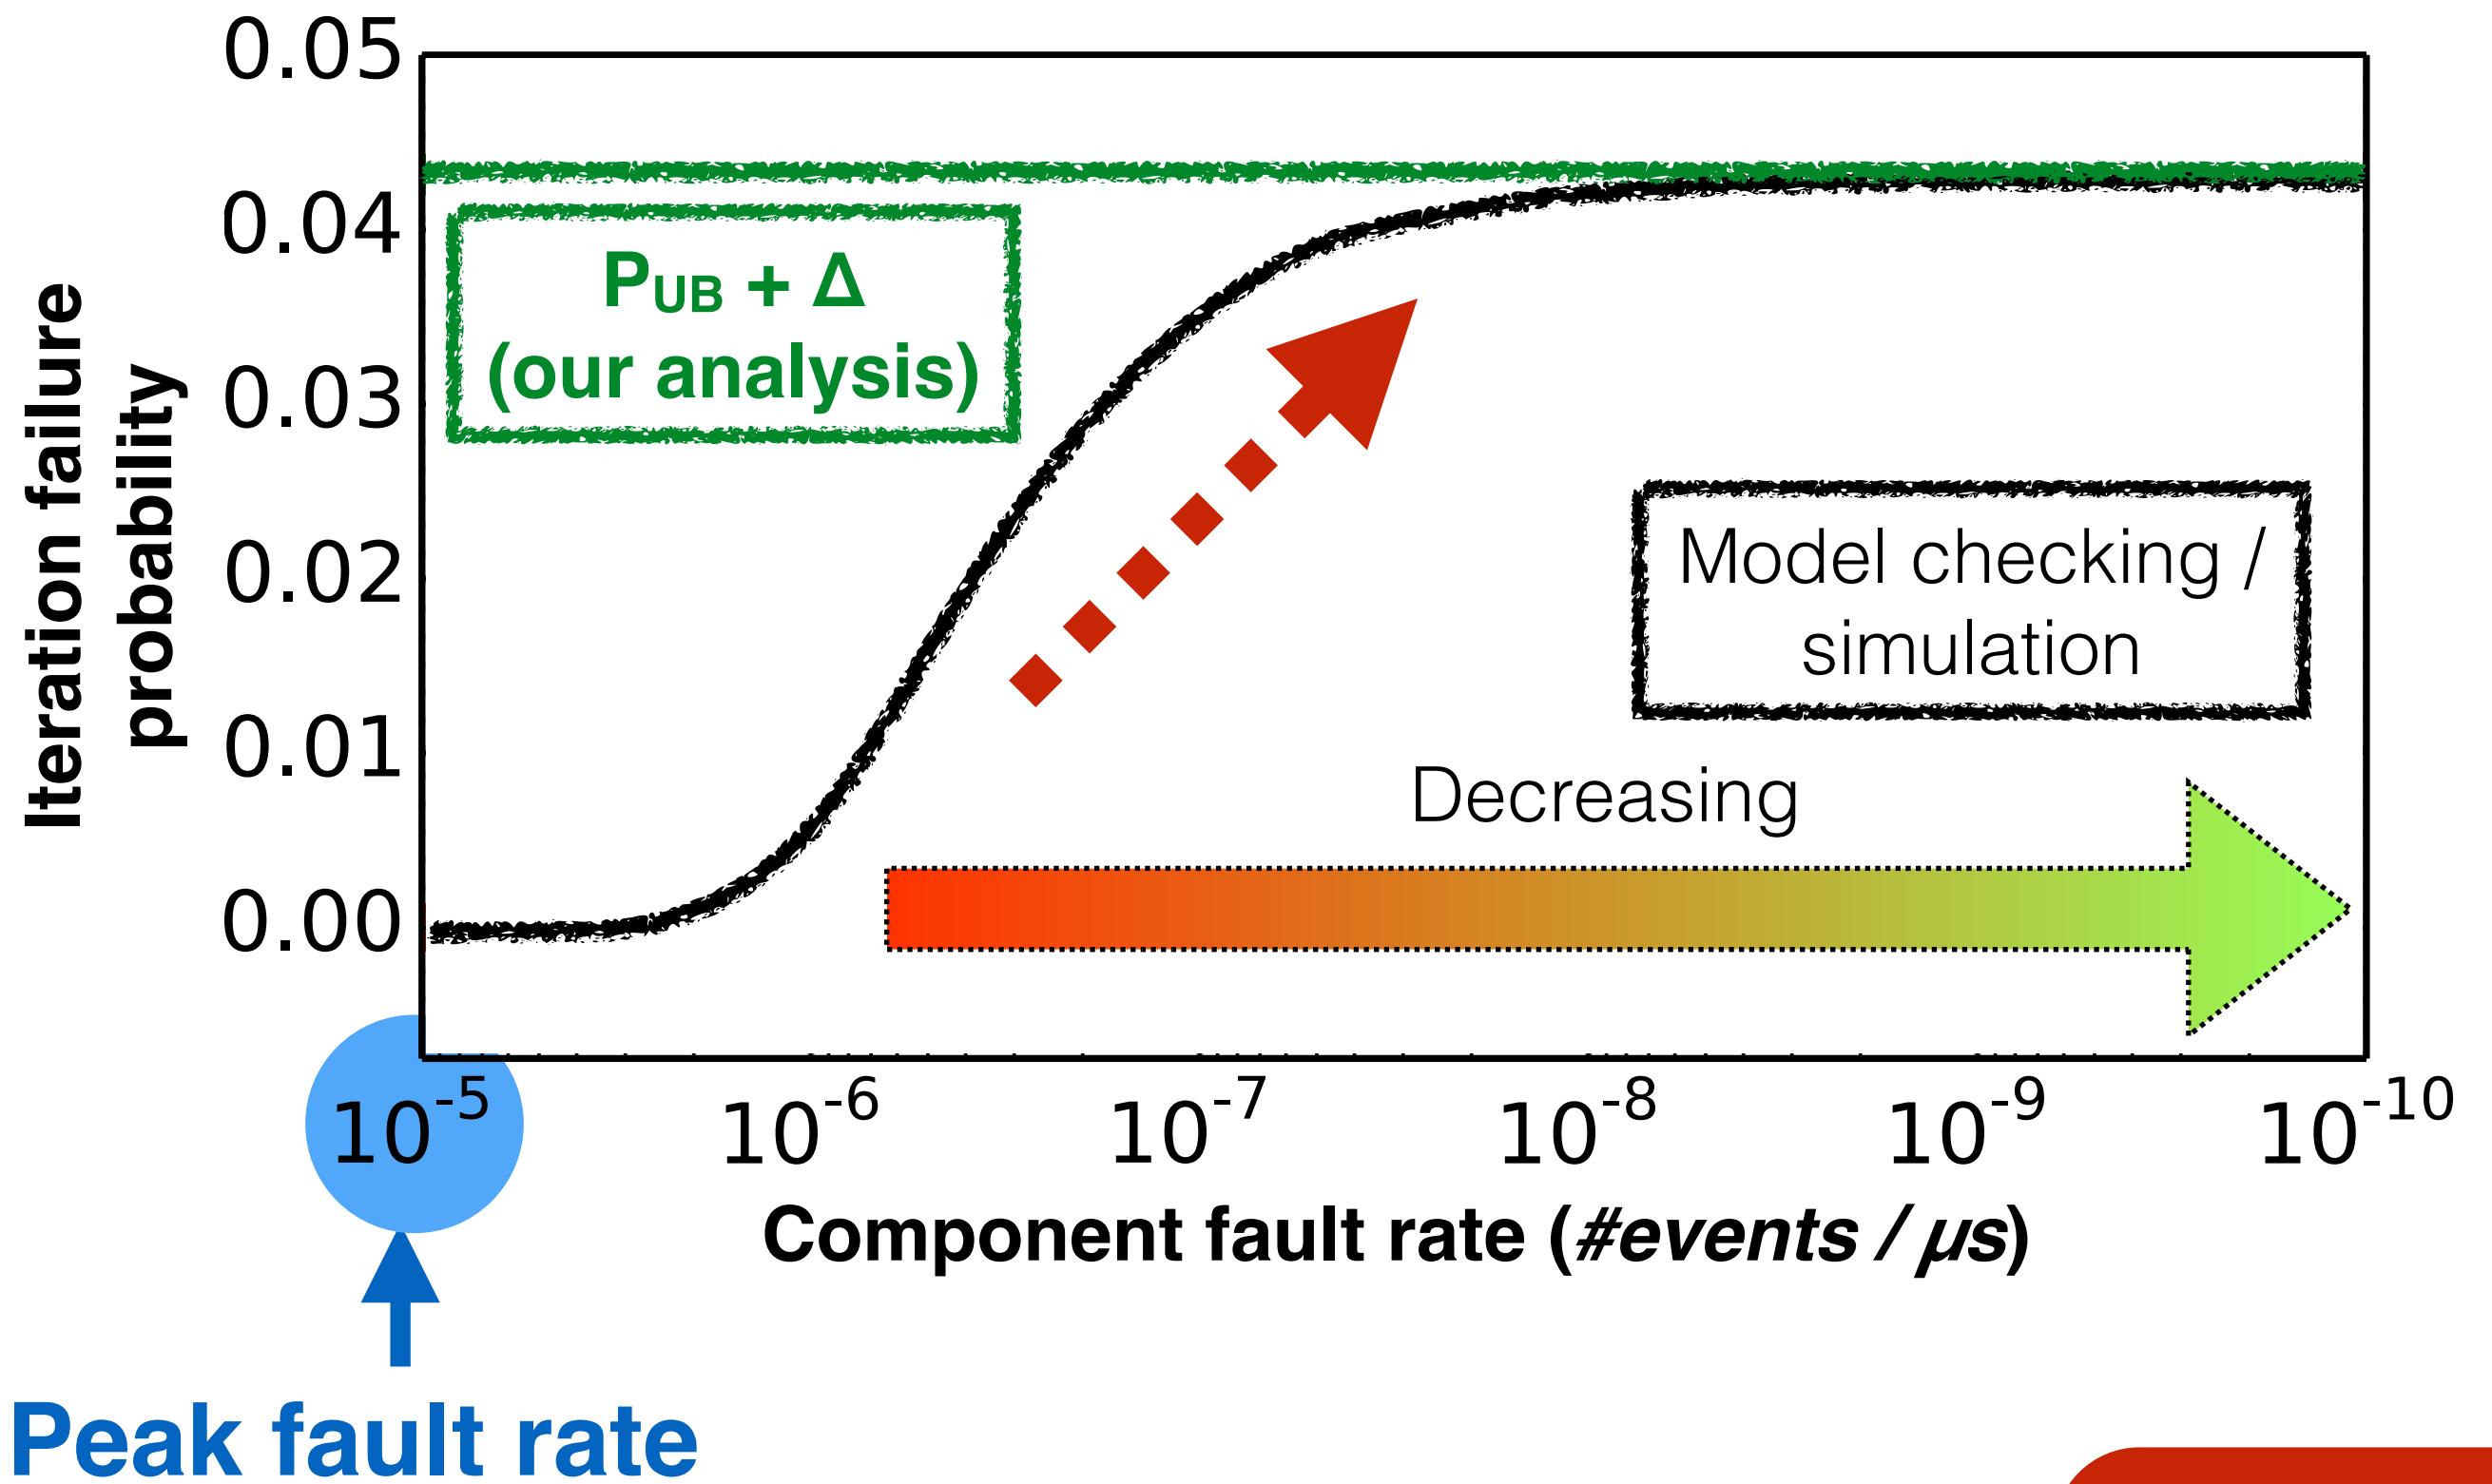


Eliminating reliability anomalies

A "fudge factor" Δ is added to P_{UB} to ensure that $P_{UB} + \Delta$ is monotonically increasing with P_1, P_2, P_3, \dots

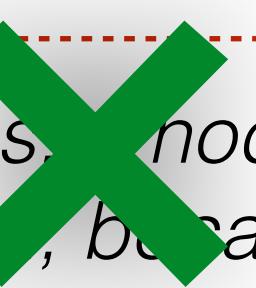


The problem of reliability anomalies



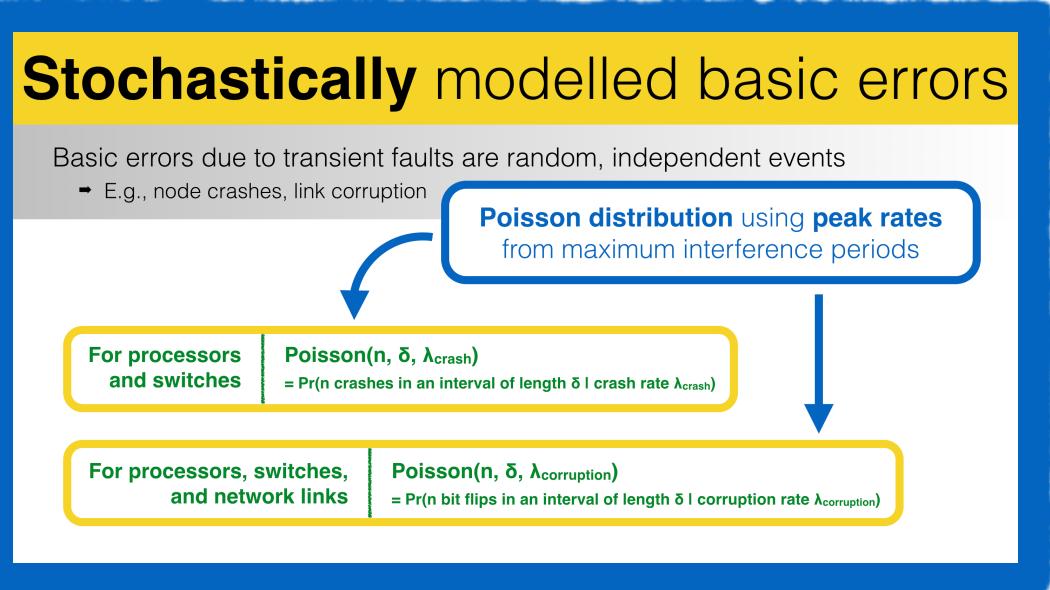
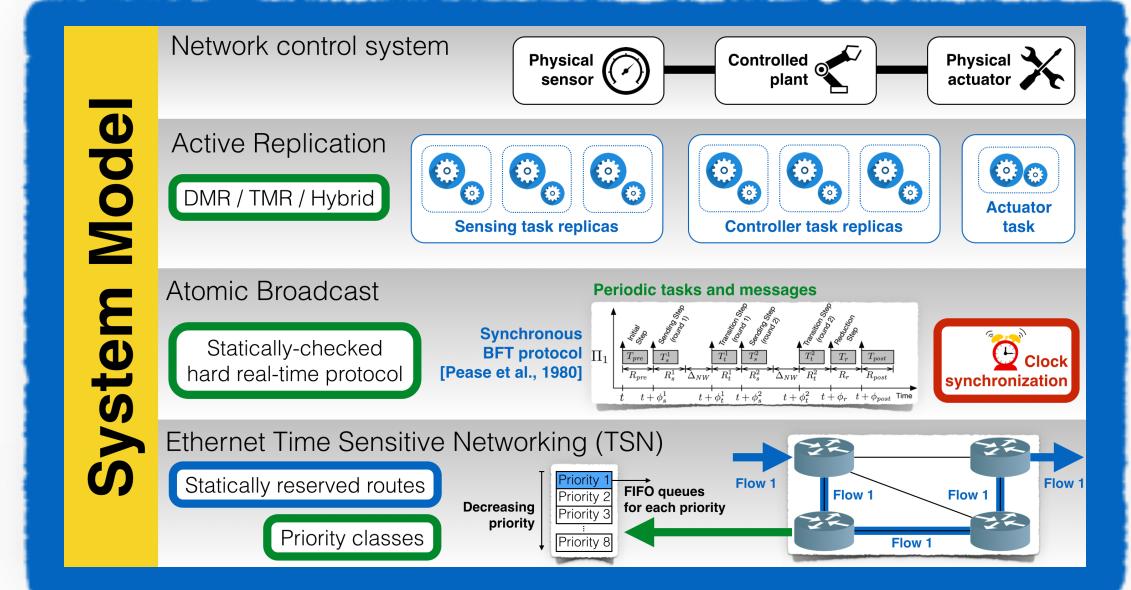
$\Pr(\text{iteration failure})$ increases despite decreasing component fault rate

Intuition: Sometimes, a node crash is good for the overall system, because it may reduce the probability of confusing a majority voting protocol in another part of the system!



For soundness, need to estimate failure probabilities for the **entire search space** $[0, 10^{-5}]$

Key idea 2: Ensure **monotonicity** to eliminate anomalies



Model checking or simulation

$\Pr(\text{ iteration failure })$

Scalability challenges

Key idea 1: Tackle scalability through abstraction and pruning

Reliability anomalies

Key idea 2: Ensure *monotonicity* to eliminate anomalies

System Model

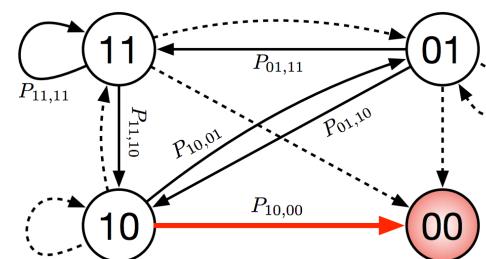
Network control system

Physical plant reliable



Q2. What is the likelihood of a **control failure**?

Periodic Weakly-Hard Systems [ECRTS '19]

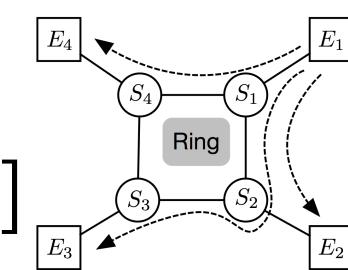


Transient
fault-induced
errors



Q1. How often does the **final actuation deviate** from an error-free scenario (iteration failure)?

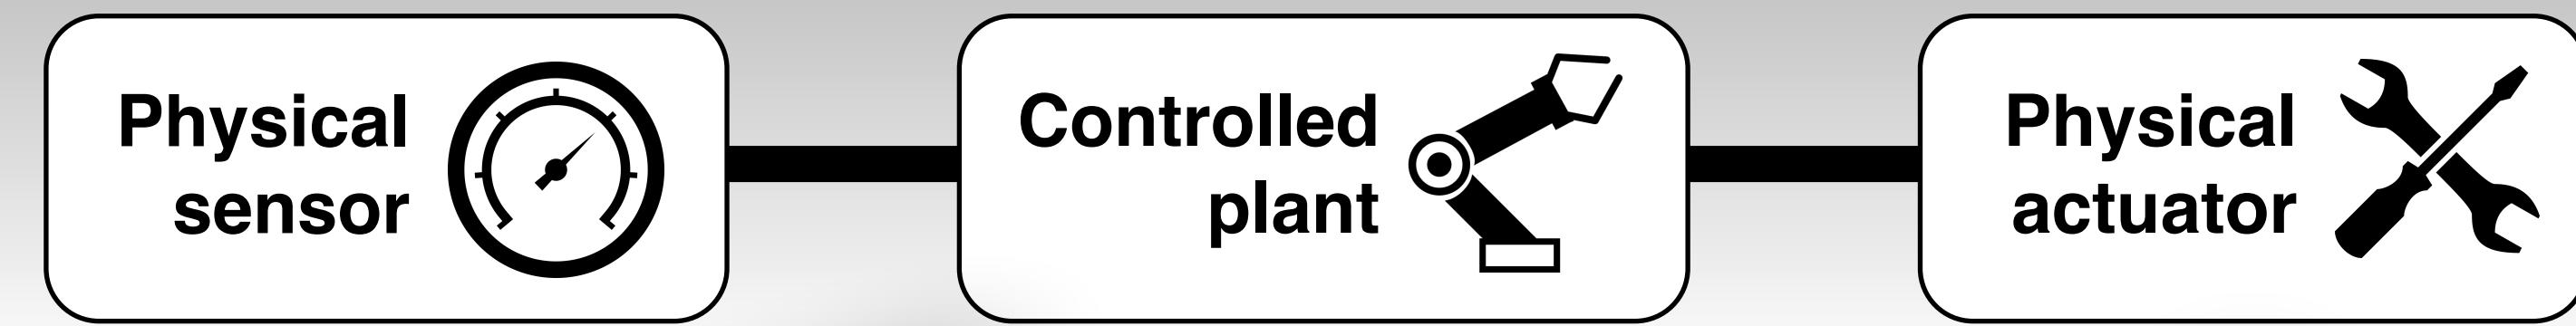
Replica Consistency over Ethernet [RTAS '20]



System Model

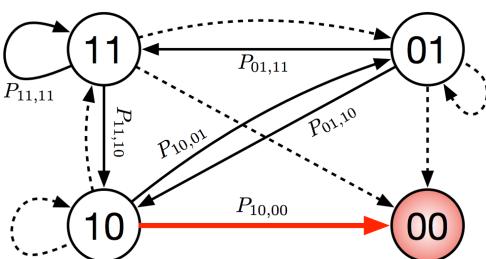
Network control system

Physical plant reliable



Q2. What is the likelihood of a **control failure**?

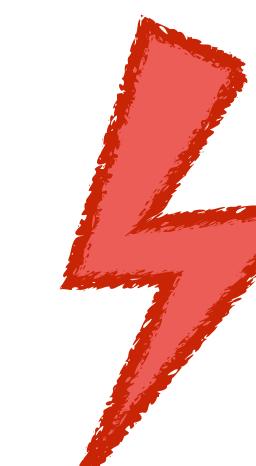
Periodic Weakly-Hard Systems [ECRTS '19]



Failures-In-Time (FIT)

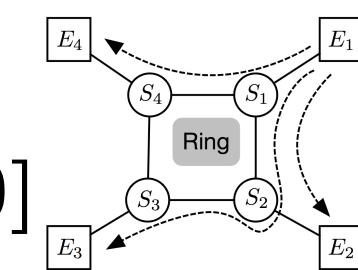
- Expected # failures in one billion operating hours

Transient fault-induced errors



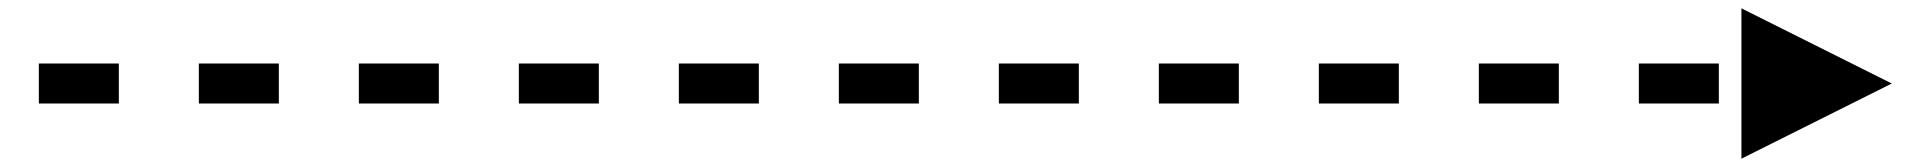
Q1. How often does the **final actuation deviate** from an error-free scenario (iteration failure)?

Replica Consistency over Ethernet [RTAS '20]



Simplistic approach

$P_{UB} + \Delta > Pr(\text{ iteration failure })$



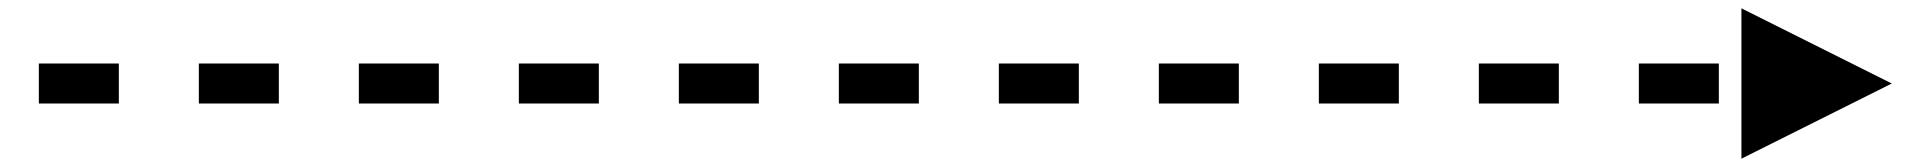
$FIT_{UB} > FIT_{\text{control-system}}$

Assumption:

1st iteration failure \rightarrow control system failure

Simplistic approach

$P_{UB} + \Delta > \Pr(\text{ iteration failure })$



$\text{FIT}_{UB} > \text{FIT}_{\text{control-system}}$

Assumption:

1st iteration failure \rightarrow control system failure

Failures-In-Time

$$\text{FIT} = \frac{10^9}{\text{MTTF} \text{ (in hours)}}$$

Simplistic approach

$$P_{UB} + \Delta > \Pr(\text{iteration failure})$$

$$FIT_{UB} > FIT_{\text{control-system}}$$

Assumption:

1st iteration failure \rightarrow control system failure

Mean Time To Failure

$$MTTF = T \times \sum_{n=0}^{\infty} n \times f(n)$$

Time-period Iteration #

Failures-In-Time

$$FIT = \frac{10^9}{MTTF \text{ (in hours)}}$$

Simplistic approach

$$P_{UB} + \Delta > \Pr(\text{ iteration failure })$$

$$FIT_{UB} > FIT_{\text{control-system}}$$

Assumption:

1st iteration failure \rightarrow control system failure

Probability density function

$$f(n) = (1 - \Pr(\text{ iteration failure }))^{n-1} \times \Pr(\text{ iteration failure })$$

Mean Time To Failure

$$MTTF = T \times \sum_{n=0}^{\infty} n \times f(n)$$

Time-period Iteration #

Failures-In-Time

$$FIT = \frac{10^9}{MTTF \text{ (in hours)}}$$

Simplistic approach

$$P_{UB} + \Delta > \Pr(\text{ iteration failure })$$

$$FIT_{UB} > FIT_{\text{control-system}}$$

Assumption:

1st iteration failure \rightarrow control system failure

Probability density function

$$f(n) = (1 - \Pr(\text{ iteration failure }))^{n-1} \times \Pr(\text{ iteration failure })$$

Mean Time To Failure

$$MTTF = T \times \sum_{n=0}^{\infty} n \times f(n)$$

Time-period Iteration #

Failures-In-Time

$$FIT = \frac{10^9}{MTTF \text{ (in hours)}}$$

Example

- ▶ Networked control system operates at **100 Hz (T = 10 ms)**
- ▶ Iteration failure probability upper-bounded by **$P_{UB} + \Delta = 10^{-10}$**

Simplistic approach

$$P_{UB} + \Delta > \Pr(\text{ iteration failure })$$

$$FIT_{UB} > FIT_{\text{control-system}}$$

Assumption:

1st iteration failure \rightarrow control system failure

Probability density function

$$f(n) = (1 - \Pr(\text{ iteration failure }))^{n-1} \times \Pr(\text{ iteration failure })$$

Mean Time To Failure

$$MTTF = T \times \sum_{n=0}^{\infty} n \times f(n)$$

Time-period Iteration #

Failures-In-Time

$$FIT = \frac{10^9}{MTTF \text{ (in hours)}}$$

Example

- Networked control system operates at **100 Hz (T = 10 ms)**
- Iteration failure probability upper-bounded by **$P_{UB} + \Delta = 10^{-10}$**

$$MTTF = 108 \text{ seconds}$$

Pessimistic!

$$FIT = 36000$$

Simplistic approach

$$P_{UB} + \Delta > \Pr(\text{iteration failure})$$

$$FIT_{UB} > FIT_{\text{control-system}}$$

Assumption:

1st iteration failure \rightarrow control system failure

Probability density function

$$f(n) = (1 - \Pr(\text{iteration failure}))^{n-1} \times \Pr(\text{iteration failure})$$

Mean Time To Failure

$$MTTF = T \times \sum_{n=0}^{\infty} n \times f(n)$$

Time-period Iteration #

Failures-In-Time

$$FIT = \frac{10^9}{MTTF \text{ (in hours)}}$$

Example

- Networked control system operates at **100 Hz (T = 10 ms)**
- Iteration failure probability upper-bounded by $P_{UB} + \Delta = 10^{-10}$

Pessimistic!

$$MTTF = 108 \text{ seconds}$$

$$FIT = 36000$$

**Well-designed
robust controllers**

The controller **tolerates** one iteration failure every four consecutive iterations

Simplistic approach

$$P_{UB} + \Delta > \Pr(\text{iteration failure})$$

$$FIT_{UB} > FIT_{\text{control-system}}$$

Assumption:

1st iteration failure \rightarrow control system failure

Probability density function

$$f(n) = (1 - \Pr(\text{iteration failure}))^{n-1} \times \Pr(\text{iteration failure})$$

Mean Time To Failure

$$MTTF = T \times \sum_{n=0}^{\infty} n \times f(n)$$

Time-period Iteration #

Failures-In-Time

$$FIT = \frac{10^9}{MTTF \text{ (in hours)}}$$

Example

- Networked control system operates at **100 Hz (T = 10 ms)**
- Iteration failure probability upper-bounded by $P_{UB} + \Delta = 10^{-10}$

Pessimistic!

$$MTTF = 108 \text{ seconds}$$

$$FIT = 36000$$

**Well-designed
robust controllers**

The controller **tolerates** one iteration failure every four consecutive iterations

$$FIT = 1.08 \times 10^{-5}$$

Simplistic approach

$$P_{UB} + \Delta > \Pr(\text{iteration failure})$$

$$FIT_{UB} > FIT_{\text{control-system}}$$

Assumption:

1st iteration failure \rightarrow control system failure

Mean Time To Failure

How to compute **FIT_{UB}** from $P_{UB} + \Delta$ while taking into account ***temporal robustness*** properties of **well-designed controllers**?

Example

- Networked control system operates at 100 Hz ($T = 10 \text{ ms}$)
- Iteration failure probability upper-bounded by $P_{UB} + \Delta = 10^{-10}$

$$\text{MTTF} = 108 \text{ seconds}$$

$$FIT = 36000$$

Pessimistic!

More accurate

$$FIT = 1.08 \times 10^{-5}$$

Well-designed
robust controllers

The controller **tolerates** one iteration
failure every four consecutive iterations

Weakly-hard constraints*

* Bernat, Burns, and Liamosi. "Weakly hard real-time systems." IEEE transactions on Computers 50.4 (2001).

Weakly-hard constraints*

Example: (m, k) constraint

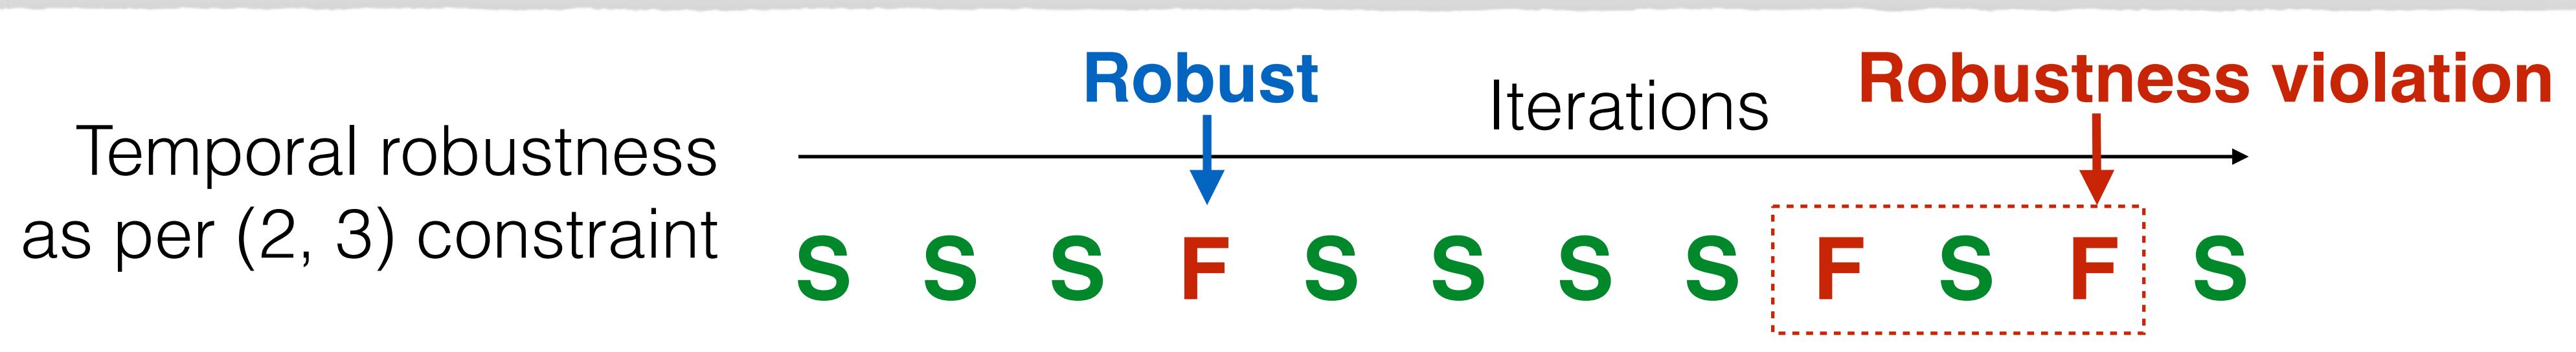
- ▶ At least m out of every k consecutive iterations must be successful

* Bernat, Burns, and Lianosi. "Weakly hard real-time systems." IEEE transactions on Computers 50.4 (2001).

Weakly-hard constraints*

Example: (m, k) constraint

- ▶ At least m out of every k consecutive iterations must be successful
- ▶ If each iteration is labeled either as a **Success** or a **Failure**



* Bernat, Burns, and Lianosi. "Weakly hard real-time systems." IEEE transactions on Computers 50.4 (2001).

Sound Approximation (SAP)

$$P_{UB} + \Delta > \Pr(\text{iteration failure})$$

$$FIT_{UB} > FIT_{\text{control-system}}$$

Assumption:

1st iteration failure \rightarrow control system failure

Probability density function

$$f(n) = (1 - \Pr(\text{iteration failure}))^{n-1} \times \Pr(\text{iteration failure})$$

Mean Time To Failure

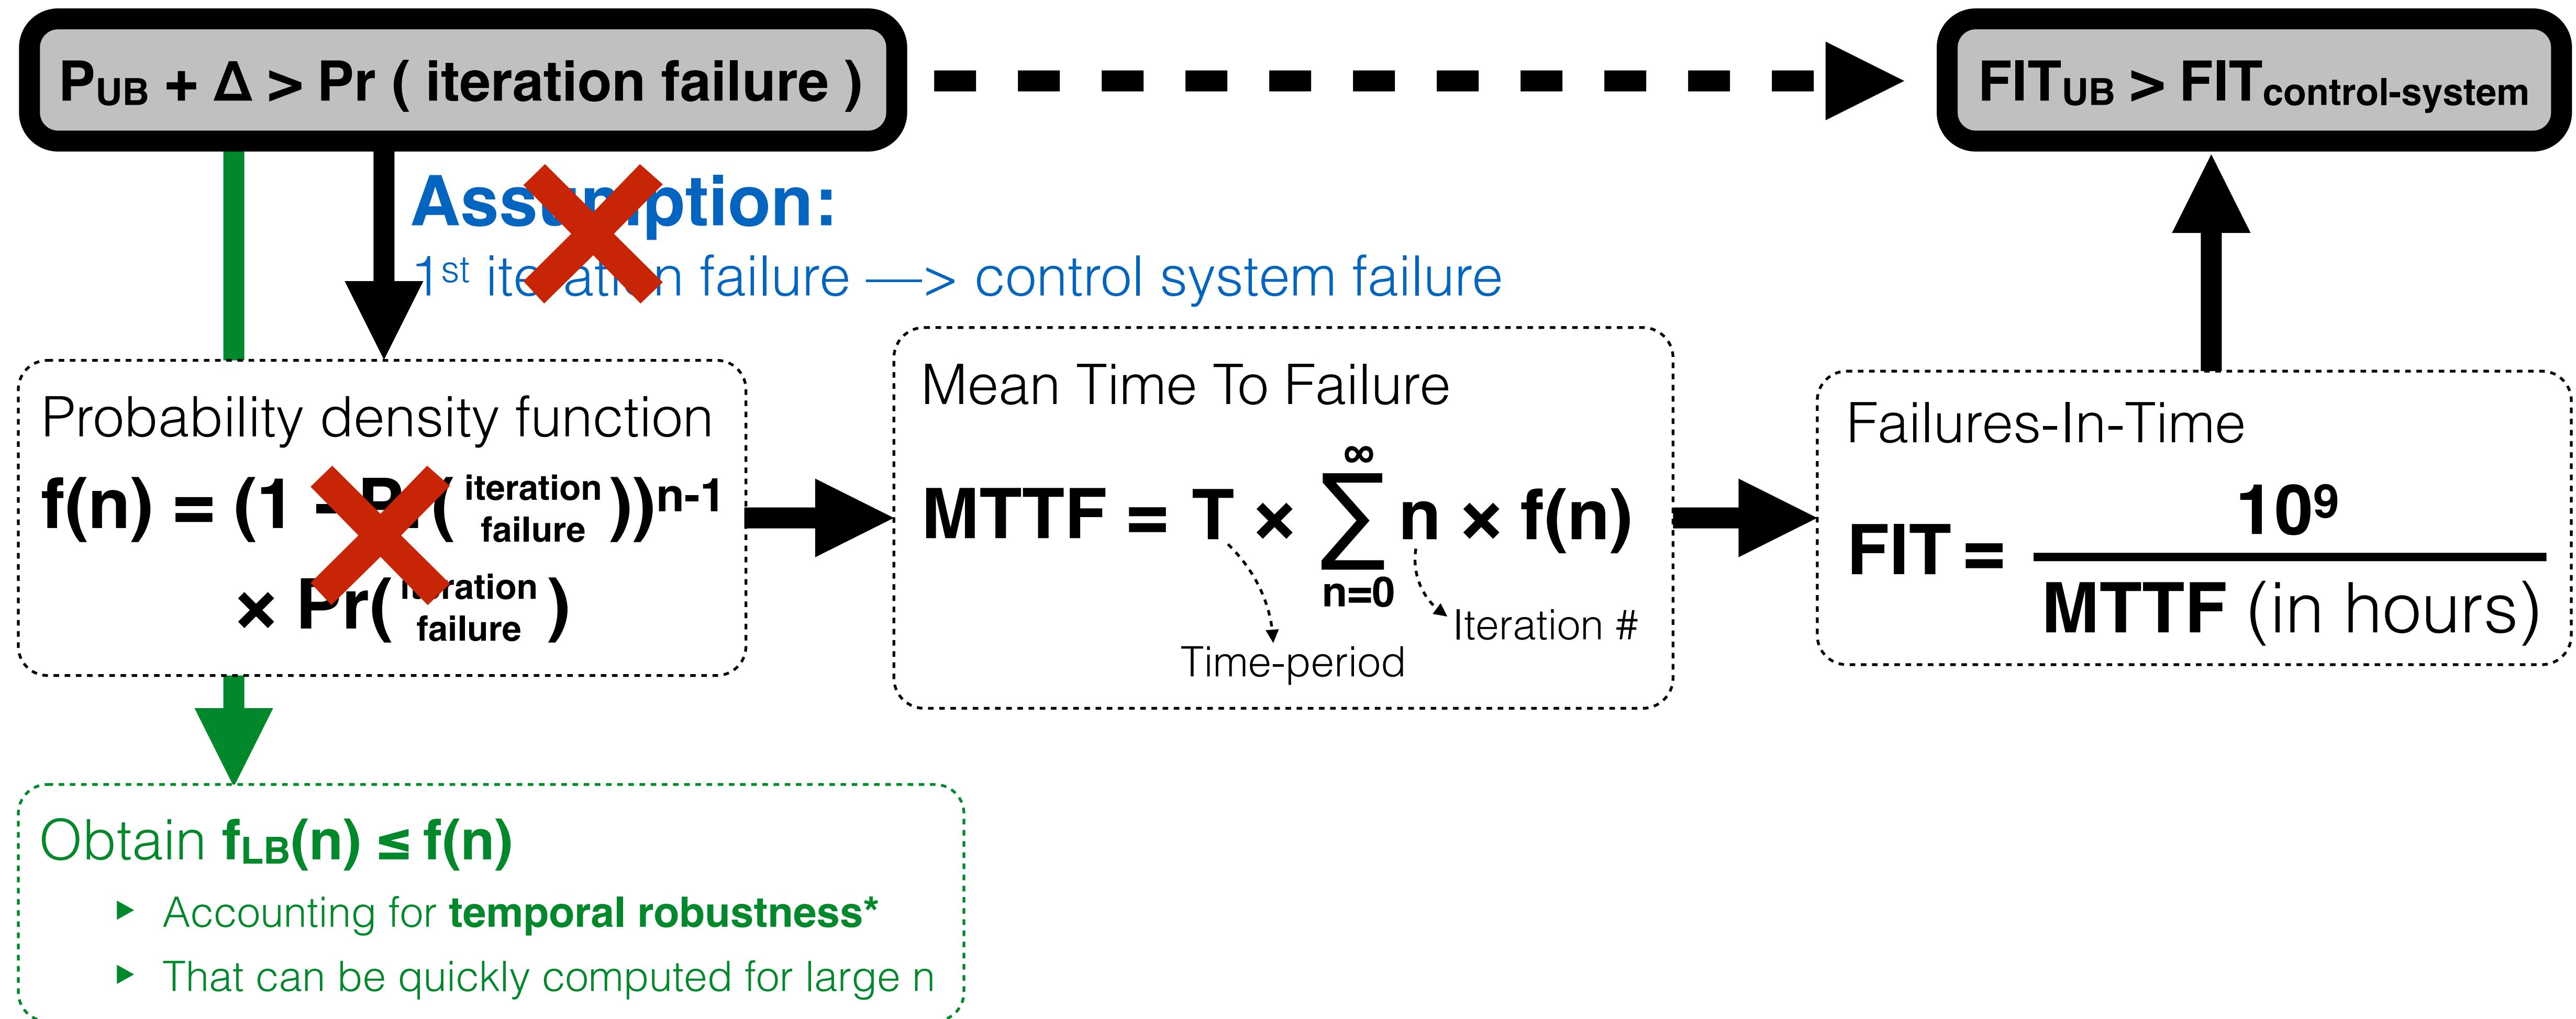
$$MTTF = T \times \sum_{n=0}^{\infty} n \times f(n)$$

Time-period Iteration #

Failures-In-Time

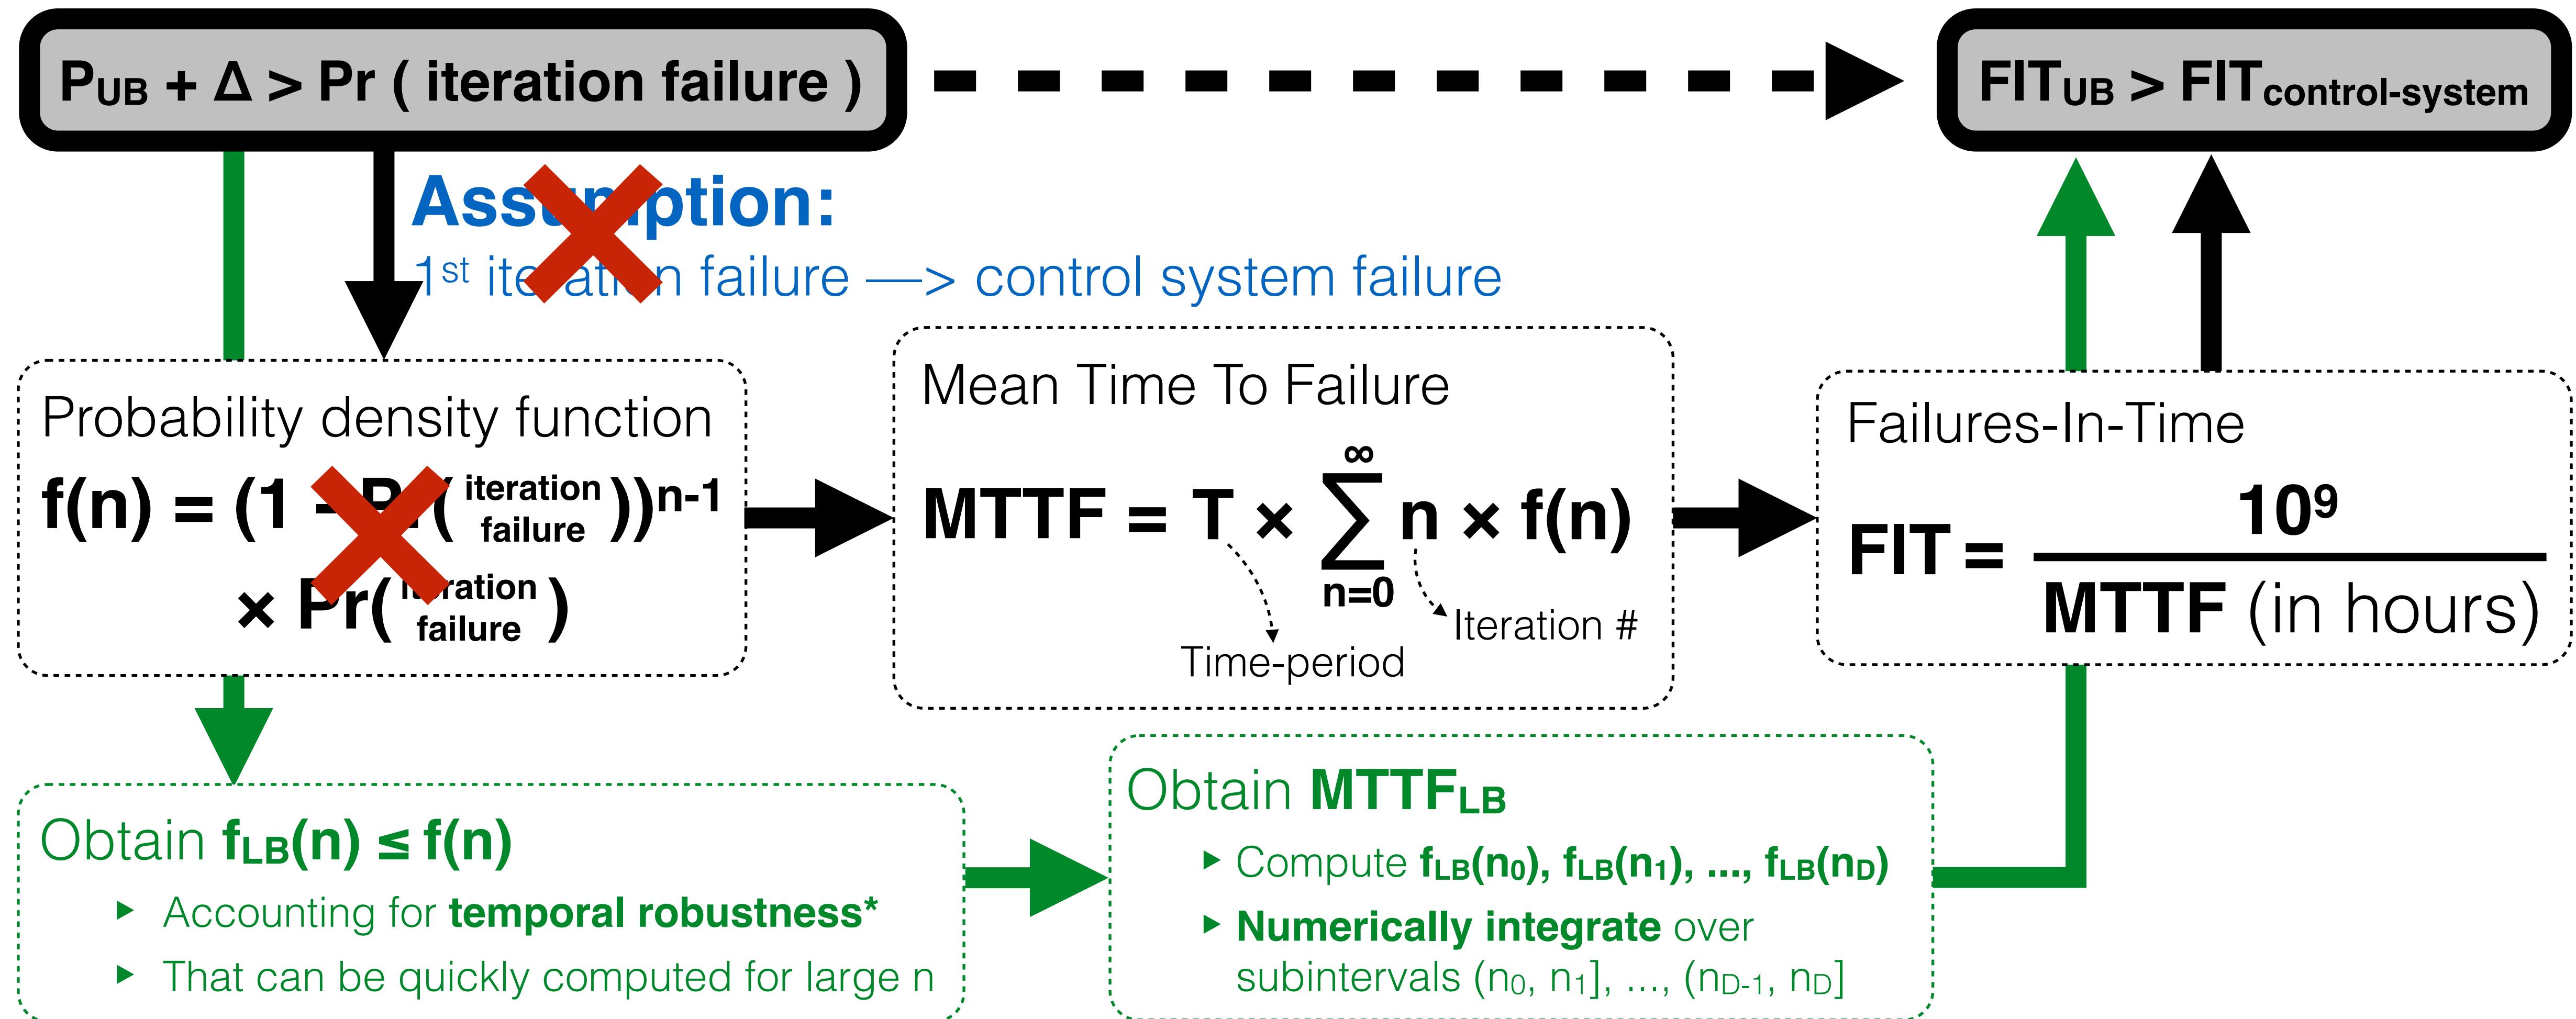
$$FIT = \frac{10^9}{MTTF \text{ (in hours)}}$$

Sound Approximation (SAP)



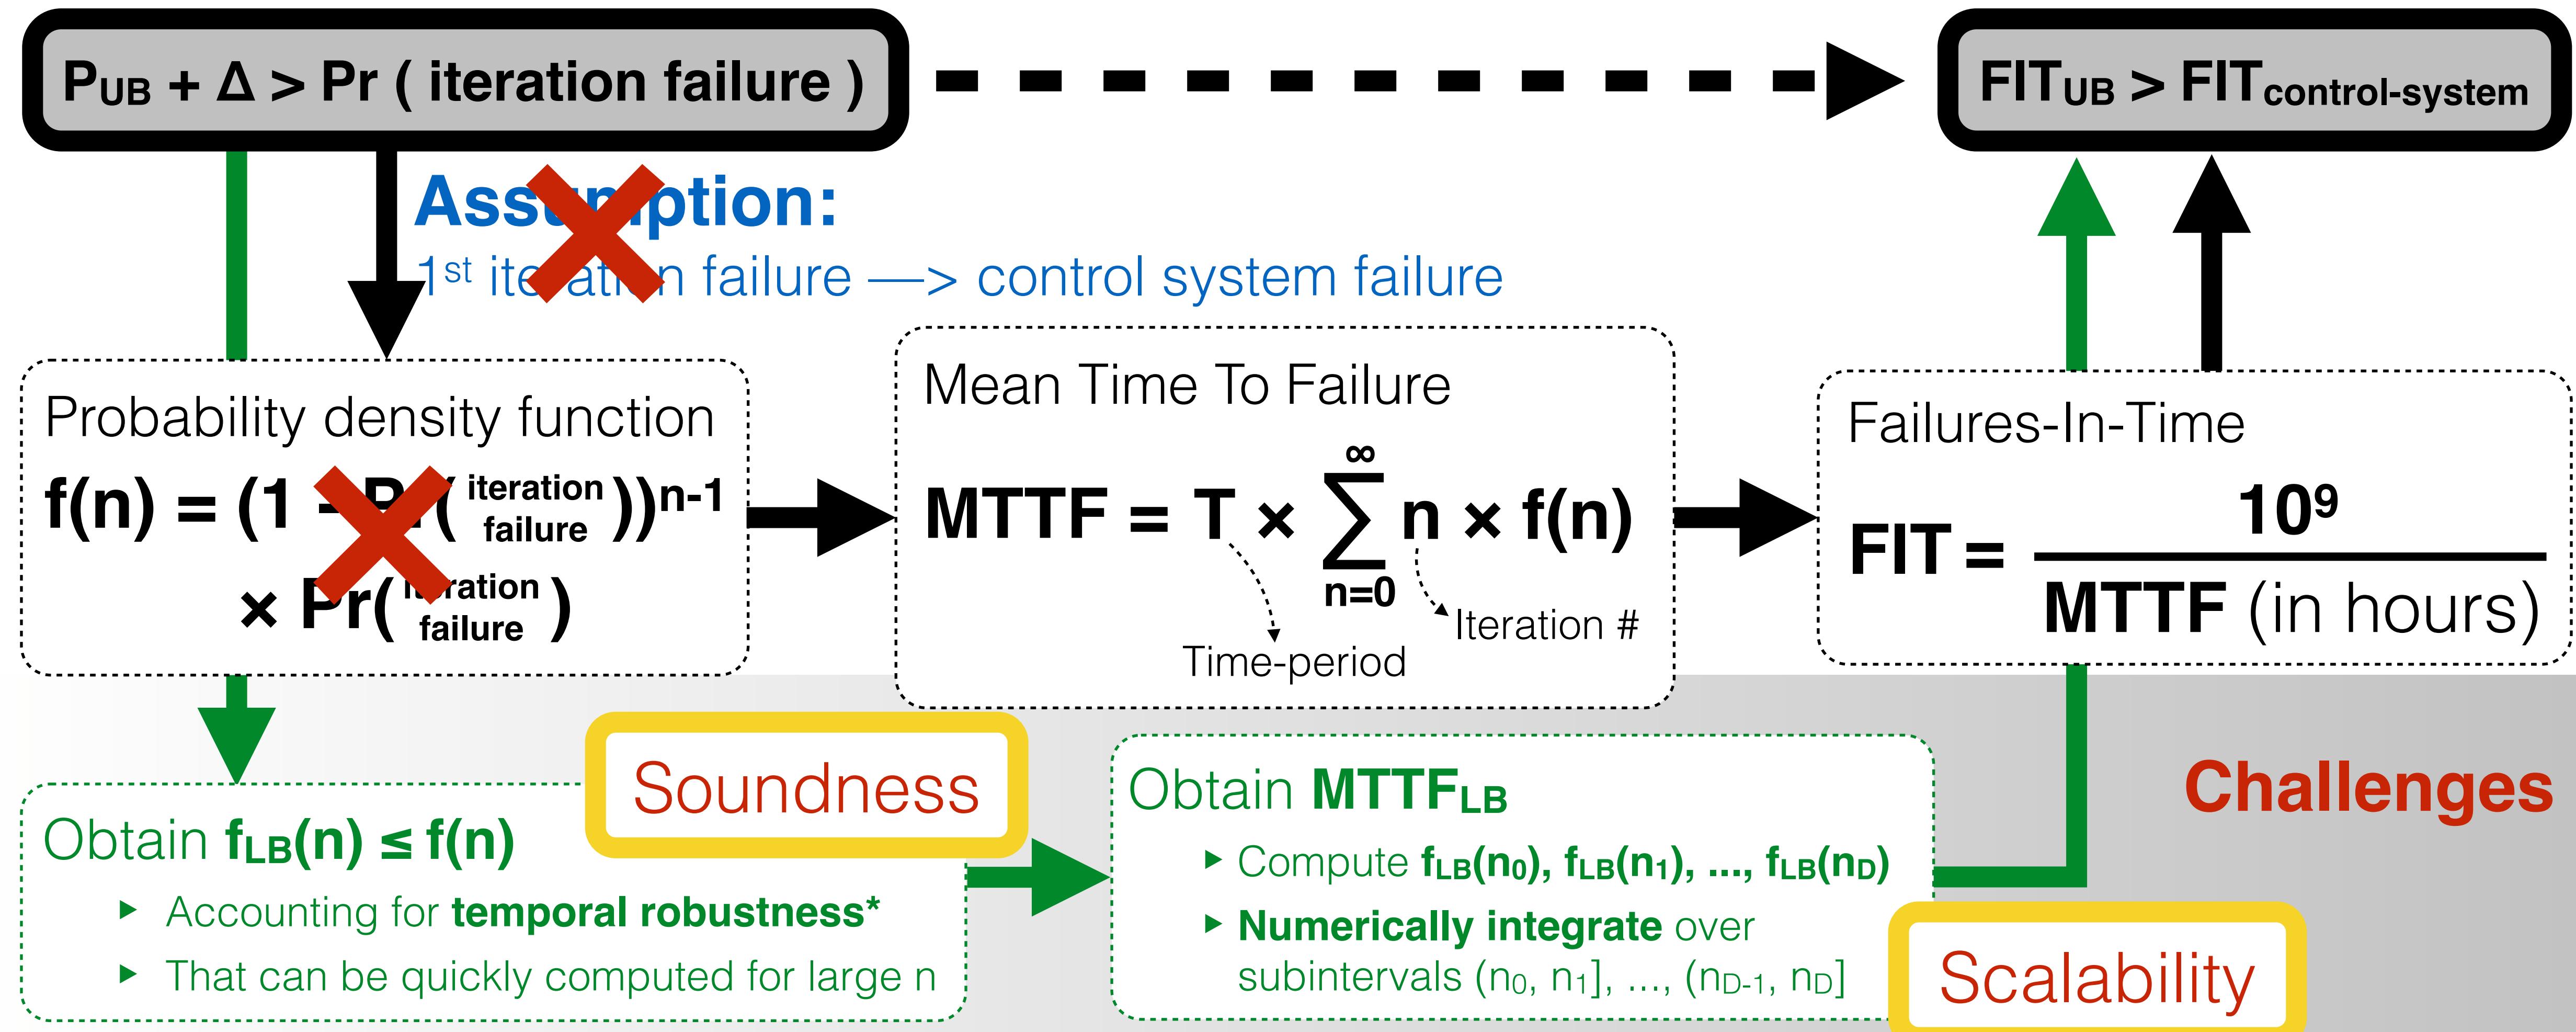
* Sfakianakis et al., "Reliability of a consecutive k-out-of-r-from-n: F system." IEEE Transactions on Reliability 41.3 (1992): 442-447.

Sound Approximation (SAP)



* Sfakianakis et al.. "Reliability of a consecutive k-out-of-r-from-n: F system." IEEE Transactions on Reliability 41.3 (1992): 442-447.

Sound Approximation (SAP)

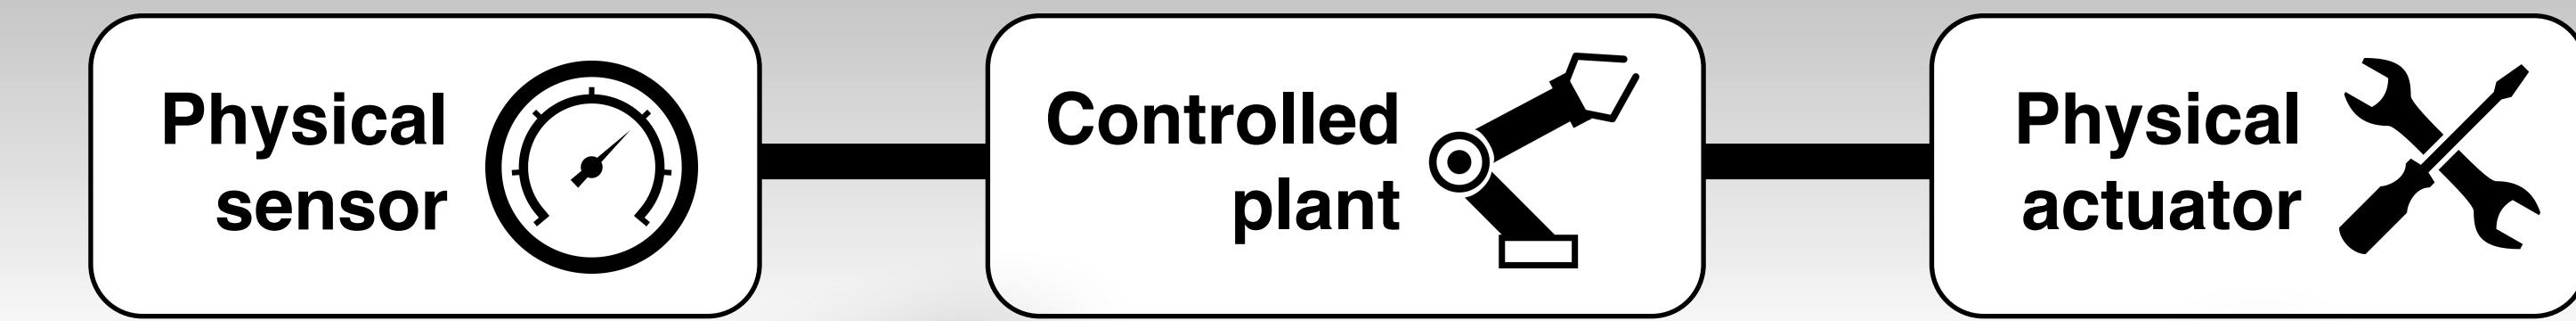


* Sfakianakis et al.. "Reliability of a consecutive k-out-of-r-from-n: F system." IEEE Transactions on Reliability 41.3 (1992): 442-447.

System Model

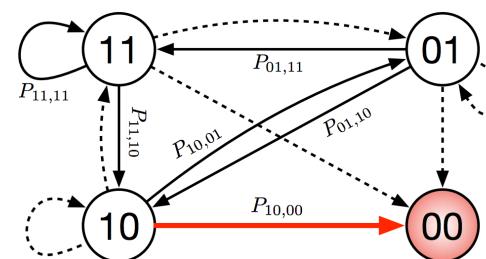
Network control system

Physical plant reliable



Q2. What is the likelihood of a **control failure**?

Periodic Weakly-Hard Systems [ECRTS '19]

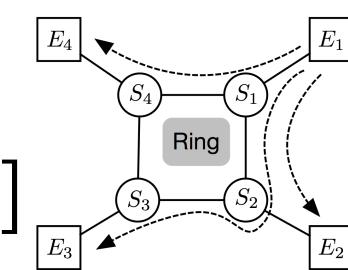


Transient
fault-induced
errors



Q1. How often does the **final actuation deviate** from an error-free scenario (iteration failure)?

Replica Consistency over Ethernet [RTAS '20]



System Model

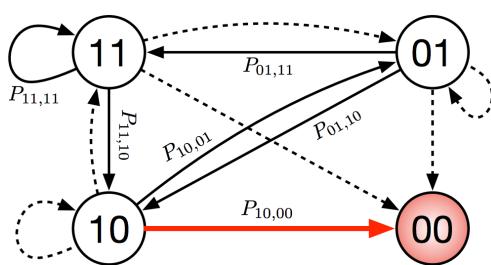
Network control system

Physical plant reliable



Q2. What is the likelihood of a **control failure**?

Periodic Weakly-Hard Systems [ECRTS '19]



Controlled plant

Physical actuator



Active Replication

Controller task replicas

Periodic tasks and messages

Statically scheduled hard real-time protocol

[Pease et al., 1980]

Ethernet Time

Statically reserved

Priority scheduling

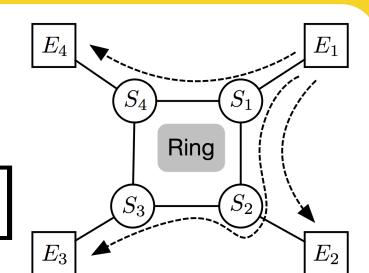
Transient fault-induced errors



Upper-bound the failure probability

- ▶ Distributed **real-time systems**
- ▶ **Reliability anomalies**
- ▶ Non-malicious **Byzantine** errors

Replica Consistency over Ethernet [RTAS '20]



System Model

Network control system

Physical plant reliable

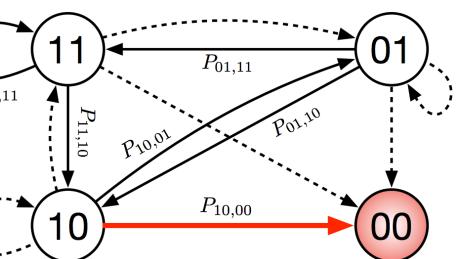
Physical sensor

Controlled plant

Physical actuator

Beyond hard constraints

Periodic Weakly-Hard Systems [ECRTS '19]



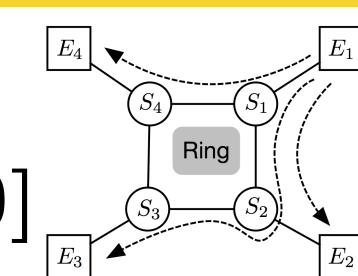
Upper-bound the failure probability

- ▶ Distributed **real-time systems**
- ▶ **Reliability anomalies**
- ▶ Non-malicious **Byzantine** errors

Transient fault-induced errors



Replica Consistency over Ethernet [RTAS '20]

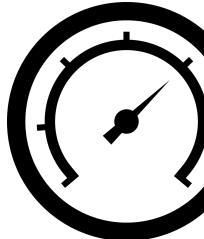


System Model

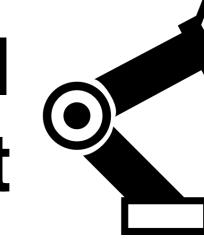
Network control system

Physical plant reliable

Physical sensor



Controlled plant

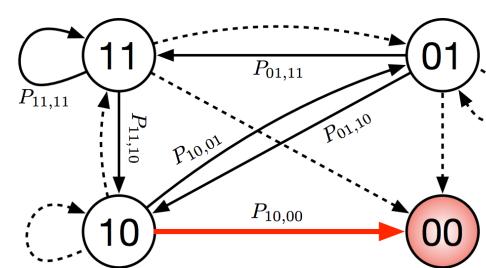


Physical actuator



Beyond hard constraints

Periodic Weakly-Hard Systems [ECRTS '19]



Towards ultra-reliable CPS

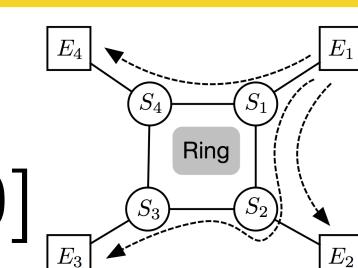
Transient fault-induced errors



Upper-bound the failure probability

- ▶ Distributed **real-time systems**
- ▶ **Reliability anomalies**
- ▶ Non-malicious **Byzantine** errors

Replica Consistency over Ethernet [RTAS '20]



Reliability Analysis of Distributed Real-Time Systems

Best Presentation Award

Controller Area
Network [ECRTS '18]

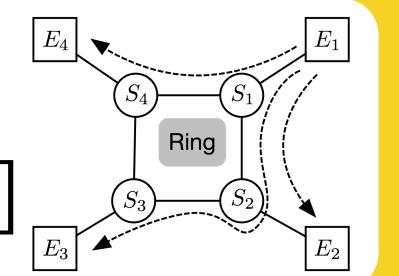


Periodic Weakly-Hard
Systems [ECRTS '19]



Distinguished Paper Award

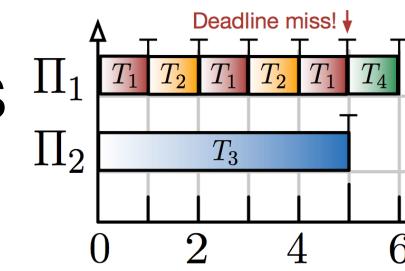
Replica Consistency
over Ethernet [RTAS '20]



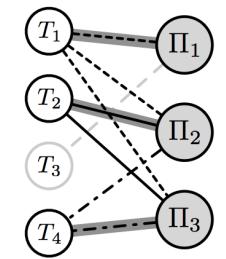
Multiprocessor Hard Real-Time Scheduling



Schedulability Analysis of Linux-like Systems
[ECRTS '13, RTS Journal '15]



Scheduling Policy for Improved Utilization [RTSS '14]



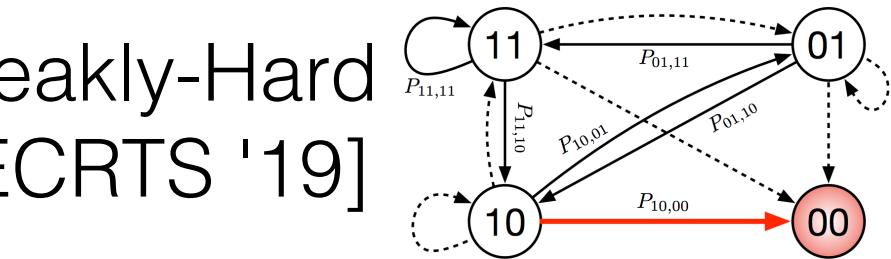
Reliability Analysis of Distributed Real-Time Systems



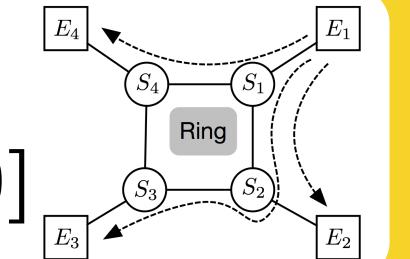
Controller Area Network [ECRTS '18]



Periodic Weakly-Hard Systems [ECRTS '19]



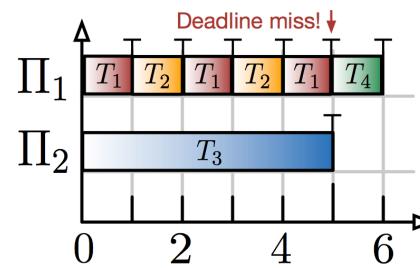
Replica Consistency over Ethernet [RTAS '20]



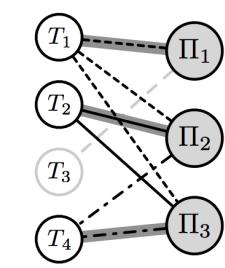
Multiprocessor Hard Real-Time Scheduling



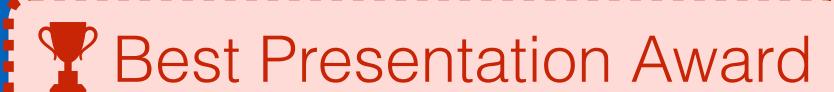
Schedulability Analysis of Linux-like Systems
[ECRTS '13, RTS Journal '15]



Scheduling Policy for Improved Utilization [RTSS '14]



Reliability Analysis of Distributed Real-Time Systems



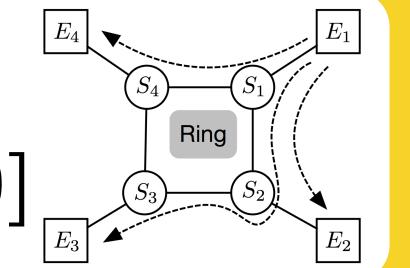
Controller Area Network [ECRTS '18]



Periodic Weakly-Hard Systems [ECRTS '19]



Replica Consistency over Ethernet [RTAS '20]



Predictable Resource Allocation



Swayam Autoscaler
[Middleware '17]

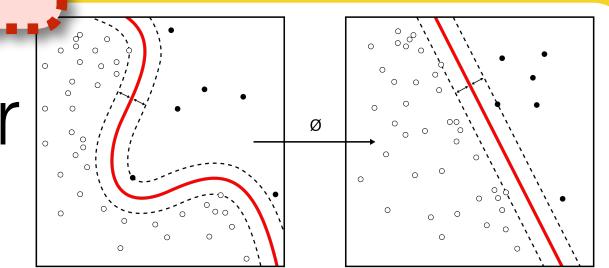
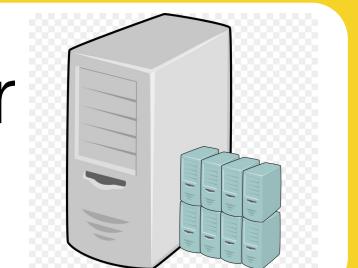


Tableau VM Scheduler
[EuroSys '18]



Clockwork for DNN Serving [OSDI '20]

