

Faster, Exact, More General Response-time Analysis for NVIDIA Holoscan Applications

Philip Schowitz, Shubhaankar

Sharma, Siddharth Balodi,

Bruce Shepherd, Arpan Gujarati

The University of British Columbia

Soham Sinha

NVIDIA



THE UNIVERSITY
OF BRITISH COLUMBIA



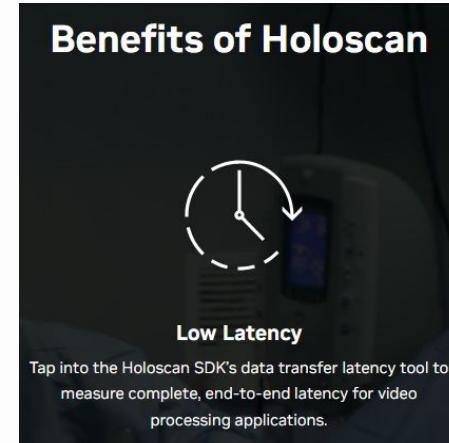
Edge Computing

- AI is fueling resource-intensive applications on the edge
- Embedded platforms become more complex
 - Harder to develop apps



Frameworks

- Frameworks like **NVIDIA Holoscan** streamline development
- Lacks hard latency guarantees
 - Potential for framework-level analysis



Why care about latency analysis?

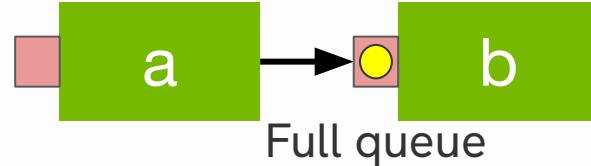
- Developers lack knowledge about framework timing properties but want low latency
 - Medical imaging, robotics
 - In the future, robotic surgery
- Create static analysis to allow informed decisions
 - E.g., what will be the timing effect of adding this node?
 - Useful during both design and development stages

Holoscan Basics



- Apps are represented as directed acyclic graphs
- App can process multiple inputs in parallel
 - Each node only processes one item at a time
 - Assume we can run all nodes at once
 - No node-to-core scheduling

Challenges



- Holoscan applications use a unique DAG model with downstream conditions
 - Node cannot begin execution if downstream queue is full

Challenges



- Holoscan applications use a unique DAG model with downstream conditions
 - Node cannot begin execution if downstream queue is full

Challenges



- Holoscan applications use a unique DAG model with downstream conditions
 - Node cannot begin execution if downstream queue is full

Challenges



- Holoscan applications use a unique DAG model with downstream conditions
 - Node cannot begin execution if downstream queue is full
 - Implements static backpressure, maintains internal consistency by preventing queue overflow

Challenges



- Holoscan applications use a unique DAG model with downstream conditions
 - Node cannot begin execution if downstream queue is full
 - Implements static backpressure, maintains internal consistency by preventing queue overflow
- This design creates timing anomalies
 - Lower node execution time → higher global response time

Prior Work

- We created a response time analysis for this model
 - Safe upper bound, scalable

Prior Work

- We created a response time analysis for this model
 - Safe upper bound, scalable
- However...
 - Pessimistic for some graphs
 - Static execution time, queue size must equal 1
 - Unclear connection to SDFG literature

Prior Work

Metric	Prior RTA
Runtime	51 μs
Pessimism	20.5%
Safe	✓
Exact	✗
Variable Execution Time	✗
Variable Queue Size	✗

Specific 9-node graph with pessimism

Contributions

Metric	Prior RTA	Synchronous Dataflow
Runtime	$51 \mu s$	$6055 s$
Pessimism	20.5%	0%
Safe	✓	✓
Exact	✗	✓
Variable Execution Time	✗	✓
Variable Queue Size	✗	✓

Specific 9-node graph with pessimism

Contributions

Metric	Prior RTA	Synchronous Dataflow	This work
Runtime	$51 \mu s$	$6055 s$	$136 ms$
Pessimism	20.5%	0%	0%
Safe	✓	✓	✓
Exact	✗	✓	✓
Variable Execution Time	✗	✓	✓
Variable Queue Size	✗	✓	✓

Specific 9-node graph with pessimism

Response-Time Analysis

Problem Statement

- What is the largest possible difference in time between source (a) starting and sink (c) starting?
 - In some iteration x, the difference is: $s(c_x) - s(a_x)$
 - Can trivially add back sink's exec to get the WCRT



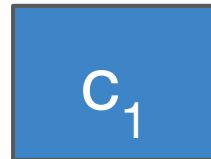
Trace Graph

First iteration
or input

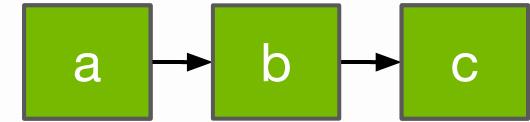
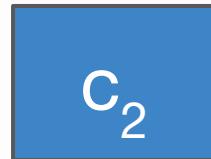
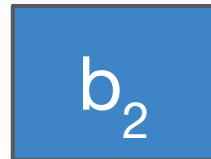
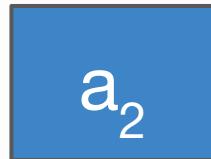


Trace Graph

First iteration
or input



a₁ / a₂ may have
different execution
times, b₁ / b₂, etc

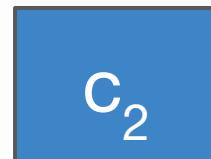
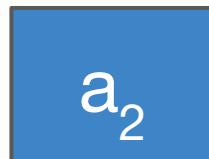


Trace Graph

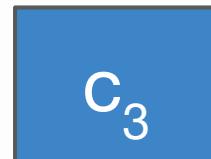
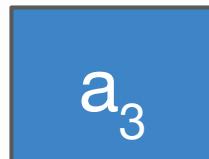
First iteration
or input



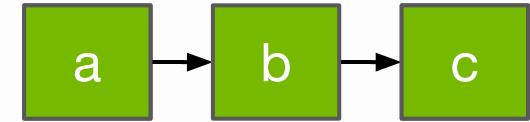
a₁ / a₂ may have
different execution
times, b₁ / b₂, etc



Infinite graph!



⋮



Trace Graph

Data-dependency edges

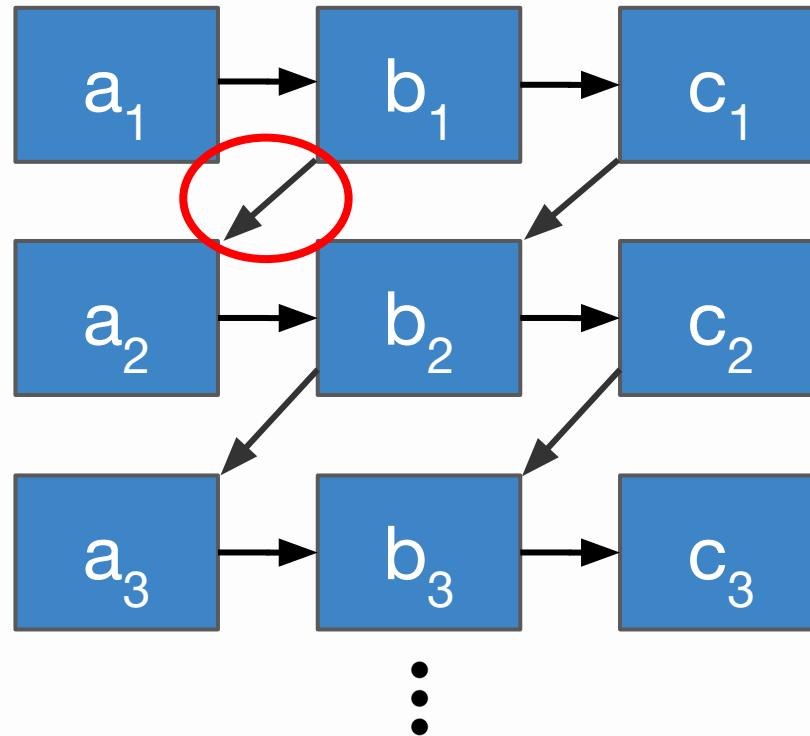
Cost = node
execution time
(finish-to-start)



Trace Graph

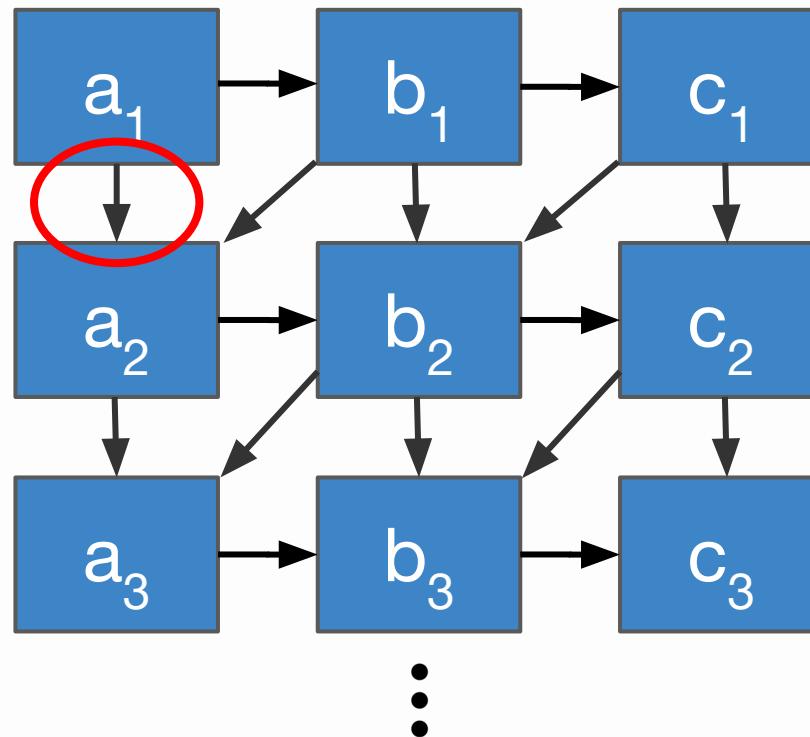
Downstream
blocking edges
(backpressure)

Cost = 0
(start-to-start)



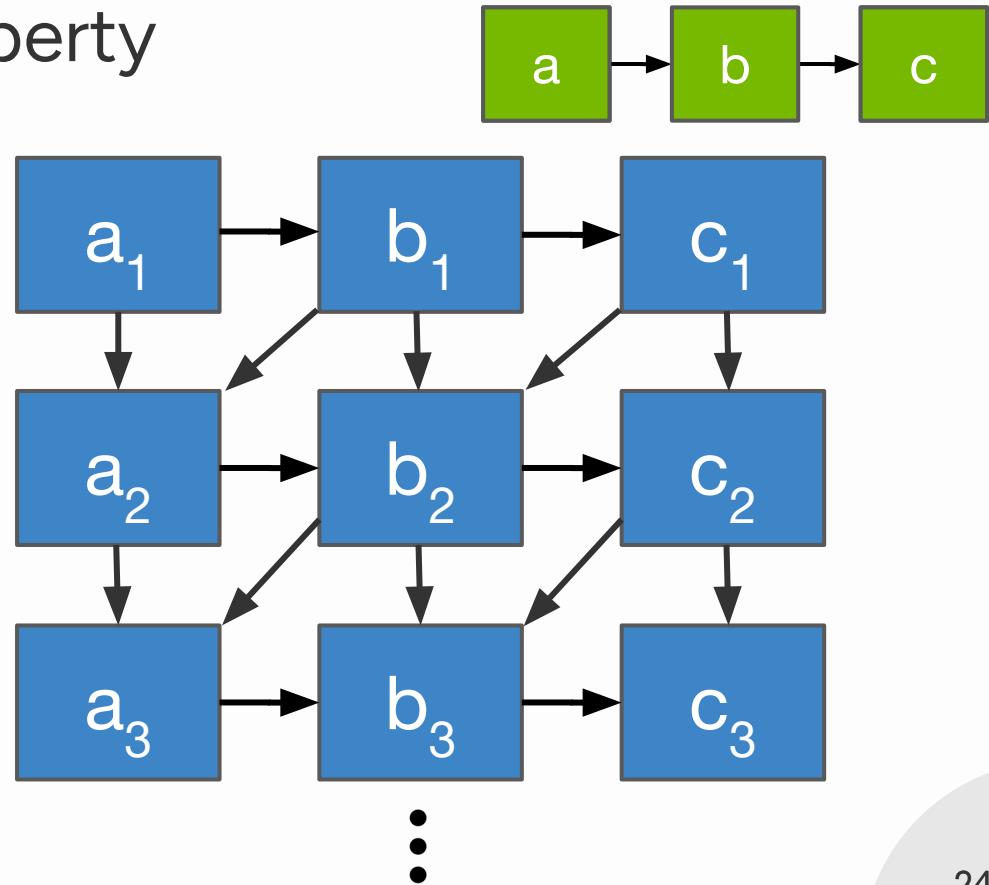
Trace Graph

Sequential-execution edges
Cost = node execution time (finish-to-start)



Key Trace Graph Property

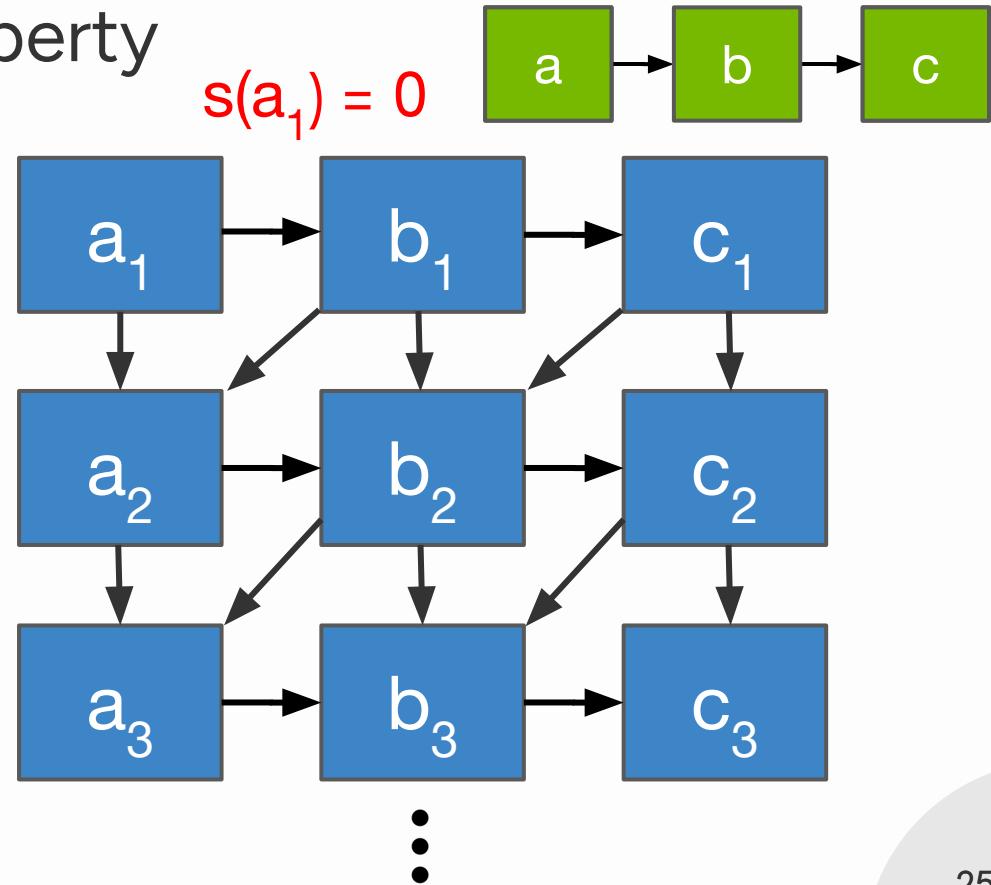
- Longest path from source (a_1) defines a node's start time
 - Preconditions must be met before node may execute



Key Trace Graph Property

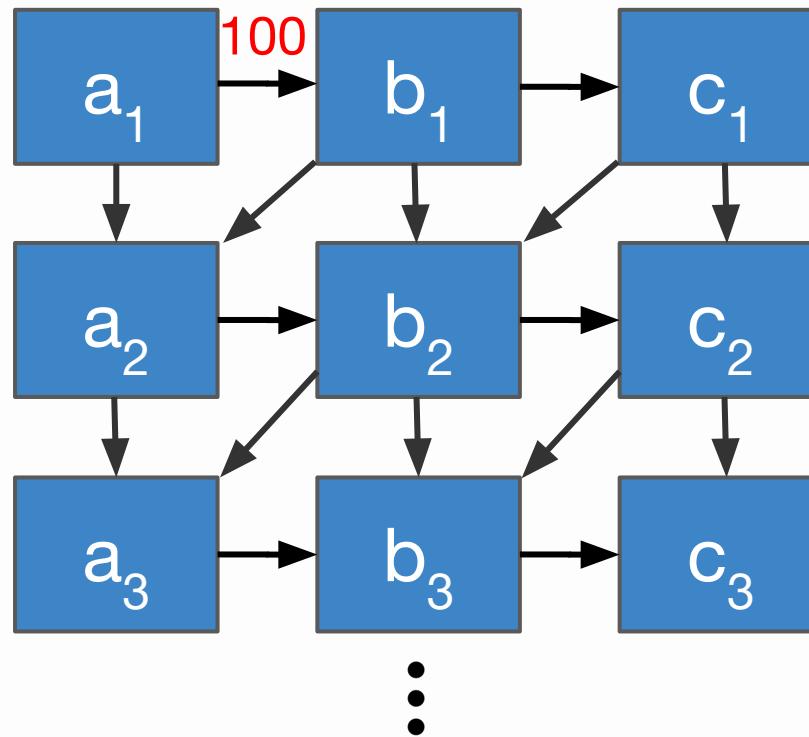
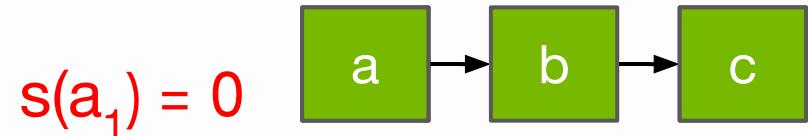
- Longest path from source (a_1) defines a node's start time
 - Preconditions must be met before node may execute

$$s(a_1) = 0$$



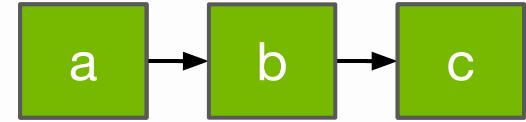
Key Trace Graph Property

- Longest path from source (a_1) defines a node's start time
 - Preconditions must be met before node may execute

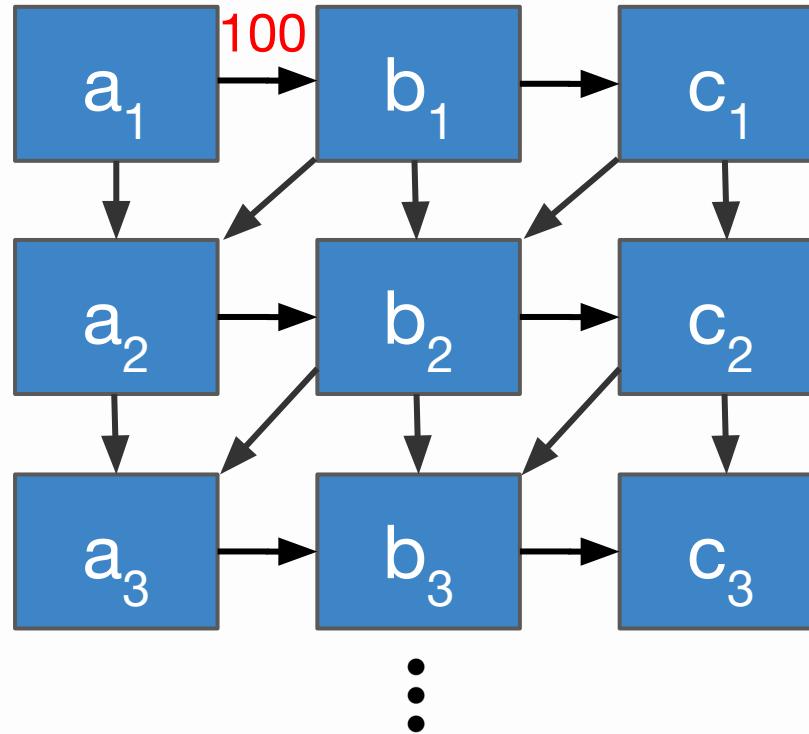


Key Trace Graph Property

$$\begin{aligned}s(b_1) &= 100 \\ s(a_1) &= 0\end{aligned}$$

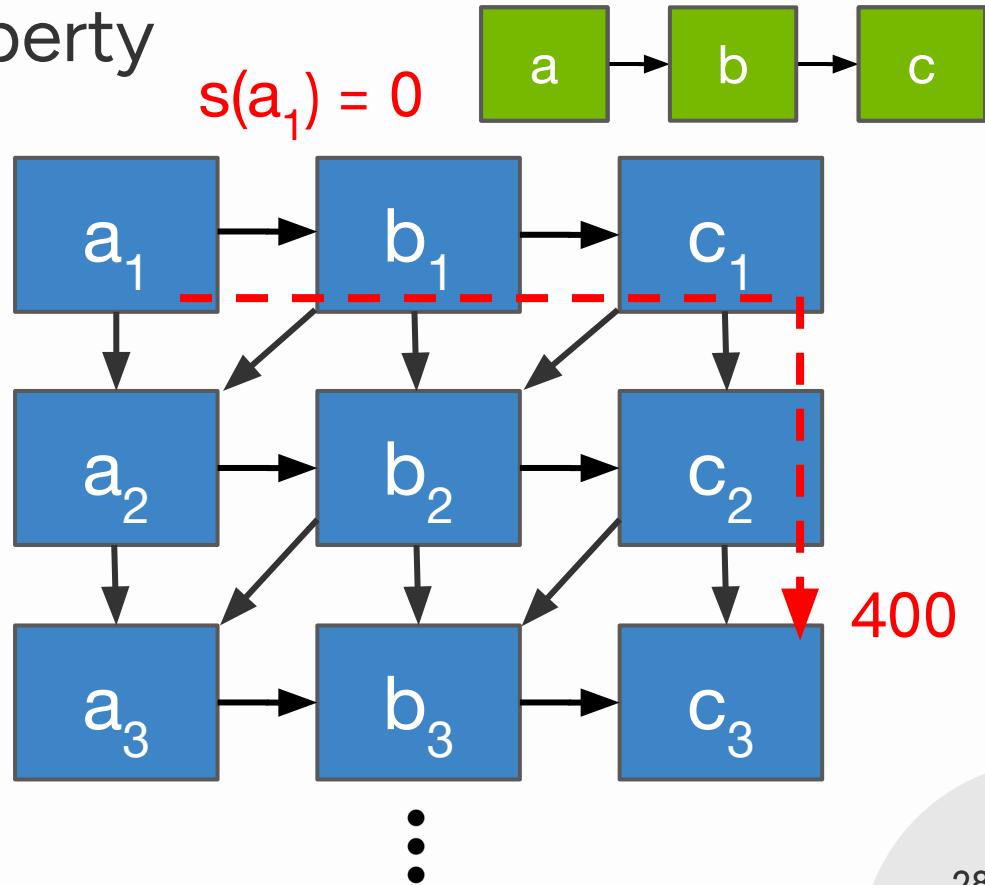


- Longest path from source (a_1) defines a node's start time
 - Preconditions must be met before node may execute



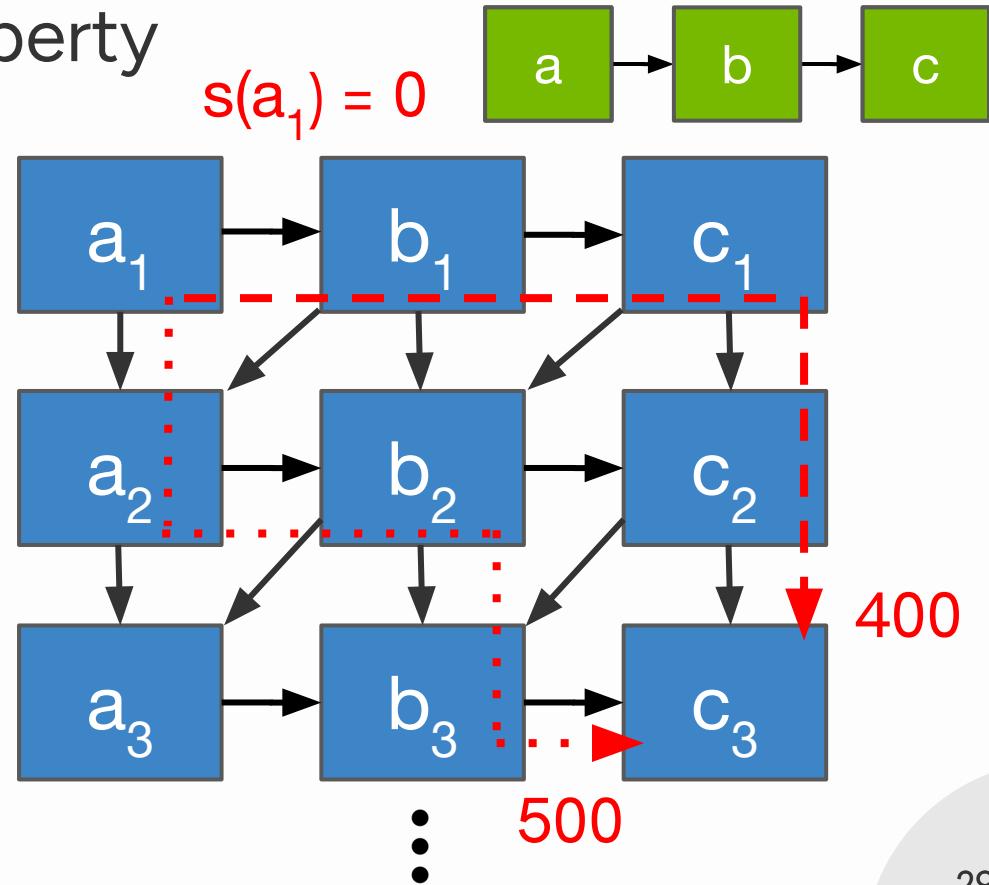
Key Trace Graph Property

- Longest path from source (a_1) defines a node's start time
 - Preconditions must be met before node may execute



Key Trace Graph Property

- Longest path from source (a_1) defines a node's start time
 - Preconditions must be met before node may execute

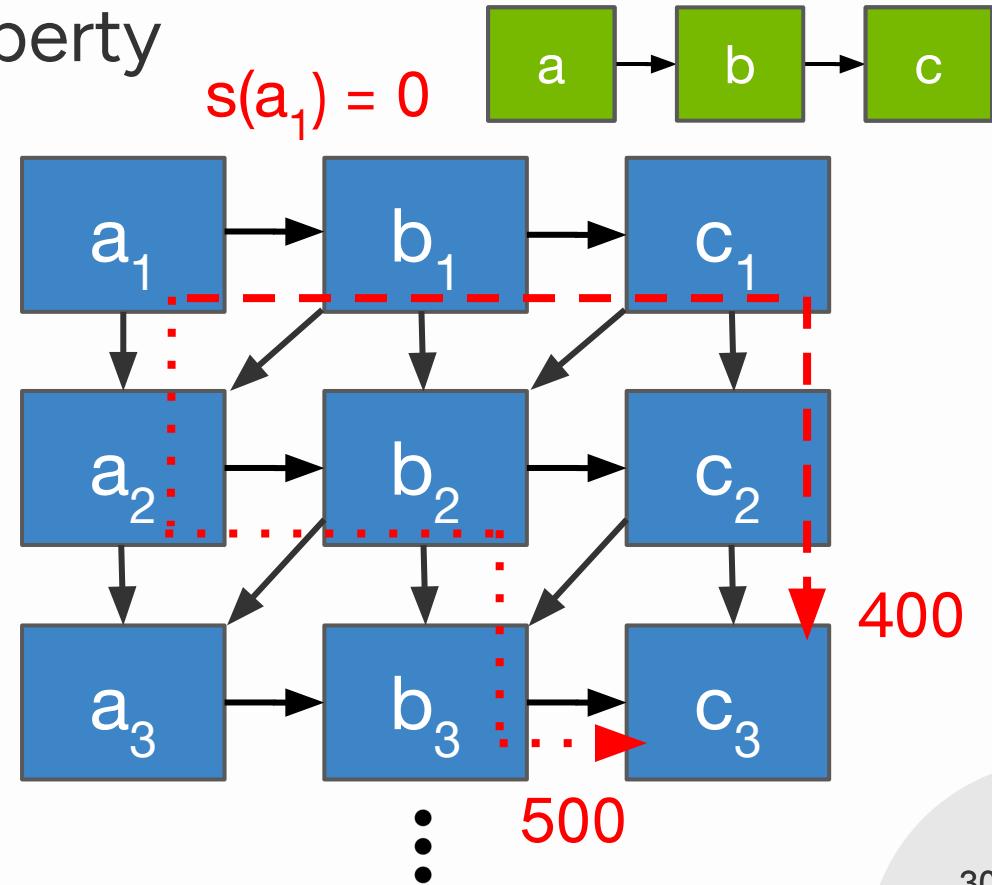


Key Trace Graph Property

- Longest path from source (a_1) defines a node's start time
 - Preconditions must be met before node may execute

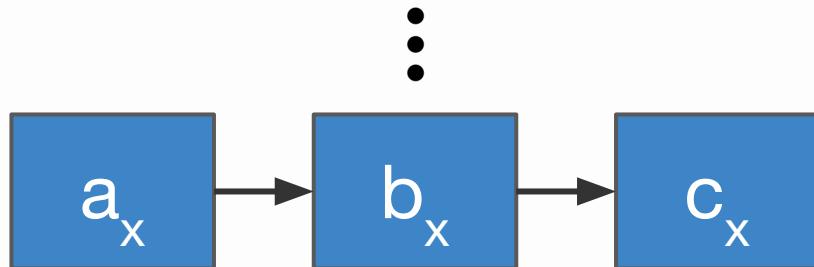
$$s(c_3) = 500$$

$$s(a_1) = 0$$



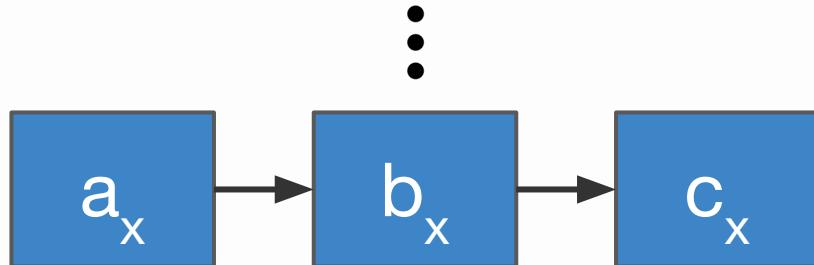
How can we leverage the trace graph for a response-time bound?

Key Idea



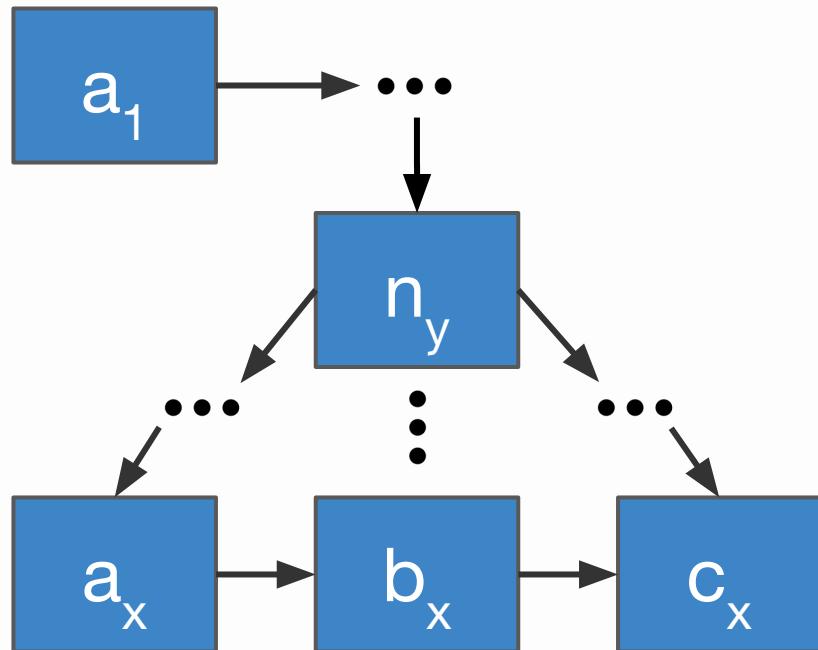
Assume the WCRT happens in iteration x

Key Idea



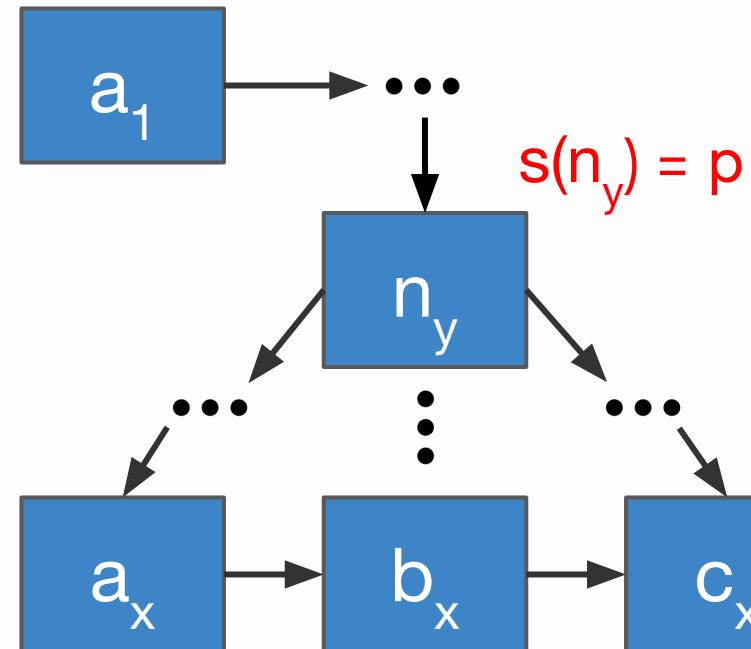
We know a_1 is somewhere in the past

Key Idea



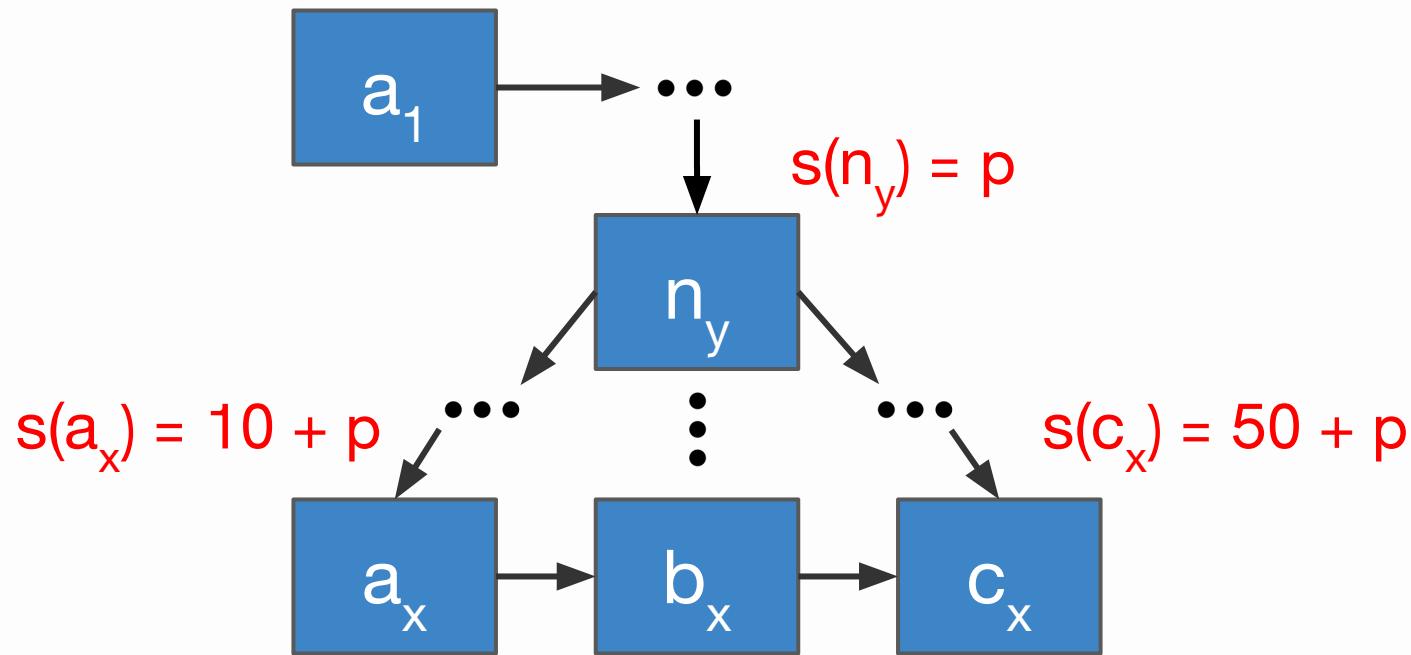
Find an n_y and assume on longest paths to a_x and c_x

Key Idea



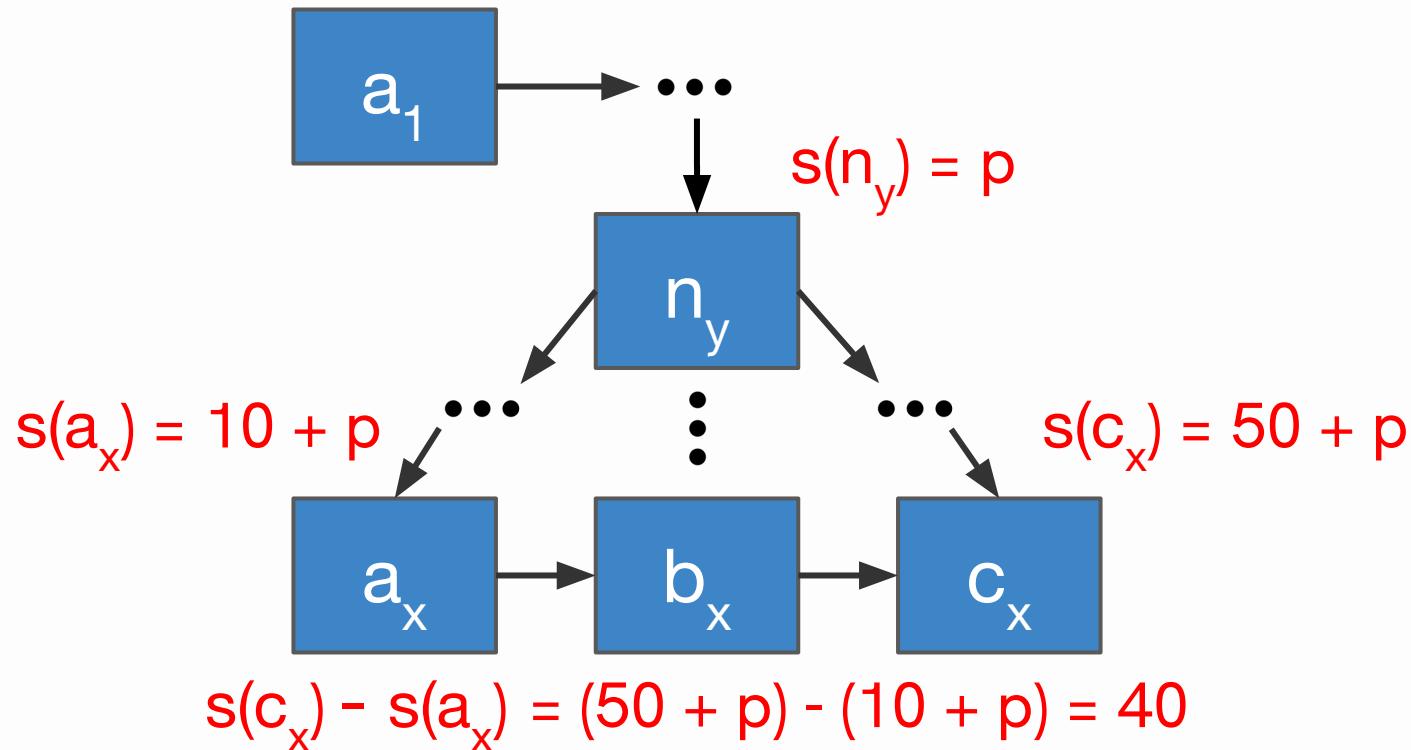
Longest path from a_1 to n_y can have any value

Key Idea



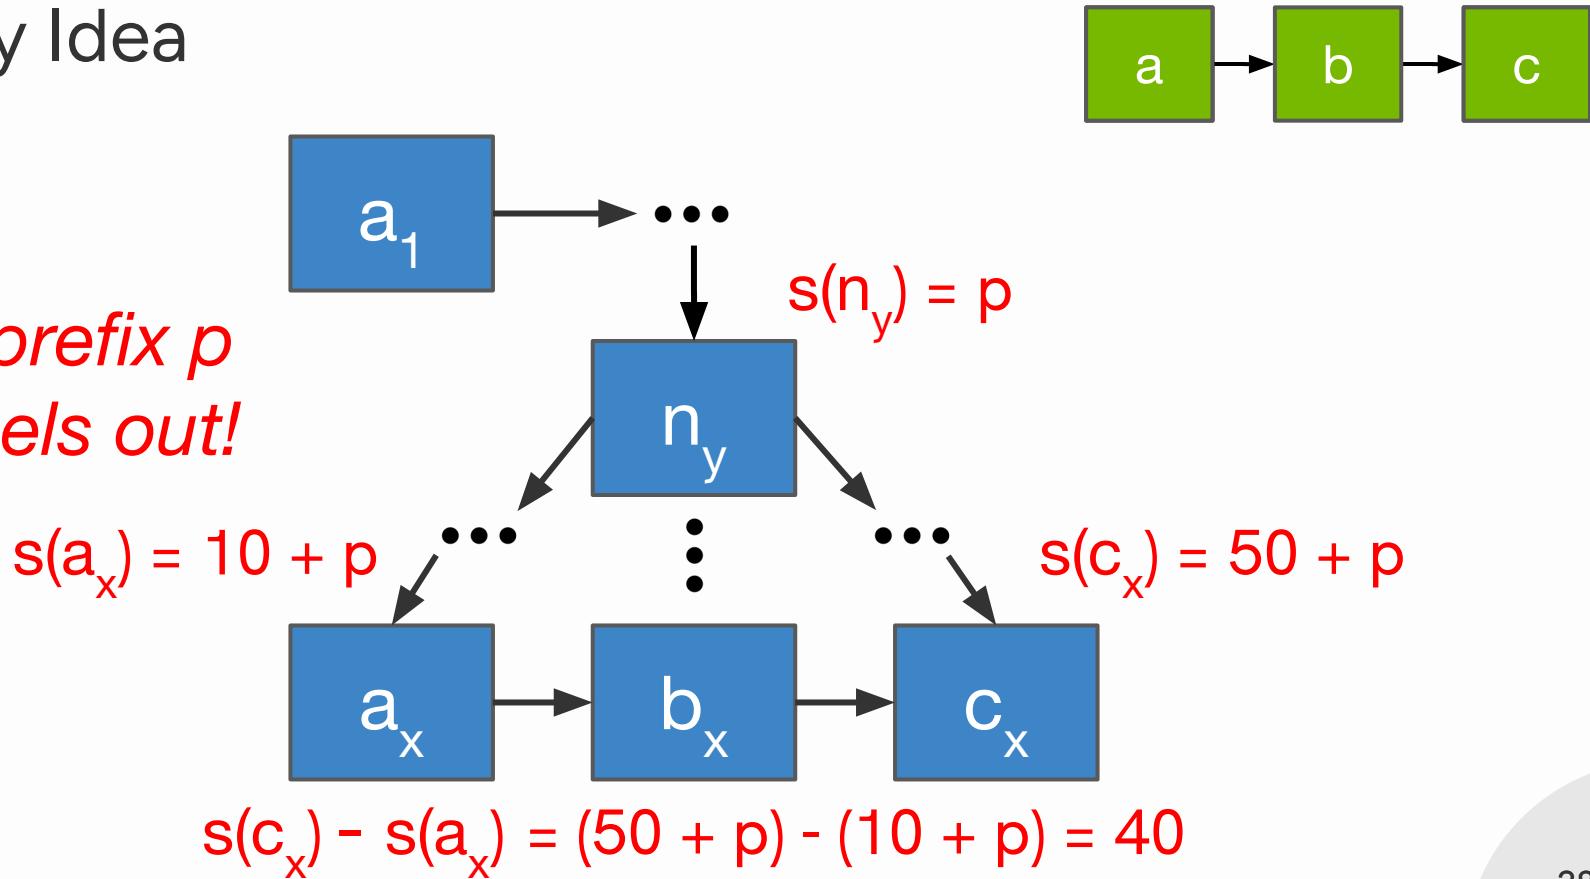
Compute path costs from n_y to a_x and c_x

Key Idea



Key Idea

The prefix p cancels out!



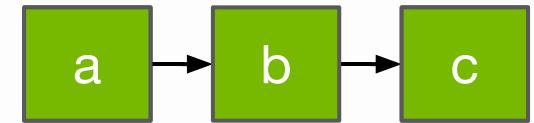
How can we leverage this insight to get a simple response-time algorithm?

- High-level overview: Start from an arbitrary iteration x , backtrack to find shared ancestors of the iteration x source and sink, and take differences of the paths from ancestor to source
 - Loop over the set of most recent shared ancestors

Algorithm

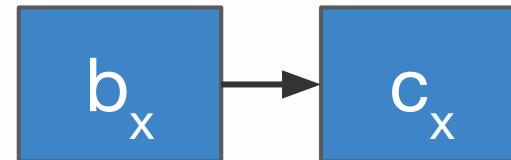
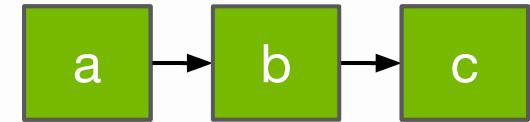
1. Find most recent shared ancestors
 - a. Lemma: all paths to c_x from a_1 have an ancestor of a_x within a bounded number of iterations from x

Search For Ancestors



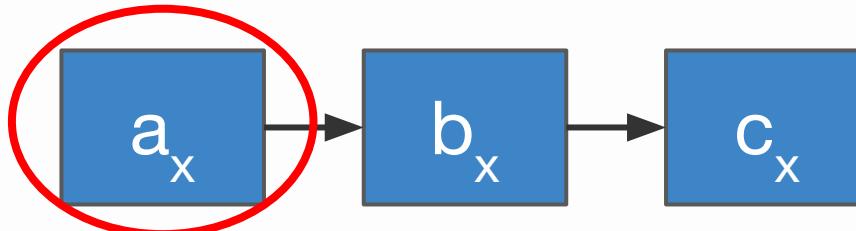
Search backwards from c_x for ancestors of a_x

Search For Ancestors



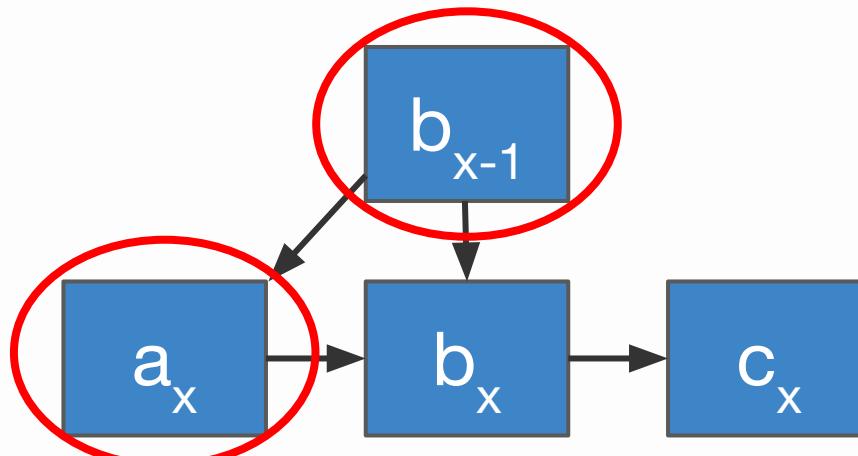
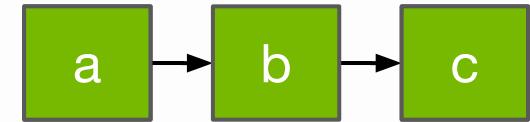
Search backwards from c_x for ancestors of a_x

Search For Ancestors



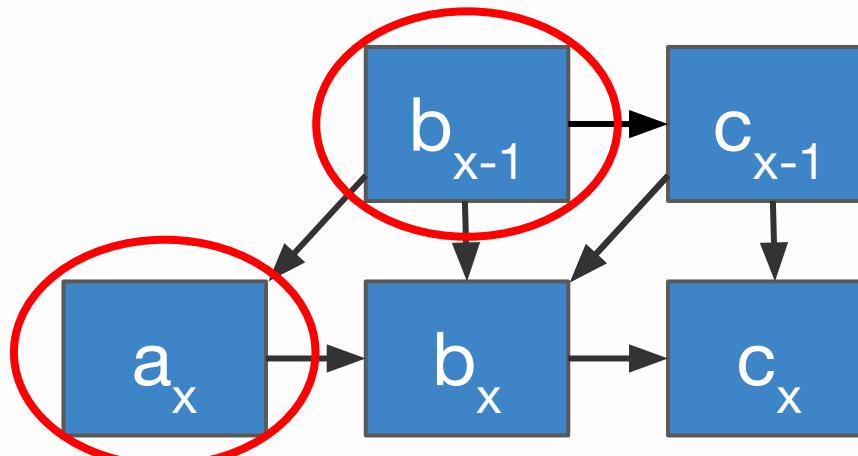
Search backwards from c_x for ancestors of a_x

Search For Ancestors



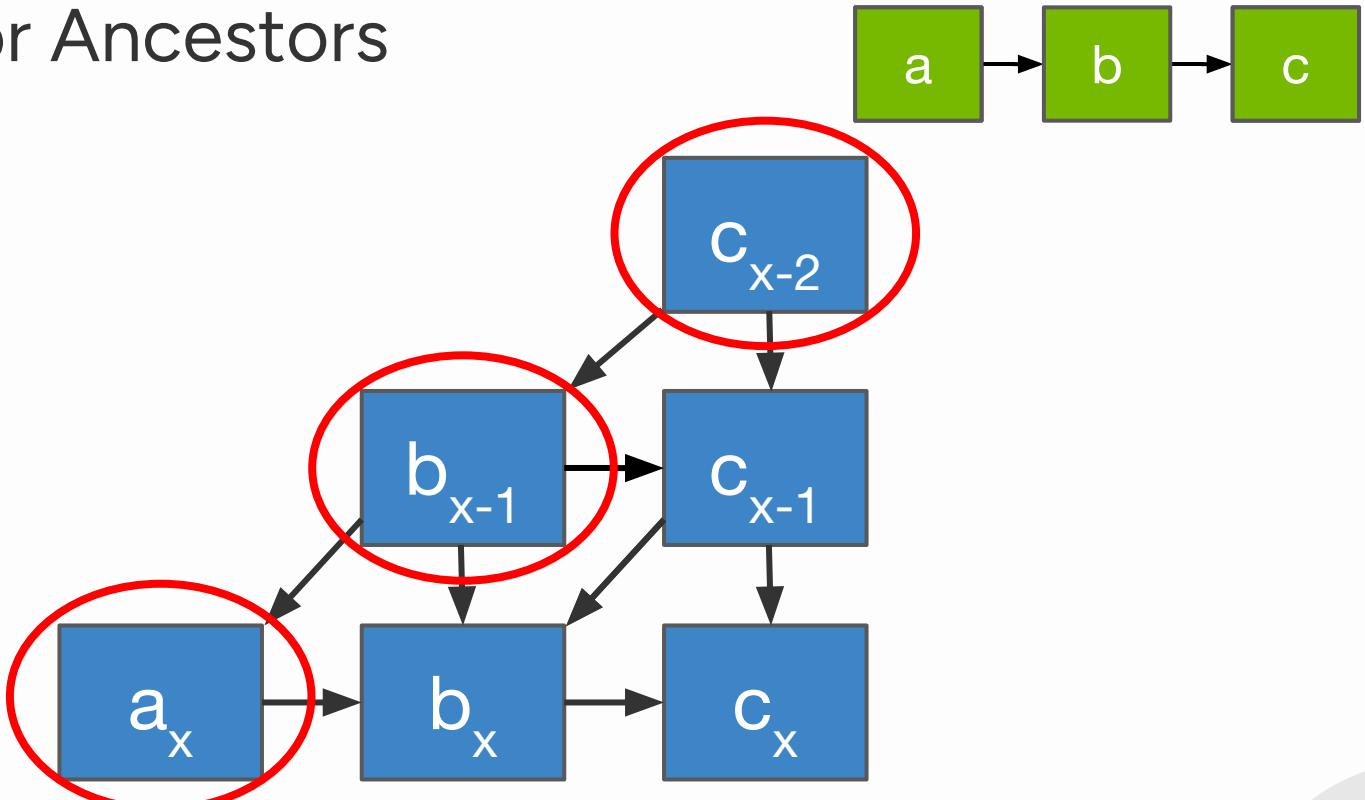
Search backwards from c_x for ancestors of a_x

Search For Ancestors



Search backwards from c_x for ancestors of a_x

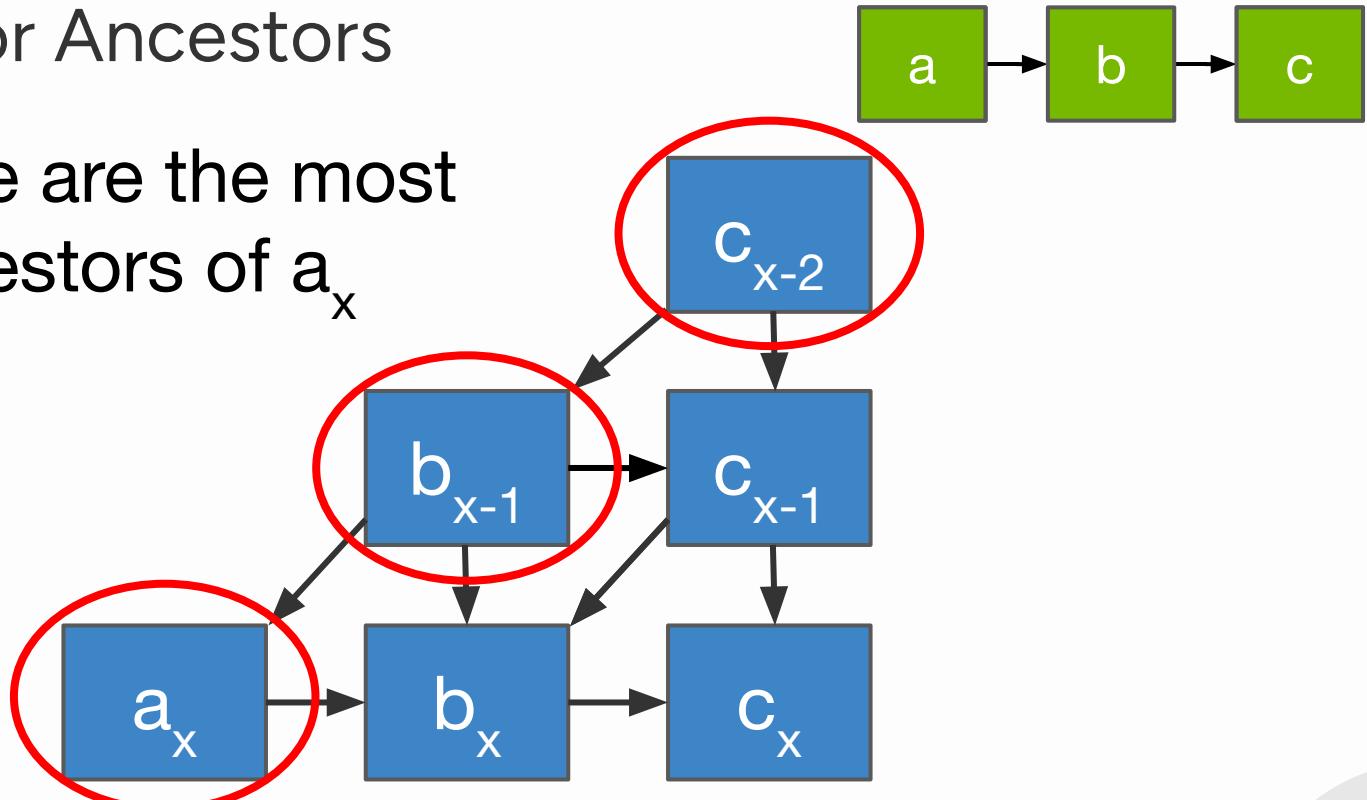
Search For Ancestors



Search backwards from c_x for ancestors of a_x

Search For Ancestors

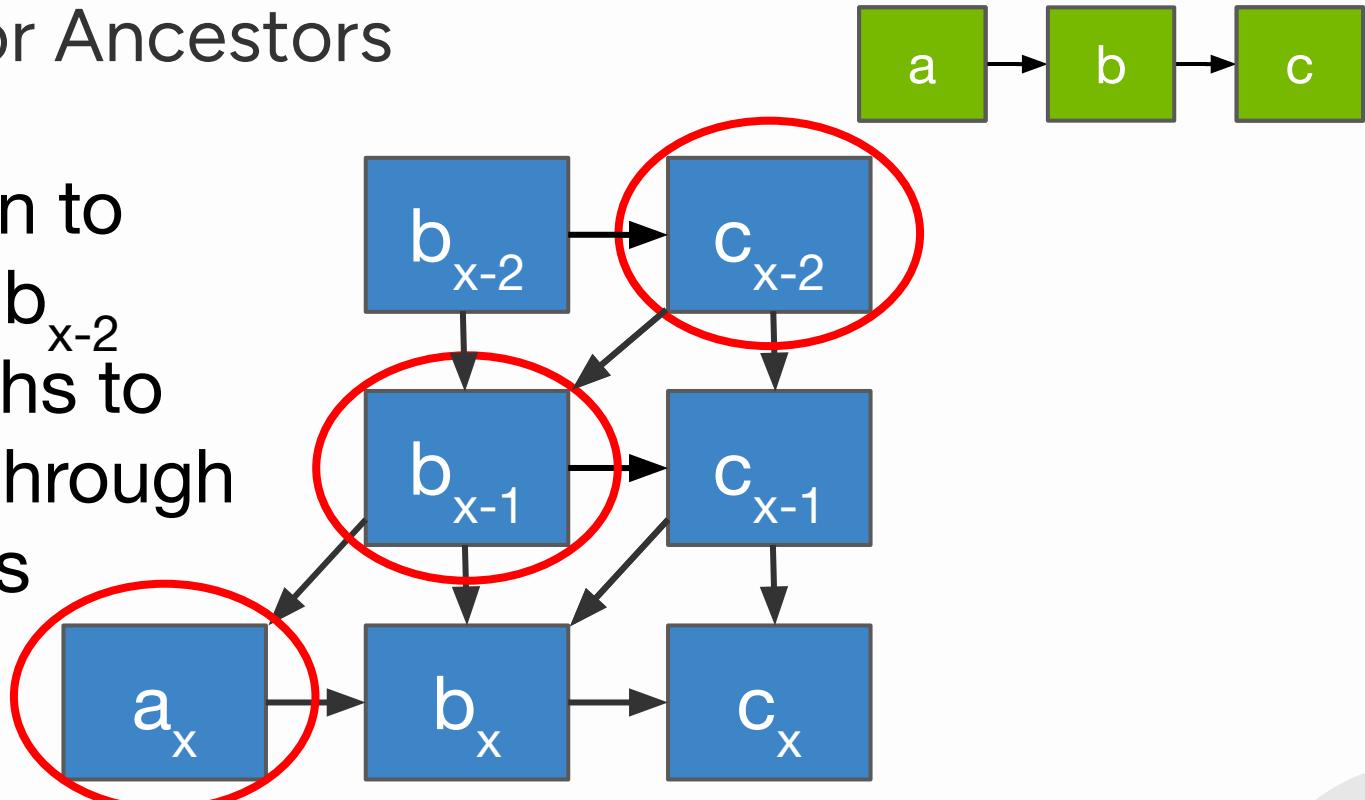
These three are the most recent ancestors of a_x



Search backwards from c_x for ancestors of a_x

Search For Ancestors

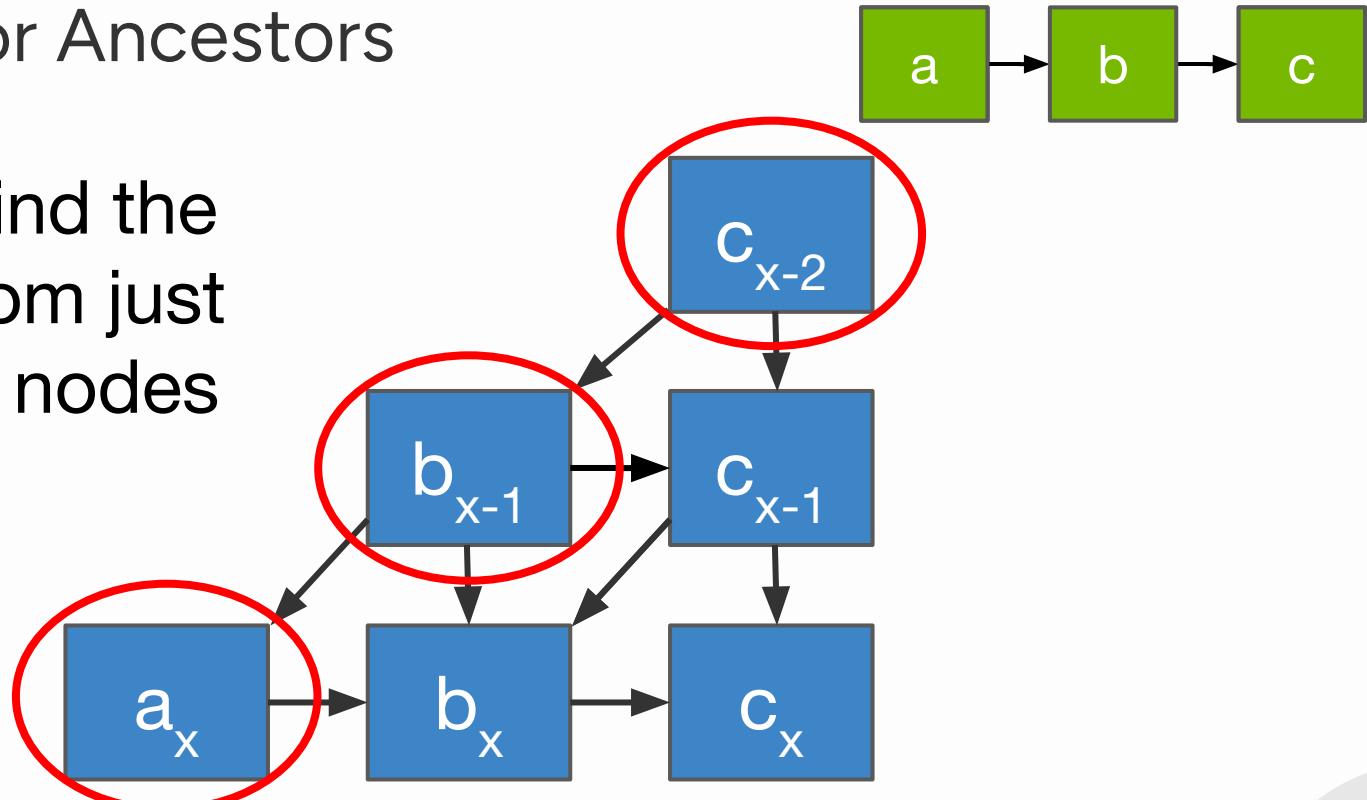
No reason to consider b_{x-2} as its paths to a_x/c_x go through the others



Search backwards from c_x for ancestors of a_x

Search For Ancestors

We can find the WCRT from just these six nodes

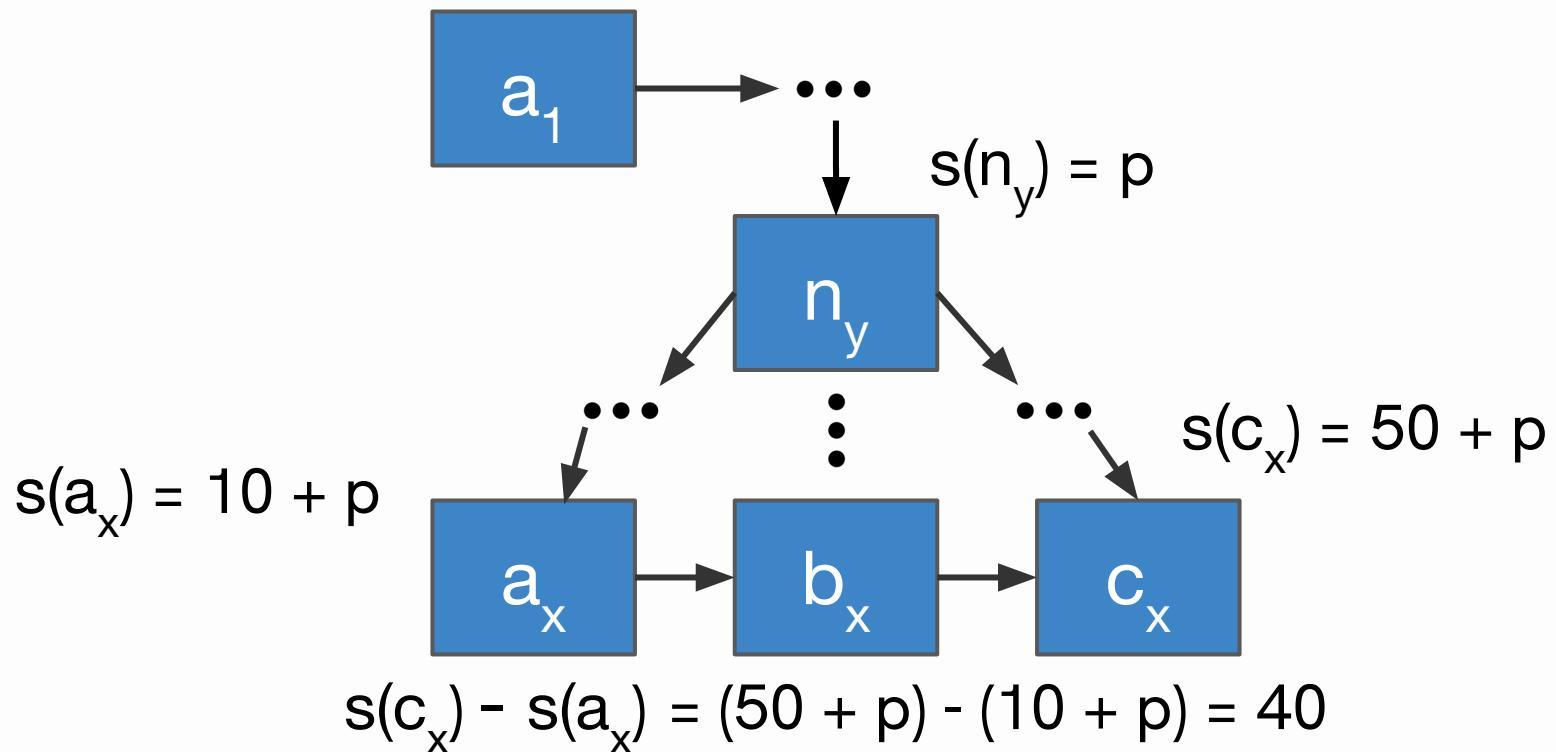


Search backwards from c_x for ancestors of a_x

Algorithm

1. Find most recent shared ancestors
 - a. Lemma: all paths to c_x from a_1 have an ancestor of a_x within a bounded number of iterations from x
2. Get the response time assuming each ancestor found in step 1 is on the longest path to c_x

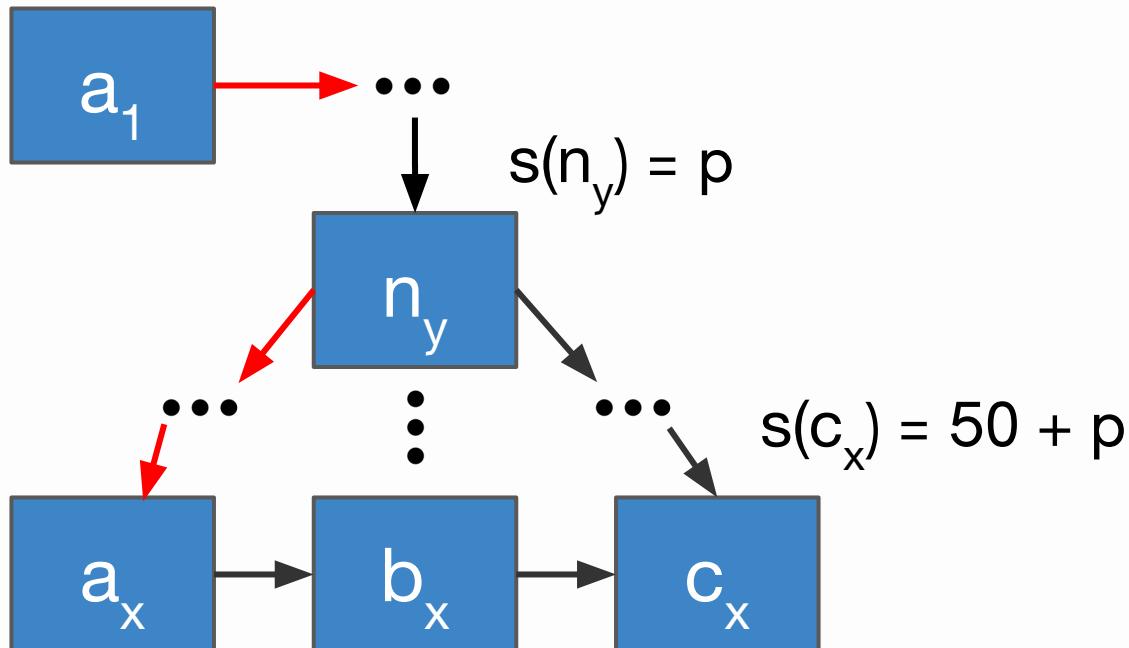
Response-Time Bound



Response-Time Bound

Assumed
longest path
to a_x came
from n_y !

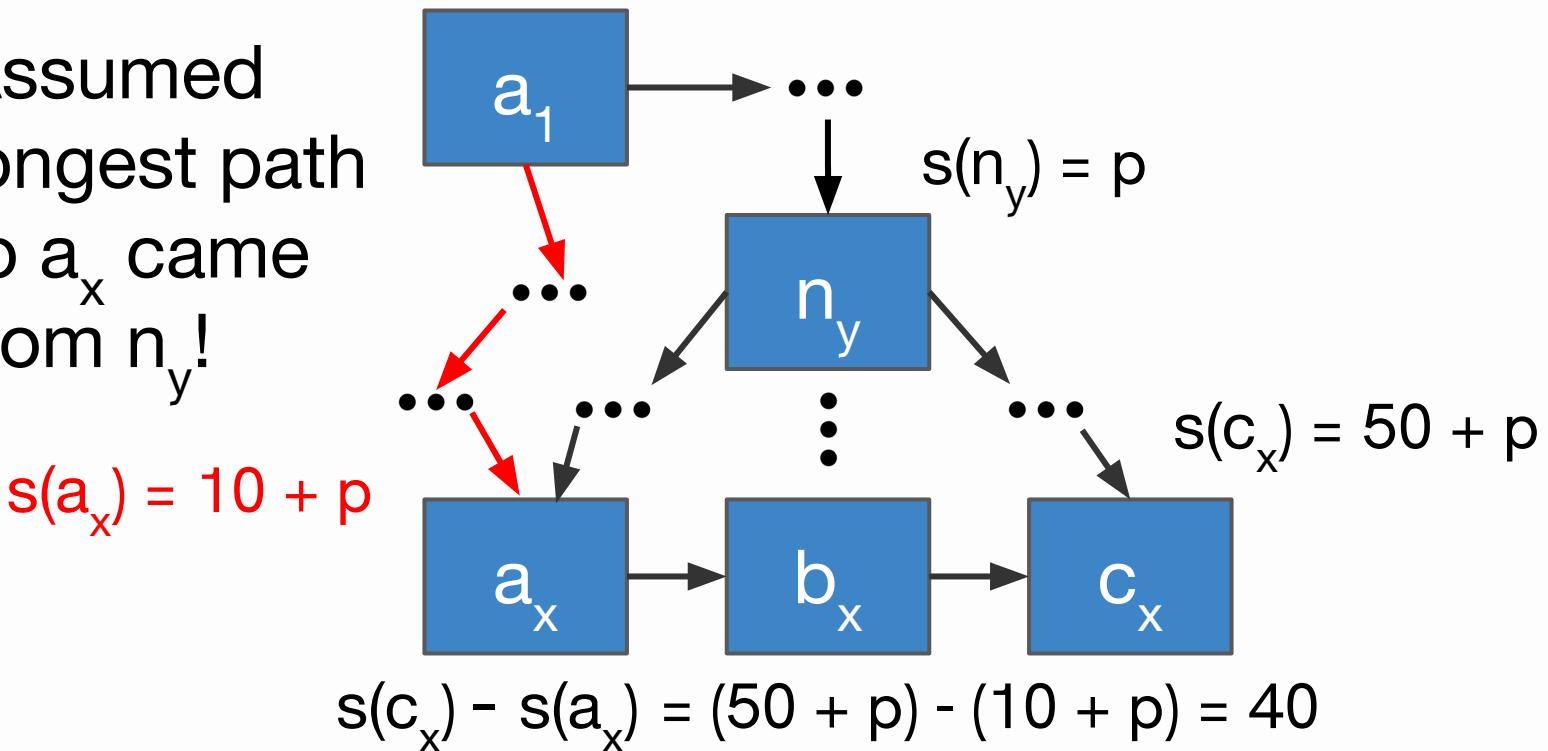
$$s(a_x) = 10 + p$$



$$s(c_x) - s(a_x) = (50 + p) - (10 + p) = 40$$

Response-Time Bound

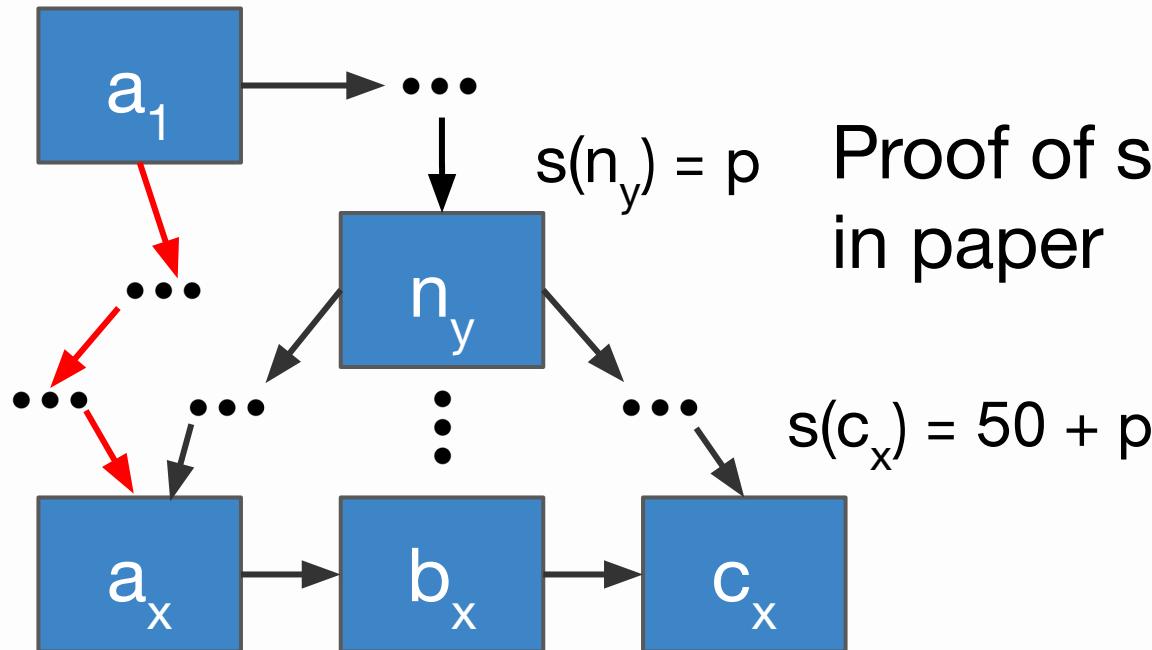
Assumed
longest path
to a_x came
from n_y !



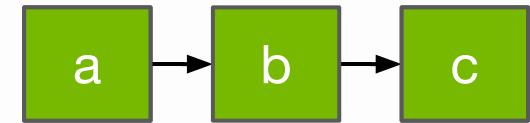
Response-Time Bound

Assumed
longest path
to a_x came
from n_y !

$$s(a_x) = 10 + p$$



$$s(c_x) - s(a_x) = (50 + p) - (10 + p) = 40$$



Proof of safety
in paper

Algorithm

1. Find most recent shared ancestors
 - a. Lemma: all paths to c_x from a_1 have an ancestor of a_x within a bounded number of iterations from x
2. Get the response time assuming each ancestor found in step 1 is on the longest path to c_x
3. Max of all candidate response times is WCRT

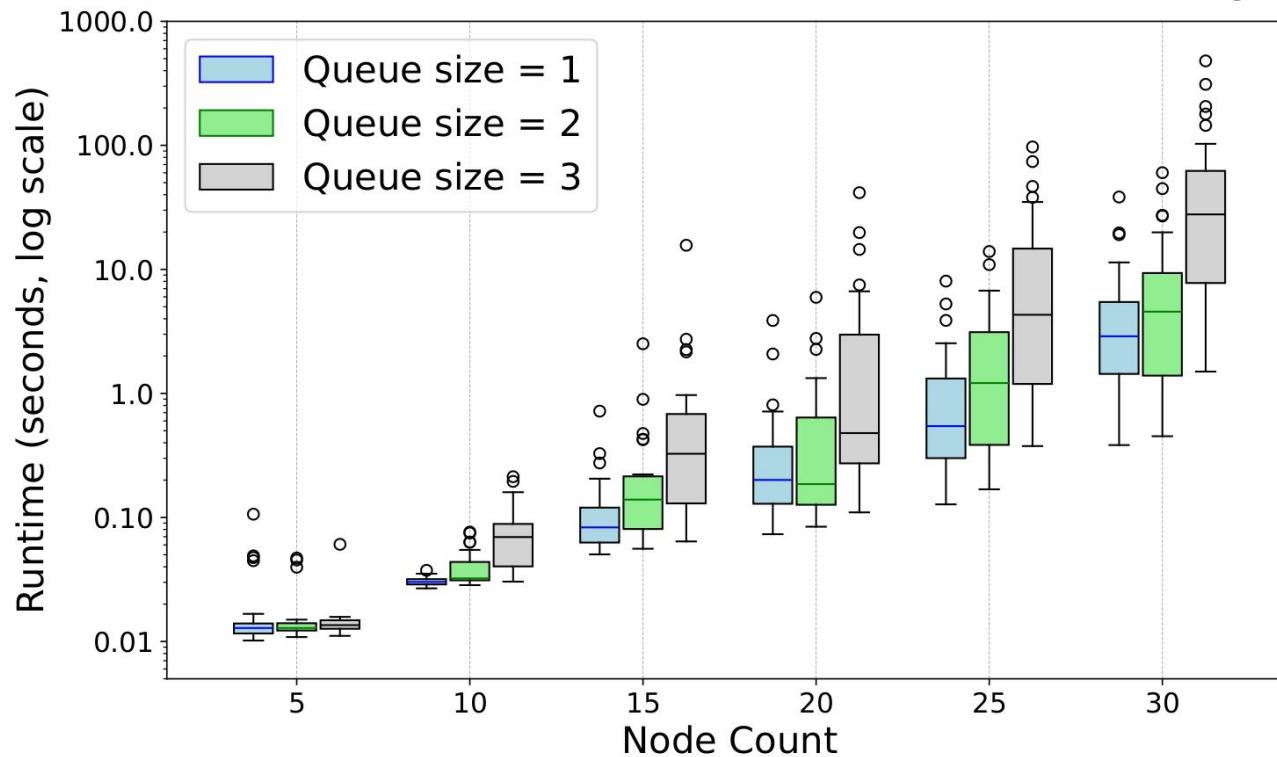
Exactness

- Our response-time bound is also exact (proof in paper)
 - We create a configuration of valid edge costs that yield the same response-time as our bound (needs 0 execution times)
- Supports variable execution time
- Supports variable queue sizes
 - This and other modifications are possible by altering the rules of how edges are drawn between nodes

Scalability

UPPAAL: 6000 seconds
for 9 node graph

Unoptimized
algorithm





Conclusion

- Novel response-time analysis for DAGs with static backpressure: exact, fast, variable execution time
- Show equivalence between DAG model and SDFGs
 - Can solve problem via model checking, but this is slow
- WiP: Incorporating GPU interactions
- Future work: Scheduling



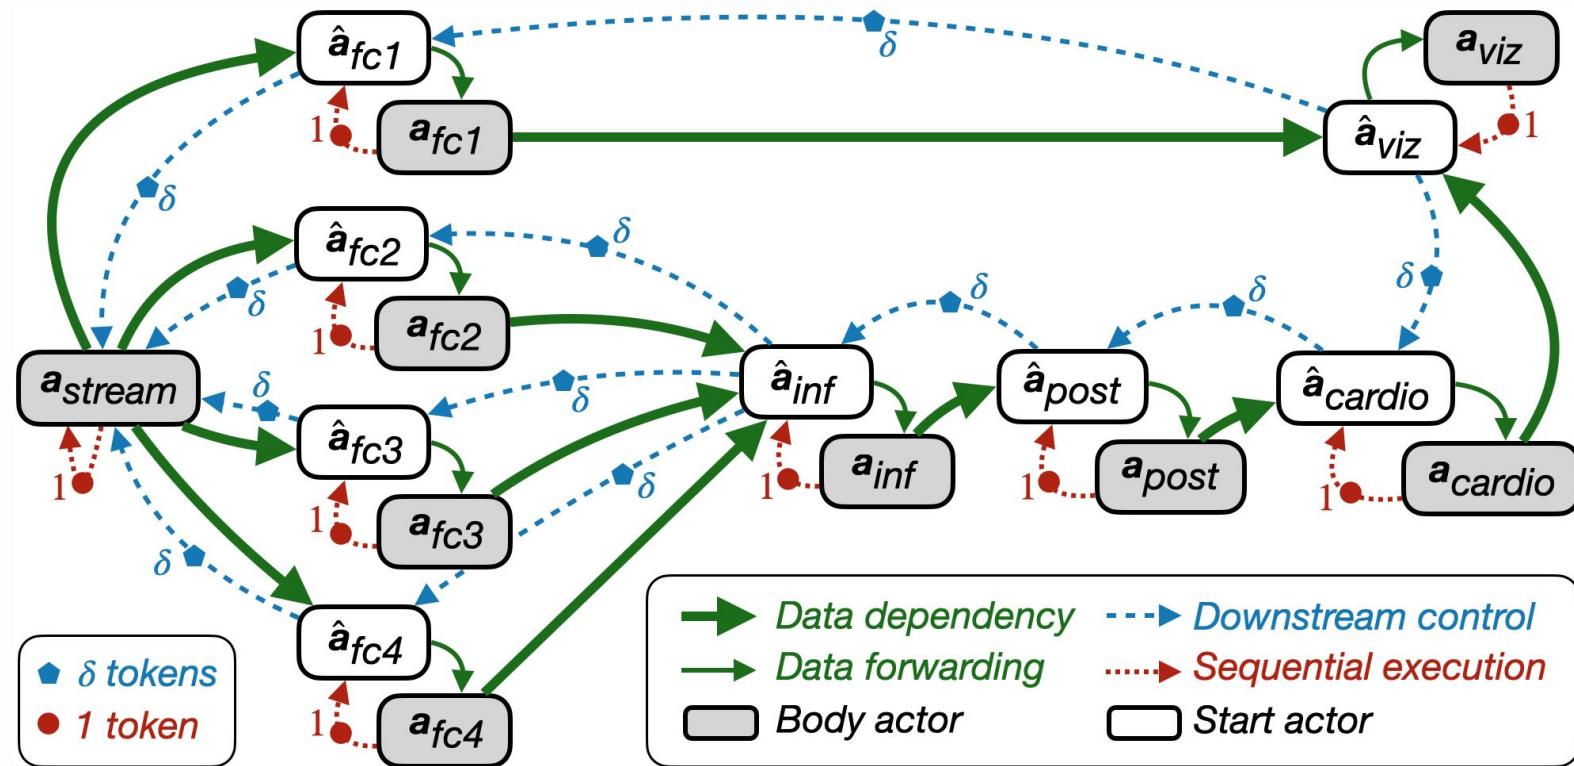
Conclusion

Thank you for listening! Questions?

- Novel response-time analysis for DAGs with static backpressure: exact, fast, variable execution time
- Show equivalence between DAG model and SDFGs
 - Can solve problem via model checking, but this is slow
- WiP: Incorporating GPU interactions
- Future work: Scheduling

Backup

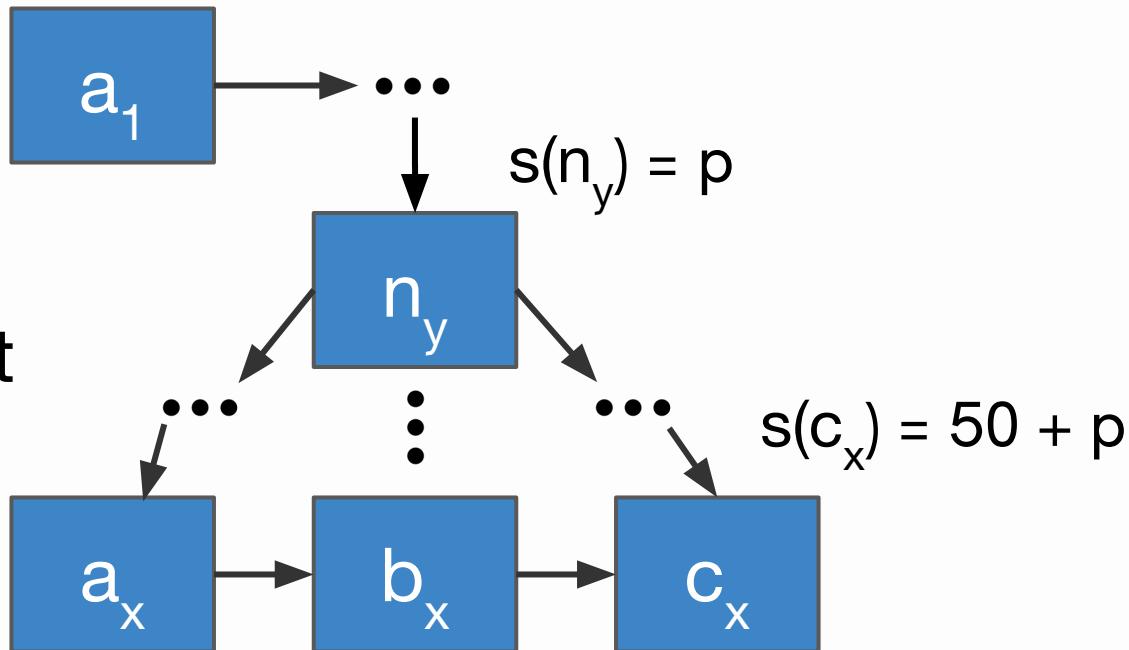
SDFG



Response-Time Bound

Two cases:
longest path
to a_x is either
from n_y or not

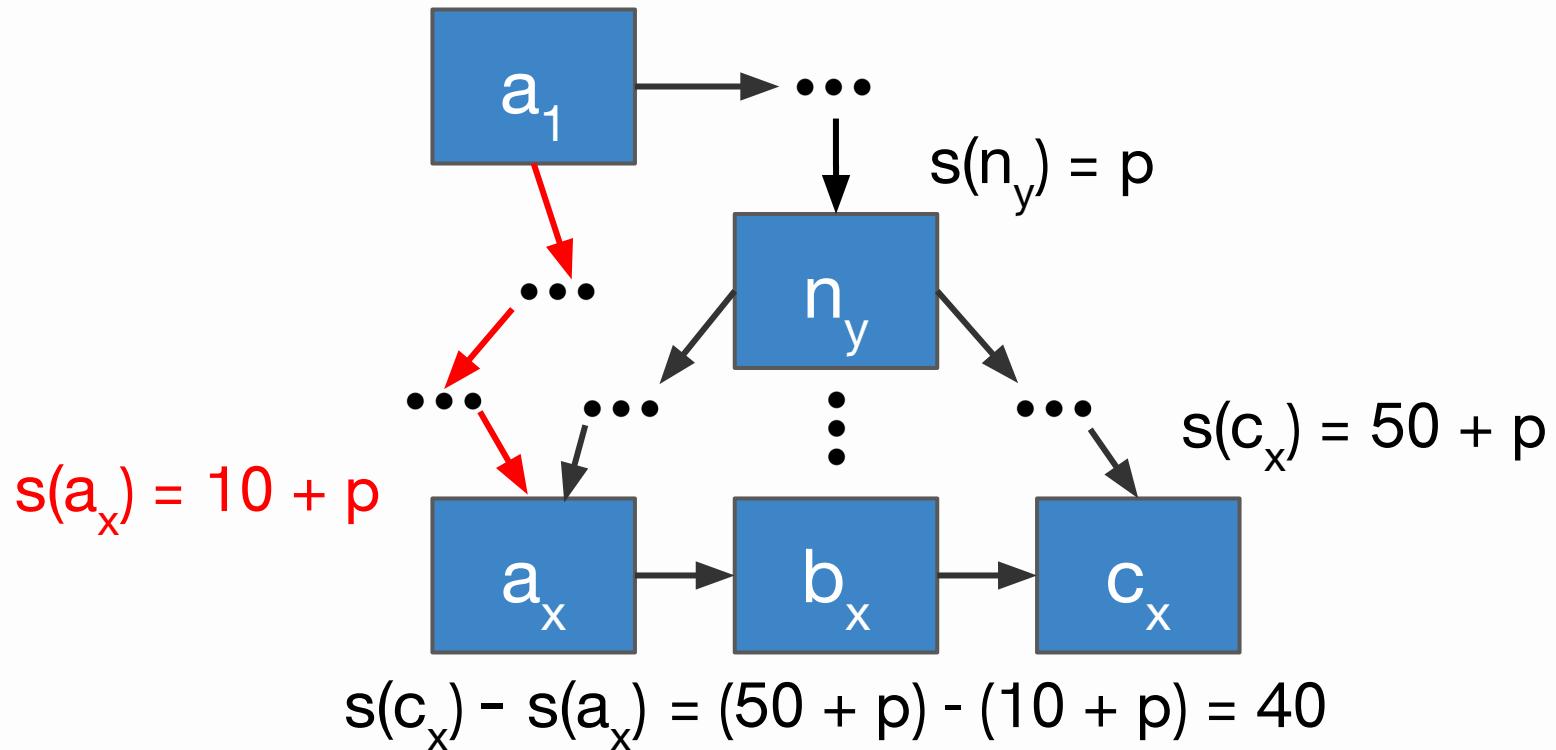
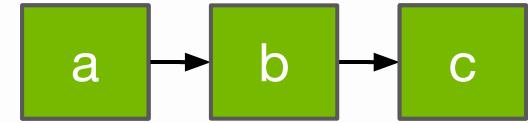
$$s(a_x) = 10 + p$$



$$s(c_x) - s(a_x) = (50 + p) - (10 + p) = 40$$

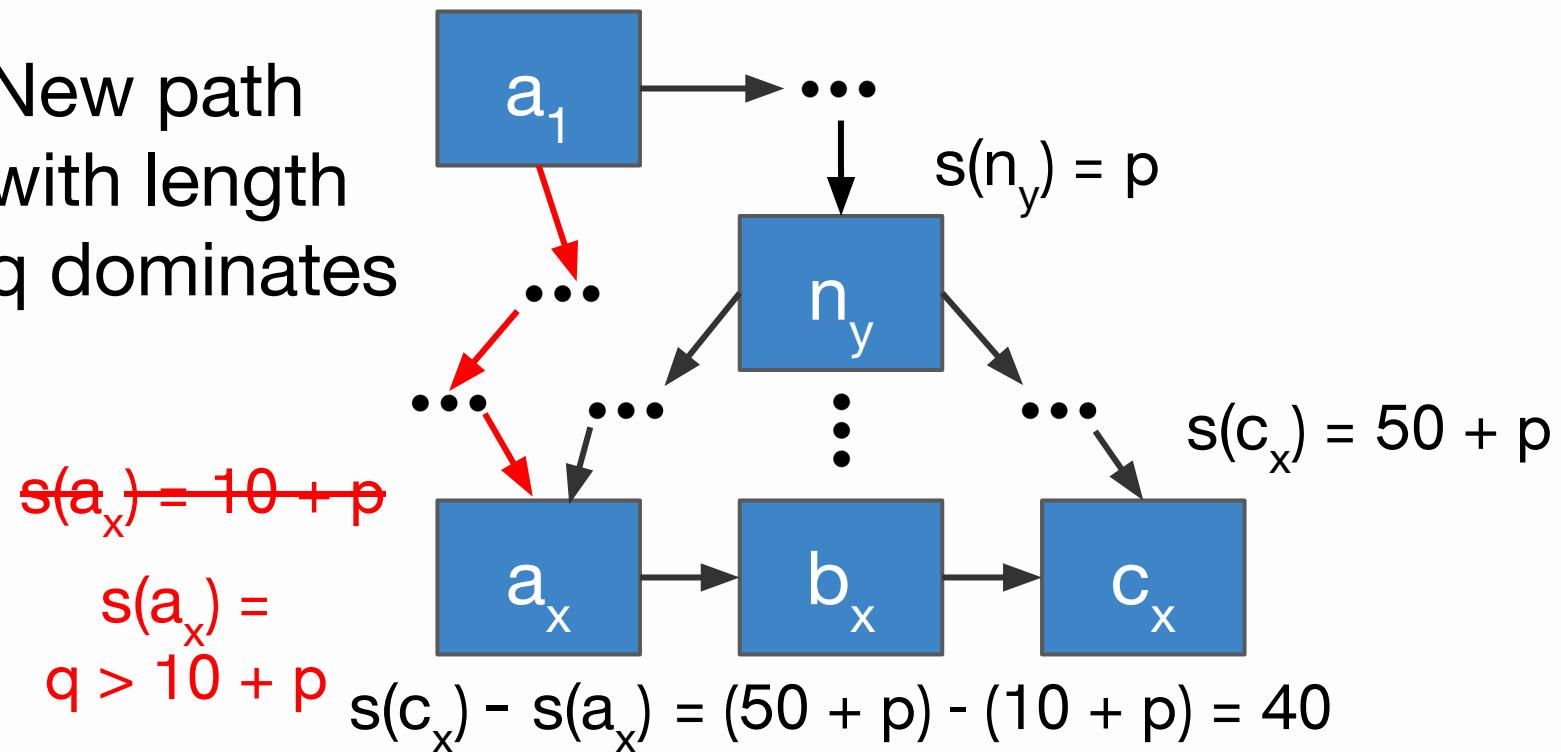


Response-Time Bound



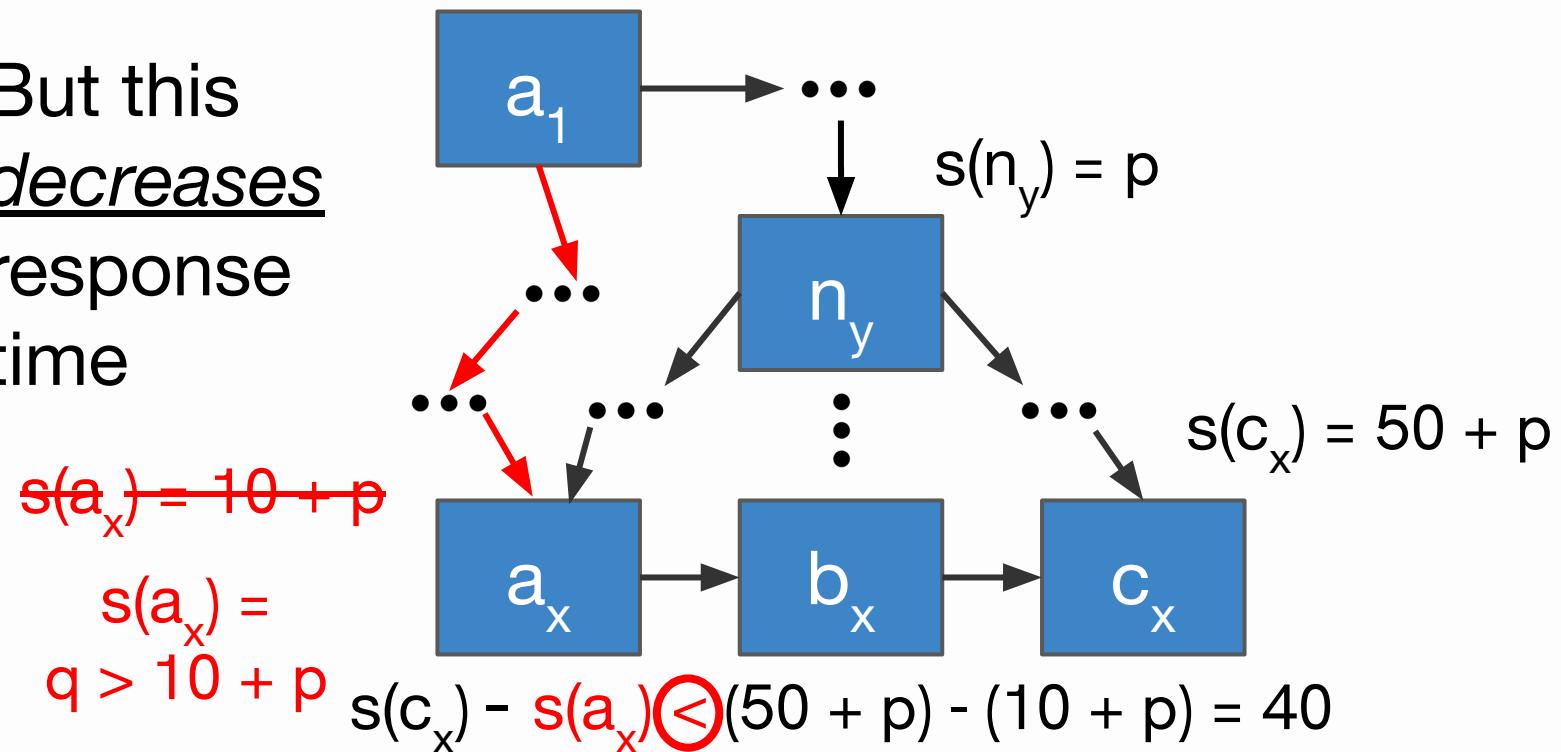
Response-Time Bound

New path
with length
 q dominates



Response-Time Bound

But this
decreases
response
time



Response-Time Bound

If n_y is on
longest path
to c_x :

$$RT \leq 40$$

$$s(a_x) = 10 + p$$

$$s(a_x) =$$

$$q > 10 + p$$

